ROSE Tutorial:
A Tool for Building
Source-to-Source Translators
Draft Tutorial
(version 0.9.3a)

Daniel Quinlan, Markus Schordan, Richard Vuduc, Qing Yi
Thomas Panas, Chunhua Liao, and Jeremiah J. Willcock
Lawrence Livermore National Laboratory
Livermore, CA 94550
925-423-2668 (office)
925-422-6278 (fax)
{dquinlan,panas2,liao6,willcock2}@llnl.gov
markus@complang.tuwien.ac.at
qingyi@cs.utsa.edu
richie@cc.gatech.edu
Project Web Page: www.rosecompiler.org
UCRL Number for ROSE User Manual: UCRL-SM-210137-DRAFT
UCRL Number for ROSE Tutorial: UCRL-SM-210032-DRAFT
UCRL Number for ROSE Source Code: UCRL-CODE-155962
ROSE User Manual (pdf)
ROSE Tutorial (pdf)
ROSE HTML Reference (html only)

October 14, 2008

October 14, 2008

Contents

List of Figures

Chapter 1

Introduction

1.1 Why you should be interested in ROSE

ROSE is a tool for building source-to-source translators. You should be interested in ROSE if you want to understand or improve any aspect of your software. ROSE makes it easy to build tools that read and operate on source code from large scale applications (millions of lines). Whole projects may be analyzed and even optimized using tools built using ROSE.

To get started immediately consult the ROSE User Manual, chapter *Getting Started* for details).

1.2 Problems that ROSE can address

ROSE is a mechanism to build source-to-source analysis or optimization tools that operate directly on the source code of large scale applications. Example tools that *have* been built include:

- Array class abstraction optimizer,
- Source-to-source instrumenter,
- Loop analyzer,
- Symbolic complexity analyzer,
- Code coverage tools,
- Inliner and outliner,
- OpenMP translator,
- and many more.

Example tools that *can* be built include:

- Custom optimization tools,

- Custom documentation generators,
- Custom analysis tools,
- Code pattern recognition tools,
- Security analysis tools,
- and many more.

1.3 Examples in this ROSE Tutorial

This tutorial lays out a set of progressively complex example programs that serve as a tutorial for the use of ROSE. Translators built using ROSE can either just analyze (and output results) or compile the input programs just like a compiler (generating object files or executables). Many of the examples in this tutorial just do simple analysis of the input source code, and a few show the full compilation of the input source code. Where the translators generate either object files of executables, the vendor's compiler is used to compile the final ROSE-generated code. Within ROSE, the call to generate source code from the AST and call the vendor's compiler is referred to as the *backend processing*. The specification of the vendor's compiler as a backend is done within the configuration step within ROSE (see options for `configure` in the ROSE User Manual).

Within the example programs below, the user can provide alternative input programs for more complex evaluation of the tutorial examples and ROSE. The end of the chapter, section 36, shows the makefiles required to compile the tutorial programs using an installed version of ROSE (compiled using `make install`). This `example_makefile` is run as part of the testing using the `make installcheck` rule.

Nearly all the chapters in this tutorial introduce a different tutorial example. Specific chapters in this tutorial include:

1. Introduction *(this chapter)*
2. Getting Started
   This chapter covers where to find ROSE documentation and how to install ROSE.
3. AST Graph Generator
   This translator reads a C or C++ application code and builds the AST, internally. The translator does not regenerate code from the AST and so does not call the backend vendor's compiler. This shows how simple it could be to build source code analysis tools; the code calls an internal ROSE function to generate a dot graph of the AST, the makefile has the details of converting the dot graph into a postscript file (also shown).
4. AST PDF Generator
   This translator reads an C or C++ application code builds the AST internally. The translator does not regenerate code from the AST and so does not call the backend vendor's compiler. This shows how simple it could be to build source code analysis tools, the code calls an internal ROSE function to generate a `pdf` file with bookmarks representing the AST. The pdf file show as output is in this case a previously generated figure of a screen shot obtained by viewing the output pdf file using `acroread`.

5. Identity Translator
   This example translator reads a C or C++ application, builds the AST internally, generates source code from the AST (unparsing), and calls the backend vendor compiler to compile the generated C or C++ application code. Thus the translator acts like and can be used to replace any compiler since it takes in source code and outputs an object code (or executable). This example also shows that the output generated from and ROSE translator is a close reproduction of the input; preserving all comments, preprocessor control structure, and most formating.

6. AST Query Library
   This example translator shows the use of the AST query library to generate a list of function declarations for any input program (and output the list of function names). It can be trivially modified to return a list of any IR node type (C or C++ language construct).

7. Introduction to AST Traversals and Attributes
   This collection of examples show the use of the simple visitor pattern for the traversal of the AST within ROSE. The simple visitor pattern permits operations to be programmed which will be invoked on different nodes within the AST. To handle communication of context information down into the AST and permit communication of analysis information up the AST, we have provided inherited and synthesized attributes (respectively). Note that an AST is most often represented as a tree with extra edges and with shared IR nodes that make the full graph (representing all edges) not a tree. We present two styles of traversal, one over the tree representing the AST (which excludes some types of IR nodes) and one over the full AST with all extra nodes and shared nodes. Extra nodes are nodes such as SgType and SgSymbol IR nodes.

   (a) AST traversals
       These traversals visit each node of the tree embedded within the AST (excluding shared SgType and SgSymbol IR nodes). These traversals visit the IR nodes is an order dependent upon the structure of the AST (the source code from which the AST is built).

       i. Classic Object-Oriented Visitor Patterns
          This example, classicObjectOrientedVisitorPatternMemoryPoolTraversal.C, show the use of a classic visitor patterns. At the moment this example uses the AST's memory pools as a basis but it is identical to a future traversal. The ROSE visitor Pattern (below) is generally more useful. The classic visitor pattern traversals are provided for completeness.

       ii. Visitor Traversal (visitorTraversal.C)
           Conventional visitor patterns without no attributes. This pattern can explicitly access global variables to provide the effect of accumulator attributes (using static data members we later show the handling of accumulator attributes).

       iii. Inherited Attributes (inheritedAttributeTraversal.C)
            Inherited attributes are used to communicate the context of any location within the AST in terms of other parent AST nodes.

       iv. Synthesized Attributes (synthesizedAttributeTraversal.C)
           Synthesized attributes are used to pass analysis results from the leaves of the AST to the parents (all the way to the root of the AST if required).

v. Accumulator Attributes (accumulatorAttributeTraversal.C)
Accumulator attributes permit the interaction of data within inherited attributes with data in synthesized attributes. In our example program we will show the use of accumulator attributes implemented as static data members. Accumulator attributes are a fancy name for what is essentially global variables (or equivalently a data structure passed by reference to all the IR nodes in the AST).

vi. Inherited and Synthesized Attributes (inheritedAndSynthesizedAttributeTraversal.C)
The combination of using inherited and synthesized attributes permits more complex analysis and is often required to compute analysis results on the AST within a specific context (e.g. number of loop nests of specific depth).

vii. Persistent Attributes (persistantAttributes.C)
Persistent attributes may be added the AST for access to stored results for later traversals of the AST. The user controls the lifetime of these persistent attributes.

viii. Nested traversals
Complex operations upon the AST can require many subordinate operations. Such subordinate operations can be accommodated using nested traversals. All traversals can operate on any subtree of the AST, and may even be nested arbitrarily. Interestingly, ROSE traversals may also be applied recursively (though care should be take using recursive traversals using accumulator attributes to avoid *over* accumulation).

(b) Memory Pool traversals
These traversals visit all IR nodes (including shared IR nodes such as SgTypes and SgSymbols). By design this traversal can visit ALL IR nodes without the worry of getting into cycles. These traversals are mostly useful for building specialized tools that operate on the AST.

i. Visit Traversal on Memory Pools
This is a similar traversal as to the Visitor Traversal over the tree in the AST.

ii. Classic Object-Oriented Visitor Pattern on Memory Pools
This is similar to the Classic Object-Oriented Visitor Pattern on the AST.

iii. IR node Type Traversal on Memory Pools
This is a specialized traversal which visits each type of IR node, but one one of each type of IR nodes. This specialized traversal is useful for building tools that call static member functions on each type or IR node. A number of memory based tools for ROSE are built using this traversal.

8. Database Support (dataBaseUsage.C)
This example shows how to use the optional (see `configure --help`) SQLite database to hold persistent program analysis results across the compilation of multiple files. This mechanism may become less critical as the only mechanism to support global analysis once we can support whole program analysis more generally within ROSE.

9. Recognizing loops within applications (loopRecognition.C)
This example program shows the use of inherited and synthesized attributes form a list of loop nests and report their depth. The inherited attributes are required to record when

the traversal is within outer loop and the synthesized attributes are required to pass the list of loop nests back up of the AST.

10. Getting the type parameters in function declaration (functionParameterTypes.C)
    This example translator builds a list to record the types used in each function. It shows an example of the sort of type information present within the AST. ROSE specifically maintains all type information.

11. Resolving overloaded functions (resolvingOverloadedFunctions.C – C++ specific)
    The AST has all type information pre-evaluated, particularly important for C++ applications where type resolution is required for determining function invocation. This example translator builds a list of functions called within each function, showing that overloaded function are fully resolved within the AST. Thus the user is not required to compute the type resolution required to identify which over loaded functions are called.

12. Getting template parameters to a templated class (templateParameters.C – C++ specific)
    All template information is saved within the AST. Templated classes and functions are separately instantiated as specializations, as such they can be transformed separately depending upon their template values. This example code shows the template types used the instantiate a specific templated class.

13. Simple Instrumentor Translator (simpleInstrumentor.C)
    This example modifies an the input application to place new code at the top and bottom of each block. The output is show with the instrumentation in place in the generated code.

14. Adding a variable declaration (addingVariableDeclaration.C)
    Here we show how to add a variable declaration to the input application. Perhaps we should show this in two ways to make it clear. This is a particularly simple use of the AST IR nodes to build an AST fragment and add it to the application's AST.

15. Adding a function (addingFunctionDeclaration.C)
    This example program shows the addition of a new function to the global scope. This example is a bit more involved than the previous example.

FIXME: *Consider a section on "Program Analysis" with subsections.*

16. Generating a CFG (buildCFG.C)
    This example shows the generation of a control flow graph within ROSE. The example is intended to be simple. Many other graphs can be built, we need to show them as well.

17. Generating a CG (buildCG.C)
    This example shows the generation of a call graph within ROSE.

18. Generating a CH (classHierarchyGraph.C)
    This example shows the generation of a class hierarchy graph within ROSE.

19. Call loop optimizer on set of loops (loopOptimization.C)
    This example program shows the optimization of a loop in C. This section contains several subsections each of which shows different sorts of optimizations. There are a large number of loop optimizations only two are shown here, we need to add more.

    (a) Optimization of Matrix Multiply

(b) Loop Fusion Optimizations

20. Calling the inliner (inlinerExample.C)
    This example shows the use of the inliner mechanism within ROSE. The function to be inlined in specified and the transformation upon the AST is done to inline the function where it is called and clean up the resulting code.

21. Calling the outliner (outlinerExample.C)
    This example shows the use of the outliner mechanism within ROSE. A segment of code is selected and a function is generated to hold the resulting code. Any required variables (including global variables) are passed through the generated function's interface. The outliner is a useful part of the empirical optimization mechanisms being developed within ROSE.

22. Program slicing (programSlicingExample.C)
    This example shows the interface to the program slicing mechanism within ROSE. Program slicing has been implemented to two ways within ROSE.

23. Recognition of Abstractions (recognitionOfAbstractions.C)
    This example program shows the automated recognition of classes and functions defined within libraries and used by applications. This is work with Brian White as part of general support for recognition and optimization of user-defined abstractions.

24. Generating Unique Names for Declarations (generatingUniqueNamesFromDeclaration.C)
    A recurring issue in the development of many tools and program analysis is the representation of unique strings from language constructs (functions, variable declarations, etc.). This example demonstrated support in ROSE for the generation of unique names. Names are unique across different ROSE tools and compilation of different files.

25. Scopes of Declarations (scopeInformation.C)
    This example shows the scopes represented by different IR nodes in the AST.

26. Symbol Table Handling (symbolTableHandling.C)
    This example shows how to use the symbol tables held within the AST for each scope.

27. AST File I/O (astFileIO_GenerateBinaryFile.C)
    This example demonstrates the file I/O for AST. This is part of ROSE support for whole program analysis.

28. Debugging Tips
    There are numerous methods ROSE provides to help debug the development of specialized source-to-source translators. This section shows some of the techniques for getting information from IR nodes and displaying it. Show how to use the PDF generator for AST's. This section may contain several subsections.

    (a) Generating the code representing any IR node

    (b) Displaying the source code position of any IR node

29. Command-line processing
    ROSE includes mechanism to simplify the processing of command-line arguments so that translators using ROSE can trivially replace compilers within makefiles. This example shows some of the many command-line handling options within ROSE and the ways in which customized options may be added.
    (a) Recognizing custom command-line options
    (b) Adding options to internal ROSE command-line driven mechanisms
30. Building custom graphs of program information
    The mechanisms used internally to build different graphs of program data is also made externally available. This section shows how new graphs of program information can be built or existing graphs customized.
31. Example Makefiles demonstrating the command lines to compile and link the example translators in this tutorial are found in `<ROSE_Compile_Tree>/tutorial/exampleMakefile`.

Other examples included come specifically from external collaborations and are more practically oriented. Each is useful as an example because each solves a specific technical problem. More of these will be included over time.

1. Tau Performance Analysis Instrumentation (tauInstrumenter.C)
   Tau currently uses an automate mechanism that modified the source code text file. This example shows the modification of the AST and the generation of the correctly instrumented files (which can otherwise be a problem when macros are used). This is part of collaborations with the Tau project.
2. Fortran promotion of constants to double precision (typeTransformation.C)
   Fortran constants are by default singe precision, and must be modified to be double precision. This is a common problem in older Fortran applications. This is part of collaborations with LANL to eventually automatically update/modify older Fortran applications.
3. Automated Runtime Library Support (charmSupport.C)
   Getting research runtime libraries into use within large scale applications requires automate mechanism to make minor changes to large amounts of code. This is part of collaborations with the Charm++ team (UIUC).
4. Code Coverage Analysis (codeCoverage.C)
   Code coverage is a useful tool by itself, but is particularly useful when combined with automated detection of bugs in programs. This is part of work with IBM, Haifa.
   (a) Shared Threaded Variable Detection Instrumentation (interveneAtVariables.C)
   Instrumentation support for variables, required to support detection of threaded bugs in applications.
   (b) Automated Modification of Function Parameters (changeFunction.C)
   This example program addresses a common problem where an applications function must be modified to include additional information. In this case each function in a threaded library is modified to include additional information to a corresponding wrapper library which instruments the library's use.

FIXME: *Add make installcheck in Makefile.am to build example translators using the installed libraries.*

FIXME: *Comments at the top of each tutorial example should match the comments in the caption for the figures that introduce them in the text.*

Part I

Stable Parts of ROSE

These parts of ROSE are tested regularly on large scale applications and are suitable for use on ROSE projects requiring the greatest levels of robustness.

FIXME: *Lay out rules and criteria for the classification of different parts of ROSE.*

Chapter 2

Getting Started

2.1 ROSE Documentation and Where To Find It

There are three forms of documentation for ROSE:

1. ROSE User Manual
   The User Manual presents how to get started with ROSE and documents features of the ROSE infrastructure. The User Manual is found in ROSE/docs/Rose directory, or at:
   ROSE User Manual (postscript version, relative link)

2. ROSE Tutorial
   The ROSE Tutorial presents a collection of examples of how to use ROSE (found in the ROSE/tutorial directory). The ROSE Tutorial documentation is found in ROSE/docs/Rose/Tutorial directory. The tutorial documentation is built in the following steps:

   (a) actual source code for each example translator in the ROSE/tutorial directory is included into the tutorial documentation

   (b) each example is compiled

   (c) inputs to the examples are taken from the ROSE/tutorial directory

   (d) output generated from running each example is placed into the tutorial documentation

   Thus the `ROSE/tutorial` directory contains the exact examples in the tutorial and each example may be modified (changing either the example translators or the inputs to the examples). The ROSE Tutorial can also be found in the `ROSE/docs/Rose/Tutorial` directory (the LaTeX document; ps or pdf file)
   : ROSE Tutorial (postscript version, relative link),

3. ROSE HTML Reference: Intermediate Representation (IR) documentation
   This web documentation presents the detail interfaces for each IR nodes (documentation generated by Doxygen). The HTML IR documentation is found in ROSE/docs/Rose directory (available as html only):
   ROSE HTML Reference (relative link)

4. ROSE Web Page
   The ROSE web pages are located at: http://www.rosecompiler.org

5. ROSE Email List
   The ROSE project maintains a mailing list (casc-rose *at* llnl *dot* gov). The email list is only used to get help and announce internal and external releases. Anyone who would like to be on the email list can send me email (dquinlan *at* llnl *dot* gov).

2.2 Using the Tutorial

First install ROSE (see ROSE User Manual, chapter *Getting Started* for details). Within the ROSE distribution at the top level is the `tutorial` directory. All of the examples in this documentation are represented there with Makefiles and sample input codes to the example translators.

Chapter 3

AST Graph Generator

What To Learn From This Example This example shows how to generate and visualize the AST from any input program. Each node of the graph in figure 3.3 shows a node of the Intermediate Representation (IR). Each edge shows the connection of the IR nodes in memory. The generated graph shows the connection of different IR nodes to form the AST. The generation of such graphs is appropriate for small input programs, chapter **??** shows a mechanism using PDF files that is more appropriate to larger programs (e.g. 100K lines of code). More information about generation of specialized AST graphs can be found in 35 and custom graph generation in 34.

```
// Example ROSE Translator: used within ROSE/tutorial

#include "rose.h"

int main( int argc, char * argv[] )
   {
  // Build the AST used by ROSE
     SgProject* project = frontend(argc,argv);

     generateDOT ( *project );
     return 0;
   }
```

Figure 3.1: Example source code to read an input program and generate an AST graph.

The program in figure 3.1 calls an internal ROSE function that traverses the AST and generates an ASCII file in `dot` format. Figure 3.2 shows an input code which is processed to generate a graph of the AST, generating a `dot` file. The `dot` file is then processed using `dot` to generate a postscript file 3.3 (within the `Makefile`). Note that a similar utility program already exists within ROSE/exampleTranslators (and includes a utility to output an alternative PDF representation (suitable for larger ASTs) as well). Figure 3.3 (../../..//tutorial/test.ps) can be found in the compile tree (in the tutorial directory) and viewed directly using ghostview or any postscript viewer to see more detail.

Figure 3.3 displays the individual C++ nodes in ROSE's intermediate representation (IR). Each circle represents a single IR node, the name of the C++ construct appears in the center of

```
// Templated class declaration used in template parameter example code
template <typename T>
class templateClass
   {
     public:
          int x;

          void foo(int);
          void foo(double);
   };

// Overloaded functions for testing overloaded function resolution
void foo(int);
void foo(double)
   {
     int x = 1;
     int y;

  // Added to allow non-trivial CFG
     if (x)
        y = 2;
     else
        y = 3;
   }
```

Figure 3.2: Example source code used as input to generate the AST graph.

the circle, with the edge numbers of the traversal on top and the number of child nodes appearing below. Internal processing to build the graph generates unique values for each IR node, a pointer address, which is displays at the bottom of each circle. The IR nodes are connected for form a tree, and abstract syntax tree (AST). Each IR node is a C++ class, see SAGE III reference for details, the edges represent the values of data members in the class (pointers which connect the IR nodes to other IR nodes). The edges are labeled with the names of the data members in the classes representing the IR nodes.

Use this first example e to explain the use of ader files (config.h and d the code to build the SgProject object.

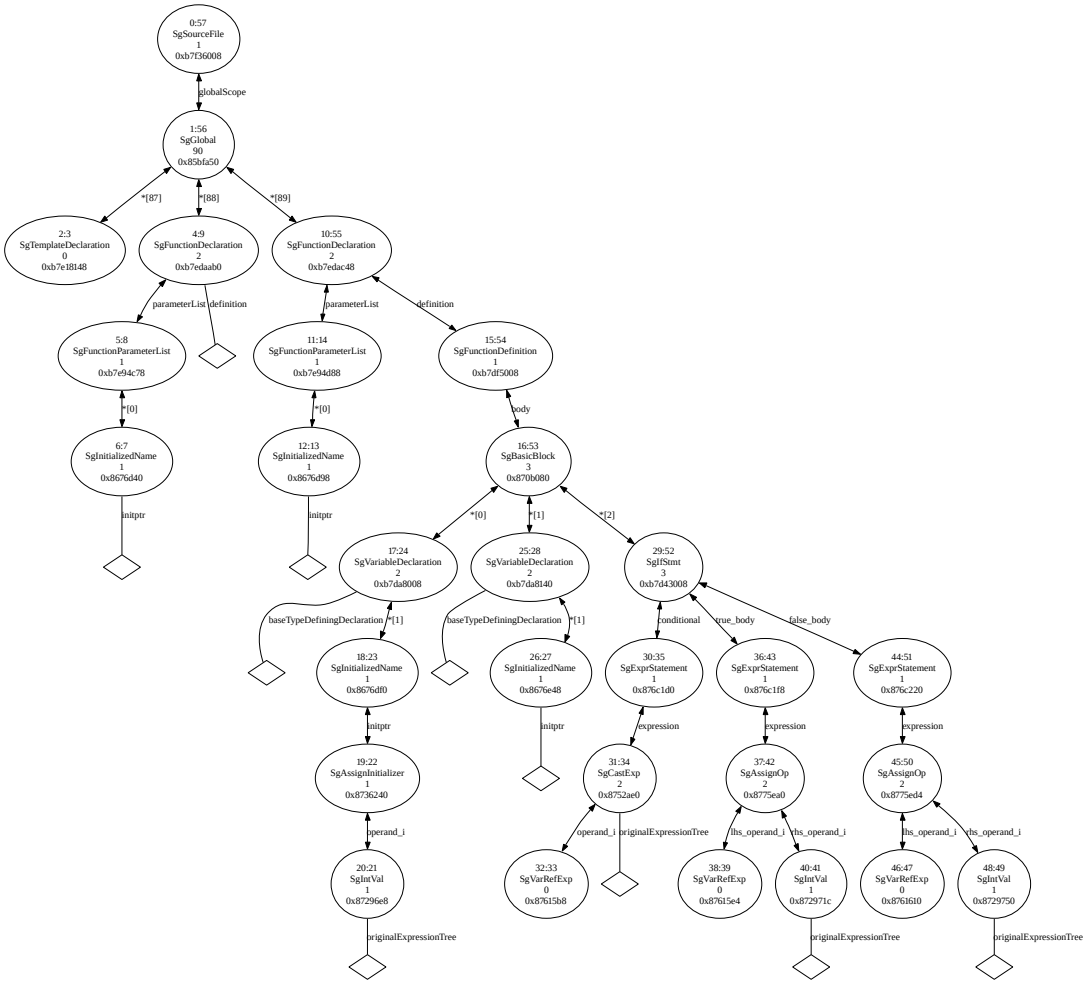

Figure 3.3: AST representing the source code file: inputCode_ASTGraphGenerator.C.

Chapter 4

AST Whole Graph Generator

What To Learn From This Example This example shows how to generate and visualize the AST from any input program. This view of the AST includes all additional IR nodes and edges that form attributes to the AST, as a result this graph is not a tree. These graphs are more complex but show significantly more detail about the AST and its additional edges and attributes. Each node of the graph in figure **??** shows a node of the Intermediate Representation (IR). Each edge shows the connection of the IR nodes in memory. The generated graph shows the connection of different IR nodes to form the AST and its additional attributes (e.g types, modifiers, etc). The generation of such graphs is appropriate for very small input programs, chapter **??** shows a mechanism using PDF files that is more appropriate to larger programs (e.g. 100K lines of code). More information about generation of specialized AST graphs can be found in 35 and custom graph generation in 34.

Viewing these dot files is best done using: **zgrviewer** at `http://zvtm.sourceforge.net/zgrviewer.html`. This tool permits zooming in and out and viewing isolated parts of even very large graphs. **Zgrviewer** permits a more natural way of understanding the AST and its additioan IR nodes than the pdf file displayed in these pages. The few lines of code used to generate the graphs can be used on any input code to better understand how the AST represents different languages and their constructs.

```
// Example ROSE Translator: used within ROSE/tutorial

#include "rose.h"

int main( int argc, char * argv[] )
   {
  // Build the AST used by ROSE
     SgProject* project = frontend(argc,argv);

  // To protect against building graphs that are too large an option is
  // provided to bound the number of IR nodes for which a graph will be
  // generated.  The layout of larger graphs is prohibitively expensive.
     const int MAX_NUMBER_OF_IR_NODES = 2000;
     generateAstGraph(project,MAX_NUMBER_OF_IR_NODES);
   }
```

Figure 4.1: Example source code to read an input program and generate a *whole* AST graph.

The program in figure 4.1 calls an internal ROSE function that traverses the AST and generates an ASCII file in `dot` format. Figure **??** shows an tiny input code which is processed to generate a graph of the AST with its attributes, generating a `dot` file. The `dot` file is then processed using `dot` to generate a pdf file 4.3 (within the `Makefile`). Note that a similar utility program already exists within ROSE/exampleTranslators (and includes a utility to output an alternative PDF representation (suitable for larger ASTs) as well). Figure **??** (../../..//tutorial/test.ps) can be found in the compile tree (in the tutorial directory) and viewed directly using any pdf or dot viewer to see more detail (**zgrviewer** working with the dot file directly is strongly advised).

Note that AST's can get very large, and that the additional IR nodes required to represent the types, modifiers, etc, can generate visually complex graphs. ROSE contains the mechanisms to traverse these graphs and do analysis on them. In one case the number of IR nodes exceeded 27 million, an analysis was done through a traversal of the graph in 10 seconds on a desktop x86 machine (the memory requirements were 6 Gig). ROSE organizes the IR in ways that permit analysis of programs that can represent rather large ASTs.

```
// Trivial function used to generate graph of AST
// with all types and additional edges shown.
// Graphs of this sort are large, and can be
// viewed using "zgrviewer" for dot files.
int foo()
   {
     return 0;
   }
```

Figure 4.2: Example tiny source code used as input to generate the small AST graph with attributes.

Figure 4.3 displays the individual C++ nodes in ROSE's intermediate representation (IR). Colors and shapes are used to represent different types or IR nodes. Although only visible using **zgrviewer** the name of the C++ construct appears in the center of each node in the graph, with the names of the data members in each IR node as edge labels. Unique pointer values are includes and printed next to the IR node name. These graphs are the single best way to develop an intuitive understanding how language constructs are organized in the AST. In these graphs, the color yellow is used for types (SgType IR nodes), the color green is used for expressions (SgExpression IR nodes), and statements are a number of different colors and shapes to make them more recognizable.

Figure 4.5 shows a graph similar to the previous graph but larger and more complex because it is from a larger code. Larger graphs of this sort are still very useful in understanding how more significant language constructs are organized and reference each other in the AST. Tools such as **zgrviewer** are essential to reviewing and understanding these graphs.

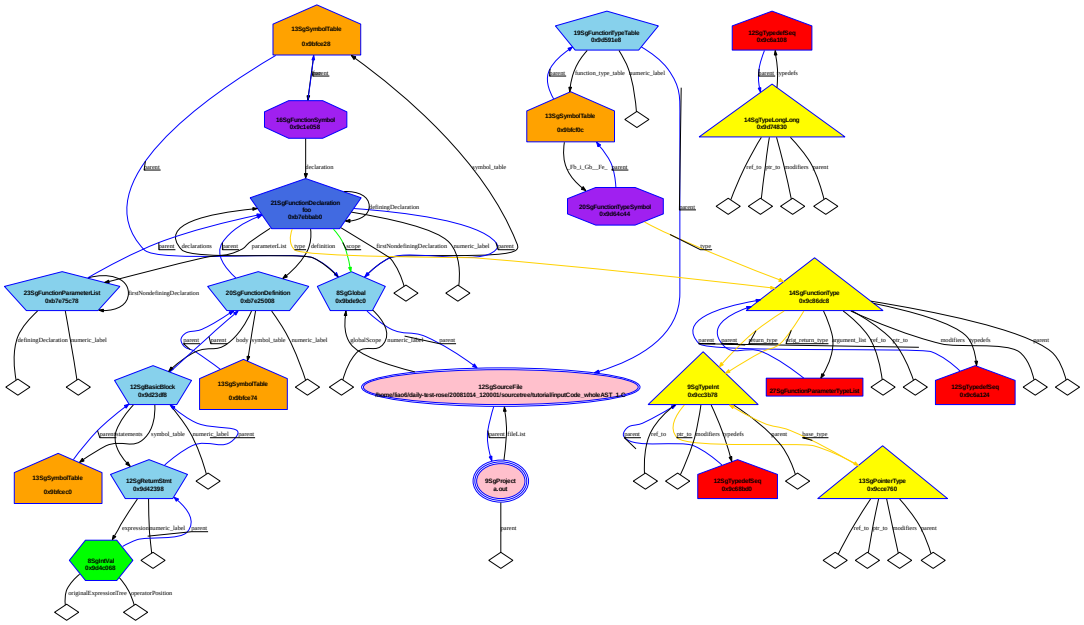

Figure 4.3: AST representing the tiny source code file: inputCode_wholeAST_1.C.

```
// Larger function used to generate graph of AST
// with all types and additional edges shown.
// Graphs of this sort are large, and can be
// viewed using "zgrviewer" for dot files.
int foo ( int x );

int globalVar = 42;

void foobar_A()
   {
     int a = 4;
     int b = a + 2;
     int c = b * globalVar;
     int x;
     x = foo (c);
     int y = x + 2;
     int z = globalVar * y;
   }


void foobar_B()
   {
     int p;
     int i = 4;
     int k = globalVar * (i+2);
     p = foo (k);
     int r = (p+2) * globalVar;
   }
```

Figure 4.4: Example source code used as input to generate a larger AST graph with attributes.

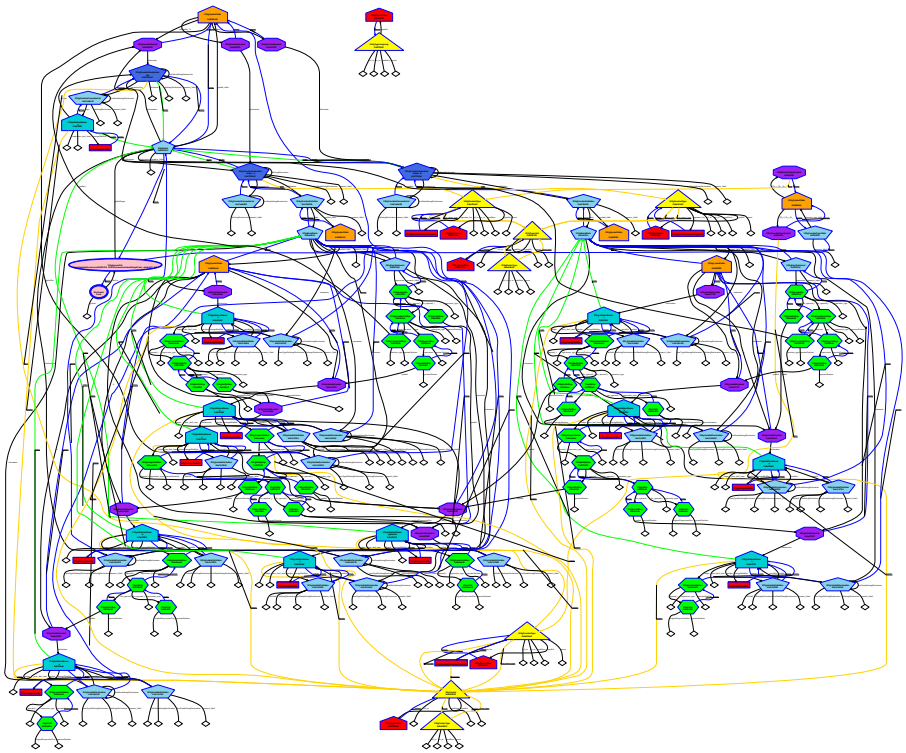

Figure 4.5: AST representing the small source code file: inputCode_wholeAST_2.C.

Chapter 5

AST PDF Generator

What To Learn From This Example This example demonstrates a mechanism for generating a visualization of the AST using pdf files. A pdf file is generated and can be viewed using `acroread`. The format is suitable for much larger input programs than the example shown in chapter **??**. This mechanism can support the visualization of input files around 100K lines of code.

```
// Example ROSE Translator: used within ROSE/tutorial

#include "rose.h"

int main( int argc, char * argv[] )
   {
  // Build the AST used by ROSE
     SgProject* project = frontend(argc,argv);

     generatePDF ( *project );
     return 0;
   }
```

Figure 5.1: Example source code to read an input program and generate a PDF file to represent the AST.

The program in figure 5.1 calls an internal ROSE function that traverses the AST and generates an ASCI file in `dot` format. Figure 3.2 shows an input code which is processed to generate a graph of the AST, generating a `pdf` file. The `pdf` file is then processed using `acroread` to generate a GUI for viewing the AST.

Figure 5.3 displays on the left hand side the individual C++ nodes in ROSE's intermediate representation (IR). The page on the right hand side shows that IR nodes member data. Pointers in boxes can be clicked on to navigate the AST (or nodes in the tree hierarchy can be clicked on jump to any location in the AST. This representation shows only the IR nodes that are traversed by the standard traversal (no SgSymbol or SgType IR nodes are presented in this view of the AST).

The output of this translator is shown in figure 5.3. The left hand side of the screen is a tree with click-able nodes to expand/collapse the subtrees. The right hand side of the screen is a

```
// Templated class declaration used in template parameter example code
template <typename T>
class templateClass
   {
     public:
          int x;

          void foo(int);
          void foo(double);
   };

// Overloaded functions for testing overloaded function resolution
void foo(int);
void foo(double)
   {
     int x = 1;
     int y;

  // Added to allow non-trivial CFG
     if (x)
        y = 2;
     else
        y = 3;
   }
```

Figure 5.2: Example source code used as input to generate the PDF file of the AST.

description of the data at a particular node in the AST (the node where the user has clicked the left mouse button). This relatively simple view of the AST is useful for debugging transformation and finding information in the AST required by specific sorts of analysis. It is also useful for developing an intuitive feel for what information is in the AST, how it is organized, and where it is stored.

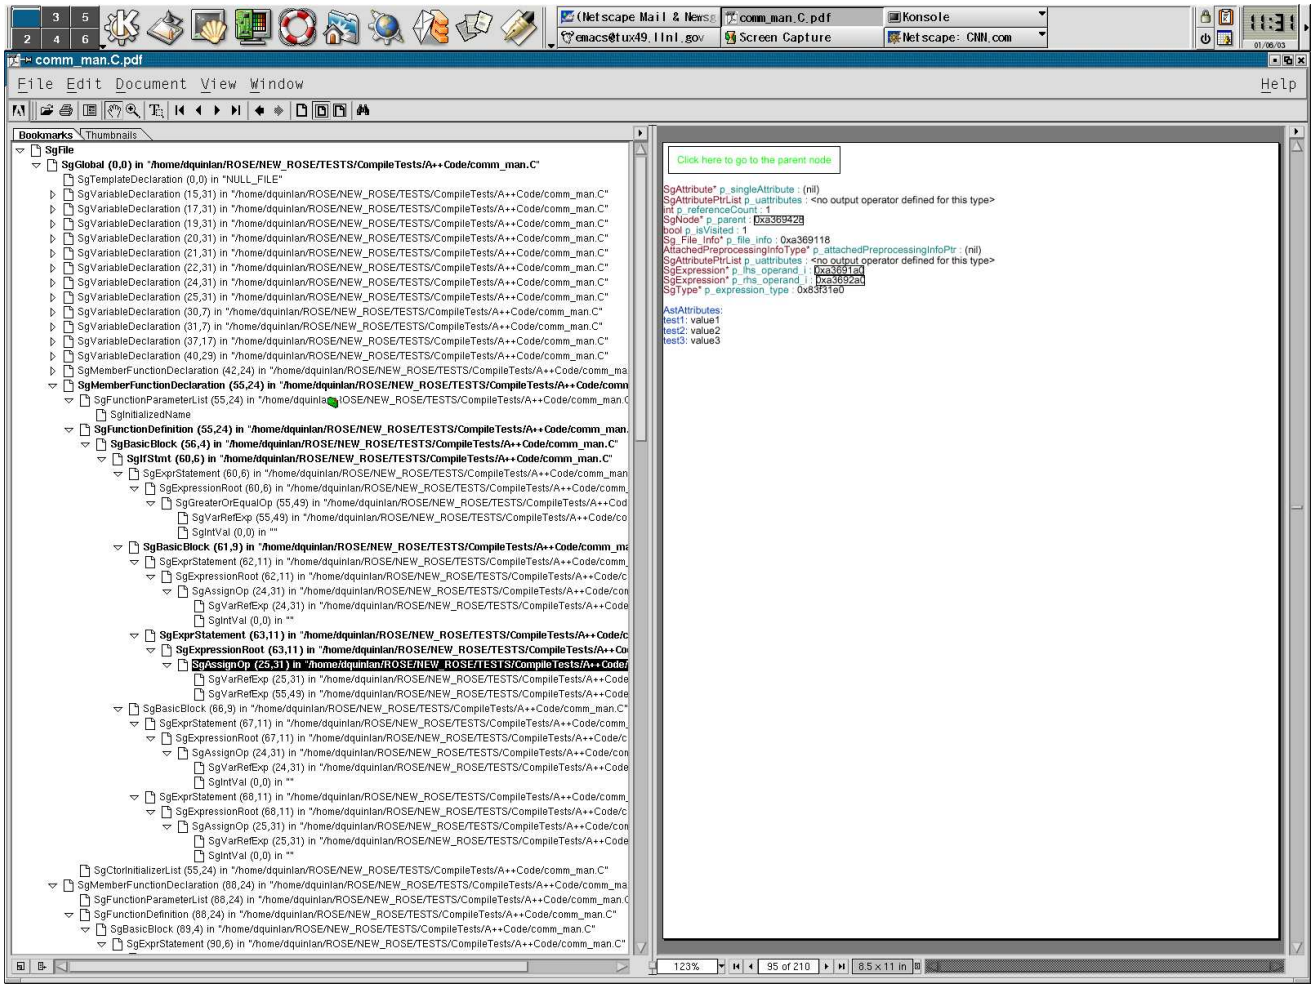

Figure 5.3: Example output from translator which outputs PDF representation of AST.

Chapter 6

Identity Translator

Using the input code in Figure 6.2 we now show a translator which builds the AST, generates the source code from the AST, and compiles the generated code using the backend vendor compiler[1]. Figure 6.1 shows the source code for this translator, the construction of the AST is identical to the previous code, but we make an explicit call to the ROSE `backend()` function.

```
// Example ROSE Translator: used for testing ROSE infrastructure

#include "rose.h"

int main( int argc, char * argv[] )
   {
  // Build the AST used by ROSE
     SgProject* project = frontend(argc,argv);

  // Run internal consistency tests on AST
     AstTests::runAllTests(project);

  // Insert your own manipulation of the AST here...

  // Generate source code from AST and call the vendor's compiler
     return backend(project);
   }
```

Figure 6.1: Source code for translator to read an input program and generate an object code (with no translation).

Figure 6.3 shows the generated code from the processing of the `identityTranslator` build using ROSE and using the input file shown in figure 6.2. This example also shows that the output generated from and ROSE translator is a close reproduction of the input; preserving all comments, preprocessor control structure, and most formating. Note that all macros are expanded in the generated code.

In this trivial case of a program in a single file, the translator compiles the application to build an executable (since `-c` was not specified on the command-line.

[1] Note: that the backend vendor compiler is selected at configuration time.

```
// Example input file for ROSE tutorial
#include <iostream>

typedef float Real;

// Main function
int main()
   {
     int x = 0;
     bool value = false;

  // for loop
     for (int i=0; i < 4; i++)
        {
          int x;
        }

     return 0;
   }
```

Figure 6.2: Example source code used as input to identity translator.

```
// Example input file for ROSE tutorial
#include <iostream>
typedef float Real;
// Main function

int main()
{
  int x = 0;
  bool value = false;
// for loop
  for (int i = 0; i < 4; i++) {
    int x;
  }
  return 0;
}
```

Figure 6.3: Generated code, from ROSE identity translator, sent to the backend (vendor) compiler.

Chapter 7

AST Query

This chapter presents a mechanism for simple queries on the AST. Such queries are typically a single line of code, instead of the class that much be declared and defined when using the traversal mechanism. While the traversal mechanism is more sophisticated and more powerful, the AST Query mechanism is particularly simple to use.

7.1 Simple Queries on the AST

This section demonstrates a simple query on the AST.

The program in figure 7.1 calls an internal ROSE Query Library. Queries of the AST using the query library are particularly simple and often are useful as nested queries within more complex analysis. More information of the ROSE AST Query Library is available within ROSE User Manual.

Using the input program in figure 7.2 the translator processes the code and generates the output in figure 7.3.

FIXME: *Put an example of composition of AST queries into the example input code.*

7.2 Nested Query

This section demonstrates a nested AST query, showing how to use composition in the construction of more elaborate queries from simple ones.

The number of traversals of the AST can be reduced by using nested queries. Nested queries permits queries on the result from a NodeQuery. Another advantage is that nested (combined) queries can be formed to query for information without writing new query, the nested query is a new query.

The program in figure 7.4 calls an internal ROSE Query Library. Two different queries are performed to find all access functions within the AST. The first query is nested, the returned list from a query is used in a traversal, and the second query queries the AST for the same nodes.

Using the input program in figure 7.5 the translator processes the code and generates the output in figure 7.6.

```
// Example ROSE Translator: used within ROSE/tutorial

#include "rose.h"

using namespace std;

int main( int argc, char * argv[] )
   {
  // Build the AST used by ROSE
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT(project != NULL);

  // Build a list of functions within the AST
     Rose_STL_Container<SgNode*> functionDeclarationList = NodeQuery::querySubTree (project,V_SgFunctionDecla

     int counter = 0;
     for (Rose_STL_Container<SgNode*>::iterator i = functionDeclarationList.begin(); i != functionDeclarationL
        {
       // Build a pointer to the current type so that we can call the get_name() member function.
          SgFunctionDeclaration* functionDeclaration = isSgFunctionDeclaration(*i);
          ROSE_ASSERT(functionDeclaration != NULL);

       // DQ (3/5/2006): Only output the non-compiler generated IR nodes
          if ( (*i)->get_file_info()->isCompilerGenerated() == false)
             {
            // output the function number and the name of the function
               printf ("Function #%2d name is %s at line %d \n",
                    counter++,functionDeclaration->get_name().str(),
                    functionDeclaration->get_file_info()->get_line());
             }
            else
             {
            // Output something about the compiler-generated builtin functions
               printf ("Compiler-generated (builtin) function #%2d name is %s \n",
                    counter++,functionDeclaration->get_name().str());
             }
        }

  // Note: Show composition of AST queries

     return 0;
   }
```

Figure 7.1: Example source code for translator to read an input program and generate a list of functions in the AST (queryLibraryExample.C).

```
// Templated class declaration used in template parameter example code
template <typename T>
class templateClass
   {
     public:
          int x;

          void foo(int);
          void foo(double);
   };

// Overloaded functions for testing overloaded function resolution
void foo(int);
void foo(double)
   {
     int x = 1;
     int y;

  // Added to allow non-trivial CFG
     if (x)
        y = 2;
     else
        y = 3;
   }

int main()
   {
     foo(42);
     foo(3.14159265);

     templateClass<char> instantiatedClass;
     instantiatedClass.foo(7);
     instantiatedClass.foo(7.0);

     for (int i=0; i < 4; i++)
        {
          int x;
        }

     return 0;
   }
```

Figure 7.2: Example source code used as input to program in figure 7.1 (queryLibraryExample.C).

```
Compiler-generated (builtin) function # 0 name is __builtin_copysign
Compiler-generated (builtin) function # 1 name is __builtin_copysignf
Compiler-generated (builtin) function # 2 name is __builtin_copysignl
Compiler-generated (builtin) function # 3 name is __builtin_acosf
Compiler-generated (builtin) function # 4 name is __builtin_acosl
Compiler-generated (builtin) function # 5 name is __builtin_asinf
Compiler-generated (builtin) function # 6 name is __builtin_asinl
Compiler-generated (builtin) function # 7 name is __builtin_atanf
Compiler-generated (builtin) function # 8 name is __builtin_atanl
Compiler-generated (builtin) function # 9 name is __builtin_atan2f
Compiler-generated (builtin) function #10 name is __builtin_atan2l
Compiler-generated (builtin) function #11 name is __builtin_ceilf
Compiler-generated (builtin) function #12 name is __builtin_ceill
Compiler-generated (builtin) function #13 name is __builtin_coshf
Compiler-generated (builtin) function #14 name is __builtin_coshl
Compiler-generated (builtin) function #15 name is __builtin_floorf
Compiler-generated (builtin) function #16 name is __builtin_floorl
Compiler-generated (builtin) function #17 name is __builtin_fmodf
Compiler-generated (builtin) function #18 name is __builtin_fmodl
Compiler-generated (builtin) function #19 name is __builtin_frexpf
Compiler-generated (builtin) function #20 name is __builtin_frexpl
Compiler-generated (builtin) function #21 name is __builtin_ldexpf
Compiler-generated (builtin) function #22 name is __builtin_ldexpl
Compiler-generated (builtin) function #23 name is __builtin_log10f
Compiler-generated (builtin) function #24 name is __builtin_log10l
Compiler-generated (builtin) function #25 name is __builtin_modff
Compiler-generated (builtin) function #26 name is __builtin_modfl
Compiler-generated (builtin) function #27 name is __builtin_powf
Compiler-generated (builtin) function #28 name is __builtin_powl
Compiler-generated (builtin) function #29 name is __builtin_sinhf
Compiler-generated (builtin) function #30 name is __builtin_sinhl
Compiler-generated (builtin) function #31 name is __builtin_tanf
Compiler-generated (builtin) function #32 name is __builtin_tanl
Compiler-generated (builtin) function #33 name is __builtin_tanhf
Compiler-generated (builtin) function #34 name is __builtin_tanhl
Compiler-generated (builtin) function #35 name is __builtin_powil
Compiler-generated (builtin) function #36 name is __builtin_powi
Compiler-generated (builtin) function #37 name is __builtin_powif
Compiler-generated (builtin) function #38 name is __builtin_strchr
Compiler-generated (builtin) function #39 name is __builtin_strrchr
Compiler-generated (builtin) function #40 name is __builtin_strpbrk
Compiler-generated (builtin) function #41 name is __builtin_strstr
Compiler-generated (builtin) function #42 name is __builtin_nansf
Compiler-generated (builtin) function #43 name is __builtin_nans
Compiler-generated (builtin) function #44 name is __builtin_nansl
Compiler-generated (builtin) function #45 name is __builtin_fabs
Compiler-generated (builtin) function #46 name is __builtin_fabsf
Compiler-generated (builtin) function #47 name is __builtin_fabsl
Compiler-generated (builtin) function #48 name is __builtin_cosf
Compiler-generated (builtin) function #49 name is __builtin_cosl
Compiler-generated (builtin) function #50 name is __builtin_sinf
Compiler-generated (builtin) function #51 name is __builtin_sinl
Compiler-generated (builtin) function #52 name is __builtin_sqrtf
Compiler-generated (builtin) function #53 name is __builtin_sqrtl
Compiler-generated (builtin) function #54 name is __builtin_return_address
Compiler-generated (builtin) function #55 name is __builtin_frame_address
Compiler-generated (builtin) function #56 name is __builtin_expect
Compiler-generated (builtin) function #57 name is __builtin_prefetch
Compiler-generated (builtin) function #58 name is __builtin_huge_val
Compiler-generated (builtin) function #59 name is __builtin_huge_valf
Compiler-generated (builtin) function #60 name is __builtin_huge_vall
Compiler-generated (builtin) function #61 name is __builtin_inf
Compiler-generated (builtin) function #62 name is __builtin_inff
Compiler-generated (builtin) function #63 name is __builtin_infl
Compiler-generated (builtin) function #64 name is __builtin_nan
Compiler-generated (builtin) function #65 name is __builtin_nanf
Compiler-generated (builtin) function #66 name is __builtin_nanl
Compiler-generated (builtin) function #67 name is __builtin_nans
Compiler-generated (builtin) function #68 name is __builtin_nansf
Compiler-generated (builtin) function #69 name is __builtin_nansl
Compiler-generated (builtin) function #70 name is __builtin_clz
Compiler-generated (builtin) function #71 name is __builtin_ctz
Compiler-generated (builtin) function #72 name is __builtin_popcount
Compiler-generated (builtin) function #73 name is __builtin_parity
Compiler-generated (builtin) function #74 name is __builtin_ffsl
Compiler-generated (builtin) function #75 name is __builtin_clzl
Compiler-generated (builtin) function #76 name is __builtin_ctzl
Compiler-generated (builtin) function #77 name is __builtin_popcountl
Compiler-generated (builtin) function #78 name is __builtin_parityl
Compiler-generated (builtin) function #79 name is __builtin_ffsll
Compiler-generated (builtin) function #80 name is __builtin_clzll
Compiler-generated (builtin) function #81 name is __builtin_ctzll
Compiler-generated (builtin) function #82 name is __builtin_popcountll
```

```
// Example ROSE Translator: used within ROSE/tutorial

#include "rose.h"

using namespace std;


 // Function querySolverAccessFunctions()
 // find access functions (function name starts with "get_" or "set_")
NodeQuerySynthesizedAttributeType
querySolverAccessFunctions (SgNode * astNode)
   {
     ROSE_ASSERT (astNode != 0);
     NodeQuerySynthesizedAttributeType returnNodeList;

     SgFunctionDeclaration* funcDecl = isSgFunctionDeclaration(astNode);

     if (funcDecl != NULL)
        {
          string functionName = funcDecl->get_name().str();
          if ( (functionName.length() >= 4) && ((functionName.substr(0,4) == "get_") || (functionName.substr(0,4) == "set
               returnNodeList.push_back (astNode);
        }

     return returnNodeList;
   }

// Function printFunctionDeclarationList will print all function names in the list
void printFunctionDeclarationList(Rose_STL_Container<SgNode*> functionDeclarationList)
   {
     int counter = 0;
     for (Rose_STL_Container<SgNode*>::iterator i = functionDeclarationList.begin(); i != functionDeclarationList.end();
        {
       // Build a pointer to the current type so that we can call the get_name() member function.
          SgFunctionDeclaration* functionDeclaration = isSgFunctionDeclaration(*i);
          ROSE_ASSERT(functionDeclaration != NULL);

       // output the function number and the name of the function
          printf ("function name #%d is %s at line %d \n",
               counter++,functionDeclaration->get_name().str(),
               functionDeclaration->get_file_info()->get_line());
        }
   }

int main( int argc, char * argv[] )
   {
  // Build the AST used by ROSE
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT(project != NULL);

  // Build a list of functions within the AST and find all access functions
  // (function name starts with "get_" or "set_")

  // Build list using a query of the whole AST
     Rose_STL_Container<SgNode*> functionDeclarationList = NodeQuery::querySubTree (project,V_SgFunctionDeclaration);

  // Build list using nested Queries (operating on return result of previous query)
     Rose_STL_Container<SgNode*> accessFunctionsList;
     accessFunctionsList = NodeQuery::queryNodeList (functionDeclarationList,&querySolverAccessFunctions);
     printFunctionDeclarationList(accessFunctionsList);

  // Alternative form of same query building the list using a query of the whole AST
     accessFunctionsList = NodeQuery::querySubTree (project,&querySolverAccessFunctions);
     printFunctionDeclarationList(accessFunctionsList);

  // Another way to query for collections of IR nodes
     VariantVector vv1 = V_SgClassDefinition;
     std::cout << "Number of class definitions in the memory pool is: " << NodeQuery::queryMemoryPool(vv1).size() << std:

  // Another way to query for collections of multiple IR nodes.
  // VariantVector(V_SgType) is internally expanded to all IR nodes derived from SgType.
     VariantVector vv2 = VariantVector(V_SgClassDefinition) + VariantVector(V_SgType);
     std::cout << "Number of class definitions AND types in the memory pool is: " << NodeQuery::queryMemoryPool(vv2).size

  // Note: Show composition of AST queries

     return 0;
   }
```

Figure 7.4: Example source code for translator to read an input program and generate a list of access functions in the AST (nestedQueryExample.C).

```
// Templated class declaration used in template parameter example code
template <typename T>
class templateClass
   {
     public:
          int x;

          void foo(int);
          void foo(double);
   };

// Overloaded functions for testing overloaded function resolution
void foo(int);
void foo(double)
   {
     int x = 1;
     int y;

  // Added to allow non-trivial CFG
     if (x)
        y = 2;
     else
        y = 3;
   }

int main()
   {
     foo(42);
     foo(3.14159265);

     templateClass<char> instantiatedClass;
     instantiatedClass.foo(7);
     instantiatedClass.foo(7.0);

     for (int i=0; i < 4; i++)
        {
          int x;
        }

     return 0;
   }
```

Figure 7.5: Example source code used as input to program in figure 7.4 (nestedQueryExample.C).

```
function name #0 is get_foo at line 0
function name #1 is set_foo at line 0
function name #2 is get_foo at line 28
function name #3 is set_foo at line 29
function name #0 is get_foo at line 0
function name #1 is set_foo at line 0
function name #2 is get_foo at line 28
function name #3 is set_foo at line 29
Number of class definitions in the memory pool is: 1
Number of class definitions AND types in the memory pool is: 165
```

Figure 7.6: Output of input file to the AST query processor (nestedQueryExample.C).

Chapter 8

Introduction to AST Traversals

An essential operation in the analysis and construction of ASTs is the definition of traversals upon the AST to gather information and modify targeted internal representation (IR) nodes. ROSE includes different sorts of traversals to address the different requirements of numerous program analysis and transformation operations. This section demonstrates the different types of traversals that are possible using ROSE.

ROSE translators most commonly introduce transformations and analysis through a traversal over the AST. Alternatives could be to generate a simpler IR that is more suitable to a specific transformation and either convert modification to that transformation specific IR into changes to the AST or generate source code from the transformation specific IR directly. These approaches are more complex than introducing changes to the AST directly, but may be better for specific transformations.

Traversals represent an essential operation on the AST and there are a number of different types of traversals. The suggested traversals for users are explained in Section 8.2. Section 8.3 introduces specialized traversals (that traverse the AST in different orders and traverse types and symbols), typically not appropriate for most translators (but perhaps appropriate for specialized tools, often internal tools within ROSE).

See the ROSE User Manual for a more complete introduction to the different types of traversals. The purpose of this tutorial is to present examples, but we focus less on the background and philosophy here than in the ROSE User Manual.

This chapter presents a number of ways of traversing the AST of any input source code. These *traversals* permit operations on the AST, which may either read or modify the AST in place. Modifications to the AST will be reflected in the source code generated when the AST is *unparsed*; the code generation phase of the source-to-source process defined by ROSE. Note that for all examples, the input code described in section 8.1 is used to generate all outputs shown with each translator.

8.1 Input For Example Traversals

The code shown in figure 8.1 shows the input code that will be used to demonstrate the traversals in this chapter. It may be modified by the user to experiment with the use of the traversals on

```
// Templated class declaration used in template parameter example code
template <typename T>
class templateClass
   {
     public:
          int x;
          void foo(int);
          void foo(double);
   };

// Overloaded functions for testing overloaded function resolution
void foo(int);

void foo(double)
   {
     int x = 1;
     int y;

     for (int i=0; i < 4; i++)
        {
          int x;
        }

  // Added to allow non-trivial CFG
     if (x)
        y = 2;
     else
        y = 3;
   }

int main()
   {
     foo(42);
     foo(3.14159265);

     templateClass<char> instantiatedClass;
     instantiatedClass.foo(7);
     instantiatedClass.foo(7.0);

     for (int i=0; i < 4; i++)
        {
          int x;
        }

     return 0;
   }
```

Figure 8.1: Example source code used as input to program in traversals shown in this chapter.

alternative input codes.

8.2 Traversals of the AST Structure

This collection of traversals operates on the AST in an order which matches the structure of the AST and the associated source code. These types of traversals are the most common traversals for users to use. A subsequent section of this chapter demonstrated more specialized traversals over all IR nodes (more than just those IR nodes in the AST representing the structure of the source code) that are suitable for some tools, mostly tools built internally within ROSE.

Because the traversals in this section traverse the structure of the source code (see the AST graph presented in the first tutorial example) they are more appropriate for most transformations of the source code. We suggest that the user focus on these traversals which represent the interface we promote for analysis and transformation of the AST, instead of the memory pools traversals which are suitable mostly for highly specialized internal tools. The simple traversals of both kinds have the same interface so the user may easily switch between them with out significant difficulty.

8.2.1 Classic Object-Oriented Visitor Pattern for the AST (Not Yet Implemented)

```
#include "rose.h"

// Classic Visitor Pattern in ROSE (implemented using the traversal over
// the elements stored in the memory pools so it has no cycles and visits
// ALL IR nodes (including all Sg_File_Info, SgSymbols, SgTypes, and the
// static builtin SgTypes).
class ClassicVisitor : public ROSE_VisitorPattern
   {
     public:
       // Override virtural function defined in base class
          void visit(SgGlobal* globalScope)
             {
               printf ("Found the SgGlobal IR node \n");
             }

          void visit(SgFunctionDeclaration* functionDeclaration)
             {
               printf ("Found a SgFunctionDeclaration IR node \n");
             }
          void visit(SgTypeInt* intType)
             {
               printf ("Found a SgTypeInt IR node \n");
             }

          void visit(SgTypeDouble* doubleType)
             {
               printf ("Found a SgTypeDouble IR node \n");
             }
   };


int
main ( int argc, char* argv[] )
   {
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT (project != NULL);

  // Classic visitor pattern over the memory pool of IR nodes
     ClassicVisitor visitor_A;
     traverseMemoryPoolVisitorPattern(visitor_A);

     return backend(project);
   }
```

Figure 8.2: Example source showing simple visitor pattern.

Figure 8.2 shows the source code for a translator using the classic object-oriented visitor pattern to traverse the AST. *This visitor pattern is not yet implemented except for the memory pool based traversal. It is however expected to appear identical to the classing visitor pattern shown for the memory pool traversal.* Figure 8.3 shows the output from this traversal using the example input source from figure 8.1.

8.2.2 Simple Traversal (no attributes)

Figure 8.4 shows the source code for a translator which traverses the AST. At each node the `visit()` function is called using only the input information represented by the current node. Note that using this simple traversal the only context information available to the visit function is what is stored in its member variables. The only option is to traverse the AST in either pre-order or postorder. The `atTraversalEnd()` function may be defined by the user to do final processing after all nodes have been visited (or to perform preparations before the nodes are visited, in the case of the corresponding `atTraversalStart()` function). Figure 8.5 shows the output from this traversal using the example input source from figure 8.1.

8.2.3 Simple Pre- and Postorder Traversal

Figure 8.6 shows the source code for a translator that traverses the AST without attributes (like the one in the previous subsection), but visiting each node twice, once in preorder (before its children) and once in postorder (after all children). Figure 8.7 shows the output from this traversal using the example input source from figure 8.1.

8.2.4 Inherited Attributes

Figure 8.8 shows the use of inherited attributes associated with each IR node. Within this traversal the attributes are managed by the traversal and exist on the stack. Thus the lifetime of the attributes is only as long as the processing of the IR node and its subtree. Attributes such as this are used to communicate context information **down** the AST and called *Inherited attributes*.

In the example the class `Inherited Attribute` is used to represent inherited attribute. Each instance of the class represents an attribute value. When the AST is traversed we obtain as output the loop nest depth at each point in the AST. The output uses the example input source from figure 8.1.

Note that inherited attributes are passed by-value down the AST. In very rare cases you might want to pass a pointer to dynamically allocated memory as an inherited attribute. In this case you can define the virtual member function `void destroyInheritedValue(SgNode *n, InheritedAttribute inheritedValue)` which is called after the last use of the inherited attribute computed at this node, i. e. after all children have been visited. You can use this function to free the memory allocated for this inherited attribute.

```
Found a SgTypeInt IR node
Found a SgTypeDouble IR node
Found the SgGlobal IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
```

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.
// rose.C: Example (default) ROSE Preprocessor: used for testing ROSE infrastructure

#include "rose.h"

class visitorTraversal : public AstSimpleProcessing
   {
     public:
          visitorTraversal();
          virtual void visit(SgNode* n);
          virtual void atTraversalEnd();
   };

visitorTraversal::visitorTraversal()
   {
   }

void visitorTraversal::visit(SgNode* n)
   {
     if (isSgForStatement(n) != NULL)
        {
          printf ("Found a for loop ... \n");
        }
   }

void visitorTraversal::atTraversalEnd()
   {
     printf ("Traversal ends here. \n");
   }

int
main ( int argc, char* argv[] )
   {
     if (SgProject::get_verbose() > 0)
          printf ("In visitorTraversal.C: main() \n");

     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT (project != NULL);

  // Build the traversal object
     visitorTraversal exampleTraversal;

  // Call the traversal starting at the project node of the AST
     exampleTraversal.traverseInputFiles(project,preorder);

     return 0;
   }
```

Figure 8.4: Example source showing simple visitor pattern.

```
Found a for loop ...
Found a for loop ...
Traversal ends here.
```

Figure 8.5: Output of input file to the visitor traversal.

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.
// rose.C: Example (default) ROSE Preprocessor: used for testing ROSE infrastructure

#include "rose.h"

class PreAndPostOrderTraversal : public AstPrePostProcessing
   {
     public:
          virtual void preOrderVisit(SgNode* n);
          virtual void postOrderVisit(SgNode* n);
   };

void PreAndPostOrderTraversal::preOrderVisit(SgNode* n)
   {
     if (isSgForStatement(n) != NULL)
        {
          printf ("Entering for loop ... \n");
        }
   }

void PreAndPostOrderTraversal::postOrderVisit(SgNode* n)
   {
     if (isSgForStatement(n) != NULL)
        {
          printf ("Leaving for loop ... \n");
        }
   }

int
main ( int argc, char* argv[] )
   {
     if (SgProject::get_verbose() > 0)
          printf ("In prePostTraversal.C: main() \n");

     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT (project != NULL);

  // Build the traversal object
     PreAndPostOrderTraversal exampleTraversal;

  // Call the traversal starting at the project node of the AST
     exampleTraversal.traverseInputFiles(project);

     return 0;
   }
```

Figure 8.6: Example source showing simple pre- and postorder pattern.

```
Entering for loop ...
Leaving for loop ...
Entering for loop ...
Leaving for loop ...
```

Figure 8.7: Output of input file to the pre- and postorder traversal.

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.
// rose.C: Example (default) ROSE Preprocessor: used for testing ROSE infrastructure

#include "rose.h"

// Build an inherited attribute for the tree traversal to test the rewrite mechanism
class InheritedAttribute
   {
     public:
      // Depth in AST
         int depth;
         int maxLinesOfOutput;

       // Specific constructors are required
         InheritedAttribute (int x) : depth(x), maxLinesOfOutput(20) {};
         InheritedAttribute ( const InheritedAttribute & X ) : depth(X.depth), maxLinesOfOutput(20){};
   };

class visitorTraversal : public AstTopDownProcessing<InheritedAttribute>
   {
     public:
       // virtual function must be defined
         virtual InheritedAttribute evaluateInheritedAttribute(SgNode* n, InheritedAttribute inheritedAttribu
   };

InheritedAttribute
visitorTraversal::evaluateInheritedAttribute(SgNode* n, InheritedAttribute inheritedAttribute)
   {
     static int linesOfOutput = 0;
     if (linesOfOutput++ < inheritedAttribute.maxLinesOfOutput)
          printf ("Depth in AST at %s = %d \n",n->sage_class_name(),inheritedAttribute.depth);

     return InheritedAttribute(inheritedAttribute.depth+1);
   }

int
main ( int argc, char* argv[] )
   {
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT (project != NULL);

  // DQ (1/18/2006): Part of debugging
     SgFile & localFile = project->get_file(0);
     localFile.get_file_info()->display("localFile information");

  // Build the inherited attribute
     InheritedAttribute inheritedAttribute(0);

  // Build the traversal object
     visitorTraversal exampleTraversal;

  // Call the traversal starting at the project node of the AST
     exampleTraversal.traverseInputFiles(project,inheritedAttribute);

  // Or the traversal over all AST IR nodes can be called!
     exampleTraversal.traverse(project,inheritedAttribute);

     return 0;
   }
```

Figure 8.8: Example source code showing use of inherited attributes (passing context information **down** the AST.

```
Inside of Sg_File_Info::display(localFile information)
     isTransformation                        = false
     isCompilerGenerated                     = false
     isOutputInCodeGeneration                = false
     isShared                                = false
     isFrontendSpecific                      = false
     isSourcePositionUnavailableInFrontend = false
     isCommentOrDirective                    = false
     isToken                                 = false
     filename = /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_7.C
     line     = 1  column = 1
Depth in AST at SgSourceFile = 0
Depth in AST at SgGlobal = 1
Depth in AST at SgTemplateDeclaration = 2
Depth in AST at SgFunctionDeclaration = 2
Depth in AST at SgFunctionParameterList = 3
Depth in AST at SgInitializedName = 4
Depth in AST at SgFunctionDeclaration = 2
Depth in AST at SgFunctionParameterList = 3
Depth in AST at SgInitializedName = 4
Depth in AST at SgFunctionDefinition = 3
Depth in AST at SgBasicBlock = 4
Depth in AST at SgVariableDeclaration = 5
Depth in AST at SgInitializedName = 6
Depth in AST at SgAssignInitializer = 7
Depth in AST at SgIntVal = 8
Depth in AST at SgVariableDeclaration = 5
Depth in AST at SgInitializedName = 6
Depth in AST at SgForStatement = 5
Depth in AST at SgForInitStatement = 6
Depth in AST at SgVariableDeclaration = 7
```

Figure 8.9: Output of input file to the inherited attribute traversal.

8.2.5 Synthesized Attributes

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.
// rose.C: Example (default) ROSE Preprocessor: used for testing ROSE infrastructure

#include "rose.h"

#include <algorithm>
#include <functional>
#include <numeric>

typedef bool SynthesizedAttribute;

class visitorTraversal : public AstBottomUpProcessing<SynthesizedAttribute>
   {
     public:
       // virtual function must be defined
          virtual SynthesizedAttribute evaluateSynthesizedAttribute (
                       SgNode* n, SynthesizedAttributesList childAttributes );
   };

SynthesizedAttribute
visitorTraversal::evaluateSynthesizedAttribute ( SgNode* n, SynthesizedAttributesList childAttributes )
   {
  // Fold up the list of child attributes using logical or, i.e. the local
  // result will be true iff one of the child attributes is true.
     SynthesizedAttribute localResult =
         std::accumulate(childAttributes.begin(), childAttributes.end(),
                         false, std::logical_or<bool>());

     if (isSgForStatement(n) != NULL)
        {
          printf ("Found a for loop ... \n");
          localResult = true;
        }

     return localResult;
   }

int
main ( int argc, char* argv[] )
   {
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT (project != NULL);

  // Build the traversal object
     visitorTraversal exampleTraversal;

  // Call the traversal starting at the project node of the AST
     SynthesizedAttribute result = exampleTraversal.traverse(project);

     if (result == true)
        {
          printf ("The program contains at least one loop!\n");
        }

     return 0;
   }
```

Figure 8.10: Example source code showing use of synthesized attributed (passing analysis information **up** the AST).

Figure 8.10 shows the use of attributes to pass information **up** the AST. The lifetime of the

```
Found a for loop ...
Found a for loop ...
The program contains at least one loop!
```

Figure 8.11: Output of input file to the synthesized attribute traversal.

attributes are similar as for inherited attributes. Attributes such as these are called synthesized attributes.

This code shows the code for a translator which does an analysis of an input source code to determine the presence of loops. It returns true if a loop exists in the input code and false otherwise. The list of synthesized attributes representing the information passed up the AST from a node's children is of type `SynthesizedAttributesList`, which is a type that behaves very similarly to `vector<SynthesizedAttribute>` (it supports iterators, can be indexed, and can be used with STL algorithms).

The example determines the existence of loops for a given program.

8.2.6 Accumulator Attributes

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.
// rose.C: Example (default) ROSE Preprocessor: used for testing ROSE infrastructure

#include "rose.h"

// Build an accumulator attribute, fancy name for what is essentially a global variable :-).
class AccumulatorAttribute
   {
     public:
          int forLoopCounter;

        // Specific constructors are optional
          AccumulatorAttribute () { forLoopCounter = 0; }
          AccumulatorAttribute ( const AccumulatorAttribute & X ) {}
          AccumulatorAttribute & operator= ( const AccumulatorAttribute & X ) { return *this; }
   };

class visitorTraversal : public AstSimpleProcessing
   {
     public:
          static AccumulatorAttribute accumulatorAttribute;
          virtual void visit(SgNode* n);
   };

// declaration required for static data member
AccumulatorAttribute visitorTraversal::accumulatorAttribute;

void visitorTraversal::visit(SgNode* n)
   {
     if (isSgForStatement(n) != NULL)
        {
          printf ("Found a for loop ... \n");
          accumulatorAttribute.forLoopCounter++;
        }
   }

int
main ( int argc, char* argv[] )
   {
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT (project != NULL);

  // Build the traversal object
     visitorTraversal exampleTraversal;

  // Call the traversal starting at the project node of the AST
  // can be specified to be preorder or postorder).
     exampleTraversal.traverseInputFiles(project,preorder);

     printf ("Number of for loops in input application = %d \n",exampleTraversal.accumulatorAttribute.forLoopC

     return 0;
   }
```

Figure 8.12: Example source code showing use of accumulator attributes (typically to count things in the AST).

Figure 8.12 shows the use of a different sort of attribute. This attribute has a lifetime equal to the lifetime of the traversal object (much longer than the traversal of any subset of IR nodes). The same attribute is accessible from each IR node. Such attributes are called *accumulator*

```
Found a for loop ...
Found a for loop ...
Number of for loops in input application = 2
```

Figure 8.13: Output of input file to the accumulator attribute traversal.

attributes and are semantically equivalent to a global variable. Accumulator attributes act as global variables which can easily be used to count application specific properties within the AST.

Note that due to the limitation that the computation of inherited attributes cannot be made dependent on the values of synthesized attributes, counting operations cannot be implemented by combining these attributes as is usually done in attribute grammars. However, the use of accumulator attributes serves well for this purpose. Therefore all counting-like operations should be implemented using accumulator attributes (= member variables of traversal or processing classes).

Although not shown in this tutorial explicitly, accumulator attributes may be easily mixed with inherited and/or synthesized attributes.

In this example we count the number of for-loops in an input program.

8.2.7 Inherited and Synthesized Attributes

Figure 8.14 shows the combined use of inherited and synthesized attributes. The example source code shows the mixed use of such attributes to list the functions containing loop. Inherited attributes are used to communicate that the traversal is in a function, which the synthesized attributes are used to pass back the existence of loops deeper within the subtrees associated with each function.

List of functions containing loops.

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.

#include "rose.h"

#include <algorithm>
#include <functional>
#include <numeric>

typedef bool InheritedAttribute;
typedef bool SynthesizedAttribute;

class Traversal : public SgTopDownBottomUpProcessing<InheritedAttribute,SynthesizedAttribute>
   {
     public:
       // Functions required
          InheritedAttribute evaluateInheritedAttribute (
             SgNode* astNode,
             InheritedAttribute inheritedAttribute );

          SynthesizedAttribute evaluateSynthesizedAttribute (
             SgNode* astNode,
             InheritedAttribute inheritedAttribute,
             SubTreeSynthesizedAttributes synthesizedAttributeList );
   };

InheritedAttribute
Traversal::evaluateInheritedAttribute (
     SgNode* astNode,
     InheritedAttribute inheritedAttribute )
   {
     if (isSgFunctionDefinition(astNode))
        {
       // The inherited attribute is true iff we are inside a function.
          return true;
        }
     return inheritedAttribute;
   }

SynthesizedAttribute
Traversal::evaluateSynthesizedAttribute (
     SgNode* astNode,
     InheritedAttribute inheritedAttribute,
     SynthesizedAttributesList childAttributes )
   {
     if (inheritedAttribute == false)
        {
       // The inherited attribute is false, i.e. we are not inside any
       // function, so there can be no loops here.
          return false;
        }
     else
        {
       // Fold up the list of child attributes using logical or, i.e. the local
       // result will be true iff one of the child attributes is true.
          SynthesizedAttribute localResult =
          std::accumulate(childAttributes.begin(), childAttributes.end(),
                          false, std::logical_or<bool>());
          if (isSgFunctionDefinition(astNode) && localResult == true)
             {
               printf ("Found a function containing a for loop ...\n");
             }
          if (isSgForStatement(astNode))
             {
               localResult = true;
             }
          return localResult;
        }
   }

int
main ( int argc, char* argv[] )
   {
  // Build the abstract syntax tree
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT (project != NULL);

  // Build the inherited attribute
     InheritedAttribute inheritedAttribute = false;

  // Define the traversal
     Traversal myTraversal;

  // Call the traversal starting at the project (root) node of the AST
     myTraversal.traverseInputFiles(project,inheritedAttribute);

  // This program only does analysis, so it need not call the backend to generate code.
     return 0;
   }
```

Figure 8.14: Example source code showing use of both inherited and synthesized attributes working together (part 1).

```
Found a function containing a for loop ...
Found a function containing a for loop ...
```

Figure 8.15: Output of input file to the inherited and synthesized attribute traversal.

8.2.8 Persistent Attributes

Figure 8.16 shows the use of another form of attribute. This attribute has a lifetime which is controlled explicitly by the user; it lives on the heap typically. These attributes are explicitly attached to the IR nodes and are not managed directly by the traversal. There attributes are called *persistent* attributes and are not required to be associated with any traversal. Persistent attributes are useful for storing information across multiple traversals (or permanently within the AST) for later traversal passes.

Persistent attributes may be used at any time and combined with other traversals (similar to accumulator attributes). Traversals may combine any or all of the types of attributes within in ROSE as needed to store, gather, or propagate information within the AST for complex program analysis.

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.
// rose.C: Example (default) ROSE Preprocessor: used for testing ROSE infrastructure

#include "rose.h"

class persistantAttribute : public AstAttribute
   {
     public:
          int value;
          persistantAttribute (int v) : value(v) {}
   };

class visitorTraversalSetAttribute : public AstSimpleProcessing
   {
     public:
          virtual void visit(SgNode* n);
   };

void visitorTraversalSetAttribute::visit(SgNode* n)
   {
     if (isSgForStatement(n) != NULL)
        {
          printf ("Found a for loop (set the attribute) ... \n");

       // Build an attribute (on the heap)
          AstAttribute* newAttribute = new persistantAttribute(5);
          ROSE_ASSERT(newAttribute != NULL);
#if 0
       // DQ (1/2/2006): Added support for new attribute interface.
       // printf ("visitorTraversalSetAttribute::visit(): using new attribute interface \n");
          if (n->get_attribute() == NULL)
             {
               AstAttributeMechanism* attributePtr = new AstAttributeMechanism();
               ROSE_ASSERT(attributePtr != NULL);
               n->set_attribute(attributePtr);
             }
#endif
       // Add it to the AST (so it can be found later in another pass over the AST)
       // n->attribute.add("MyNewAttribute",newAttribute);
       // n->attribute().add("MyNewAttribute",newAttribute);
          n->addNewAttribute("MyNewAttribute",newAttribute);
        }
   }

class visitorTraversalReadAttribute : public AstSimpleProcessing
   {
     public:
          virtual void visit(SgNode* n);
   };

void visitorTraversalReadAttribute::visit(SgNode* n)
   {
     if (isSgForStatement(n) != NULL)
        {
          printf ("Found a for loop (read the attribute) ... \n");

       // Add it to the AST (so it can be found later in another pass over the AST)
       // AstAttribute* existingAttribute = n->attribute["MyNewAttribute"];
       // DQ (1/2/2006): Added support for new attribute interface.
       // printf ("visitorTraversalReadAttribute::visit(): using new attribute interface \n");
       // AstAttribute* existingAttribute = n->attribute()["MyNewAttribute"];
          AstAttribute* existingAttribute = n->getAttribute("MyNewAttribute");
          ROSE_ASSERT(existingAttribute != NULL);

          printf ("Existing attribute at %p value = %d \n",n,dynamic_cast<persistantAttribute*>(existingAttribute)->value
        }
   }

int
main ( int argc, char* argv[] )
   {
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT (project != NULL);

  // Build the traversal object to set persistant AST attributes
     visitorTraversalSetAttribute exampleTraversalSettingAttribute;

  // Call the traversal starting at the project node of the AST
     exampleTraversalSettingAttribute.traverseInputFiles(project,preorder);

  // Build the traversal object to read any existing AST attributes
     visitorTraversalReadAttribute exampleTraversalReadingAtribute;
```

```
Found a for loop (set the attribute) ...
Found a for loop (set the attribute) ...
Found a for loop (read the attribute) ...
Existing attribute at 0xb7cca008 value = 5
Found a for loop (read the attribute) ...
Existing attribute at 0xb7cca08c value = 5
```

Figure 8.17: Output of input file to the persistent attribute traversal showing the passing of information from one AST traversal to a second AST traversal.

8.2.9 Nested Traversals

```
// Example ROSE Translator: used within ROSE/tutorial

#include "rose.h"

class visitorTraversal : public AstSimpleProcessing
   {
     public:
          virtual void visit(SgNode* n);
   };

class nestedVisitorTraversal : public AstSimpleProcessing
   {
     public:
          virtual void visit(SgNode* n);
   };

void visitorTraversal::visit(SgNode* n)
   {
     if (isSgFunctionDeclaration(n) != NULL)
        {
          printf ("Found a function declaration ... \n");

       // Build the nested traversal object
          nestedVisitorTraversal exampleTraversal;

       // Call the traversal starting at the project node of the AST (traverse in postorder just to be different)
       // Note that we call the traverse function instead of traverseInputFiles, because we are not starting at
       // the AST root.
          exampleTraversal.traverse(n,postorder);
        }
   }

void nestedVisitorTraversal::visit(SgNode* n)
   {
     if (isSgFunctionDefinition(n) != NULL)
        {
          printf ("Found a function definition within the function declaration ... \n");
        }
   }

int
main ( int argc, char* argv[] )
   {
     if (SgProject::get_verbose() > 0)
          printf ("In visitorTraversal.C: main() \n");

     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT (project != NULL);

  // Build the traversal object
     visitorTraversal exampleTraversal;

  // Call the traversal starting at the project node of the AST
     exampleTraversal.traverseInputFiles(project,preorder);

     return 0;
   }
```

Figure 8.18: Example source code showing use nested traversals.

Figure 8.18 shows the use of multiple traversals in composition. Figure 8.19 shows the output

```
Found a function declaration ...
Found a function declaration ...
Found a function definition within the function declaration ...
Found a function declaration ...
Found a function declaration ...
Found a function declaration ...
Found a function declaration ...
Found a function definition within the function declaration ...
```

Figure 8.19: Output of input file to the nested traversal example.

of the nested traversal.

8.2.10 Combining all Attributes and Using Primitive Types

```
int main() {
  int x=1;
  for(int i=1; i<10;i++)
    for(int j=i;j<10;j++)
      for(int k=i;k<10;k++)
        for(int l=i;l<10;l++)
          for(int m=i;m<10;m++)
            x++;

  int i=5,j=7;
  while(i>0) {
    while(j>0) {
      x++;
      j--;
      i--;
    }
  }

  i=10;
  do {
    x++;
    i--;
  } while (i>0);

  return x;
}
```

Figure 8.20: Input code with nested loops for nesting info processing

The previous examples have shown cases where attributes were classes, alternatively attributes can be any primitive type (int, bool, etc.). This example demonstrates how to use AstTopDownBottomUpProcessing to compute inherited and synthesized attributes, generate pdf and dot output, how to accumulate information, and how to attach attributes to AST nodes in the same pass.

The attributes are used to compute the nesting level and the nesting depth of for/while/do-while loops: The nesting level is computed using an inherited attribute. It holds that $nesting-level(innerloop) = nesting-level(outerloop)+1$ starting with 1 at the outer most loop. The nesting depth is computed using a synthesized attribute. It holds that $nesting{-}depth(innerloop) = nesting-level(outerloop)-1$ starting with 1 at the inner most loop.

To compute the values we use a primitive type (unsigned int). This example also shows how to use defaultSynthesizedAttribute to initialize a synthesized attribute of primitive type. The values of the attributes are attached to the AST using AstAttribute and the AST node attribute mechanism available at every AST node (which can be accessed with `node->attribute`). (see loopNestingInfoProcessing.C)

For the entire program the maximum nesting level (= max nesting depth) is computed as accumulated value using member variable `_maxNestingLevel` of class LoopNestingInfoProcessing. We also demonstrate how to customize an AstAttribute such that the value of the attribute is printed in a pdf output. (by overriding toString, see LoopNestingInfo class)

In the generated pdf file (for some C++ input file) the values of the attributes can be viewed for each node (see printLoopInfo implementation). Further more we also generate a dot file, to visualize the tree using the graph visualization tool dot. The generated file can be converted to

postscript (using dot) and viewed with gv.

8.2.11 Combined Traversals

Performing a large number of program analyses as separate traversals of the AST can be somewhat inefficient as there is some overhead associated with visiting every node several times. ROSE therefore provides a mechanism for combining traversal objects of the same base type and evaluating them in a single traversal of the AST. This is entirely transparent to the individual traversal object, so existing code can be reused with the combination mechanism, and analyzers can be developed and tested in isolation and combined when needed.

The one requirement that is placed on traversals to be combined is that they be independent of each other; in particular, this means that they should not modify the AST or any shared global data. Any output produced by the analyzers will be interleaved.

Figure 8.24 shows the source code for a translator that combines three different analyzers into one traversal, each one counting the occurrences of a different type of AST node (as determined by a VariantT value). First three traversals are run after each other, as usual; then three traversal objects are passed (by pointer) to an object of type `AstCombinedSimpleProcessing` using its `addTraversal` method. One then invokes one of the usual traverse methods on this combined object with the same effect as if it had been called for each of the traversal objects individually.

Any operation on the list of analyzers is possible using the `get_traversalPtrListRef` method of the combined processing class that returns a reference to its internal list of analyzers (an object of type `vector<AstSimpleProcessing *>`). Any changes made through this reference will be reflected in further traversals.

In addition to `AstCombinedSimpleProcessing`, there is also a combined class for each of the other types of traversals discussed above: `AstCombinedTopDownProcessing`, `AstCombinedBottomUpProcessing`, etc. Where traversals using attributes are combined, all of the combined traversals must have the same attribute types (i. e. the same template parameters). Attributes are passed to and returned from the combined traversal as a vector.

```
// Author: Markus Schordan, Vienna University of Technology, 2004.
// $Id: loopNestingInfoProcessing.C,v 1.1 2006/04/24 00:22:00 dquinlan Exp $

// #include <string>
// #include <iostream>

#include "rose.h"

using namespace std;

typedef unsigned int NestingLevel;
typedef unsigned int NestingDepth;
typedef NestingLevel InhNestingLevel;
typedef NestingDepth SynNestingDepth;

/*! This class is used to attach information to AST nodes.
    Method 'toString' is overridden and
    called when a pdf file is generated. This allows to display
    the value of an AST node attribute (annotation) in a pdf file.
*/
class NestingLevelAnnotation : public AstAttribute {
public:
  NestingLevelAnnotation(NestingLevel n,NestingDepth d)
    : _nestingLevel(n),_nestingDepth(d) {}
  NestingLevel getNestingLevel() { return _nestingLevel; }
  NestingDepth getNestingDepth() { return _nestingDepth; }
  string toString() {
    ostringstream ss; ss<<_nestingLevel<<","<<_nestingDepth;
    return ss.str();
  }
private:
  NestingLevel _nestingLevel;
  NestingDepth _nestingDepth;
};

/* !
The loop nesting level and nesting depth for each while/dowhile/for
loop nest is computed. It is attached to the AST as annotation and
can be accessed as node->attribute["loopNestingInfo"]
after the processing has been performed.
The maximum nesting level of the whole AST is computed as
"accumulated" value in a member variable and can be accessed with
getMaxNestingLevel().
*/
class LoopLevelProcessing : public AstTopDownBottomUpProcessing<InhNestingLevel,SynNestingDepth> {
public:
  LoopLevelProcessing():_maxNestingLevel(0) {}

  /*! Performs a traversal of the AST and computes loop-nesting information by using
      inherited and synthesized attributes. The results are attached to the AST as
      annotation.
  */
  void attachLoopNestingAnnotaton(SgProject* node) { traverseInputFiles(node,0); }

  /*! Returns the maximum nesting level of the entire AST (of the input file).
      Requires attachLoopNestingAnnotation (to be called before)
  */
  NestingLevel getMaxNestingLevel() { return _maxNestingLevel; }

protected:
  //! computes the nesting level
  InhNestingLevel evaluateInheritedAttribute(SgNode*,InhNestingLevel);
  //! computes the nesting depth
  SynNestingDepth evaluateSynthesizedAttribute(SgNode*,InhNestingLevel,SynthesizedAttributesList);
  //! provides the default value 0 for the nesting depth
  SynNestingDepth defaultSynthesizedAttribute(InhNestingLevel inh);
private:
  NestingLevel _maxNestingLevel;
};


NestingLevel
LoopLevelProcessing::evaluateInheritedAttribute(SgNode* node, NestingLevel loopNestingLevel) {

  //! compute maximum nesting level of entire program in accumulator (member variable)
  if(loopNestingLevel>_maxNestingLevel)
    _maxNestingLevel=loopNestingLevel;

  switch(node->variantT()) {
  case V_SgGotoStatement:
```

Figure 8.21: Example source code showing use of inherited, synthesized, accumulator, and per

```
    cout << "WARNING: Goto statement found. We do not consider goto loops.\n";
 // DQ (11/17/2005): Added return statment to avoid g++ warning: control reaches end of non-void function
    return loopNestingLevel;
    break;
  case V_SgDoWhileStmt:
  case V_SgForStatement:
  case V_SgWhileStmt:
    return loopNestingLevel+1;
  default:
    return loopNestingLevel;
  }
}

SynNestingDepth
LoopLevelProcessing::defaultSynthesizedAttribute(InhNestingLevel inh) {
  /*! we do not need the inherited attribute here
      as default value for synthesized attribute we set 0, representing nesting depth 0.
  */
  return 0;
}

SynNestingDepth
LoopLevelProcessing::evaluateSynthesizedAttribute(SgNode* node,InhNestingLevel nestingLevel,SynthesizedAttribu
   {
     if (nestingLevel>_maxNestingLevel)
          _maxNestingLevel=nestingLevel;

  // compute maximum nesting depth of synthesized attributes
     SynNestingDepth nestingDepth=0;
     for(SynthesizedAttributesList::iterator i=l.begin();i!=l.end();i++)
        {
          if(*i>nestingDepth) nestingDepth=*i;
        }

     switch(node->variantT())
        {
          case V_SgDoWhileStmt:
          case V_SgForStatement:
          case V_SgWhileStmt:
             {
               nestingDepth++;
               cout << "Nesting level:" << nestingLevel << ", nesting depth:" << nestingDepth << endl;
               break;
             }

          default:
             {
            // DQ (11/17/2005): Nothing to do here, but explicit default in switch avoids lots of warnings.
             }
        }

  // add loop nesting level as annotation to AST
     NestingLevelAnnotation* nla = new NestingLevelAnnotation(nestingLevel,nestingDepth);
     ROSE_ASSERT(nla != NULL);

  // DQ (1/2/2006): Added support for new attribute interface.
  // printf ("LoopLevelProcessing::evaluateSynthesizedAttribute(): using new attribute interface \n");
#if 0
     if (node->get_attribute() == NULL)
        {
          AstAttributeMechanism* attributePtr = new AstAttributeMechanism();
          ROSE_ASSERT(attributePtr != NULL);
          node->set_attribute(attributePtr);
        }
#endif
// node->attribute.add("loopNestingInfo",nla);
// node->attribute().add("loopNestingInfo",nla);
   node->addNewAttribute("loopNestingInfo",nla);

  //! return the maximum nesting depth as synthesized attribute
  return nestingDepth;
}

int main ( int argc, char** argv) {

   // command line parameters are passed to EDG
   // non-EDG parameters are passed (through) to ROSE (and the vendor compiler)
   SgProject* root=frontend(argc,argv);
   LoopLevelProcessing t;
```

Figure 8.22: Example source code showing use of inherited, synthesized, accumulator, and per-

```
Output:
Nesting level:5, nesting depth:1
Nesting level:4, nesting depth:2
Nesting level:3, nesting depth:3
Nesting level:2, nesting depth:4
Nesting level:1, nesting depth:5
Nesting level:2, nesting depth:1
Nesting level:1, nesting depth:2
Nesting level:1, nesting depth:1
Max loop nesting level: 5
```

Figure 8.23: Output code showing the result of using inherited, synthesized, and accumulator attributes.

```
#include <rose.h>

class NodeTypeCounter: public AstSimpleProcessing {
public:
    NodeTypeCounter(enum VariantT variant, std::string typeName)
        : myVariant(variant), typeName(typeName), count(0) {
    }

protected:
    virtual void visit(SgNode *node) {
        if (node->variantT() == myVariant) {
            std::cout << "Found " << typeName << std::endl;
            count++;
        }
    }

    virtual void atTraversalEnd() {
        std::cout << typeName << " total: " << count << std::endl;
    }

private:
    enum VariantT myVariant;
    std::string typeName;
    unsigned int count;
};

int main(int argc, char **argv) {
    SgProject *project = frontend(argc, argv);

    std::cout << "sequential execution of traversals" << std::endl;
    NodeTypeCounter forStatementCounter(V_SgForStatement, "for loop");
    NodeTypeCounter intValueCounter(V_SgIntVal, "int constant");
    NodeTypeCounter varDeclCounter(V_SgVariableDeclaration, "variable declaration");
    // three calls to traverse, executed sequentially
    forStatementCounter.traverseInputFiles(project, preorder);
    intValueCounter.traverseInputFiles(project, preorder);
    varDeclCounter.traverseInputFiles(project, preorder);
    std::cout << std::endl;

    std::cout << "combined execution of traversals" << std::endl;
    AstCombinedSimpleProcessing combinedTraversal;
    combinedTraversal.addTraversal(new NodeTypeCounter(V_SgForStatement, "for loop"));
    combinedTraversal.addTraversal(new NodeTypeCounter(V_SgIntVal, "int constant"));
    combinedTraversal.addTraversal(new NodeTypeCounter(V_SgVariableDeclaration, "variable declaration"));
    // one call to traverse, execution is interleaved
    combinedTraversal.traverseInputFiles(project, preorder);
}
```

Figure 8.24: Example source showing the combination of traversals.

```
sequential execution of traversals
Found for loop
Found for loop
for loop total: 2
Found int constant
Found int constant
Found int constant
Found int constant
Found int constant
Found int constant
Found int constant
Found int constant
Found int constant
Found int constant
int constant total: 10
Found variable declaration
Found variable declaration
Found variable declaration
Found variable declaration
Found variable declaration
Found variable declaration
Found variable declaration
Found variable declaration
variable declaration total: 8

combined execution of traversals
Found variable declaration
Found int constant
Found variable declaration
Found for loop
Found variable declaration
Found int constant
Found int constant
Found variable declaration
Found int constant
Found int constant
Found variable declaration
Found int constant
Found variable declaration
Found int constant
Found for loop
Found variable declaration
Found int constant
Found int constant
Found variable declaration
Found int constant
for loop total: 2
int constant total: 10
variable declaration total: 8
```

Figure 8.25: Output of input file to the combined traversals. Note that the order of outputs changes as execution of several analyzers is interleaved.

8.2.12 Short-Circuiting Traversals

The traversal short-circuit mechanism is a simple way to cut short the traversal of a large AST once specific information has been obtained. It is purely an optimization mechanism, and a bit of a hack, but common within the C++ Boost community. Since the technique works we present it as a way of permitting users to avoid the full traversal of an AST that they might deam to be redundant of inappropriate. We don't expect that this mechanism will be particularly useful to most users and we don't recommend it. It may even at some point not be supported. However, we present it because it is a common technique used in the C++ Boost community and it happens to work (at one point it didn't work and so we have no idea what we fixed that permitted it to work now). We have regarded this technique as a rather ugly hack. It is presented in case you really need it. It is, we think, better than the direct use of lower level mechanisms that are used to support the AST traversal.

```
// Input for translator to show exception-based exiting from a translator.

namespace A
{
  int __go__;
  struct B
  {
    static int __stop__;
  };
};

void foo (void)
{
  extern void bar (int);
  bar (A::__go__);
  bar (A::B::__stop__);
}
```

Figure 8.26: Input code with used to demonstrate the traversal short-circuit mechanism.

Figure 8.27 shows the example code demonstrating a traversal setup to support the short-circuit mechanism (a conventional mechanism used often within the C++ Boost community). The input code shown in figure 8.26 is compiled using the example translator, the output is shown in figure 8.28.

The output shown in figure 8.28 demonstrates the initiation of a traversal over the AST and that traversal being short-circuited after a specific point in the evaluation. The result is that there is no further traversal of the AST after that point where it is short-circuited.

```
// Example of an AST traversal that uses the Boost idiom of throwing
// an exception to exit the traversal early.

#include <rose.h>
#include <string>
#include <iostream>

using namespace std;

// Exception to indicate an early exit from a traversal at some node.
class StopEarly
{
public:
  StopEarly (const SgNode* n) : exit_node_ (n) {}
  StopEarly (const StopEarly& e) : exit_node_ (e.exit_node_) {}

  // Prints information about the exit node.
  void print (ostream& o) const
  {
    if (exit_node_) {
      o << '\t' << (const void *)exit_node_ << ":" << exit_node_->class_name () << endl;
      const SgLocatedNode* loc_n = isSgLocatedNode (exit_node_);
      if (loc_n) {
        const Sg_File_Info* info = loc_n->get_startOfConstruct ();
        ROSE_ASSERT (info);
        o << "\tAt " << info->get_filename () << ":" << info->get_line () << endl;
      }
    }
  }

private:
  const SgNode* exit_node_; // Node at early exit from traversal
};

// Preorder traversal to find the first SgVarRefExp of a particular name.
class VarRefFinderTraversal : public AstSimpleProcessing
{
public:
  // Initiate traversal to find 'target' in 'proj'.
  void find (SgProject* proj, const string& target)
  {
    target_ = target;
    traverseInputFiles (proj, preorder);
  }

  void visit (SgNode* node)
  {
    const SgVarRefExp* ref = isSgVarRefExp (node);
    if (ref) {
      const SgVariableSymbol* sym = ref->get_symbol ();
      ROSE_ASSERT (sym);
      cout << "Visiting SgVarRef '" << sym->get_name ().str () << "'" << endl;
      if (sym->get_name ().str () == target_) // Early exit at first match.
        throw StopEarly (ref);
    }
  }

private:
  string target_; // Symbol reference name to find.
};

int main (int argc, char* argv[])
{
  SgProject* proj = frontend (argc, argv);
  VarRefFinderTraversal finder;

  // Look for a reference to "__stop__".
  try {
    finder.find (proj, "__stop__");
    cout << "*** Reference to a symbol '__stop__' not found. ***" << endl;
  } catch (StopEarly& stop) {
    cout << "*** Reference to a symbol '__stop__' found. ***" << endl;
    stop.print (cout);
  }

  // Look for a reference to "__go__".
  try {
    finder.find (proj, "__go__");
    cout << "*** Reference to a symbol '__go__' not found. ***" << endl;
  } catch (StopEarly& go) {
    cout << "*** Reference to a symbol '__go__' found. ***" << endl;
    go.print (cout);
  }
```

```
Visiting SgVarRef '__go__'
Visiting SgVarRef '__stop__'
*** Reference to a symbol '__stop__' found. ***
        0x9717834:SgVarRefExp
        At /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_traversalShortCircuit.C:1
Visiting SgVarRef '__go__'
*** Reference to a symbol '__go__' found. ***
        0x9717808:SgVarRefExp
        At /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_traversalShortCircuit.C:1
```

Figure 8.28: Output code showing the result of short-circuiting the traversal.

8.3 Memory Pool Traversals

Allocation of IR nodes in ROSE is made more efficient through the use of specialized allocators implemented at member function new operators for each class of the IR in Sage III. Such specialized memory allocators avoid significant fragmentation of memory, provide more efficient packing of memory, improve performance of allocation of memory and IR node access, and additionally provide a secondary mechanism to accessing all the IR nodes. Each IR node has a memory pool which is an STL vector of blocks (a fixed or variable sized array of contiguously stored IR nodes).

The three types of traversals are:

1. ROSE Memory Pool Visit Traversal
   This traversal is similar to the one provided by the SimpleProcessing Class (using the visit() function and no inherited or synthesized attributes).

2. Classic Object-Oriented Visitor Pattern for Memory Pool
   This is a classic object-oriented visitor pattern.

3. IR node type traversal, visits one type of IR node for all IR types in the AST. This is useful for building specialized tools.

8.3.1 ROSE Memory Pool Visit Traversal

Figure 8.29 shows the source code for a translator which traverses the memory pool containing the AST. At each node the `visit()` function is called using only the input information represented by the current node. Note that using this simple traversal no context information is available to the visit function. All the IR nodes for a given memory pool are iterated over at one time. The order of the traversal of the different memory pools is random but fixed. Thus the order of the traversal of the IR nodes is in no way connected to the structure of the AST (unlike the previous non-memory pool traversals that were very much tied to the structure of the AST and which matched the structure of the original input source code being compiled).

```
#include "rose.h"

// ROSE Visit Traversal (similar interface as Markus's visit traversal)
// in ROSE (implemented using the traversal over
// the elements stored in the memory pools so it has no cycles and visits
// ALL IR nodes (including all Sg_File_Info, SgSymbols, SgTypes, and the
// static builtin SgTypes).
class RoseVisitor : public ROSE_VisitTraversal
   {
     public:
          int counter;
          void visit ( SgNode* node);

          RoseVisitor() : counter(0) {}
   };


void RoseVisitor::visit ( SgNode* node)
   {
  // printf ("roseVisitor::visit: counter %4d node = %s \n",counter,node->class_name().c_str());
     counter++;
   }

int
main ( int argc, char* argv[] )
   {
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT (project != NULL);

  // ROSE visit traversal
     RoseVisitor visitor;
     visitor.traverseMemoryPool();

     printf ("Number of IR nodes in AST = %d \n",visitor.counter);

     return backend(project);
   }
```

Figure 8.29: Example source showing simple visit traversal over the memory pools.

```
Number of IR nodes in AST = 1794
```

Figure 8.30: Output of input file to the visitor traversal over the memory pool.

8.3.2 Classic Object-Oriented Visitor Pattern for Memory Pool

Figure 8.31 shows the source code for a translator which traverses the memory pools containing the AST. At each node the `visit()` function is called using only the input information represented by the current node. Note that using this simple traversal no context information is available to the visit function. The traversal order is the same as in the 8.29.

```
#include "rose.h"

// Classic Visitor Pattern in ROSE (implemented using the traversal over
// the elements stored in the memory pools so it has no cycles and visits
// ALL IR nodes (including all Sg_File_Info, SgSymbols, SgTypes, and the
// static builtin SgTypes).
class ClassicVisitor : public ROSE_VisitorPattern
   {
     public:
       // Override virtural function defined in base class
          void visit(SgGlobal* globalScope)
             {
               printf ("Found the SgGlobal IR node \n");
             }

          void visit(SgFunctionDeclaration* functionDeclaration)
             {
               printf ("Found a SgFunctionDeclaration IR node \n");
             }
          void visit(SgTypeInt* intType)
             {
               printf ("Found a SgTypeInt IR node \n");
             }

          void visit(SgTypeDouble* doubleType)
             {
               printf ("Found a SgTypeDouble IR node \n");
             }
   };


int
main ( int argc, char* argv[] )
   {
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT (project != NULL);

  // Classic visitor pattern over the memory pool of IR nodes
     ClassicVisitor visitor_A;
     traverseMemoryPoolVisitorPattern(visitor_A);

     return backend(project);
   }
```

Figure 8.31: Example source showing simple visitor pattern.

```
Found a SgTypeInt IR node
Found a SgTypeDouble IR node
Found the SgGlobal IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
Found a SgFunctionDeclaration IR node
```

8.3.3 ROSE IR Type Traversal (uses Memory Pools)

Figure 8.33 shows the source code for a translator which traverses only one type of IR node using the memory pool containing the AST. This traversal is useful for building specialized tools (often tools which only call static functions on each type of IR node).

This example shows the use of an alternative traversal which traverses a representative of each type or IR node just one, but only if it exists in the AST (memory pools). This sort of traversal is useful for building tools that need only operate on static member functions of the IR nodes or need only sample one of each type or IR node present in the AST. this specific example also appears in: *ROSE/src/midend/astDiagnostics/AstStatistics.C*.

The user's use of the traversal is the same as for other ROSE AST traversals except that the ROSE_VisitTraversal::traverseRepresentativeIRnodes() member function is called instead of ROSE_VisitTraversal::traverseMemoryPool().

This mechanism can be used to generate more complete reports of the memory consumption of the AST, which is reported on if *-rose:verbose 2* is used. Figure 8.35 shows a partial snapshot of current IR node frequency and memory consumption for a moderate 40,000 line source code file (one file calling a number of header files), sorted by memory consumption. The AST contains approximately 280K IR nodes. Note that the Sg_File_Info IR nodes is most frequent and consumes the greatest amount of memory, this reflects our bias toward preserving significant information about the mapping of language constructs back to the positions in the source file to support a rich set of source-to-source functionality. *Note: more complete information about the memory use of the AST in in the ROSE User Manual appendix.*

```
// This example code shows the traversal of IR types not available using the other traversal mechanism.
#include "rose.h"
using namespace std;

// CPP Macro to implement case for each IR node (we could alternatively use a visitor pattern and a function t
#define IR_NODE_VISIT_CASE(X) \
          case V_##X: \
             { \
               X* castNode = is##X(node); \
               int numberOfNodes    = castNode->numberOfNodes(); \
               int memoryFootprint = castNode->memoryUsage(); \
               printf ("count = %7d, memory use = %7d bytes, node name = %s \n",numberOfNodes,memoryFootprint,
               break; \
             }

class RoseIRnodeVisitor : public ROSE_VisitTraversal {
     public:
          int counter;
          void visit ( SgNode* node);
          RoseIRnodeVisitor() : counter(0) {}
   };

void RoseIRnodeVisitor::visit ( SgNode* node)
   {
  // Using a classic visitor pattern should avoid all this casting,
  // but each function must be created separately (so it is wash if
  // we want to do all IR nodes, as we do here).
     switch(node->variantT())
        {
          IR_NODE_VISIT_CASE(Sg_File_Info)
          IR_NODE_VISIT_CASE(SgPartialFunctionType)
          IR_NODE_VISIT_CASE(SgFunctionType)
          IR_NODE_VISIT_CASE(SgPointerType)
          IR_NODE_VISIT_CASE(SgFunctionDeclaration)
          IR_NODE_VISIT_CASE(SgFunctionSymbol)
          IR_NODE_VISIT_CASE(SgSymbolTable)
          IR_NODE_VISIT_CASE(SgInitializedName)
          IR_NODE_VISIT_CASE(SgStorageModifier)
          IR_NODE_VISIT_CASE(SgForStatement)
          IR_NODE_VISIT_CASE(SgForInitStatement)
          IR_NODE_VISIT_CASE(SgCtorInitializerList)
          IR_NODE_VISIT_CASE(SgIfStmt)
          IR_NODE_VISIT_CASE(SgExprStatement)
          IR_NODE_VISIT_CASE(SgTemplateDeclaration)
          IR_NODE_VISIT_CASE(SgTemplateInstantiationDecl)
          IR_NODE_VISIT_CASE(SgTemplateInstantiationDefn)
          IR_NODE_VISIT_CASE(SgTemplateInstantiationMemberFunctionDecl)
          IR_NODE_VISIT_CASE(SgClassSymbol)
          IR_NODE_VISIT_CASE(SgTemplateSymbol)
          IR_NODE_VISIT_CASE(SgMemberFunctionSymbol)

          default:
             {
#if 0
               printf ("Case not handled: %s \n",node->class_name().c_str());
#endif
             }
        }
   }


int
main ( int argc, char* argv[] )
   {
  // ROSE visit traversal
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT (project != NULL);

  // ROSE visit traversal
     RoseIRnodeVisitor visitor;
     visitor.traverseRepresentativeIRnodes();
     printf ("Number of types of IR nodes (after building AST) = %d \n",visitor.counter);

#if 1
  // IR nodes statistics
     if (project->get_verbose() > 1)
          cout << AstNodeStatistics::IRnodeUsageStatistics();
#endif

     int errorCode = 0;
     errorCode = backend(project);

     return errorCode;
```

```
count =      12, memory use =     912 bytes, node name = SgSymbolTable
count =     121, memory use =   10648 bytes, node name = SgInitializedName
count =     121, memory use =    2904 bytes, node name = SgStorageModifier
count =     709, memory use =   31196 bytes, node name = Sg_File_Info
No representative for SgPartialFunctionType found in memory pools
count =      84, memory use =    4704 bytes, node name = SgFunctionType
count =      34, memory use =    1496 bytes, node name = SgPointerType
count =       2, memory use =     264 bytes, node name = SgForStatement
count =       2, memory use =      96 bytes, node name = SgForInitStatement
count =       4, memory use =    1088 bytes, node name = SgCtorInitializerList
count =       1, memory use =     140 bytes, node name = SgIfStmt
count =       9, memory use =     360 bytes, node name = SgExprStatement
count =       6, memory use =    1920 bytes, node name = SgTemplateDeclaration
count =       5, memory use =    1820 bytes, node name = SgTemplateInstantiationDecl
count =       1, memory use =     144 bytes, node name = SgTemplateInstantiationDefn
count =       3, memory use =    1368 bytes, node name = SgTemplateInstantiationMemberFunctionDecl
count =      90, memory use =   36720 bytes, node name = SgFunctionDeclaration
count =       1, memory use =      24 bytes, node name = SgClassSymbol
count =       3, memory use =      72 bytes, node name = SgTemplateSymbol
count =       2, memory use =      48 bytes, node name = SgMemberFunctionSymbol
count =      84, memory use =    2016 bytes, node name = SgFunctionSymbol
Number of types of IR nodes (after building AST) = 0
```

Figure 8.34: Output of input file to the IR Type traversal over the memory pool.

```
AST Memory Pool Statistics: numberOfNodes = 114081 memory consumption = 5019564 node = Sg_File_Info
AST Memory Pool Statistics: numberOfNodes =  31403 memory consumption =  628060 node = SgTypedefSeq
AST Memory Pool Statistics: numberOfNodes =  14254 memory consumption =  285080 node = SgStorageModifier
AST Memory Pool Statistics: numberOfNodes =  14254 memory consumption = 1140320 node = SgInitializedName
AST Memory Pool Statistics: numberOfNodes =   8458 memory consumption =  169160 node = SgFunctionParameterTypeList
AST Memory Pool Statistics: numberOfNodes =   7868 memory consumption = 1101520 node = SgModifierType
AST Memory Pool Statistics: numberOfNodes =   7657 memory consumption =  398164 node = SgClassType
AST Memory Pool Statistics: numberOfNodes =   7507 memory consumption = 2071932 node = SgClassDeclaration
AST Memory Pool Statistics: numberOfNodes =   7060 memory consumption =  282400 node = SgTemplateArgument
AST Memory Pool Statistics: numberOfNodes =   6024 memory consumption =  385536 node = SgPartialFunctionType
AST Memory Pool Statistics: numberOfNodes =   5985 memory consumption = 1388520 node = SgFunctionParameterList
AST Memory Pool Statistics: numberOfNodes =   4505 memory consumption = 1477640 node = SgTemplateInstantiationDecl
AST Memory Pool Statistics: numberOfNodes =   3697 memory consumption =  162668 node = SgReferenceType
AST Memory Pool Statistics: numberOfNodes =   3270 memory consumption =  758640 node = SgCtorInitializerList
AST Memory Pool Statistics: numberOfNodes =   3178 memory consumption =   76272 node = SgMemberFunctionSymbol
AST Memory Pool Statistics: numberOfNodes =   2713 memory consumption =  119372 node = SgPointerType
AST Memory Pool Statistics: numberOfNodes =   2688 memory consumption =  161280 node = SgThrowOp
AST Memory Pool Statistics: numberOfNodes =   2503 memory consumption =   60072 node = SgFunctionSymbol
AST Memory Pool Statistics: numberOfNodes =   2434 memory consumption =  107096 node = SgFunctionTypeSymbol
AST Memory Pool Statistics: numberOfNodes =   2418 memory consumption =  831792 node = SgFunctionDeclaration
AST Memory Pool Statistics: numberOfNodes =   2304 memory consumption =   55296 node = SgVariableSymbol
AST Memory Pool Statistics: numberOfNodes =   2298 memory consumption =  101112 node = SgVarRefExp
AST Memory Pool Statistics: numberOfNodes =   2195 memory consumption =  114140 node = SgSymbolTable
AST Memory Pool Statistics: numberOfNodes =   2072 memory consumption =  721056 node = SgMemberFunctionDeclaration
AST Memory Pool Statistics: numberOfNodes =   1668 memory consumption =  400320 node = SgVariableDeclaration
AST Memory Pool Statistics: numberOfNodes =   1667 memory consumption =  393412 node = SgVariableDefinition
AST Memory Pool Statistics: numberOfNodes =   1579 memory consumption =  101056 node = SgMemberFunctionType
AST Memory Pool Statistics: numberOfNodes =   1301 memory consumption =   31224 node = SgTemplateSymbol
AST Memory Pool Statistics: numberOfNodes =   1300 memory consumption =  364000 node = SgTemplateDeclaration
AST Memory Pool Statistics: numberOfNodes =   1198 memory consumption =  455240 node = SgTemplateInstantiationMemberFunctionDecl
AST Memory Pool Statistics: numberOfNodes =   1129 memory consumption =   54192 node = SgIntVal
AST Memory Pool Statistics: numberOfNodes =   1092 memory consumption =   56784 node = SgAssignInitializer
AST Memory Pool Statistics: numberOfNodes =   1006 memory consumption =   52312 node = SgExpressionRoot

Truncated results presented ...
```

Figure 8.35: Example of output using -rose:verbose 2 (memory use report for AST).

Part II

Testing Parts of ROSE

These parts of ROSE are moderately well tested but not tested at a scale sufficent for our own use on large scale applications.

FIXME: *Lay out rules and criteria for the classification of different parts of ROSE.*

Chapter 9

Database Support

This chapter is specific to support in ROSE for persistent storage. ROSE uses the SQLite database and makes it simple to store data in the database for retrieval in later phases of processing large multiple file projects.

FIXME: *Need more information here.*

9.1 ROSE DB Support for Persistent Analysis

This section presents figure 9.3, a simple C++ source code using a template. It is used as a basis for showing how template instantiations are handled within ROSE. An example translator using a database connection to store function information is shown in Fig.9.1 and Fig.9.2. The output by the translator operating on the C++ source code is shown in Fig. 9.4.

9.2 Call Graph for Multi-file Application

This section shows an example of the use of the ROSE Database mechanism where information is stored after processing each file as part of generating the call graph for a project consisting of multiple files. The separate files are show in figures 9.3 and **??**. These files are processed using the translator in figure **??** to generate the final project call graph shown in figure **??**.

FIXME: *This example still needs to be implemented to use the new ROSE call graph generator.*

9.3 Class Hierarchy Graph

This section presents a translator in figure **??**, to generate the class hierarchy graph of the example shown in figure **??**. The input is a multi-file application show in figure **??** and figure **??**. *This example is incomplete.*

FIXME: *This example is still incomplete.*

```
// Example ROSE Translator: used for testing ROSE infrastructure

#include "rose.h"

using namespace std;

// DQ (9/9/2005): Don't include the database by default
#ifdef HAVE_MYSQL
   #include "GlobalDatabaseConnection.h"
#endif

int main( int argc, char * argv[] )
   {
#ifdef HAVE_MYSQL
  // Build the Data base
     GlobalDatabaseConnection *gDB;
     gDB = new GlobalDatabaseConnection( "functionNameDataBase" );
     gDB->initialize();
     string command = "";
     command = command + "CREATE TABLE Functions ( name TEXT, counter );";

     Query *q = gDB->getQuery();
     q->set( command );
     q->execute();

     if ( q->success() != 0 )
          cout << "Error creating schema: " << q->error() << "\n";
  // Alternative syntax, but does not permit access to error messages and exit codes
  // gDB->execute(command.c_str());
#endif

  // Build the AST used by ROSE
     SgProject* project = frontend(argc,argv);

  // Run internal consistency tests on AST
     AstTests::runAllTests(project);

  // Build a list of functions within the AST
     Rose_STL_Container<SgNode*> functionDeclarationList =
          NodeQuery::querySubTree (project,V_SgFunctionDeclaration);

     int counter = 0;
     for (Rose_STL_Container<SgNode*>::iterator i = functionDeclarationList.begin();
               i != functionDeclarationList.end(); i++)
        {
       // Build a pointer to the current type so that we can call
       // the get_name() member function.
          SgFunctionDeclaration* functionDeclaration = isSgFunctionDeclaration(*i);
          ROSE_ASSERT(functionDeclaration != NULL);

          SgName func_name = functionDeclaration->get_name();
      // Skip builtin functions for shorter output, Liao 4/28/2008
          if (func_name.getString().find("__builtin",0)==0)
            continue;

       // output the function number and the name of the function
          printf ("function name#%d is %s at line %d \n",
               counter++,func_name.str(),
               functionDeclaration->get_file_info()->get_line());

          string functionName = functionDeclaration->get_qualified_name().str();

#ifdef HAVE_MYSQL
          command = "INSERT INTO Functions values(\"" + functionName + "\"," +
```

Figure 9.1: Example translator (part 1) using database connection to store function names.

```
                StringUtility::numberToString(counter) + ");";
       // Alternative interface
       // q->set( command );
       // cout << "Executing: " << q->preview() << "\n";
       // q->execute();
          gDB->execute(command.c_str());
#endif
        }

#ifdef HAVE_MYSQL
     command = "SELECT * from Functions;";

  // Alternative Interface (using query objects)
  // q << command;
     q->set(command);
     cout << "Executing: " << q->preview() << "\n";

  // execute and return result (alternative usage: "gDB->select()")
     Result *res = q->store();
     if ( q->success() != 0 )
          cout << "Error reading values: " << q->error() << "\n";
     else
      {
     // Read the table returned from the query
     // res->showResult();
        for ( Result::iterator i = res->begin(); i != res->end(); i++ )
           {
          // Alternative syntax is possible: "Row r = *i;"
             string functionName = (*i)[0].get_string();
             int counter = (*i)[1];
             printf ("functionName = %s counter = %d \n",functionName.c_str(),counter);
           }
      }

     gDB->shutdown();
#else
     printf ("Program compiled without data base connection support (add using ROSE configure option) \n");
#endif

     return 0;
   }
```

Figure 9.2: Example translator (part 2) using database connection to store function names.

```
// This example code is used to record names of functions into the data base.

class A
   {
     public:
          virtual int f1() = 0;
          virtual int f2() {}
          int f3();
          virtual int f4();
   };

int A::f3() { f1(); return f3();}
int A::f4() {}

class B : public A
   {
     public:
          virtual int f1();
          virtual int f2() {}
   };

int B::f1() {}

class C : public A
   {
     public:
          virtual int f1() {}
          int f3() {}
   };

class D : public B
   {
     public:
          virtual int f2() {}
   };

class E : public D
   {
     public:
          virtual int f1() { return 5; }
   };

class G : public E
   {
     public:
          virtual int f1();
   };

int G::f1() {}

class F : public D
   {
     public:
          virtual int f1() {}
          virtual int f2() {return 5;}
          int f3() {return 2;}
   };

class H : public C
   {
     public:
          virtual int f1() {}
          virtual int f2() {}
          int f3() {}
   };
```

Figure 9.3: Example source code used as input to database example.

```
function name #0 is f1 at line 6
function name #1 is f2 at line 7
function name #2 is f3 at line 8
function name #3 is f4 at line 9
function name #4 is f3 at line 12
function name #5 is f4 at line 13
function name #6 is f1 at line 18
function name #7 is f2 at line 19
function name #8 is f1 at line 22
function name #9 is f1 at line 27
function name #10 is f3 at line 28
function name #11 is f2 at line 34
function name #12 is f1 at line 40
function name #13 is f1 at line 46
function name #14 is f1 at line 49
function name #15 is f1 at line 54
function name #16 is f2 at line 55
function name #17 is f3 at line 56
function name #18 is f1 at line 62
function name #19 is f2 at line 63
function name #20 is f3 at line 64
Program compiled without data base connection support (add using ROSE configure option)
```

Figure 9.4: Output from processing input code through database example dataBaseTranslator9.1.

Chapter 10

Recognizing Loops

Figures 10.1 and 10.2 show a translator which reads an application and gathers a list of loop nests. At the end of the traversal it reports information about each loop nest, including the function where it occurred and the depth of the loop nest.

Using this translator we can compile the code shown in figure 10.3. The output is shown in figure 10.4.

FIXME: *This example program is unfinished. It will output a list of objects representing information about perfectly nested loops.*

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.
// rose.C: Example (default) ROSE Preprocessor: used for testing ROSE infrastructure

#include "rose.h"

class InheritedAttribute
   {
     public:
          int loopNestDepth;

          InheritedAttribute () : loopNestDepth(0) {};
          InheritedAttribute ( const InheritedAttribute & X ) {};
   };

class SynthesizedAttribute
   {
     public:
         SynthesizedAttribute() {};
   };

class Traversal : public SgTopDownBottomUpProcessing<InheritedAttribute,SynthesizedAttribute>
   {
     public:
       // Functions required
          InheritedAttribute evaluateInheritedAttribute (
             SgNode* astNode,
             InheritedAttribute inheritedAttribute );

          SynthesizedAttribute evaluateSynthesizedAttribute (
             SgNode* astNode,
             InheritedAttribute inheritedAttribute,
             SubTreeSynthesizedAttributes synthesizedAttributeList );
   };

InheritedAttribute
Traversal::evaluateInheritedAttribute (
     SgNode* astNode,
     InheritedAttribute inheritedAttribute )
   {
     switch(astNode->variantT())
        {
          case V_SgForStatement:
              {
                printf ("Found a SgForStatement \n");

            // This loop is one deepper than the depth of the parent's inherited attribute
                inheritedAttribute.loopNestDepth++;

                break;
              }

          default:
             {
            // g++ needs a block here
             }
        }

     return inheritedAttribute;
   }

SynthesizedAttribute
Traversal::evaluateSynthesizedAttribute (
     SgNode* astNode,
     InheritedAttribute inheritedAttribute,
```

Figure 10.1: Example source code showing loop recognition (part 1).

```
     SubTreeSynthesizedAttributes synthesizedAttributeList )
   {
     SynthesizedAttribute returnAttribute;

     switch(astNode->variantT())
        {
          case V_SgForStatement:
             {

               break;
             }

          default:
             {
            // g++ needs a block here
             }
        }

     return returnAttribute;
   }


int
main ( int argc, char* argv[] )
   {
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT (project != NULL);

  // Build the inherited attribute
     InheritedAttribute inheritedAttribute;

     Traversal myTraversal;

  // Call the traversal starting at the sageProject node of the AST
     myTraversal.traverseInputFiles(project,inheritedAttribute);

     return 0;
   }
```

Figure 10.2: Example source code showing loop recognition (part 2).

```
int main()
   {
     int x[4];
     int y[4][4];

     for (int i=0; i < 4; i++)
        {
          x[i] = 7;
        }

     for (int i=0; i < 4; i++)
        {
          for (int j=0; j < 4; j++)
             {
               y[i][j] = 42;
             }
        }

     return 0;
   }
```

Figure 10.3: Example source code used as input to loop recognition processor.

```
Found a SgForStatement
Found a SgForStatement
Found a SgForStatement
```

Figure 10.4: Output of input to loop recognition processor.

Chapter 11

Function Parameter Types

The analysis of functions often requires the query of the function types. This tutorial example shows how to obtain the function parameter types for any function. Note that functions also have a type which is based on their signature, a combination of their return type and functions parameter types. Any functions sharing the same return type and function parameter types have the same function type (the function type, a SgFunctionType IR node, will be shared between such functions).

Figure 11.1 shows a translator which reads an application (shown in figure 11.2) and outputs information about the function parameter types for each function, shown in figure 11.3. This information includes the order of the function declaration in the global scope, and name of the function, and the types of each parameter declared in the function declaration.

Note that there are a number of builtin functions defined as part of the GNU g++ and gcc compatibility and these are output as well. These are marked as compiler generated functions within ROSE. The code shows how to differentiate between the two different types. Notice also that instantiated template functions are classified as *compiler generated*.

```
// Example ROSE Translator: used within ROSE/tutorial

#include "rose.h"

using namespace std;

int main( int argc, char * argv[] )
   {
  // Build the AST used by ROSE
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT(project != NULL);

  // Build a list of functions within the AST
     Rose_STL_Container<SgNode*> functionDeclarationList = NodeQuery::querySubTree (project,V_SgFunctionDecla

     int functionCounter = 0;
     for (Rose_STL_Container<SgNode*>::iterator i = functionDeclarationList.begin(); i != functionDeclarationI
        {
       // Build a pointer to the current type so that we can call the get_name() member function.
          SgFunctionDeclaration* functionDeclaration = isSgFunctionDeclaration(*i);
          ROSE_ASSERT(functionDeclaration != NULL);

       // DQ (3/5/2006): Only output the non-compiler generated IR nodes
          if ( (*i)->get_file_info()->isCompilerGenerated() == false)
             {
               SgFunctionParameterList* functionParameters = functionDeclaration->get_parameterList();
               ROSE_ASSERT(functionDeclaration != NULL);

             // output the function number and the name of the function
               printf ("Non-compiler generated function name #%3d is %s \n",functionCounter++,functionDeclarat

               SgInitializedNamePtrList & parameterList = functionParameters->get_args();
               int parameterCounter = 0;
               for (SgInitializedNamePtrList::iterator j = parameterList.begin(); j != parameterList.end(); j+
                  {
                    SgType* parameterType = (*j)->get_type();
                    printf ("    parameterType #%2d = %s \n",parameterCounter++,parameterType->unparseToString(
                  }
             }
            else
             {
               printf ("Compiler generated function name #%3d is %s \n",functionCounter++,functionDeclaration-
             }
        }

     return 0;
   }
```

Figure 11.1: Example source code showing how to get type information from function parameters.

```
// Templated class declaration used in template parameter example code
template <typename T>
class templateClass
   {
     public:
          int x;

          void foo(int);
          void foo(double);
   };

// Overloaded functions for testing overloaded function resolution
void foo(int);
void foo(double)
   {
     int x = 1;
     int y;

  // Added to allow non-trivial CFG
     if (x)
        y = 2;
     else
        y = 3;
   }

int main ( int argc, char* argv[] )
   {
     foo(42);
     foo(3.14159265);

     templateClass<char> instantiatedClass;
     instantiatedClass.foo(7);
     instantiatedClass.foo(7.0);

     for (int i=0; i < 4; i++)
        {
          int x;
        }

     return 0;
   }
```

Figure 11.2: Example source code used as input to typeInfoFromFunctionParameters.C.

```
Compiler generated function name #  0 is __builtin_copysign
Compiler generated function name #  1 is __builtin_copysignf
Compiler generated function name #  2 is __builtin_copysignl
Compiler generated function name #  3 is __builtin_acosf
Compiler generated function name #  4 is __builtin_acosl
Compiler generated function name #  5 is __builtin_asinf
Compiler generated function name #  6 is __builtin_asinl
Compiler generated function name #  7 is __builtin_atanf
Compiler generated function name #  8 is __builtin_atanl
Compiler generated function name #  9 is __builtin_atan2f
Compiler generated function name # 10 is __builtin_atan2l
Compiler generated function name # 11 is __builtin_ceilf
Compiler generated function name # 12 is __builtin_ceill
Compiler generated function name # 13 is __builtin_coshf
Compiler generated function name # 14 is __builtin_coshl
Compiler generated function name # 15 is __builtin_floorf
Compiler generated function name # 16 is __builtin_floorl
Compiler generated function name # 17 is __builtin_fmodf
Compiler generated function name # 18 is __builtin_fmodl
Compiler generated function name # 19 is __builtin_frexpf
Compiler generated function name # 20 is __builtin_frexpl
Compiler generated function name # 21 is __builtin_ldexpf
Compiler generated function name # 22 is __builtin_ldexpl
Compiler generated function name # 23 is __builtin_log10f
Compiler generated function name # 24 is __builtin_log10l
Compiler generated function name # 25 is __builtin_modff
Compiler generated function name # 26 is __builtin_modfl
Compiler generated function name # 27 is __builtin_powf
Compiler generated function name # 28 is __builtin_powl
Compiler generated function name # 29 is __builtin_sinhf
Compiler generated function name # 30 is __builtin_sinhl
Compiler generated function name # 31 is __builtin_tanf
Compiler generated function name # 32 is __builtin_tanl
Compiler generated function name # 33 is __builtin_tanhf
Compiler generated function name # 34 is __builtin_tanhl
Compiler generated function name # 35 is __builtin_powil
Compiler generated function name # 36 is __builtin_powi
Compiler generated function name # 37 is __builtin_powif
Compiler generated function name # 38 is __builtin_strchr
Compiler generated function name # 39 is __builtin_strrchr
Compiler generated function name # 40 is __builtin_strpbrk
Compiler generated function name # 41 is __builtin_strstr
Compiler generated function name # 42 is __builtin_nansf
Compiler generated function name # 43 is __builtin_nans
Compiler generated function name # 44 is __builtin_nansl
Compiler generated function name # 45 is __builtin_fabs
Compiler generated function name # 46 is __builtin_fabsf
Compiler generated function name # 47 is __builtin_fabsl
Compiler generated function name # 48 is __builtin_cosf
Compiler generated function name # 49 is __builtin_cosl
Compiler generated function name # 50 is __builtin_sinf
Compiler generated function name # 51 is __builtin_sinl
Compiler generated function name # 52 is __builtin_sqrtf
Compiler generated function name # 53 is __builtin_sqrtl
Compiler generated function name # 54 is __builtin_return_address
Compiler generated function name # 55 is __builtin_frame_address
Compiler generated function name # 56 is __builtin_expect
Compiler generated function name # 57 is __builtin_prefetch
Compiler generated function name # 58 is __builtin_huge_val
Compiler generated function name # 59 is __builtin_huge_valf
Compiler generated function name # 60 is __builtin_huge_vall
Compiler generated function name # 61 is __builtin_inf
Compiler generated function name # 62 is __builtin_inff
Compiler generated function name # 63 is __builtin_infl
Compiler generated function name # 64 is __builtin_nan
Compiler generated function name # 65 is __builtin_nanf
Compiler generated function name # 66 is __builtin_nanl
Compiler generated function name # 67 is __builtin_nans
Compiler generated function name # 68 is __builtin_nansf
Compiler generated function name # 69 is __builtin_nansl
Compiler generated function name # 70 is __builtin_clz
Compiler generated function name # 71 is __builtin_ctz
Compiler generated function name # 72 is __builtin_popcount
Compiler generated function name # 73 is __builtin_parity
Compiler generated function name # 74 is __builtin_ffsl
Compiler generated function name # 75 is __builtin_clzl
Compiler generated function name # 76 is __builtin_ctzl
Compiler generated function name # 77 is __builtin_popcountl
Compiler generated function name # 78 is __builtin_parityl
Compiler generated function name # 79 is __builtin_ffsll
Compiler generated function name # 80 is __builtin_clzll
Compiler generated function name # 81 is __builtin_ctzll
Compiler generated function name # 82 is __builtin_popcountll
```

Chapter 12

Resolving Overloaded Functions

Figure 12.1 shows a translator which reads an application and reposts on the mapping between function calls and function declarations. This is trivial since all overloaded function resolution is done within the frontend and so need not be computed (this is because all type resolution is done in the frontend and stored in the AST explicitly). Other compiler infrastructures often require this to be figured out from the AST, when type resolution is unavailable, and while not too hard for C, this is particularly complex for C++ (due to overloading and type promotion within function arguments).

Figure 12.2 shows the input code used to get the translator. Figure 12.3 shows the resulting output.

```
// Example ROSE Translator: used within ROSE/tutorial

#include "rose.h"

using namespace std;

int main( int argc, char * argv[] )
   {
  // Build the AST used by ROSE
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT(project != NULL);

  // Build a list of functions within the AST
     Rose_STL_Container<SgNode*> functionCallList = NodeQuery::querySubTree (project,V_SgFunctionCallExp);

     int functionCounter = 0;
     for (Rose_STL_Container<SgNode*>::iterator i = functionCallList.begin(); i != functionCallList.end(); i++
        {
          SgExpression* functionExpression = isSgFunctionCallExp(*i)->get_function();
          ROSE_ASSERT(functionExpression != NULL);

          SgFunctionRefExp* functionRefExp = isSgFunctionRefExp(functionExpression);

          SgFunctionSymbol* functionSymbol = NULL;
          if (functionRefExp != NULL)
             {
            // Case of non-member function
               functionSymbol = functionRefExp->get_symbol();
             }
            else
             {
            // Case of member function (hidden in rhs of binary dot operator expression)
               SgDotExp* dotExp = isSgDotExp(functionExpression);
               ROSE_ASSERT(dotExp != NULL);

               functionExpression = dotExp->get_rhs_operand();
               SgMemberFunctionRefExp* memberFunctionRefExp = isSgMemberFunctionRefExp(functionExpression);
               ROSE_ASSERT(memberFunctionRefExp != NULL);

               functionSymbol = memberFunctionRefExp->get_symbol();
             }

          ROSE_ASSERT(functionSymbol != NULL);

          SgFunctionDeclaration* functionDeclaration = functionSymbol->get_declaration();
          ROSE_ASSERT(functionDeclaration != NULL);

       // Output mapping of function calls to function declarations
          printf ("Location of function call #%d at line %d resolved by overloaded function declared at line %
               functionCounter++,
               isSgFunctionCallExp(*i)->get_file_info()->get_line(),
               functionDeclaration->get_file_info()->get_line());
        }

     return 0;
   }
```

Figure 12.1: Example source code showing mapping of function calls to overloaded function declarations.

```
// Templated class declaration used in template parameter example code
template <typename T>
class templateClass
   {
     public:
          int x;

          void foo(int);
          void foo(double);
   };

// Overloaded functions for testing overloaded function resolution
void foo(int);
void foo(double)
   {
     int x = 1;
     int y;

  // Added to allow non-trivial CFG
     if (x)
        y = 2;
     else
        y = 3;
   }

int main()
   {
     foo(42);
     foo(3.14159265);

     templateClass<char> instantiatedClass;
     instantiatedClass.foo(7);
     instantiatedClass.foo(7.0);

     for (int i=0; i < 4; i++)
        {
          int x;
        }

     return 0;
   }
```

Figure 12.2: Example source code used as input to resolveOverloadedFunction.C.

```
Location of function call #0 at line 29 resolved by overloaded function declared at line 14
Location of function call #1 at line 30 resolved by overloaded function declared at line 15
Location of function call #2 at line 33 resolved by overloaded function declared at line 0
Location of function call #3 at line 34 resolved by overloaded function declared at line 0
```

Figure 12.3: Output of input to resolveOverloadedFunction.C.

Chapter 13

Template Parameter Extraction

```
// Example ROSE Translator: used within ROSE/tutorial

#include "rose.h"

using namespace std;

int main( int argc, char * argv[] )
   {
  // Build the AST used by ROSE
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT(project != NULL);

  // Build a list of functions within the AST
     Rose_STL_Container<SgNode*> templateInstantiationDeclList =
          NodeQuery::querySubTree (project,V_SgTemplateInstantiationDecl);

     int classTemplateCounter = 0;
     for (Rose_STL_Container<SgNode*>::iterator i = templateInstantiationDeclList.begin();
          i != templateInstantiationDeclList.end(); i++)
        {
          SgTemplateInstantiationDecl* instantiatedTemplateClass = isSgTemplateInstantiationDecl(*i);
          ROSE_ASSERT(instantiatedTemplateClass != NULL);

        // output the function number and the name of the function
          printf ("Class name #%d is %s \n",
               classTemplateCounter++,
               instantiatedTemplateClass->get_templateName().str());

          const SgTemplateArgumentPtrList& templateParameterList = instantiatedTemplateClass->get_templateArguments();
          int parameterCounter = 0;
          for (SgTemplateArgumentPtrList::const_iterator j = templateParameterList.begin();
               j != templateParameterList.end(); j++)
             {
               printf ("     TemplateArgument #%d = %s \n",parameterCounter++,(*j)->unparseToString().c_str());
             }
        }

     return 0;
   }
```

Figure 13.1: Example source code used to extract template parameter information.

Figure 13.1 shows a translator which reads an application and gathers a list of loop nests. At the end of the traversal it reports information about each instantiated template, including the template arguments.

Figure 13.2 shows the input code used to get the translator. Figure 13.3 shows the resulting output.

```
// Templated class declaration used in template parameter example code
template <typename T>
class templateClass
   {
     public:
          int x;

          void foo(int);
          void foo(double);
   };


int main()
   {
     templateClass<char> instantiatedClass;
     instantiatedClass.foo(7);
     instantiatedClass.foo(7.0);

     templateClass<int> instantiatedClassInt;
     templateClass<float> instantiatedClassFloat;
     templateClass<templateClass<char> > instantiatedClassNestedChar;

     for (int i=0; i < 4; i++)
        {
          int x;
        }

     return 0;
   }
```

Figure 13.2: Example source code used as input to templateParameter.C.

```
Class name #0 is templateClass
   TemplateArgument #0 = char
Class name #1 is templateClass
   TemplateArgument #0 = int
Class name #2 is templateClass
   TemplateArgument #0 = float
Class name #3 is templateClass
   TemplateArgument #0 = templateClass< char >
```

Figure 13.3: Output of input to templateParameter.C.

Chapter 14

Template Support

This chapter is specific to demonstrating the C++ template support in ROSE. *This section is not an introduction to the general subject of C++ templates.* ROSE provides special handling for C++ templates because template instantiation must be controlled by the compiler.

Templates that require instantiation are instantiated by ROSE and can be seen in the traversal of the AST (and transformed). Any templates that can be instantiated by the backend compiler **and** *are not transformed* are not output within the code generation phase.

FIXME: *Provide a list of when templates are generated internally in the AST and when template instantiations are output.*

14.1 Example Template Code #1

This section presents figure 14.4, a simple C++ source code using a template. It is used as a basis for showing how template instantiations are handled within ROSE.

```
template <typename T>
class X
   {
     public:
          void foo();
   };

X<int> x;

void X<int>::foo()
   {
   }
```

Figure 14.1: Example source code showing use of a C++ template.

14.2 Example Template Code #2

This section presents figure 14.4, a simple C++ source code using a template function. It is used as a basis for showing how template instantiations are handled within ROSE.

```
template < typename T >
class X
   {
     public :
         void foo ( );
   };
class X< int  > x;

template<> void X < int  > ::foo()
{
}
```

Figure 14.2: Example source code after processing using identityTranslator (shown in figure 6.1).

```
// template function
template <typename T>
void foo( T t )
   {
   }

// Specialization from user
template<> void foo<int>(int x) {}

int main()
   {
     foo(1);
   }
```

Figure 14.3: Example source code showing use of a C++ template.

```
// template function
template < typename T >
void foo ( T t )
   {
   }
template<> void foo < int > (int t);
// Specialization from user

template<> void foo < int > (int x)
{
}


int main()
{
  foo < int > (1);
  return 0;
}
```

Figure 14.4: Example source code after processing using identityTranslator (shown in figure 6.1).

Chapter 15

AST Construction

AST construction is a fundamental operation needed for building ROSE source-to-source translators. Several levels of interfaces are available in ROSE for users to build AST from scratch. High level interfaces are recommended to use whenever possible for their simplicity. Low level interfaces can give users the maximum freedom to manipulate some details in AST trees.

This chapter uses several examples to demonstrate how to create AST fragments for common language constructs (such as variable declarations, functions, function calls, etc.) and how to insert them into an existing AST tree. More examples of constructing AST using high level interfaces can be found at *rose/tests/roseTests/astInterfaceTests*.

15.1 Variable Declarations

What To Learn Two examples are given to show how to construct a SAGE III AST subtree for a variable declaration and its insertion into the existing AST tree.

- Example 1. Building a variable declaration using the high level AST construction and manipulation interfaces defined in namespace SageBuilder and SageInterface.

  Figure 15.1 shows the high level construction of an AST fragment (a variable declaration) and its insertion into the AST at the top of each block. buildVariableDeclaration() takes the name and type to build a variable declaration node. prependStatement() inserts the declaration at the top of a basic block node. Details for parent and scope pointers, symbol tables, file information and so on are handled transparently.

- Example 2. Building the variable declaration using low level member functions of SAGE III node classes.

  Figure 15.2 shows the low level construction of the same AST fragment (for the same variable declaration) and its insertion into the AST at the top of each block. SgNode constructors and their member functions are used. Side effects for scope, parent pointers and symbol tables have to be handled by programmers explicitly.

Figure 15.3 shows the input code used to test the translator. Figure 15.4 shows the resulting output.

```
// SageBuilder contains all high level buildXXX() functions,
// such as buildVariableDeclaration(), buildLabelStatement() etc.
// SageInterface contains high level AST manipulation and utility functions,
// e.g. appendStatement(), lookupFunctionSymbolInParentScopes() etc.
#include "rose.h"
using namespace SageBuilder;
using namespace SageInterface;

class SimpleInstrumentation:public SgSimpleProcessing
{
public:
  void visit (SgNode * astNode);
};

void
SimpleInstrumentation::visit (SgNode * astNode)
{
  SgBasicBlock *block = isSgBasicBlock (astNode);
  if (block != NULL)
    {
      SgVariableDeclaration *variableDeclaration =
                buildVariableDeclaration ("newVariable", buildIntType ());
      prependStatement (variableDeclaration, block);
    }
}

int
main (int argc, char *argv[])
{
  SgProject *project = frontend (argc, argv);
  ROSE_ASSERT (project != NULL);

  SimpleInstrumentation treeTraversal;
  treeTraversal.traverseInputFiles (project, preorder);

  AstTests::runAllTests (project);
  return backend (project);
}
```

Figure 15.1: AST construction and insertion for a variable using the high level interfaces

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.
// Specifically it shows the design of a transformation to instrument source code, placing source code
// at the top and bottom of each basic block.
// Member functions of SAGE III AST node classes are directly used.
// So all details for Sg_File_Info, scope, parent, symbol tables have to be explicitly handled.

#include "rose.h"

class SimpleInstrumentation : public SgSimpleProcessing
   {
     public:
          void visit ( SgNode* astNode );
   };

void
SimpleInstrumentation::visit ( SgNode* astNode )
   {
     SgBasicBlock* block = isSgBasicBlock(astNode);
     if (block != NULL)
        {
       // Mark this as a transformation (required)
          Sg_File_Info* sourceLocation = Sg_File_Info::generateDefaultFileInfoForTransformationNode();
          ROSE_ASSERT(sourceLocation != NULL);

          SgType* type = new SgTypeInt();
          ROSE_ASSERT(type != NULL);

          SgName name = "newVariable";

          SgVariableDeclaration* variableDeclaration = new SgVariableDeclaration(sourceLocation,name,type);
          ROSE_ASSERT(variableDeclaration != NULL);

          SgInitializedName* initializedName = *(variableDeclaration->get_variables().begin());

       // DQ (6/18/2007): The unparser requires that the scope be set (for name qualification to work).
          initializedName->set_scope(block);

       // Liao (2/13/2008): AstTests requires this to be set
          variableDeclaration->set_firstNondefiningDeclaration(variableDeclaration);

          ROSE_ASSERT(block->get_statements().size() > 0);

          block->get_statements().insert(block->get_statements().begin(),variableDeclaration);
          variableDeclaration->set_parent(block);

       // Add a symbol to the sybol table for the new variable
          SgVariableSymbol* variableSymbol = new SgVariableSymbol(initializedName);
          block->insert_symbol(name,variableSymbol);
        }
   }

int
main ( int argc, char * argv[] )
   {
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT(project != NULL);

     SimpleInstrumentation treeTraversal;
     treeTraversal.traverseInputFiles ( project, preorder );

     AstTests::runAllTests(project);
     return backend(project);
   }
```

Figure 15.2: Example source code to read an input program and add a new variable declaration at the top of each block.

```
int main()
   {
     for (int i=0; i < 4; i++)
        {
          int x;
        }

     return 0;
   }
```

Figure 15.3: Example source code used as input to the translators adding new variable.

```
int main()
{
  int newVariable;
  for (int i = 0; i < 4; i++) {
    int newVariable;
    int x;
  }
  return 0;
}
```

Figure 15.4: Output of input to the translators adding new variable.

15.2 Expressions

Figure 15.5 shows a translator using the high level AST builder interface to add an assignment statement right before the last statement in a main() function.

Figure 15.6 shows the input code used to test the translator. Figure 15.7 shows the resulting output.

```
// Expressions can be built using both bottomup (recommended ) and topdown orders.
// Bottomup: build operands first , operation later
// Topdown: build operation first , set operands later on.

#include "rose.h"
using namespace SageBuilder;
using namespace SageInterface;

int main (int argc, char *argv[])
{
  SgProject *project = frontend (argc, argv);
  // go to the function body
  SgFunctionDeclaration* mainFunc= findMain(project);

  SgBasicBlock* body= mainFunc->get_definition()->get_body();
  pushScopeStack(body);

 // bottomup: build operands first , create expression later on
 //  double result = 2 * (1 - gama * gama);
  SgExpression * init_exp =
            buildMultiplyOp(buildDoubleVal(2.0),
                 buildSubtractOp(buildDoubleVal(1.0),
                      buildMultiplyOp (buildVarRefExp("gama"),buildVarRefExp("gama")
                                       )));
  SgVariableDeclaration* decl = buildVariableDeclaration("result",buildDoubleType(),buildAssignInitializer(init_exp));

  SgStatement* laststmt = getLastStatement(topScopeStack());
  insertStatementBefore(laststmt,decl);

 // topdown: build expression first , set operands later on
 // double result2 = alpha * beta;
  SgExpression * init_exp2 = buildMultiplyOp();
  setLhsOperand(init_exp2 ,buildVarRefExp("alpha"));
  setRhsOperand(init_exp2 ,buildVarRefExp("beta"));

  SgVariableDeclaration* decl2 = buildVariableDeclaration("result2",buildDoubleType(),buildAssignInitializer(init_exp2));
  laststmt = getLastStatement(topScopeStack());
  insertStatementBefore(laststmt,decl2);

  popScopeStack();
  AstTests::runAllTests(project);

  //invoke backend compiler to generate object/binary files
   return backend (project);

}
```

Figure 15.5: Example translator to add expressions

```
int main()
{
  double alpha= 0.5;
  double beta = 0.1;
  double gama = 0.7;

  return 0;
}
```

Figure 15.6: Example source code used as input

```
int main()
{
  double alpha = 0.5;
  double beta = 0.1;
  double gama = 0.7;
  double result = 2 * (1 - gama * gama);
  double result2 = alpha * beta;
  return 0;
}
```

Figure 15.7: Output of the input

15.3 Assignment Statements

Figure 15.8 shows a translator using the high level AST builder interface to add an assignment statement right before the last statement in a main() function.

Figure 15.9 shows the input code used to test the translator. Figure 15.10 shows the resulting output.

```
// SageBuilder contains all high level buildXXX() functions,
// such as buildVariableDeclaration(), buildLabelStatement() etc.
// SageInterface contains high level AST manipulation and utility functions,
// e.g. appendStatement(), lookupFunctionSymbolInParentScopes() etc.
#include "rose.h"
using namespace SageBuilder;
using namespace SageInterface;

int main (int argc, char *argv[])
{
  SgProject *project = frontend (argc, argv);

  // go to the function body of main()
  // and push it to the scope stack
  SgFunctionDeclaration* mainFunc= findMain(project);
  SgBasicBlock* body= mainFunc->get_definition()->get_body();
  pushScopeStack(body);

  // build a variable assignment statement: i=9;
  // buildVarRefExp(string varName) will automatically search for a matching variable symbol starting
  // from the current scope to the global scope.
  SgExprStatement* assignStmt = buildAssignStatement(buildVarRefExp("i"),buildIntVal(9));

 // insert it before the last return statement
  SgStatement* lastStmt = getLastStatement(topScopeStack());
  insertStatementBefore(lastStmt,assignStmt);

  popScopeStack();

  //AstTests ensures there is no dangling SgVarRefExp without a mathing symbol
  AstTests::runAllTests(project);
  return backend (project);
}
```

Figure 15.8: Example source code to add an assignment statement

```
int main(int argc, char* argv[])
{
  int i;
  return 0;
}
```

Figure 15.9: Example source code used as input

```
int main(int argc,char *argv[])
{
  int i;
  i = 9;
  return 0;
}
```

Figure 15.10: Output of the input

15.4 Functions

This section shows how to add a function at the top of a global scope in a file. Again, examples for both high level and low level constructions of AST are given.

- Figure 15.11 shows the high level construction of a defining function (a function with a function body). Scope information is passed to builder functions explicitly when it is needed.

```
// This example shows how to construct a defining function (with a function body)
// using high level AST construction interfaces.
//
#include "rose.h"
using namespace SageBuilder;
using namespace SageInterface;

class SimpleInstrumentation : public SgSimpleProcessing
   {
     public:
          void visit ( SgNode* astNode );
   };

void
SimpleInstrumentation::visit ( SgNode* astNode )
   {
     SgGlobal* globalScope = isSgGlobal(astNode);
     if (globalScope != NULL)
        {
       // ************************************************************
       // Create a parameter list with a parameter
       // ************************************************************
          SgName var1_name = "var_name";
          SgReferenceType *ref_type = buildReferenceType(buildIntType());
          SgInitializedName *var1_init_name = buildInitializedName(var1_name, ref_type);
          SgFunctionParameterList* parameterList = buildFunctionParameterList();
          appendArg(parameterList,var1_init_name);

       // ****************************************************
       // Create a defining functionDeclaration (with a function body)
       // ****************************************************
          SgName func_name                    = "my_function";
          SgFunctionDeclaration * func        = buildDefiningFunctionDeclaration
                    (func_name, buildIntType(), parameterList,globalScope);
          SgBasicBlock*  func_body    = func->get_definition()->get_body();

       // ******************************************************
       // Insert a statement in the function body
       // ******************************************************

          SgVarRefExp *var_ref = buildVarRefExp(var1_name,func_body);
          SgPlusPlusOp *pp_expression = buildPlusPlusOp(var_ref);
          SgExprStatement* new_stmt = buildExprStatement(pp_expression);

       // insert a statement into the function body
          prependStatement(new_stmt,func_body);
          prependStatement(func,globalScope);

        }
   }

int
main ( int argc, char * argv[] )
   {
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT(project != NULL);

     SimpleInstrumentation treeTraversal;
     treeTraversal.traverseInputFiles ( project, preorder );

     AstTests::runAllTests(project);
     return backend(project);
   }
```

Figure 15.11: Addition of function to global scope using high level interfaces

- Figure 15.12 shows almost the same high level construction of the defining function, but

with an additional scope stack. Scope information is passed to builder functions implicitly when it is needed.

```
// This example shows how to construct a defining function (with a function body)
// using high level AST construction interfaces.
// A scope stack is used to pass scope information implicitly to some builder functions
#include "rose.h"
using namespace SageBuilder;
using namespace SageInterface;

int
main ( int argc, char * argv[] )
   {
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT(project != NULL);
     SgGlobal *globalScope = getFirstGlobalScope (project);

 //push global scope into stack
     pushScopeStack (isSgScopeStatement (globalScope));

 // Create a parameter list with a parameter
    SgName var1_name = "var_name";
    SgReferenceType *ref_type = buildReferenceType(buildIntType());
    SgInitializedName *var1_init_name = buildInitializedName(var1_name, ref_type);
    SgFunctionParameterList* parameterList = buildFunctionParameterList();
    appendArg(parameterList,var1_init_name);

 // Create a defining functionDeclaration (with a function body)
    SgName func_name                    = "my_function";
    SgFunctionDeclaration * func        = buildDefiningFunctionDeclaration
                  (func_name, buildIntType(), parameterList);
    SgBasicBlock*  func_body    = func->get_definition()->get_body();

 // push function body scope into stack
    pushScopeStack(isSgScopeStatement(func_body));

 // build a statement in the function body
    SgVarRefExp *var_ref = buildVarRefExp(var1_name);
    SgPlusPlusOp *pp_expression = buildPlusPlusOp(var_ref);
    SgExprStatement* new_stmt = buildExprStatement(pp_expression);

 // insert a statement into the function body
    appendStatement(new_stmt);
//  pop function body off the stack
    popScopeStack();

//  insert the function declaration into the scope at the top of the scope stack
    prependStatement(func);
    popScopeStack();

    AstTests::runAllTests(project);
    return backend(project);
 }
```

Figure 15.12: Addition of function to global scope using high level interfaces and a scope stack

- The low level construction of the AST fragment of the same function declaration and its insertion is separated into two portions and shown in two figures (Figure 15.13 and Figure 15.14).

Figure 39.6 and Figure 39.7 give the input code and output result for the translators above.

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.
// Specifically it shows the design of a transformation to instrument source code, placing source code
// at the top of the source file.

#include "rose.h"

#define TRANSFORMATION_FILE_INFO Sg_File_Info::generateDefaultFileInfoForTransformationNode()

class SimpleInstrumentation : public SgSimpleProcessing
   {
     public:
          void visit ( SgNode* astNode );
   };

void
SimpleInstrumentation::visit ( SgNode* astNode )
   {
     SgGlobal* globalScope = isSgGlobal(astNode);
     if (globalScope != NULL)
        {
       // ************************************************
       // Create the functionDeclaration
       // ************************************************
          SgType * func_return_type           = new SgTypeInt();
          SgName func_name                    = "my_function";
          SgFunctionType * func_type          = new SgFunctionType(func_return_type,false);
          SgFunctionDeclaration * func        = new SgFunctionDeclaration(TRANSFORMATION_FILE_INFO, func_name, func_type);
          SgFunctionDefinition *  func_def    = new SgFunctionDefinition(TRANSFORMATION_FILE_INFO, func);
          SgBasicBlock          * func_body   = new SgBasicBlock(TRANSFORMATION_FILE_INFO);

       // set the end source position as transformation generated
       // since the constructors only set the beginning source position by default
          func->set_endOfConstruct(TRANSFORMATION_FILE_INFO);
          func->get_endOfConstruct()->set_parent(func);

          func_def->set_endOfConstruct(TRANSFORMATION_FILE_INFO);
          func_def->get_endOfConstruct()->set_parent(func_def);

          func_body->set_endOfConstruct(TRANSFORMATION_FILE_INFO);
          func_body->get_endOfConstruct()->set_parent(func_body);

       // Sets the body into the definition
          func_def->set_body(func_body);
       // Sets the defintion's parent to the declaration
          func_def->set_parent(func);

       // DQ (9/8/2007): Fixup the defining and non-defining declarations
          ROSE_ASSERT(func->get_definingDeclaration() == NULL);
          func->set_definingDeclaration(func);
          ROSE_ASSERT(func->get_definingDeclaration()         != NULL);
          ROSE_ASSERT(func->get_firstNondefiningDeclaration() != func);

       // DQ (9/8/2007): We have not build a non-defining declaration, so this should be NULL.
          ROSE_ASSERT(func->get_firstNondefiningDeclaration() == NULL);

       // DQ (9/8/2007): Need to add function symbol to global scope!
          //printf ("Fixing up the symbol table in scope = %p = %s for function = %p = %s \n",globalScope,globalScope->class_name().c_str(),fu
          SgFunctionSymbol* functionSymbol = new SgFunctionSymbol(func);
          globalScope->insert_symbol(func->get_name(),functionSymbol);
          ROSE_ASSERT(globalScope->lookup_function_symbol(func->get_name()) != NULL);

       // ************************************************************
       // Create the InitializedName for a parameter within the parameter list
       // ************************************************************
          SgName var1_name = "var_name";

          SgTypeInt * var1_type      = new SgTypeInt();
          SgReferenceType *ref_type = new SgReferenceType(var1_type);
          SgInitializer * var1_initializer = NULL;

          SgInitializedName *var1_init_name = new SgInitializedName(var1_name, ref_type, var1_initializer, NULL);
          var1_init_name->set_file_info(TRANSFORMATION_FILE_INFO);

       // DQ (9/8/2007): We now test this, so it has to be set explicitly.
          var1_init_name->set_scope(func_def);

       // DQ (9/8/2007): Need to add variable symbol to global scope!
          //printf ("Fixing up the symbol table in scope = %p = %s for SgInitializedName = %p = %s \n",globalScope,globalScope->class_name().c
          SgVariableSymbol *var_symbol = new SgVariableSymbol(var1_init_name);
          func_def->insert_symbol(var1_init_name->get_name(),var_symbol);
```

Figure 15.13: Example source code shows addition of function to global scope (part 1).

```
          ROSE_ASSERT(func_def->lookup_variable_symbol(var1_init_name->get_name()) != NULL);
          ROSE_ASSERT(var1_init_name->get_symbol_from_symbol_table() != NULL);

        // Done constructing the InitializedName variable

        // Insert argument in function parameter list
          ROSE_ASSERT(func != NULL);
        // Sg_File_Info * parameterListFileInfo    = new Sg_File_Info();
        // Sg_File_Info* parameterListFileInfo = Sg_File_Info::generateDefaultFileInfoForTransformationNode();
          SgFunctionParameterList* parameterList = new SgFunctionParameterList(TRANSFORMATION_FILE_INFO);
          ROSE_ASSERT(parameterList != NULL);
          func->set_parameterList(parameterList);
          ROSE_ASSERT(func->get_parameterList() != NULL);
          func->get_parameterList()->append_arg(var1_init_name);

        // *****************************************************
        // Insert a statement in the function body
        // *****************************************************

        // create a VarRefExp
        // SgVariableSymbol *var_symbol = new SgVariableSymbol(var1_init_name);
          SgVarRefExp *var_ref = new SgVarRefExp(TRANSFORMATION_FILE_INFO,var_symbol);
          var_ref->set_endOfConstruct(TRANSFORMATION_FILE_INFO);
          var_ref->get_endOfConstruct()->set_parent(var_ref);

        // create a ++ expression, 0 for prefix ++
          SgPlusPlusOp *pp_expression = new SgPlusPlusOp(TRANSFORMATION_FILE_INFO,var_ref,0);
          pp_expression->set_endOfConstruct(TRANSFORMATION_FILE_INFO);
          pp_expression->get_endOfConstruct()->set_parent(pp_expression);

        // create an expression statement
          SgExprStatement* new_stmt = new SgExprStatement(TRANSFORMATION_FILE_INFO,pp_expression);
          new_stmt->set_endOfConstruct(TRANSFORMATION_FILE_INFO);
          new_stmt->get_endOfConstruct()->set_parent(new_stmt);

#if 0
        // DQ (9/8/2007): This is no longer required, SgExpressionRoot is not longer used in the ROSE IR.
        // create an expression type
          SgTypeInt* expr_type = new SgTypeInt();

        // create an expression root
          SgExpressionRoot * expr_root = new SgExpressionRoot(TRANSFORMATION_FILE_INFO,pp_expression,expr_type);
          expr_root->set_parent(new_stmt);

        // DQ (11/8/2006): Modified to reflect use of SgExpression instead of SgExpressionRoot
          new_stmt->set_expression(expr_root);

          pp_expression->set_parent(new_stmt->get_expression());
#endif
          pp_expression->set_parent(new_stmt);

        // insert a statement into the function body
          func_body->prepend_statement(new_stmt);

        // setting the parent explicitly is not required since it would be done within AST post-processing
          func->set_parent(globalScope);

        // scopes of statments must be set explicitly since within C++ they are not guaranteed
        // to be the same as that indicated by the parent (see ChangeLog for Spring 2005).
          func->set_scope(globalScope);

        // *****************************************************
        // Insert the function declaration in the code
        // *****************************************************
          globalScope->prepend_declaration(func);

        // Required post processing of AST required to set parent pointers and fixup template names, etc.
        // temporaryAstFixes(globalScope);
          AstPostProcessing(globalScope);
         }
    }

int
main ( int argc, char * argv[] )
    {
      SgProject* project = frontend(argc,argv);
      ROSE_ASSERT(project != NULL);

      SimpleInstrumentation treeTraversal;
      treeTraversal.traverseInputFiles ( project, preorder );
```

Figure 15.14: Example source code shows addition of function to global scope (part 2).

```
int main()
   {
     for (int i=0; i < 4; i++)
        {
          int x;
        }

     return 0;
   }
```

Figure 15.15: Example source code used as input to translator adding new function.

```
int my_function(int &var_name)
{
  ++var_name;
}


int main()
{
  for (int i = 0; i < 4; i++) {
    int x;
  }
  return 0;
}
```

Figure 15.16: Output of input to translator adding new function.

15.5 Function Calls

Adding functions calls is a typical task for instrumentation translator.

- Figure 15.17 shows the use of the AST string based rewrite mechanism to add function calls to the top and bottom of each block within the AST.
- Figure 15.18 shows the use of the AST builder interface to do the same instrumentation work.

Figure 15.19 shows the input code used to get the translator. Figure 15.20 shows the resulting output.

Another example shows how to add a function call at the end of each function body. A utility function, *instrumentEndOfFunction()*, from SageInterface name space is used. The interface tries to locate all return statements of a target function and rewriting return expressions with side effects, if there are any. Figure 15.21 shows the translator code. Figure 15.22 shows the input code. The instrumented code is shown in Figure 15.23.

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.
// Specifically it shows the design of a transformation to instrument source code, placing source code
// at the top and bottome of each basic block.

#include "rose.h"

using namespace std;

class SimpleInstrumentation : public SgSimpleProcessing
   {
     public:
          void visit ( SgNode* astNode );
   };

void
SimpleInstrumentation::visit ( SgNode* astNode )
   {
     SgBasicBlock* block = isSgBasicBlock(astNode);
     if (block != NULL)
        {
          const unsigned int SIZE_OF_BLOCK = 1;
          if (block->get_statements().size() > SIZE_OF_BLOCK)
             {
            // It is up to the user to link the implementations of these functions link time
               string codeAtTopOfBlock    = "void myTimerFunctionStart(); myTimerFunctionStart();";
               string codeAtBottomOfBlock = "void myTimerFunctionEnd(); myTimerFunctionEnd();";

            // Insert new code into the scope represented by the statement (applies to SgScopeStatements)
               MiddleLevelRewrite::ScopeIdentifierEnum scope = MidLevelCollectionTypedefs::StatementScope;

            // Insert the new code at the top and bottom of the scope represented by block
               MiddleLevelRewrite::insert(block,codeAtTopOfBlock,scope,
                                          MidLevelCollectionTypedefs::TopOfCurrentScope);
               MiddleLevelRewrite::insert(block,codeAtBottomOfBlock,scope,
                                          MidLevelCollectionTypedefs::BottomOfCurrentScope);
             }
        }
   }

int
main ( int argc, char * argv[] )
   {
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT(project != NULL);

     SimpleInstrumentation treeTraversal;
     treeTraversal.traverseInputFiles ( project, preorder );

     AstTests::runAllTests(project);
     return backend(project);
   }
```

Figure 15.17: Example source code to instrument any input program.

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.
// Specifically it shows the design of a transformation to instrument source code, placing source code
// at the top and bottom of each basic block.

#include "rose.h"
using namespace std;
using namespace SageInterface;
using namespace SageBuilder;

class SimpleInstrumentation:public SgSimpleProcessing
{
public:
  void visit (SgNode * astNode);
};

void
SimpleInstrumentation::visit (SgNode * astNode)
{
  SgBasicBlock *block = isSgBasicBlock (astNode);
  if (block != NULL)
    {
      const unsigned int SIZE_OF_BLOCK = 1;
      if (block->get_statements ().size () > SIZE_OF_BLOCK)
        {
          SgName name1("myTimerFunctionStart");
          // It is up to the user to link the implementations of these functions link time
          SgFunctionDeclaration *decl_1 = buildNondefiningFunctionDeclaration
              (name1,buildVoidType(),buildFunctionParameterList(),block);
          ((decl_1->get_declarationModifier()).get_storageModifier()).setExtern();

          SgExprStatement* callStmt_1 = buildFunctionCallStmt
              (name1,buildVoidType(), buildExprListExp(),block);

          prependStatement(callStmt_1,block);
          prependStatement(decl_1,block);

          SgName name2("myTimerFunctionEnd");
          // It is up to the user to link the implementations of these functions link time
          SgFunctionDeclaration *decl_2 = buildNondefiningFunctionDeclaration
              (name2,buildVoidType(),buildFunctionParameterList(),block);
          ((decl_2->get_declarationModifier()).get_storageModifier()).setExtern();

          SgExprStatement* callStmt_2 = buildFunctionCallStmt
              (name2,buildVoidType(), buildExprListExp(),block);

          appendStatement(decl_2,block);
          appendStatement(callStmt_2,block);
        }
    }
}

int
main (int argc, char *argv[])
{
  SgProject *project = frontend (argc, argv);
  ROSE_ASSERT (project != NULL);

  SimpleInstrumentation treeTraversal;
  treeTraversal.traverseInputFiles (project, preorder);

  AstTests::runAllTests (project);
  return backend (project);
}
```

Figure 15.18: Example source code using the high level interfaces

```
// Overloaded functions for testing overloaded function resolution
void foo(double)
   {
     int x = 1;
     int y;

  // I think that this case fails currently
  // if (x) y = 1; else y = 2;
   }
```

Figure 15.19: Example source code used as input to instrumenting translator.

```
// Overloaded functions for testing overloaded function resolution

void foo(double )
{
  extern void myTimerFunctionStart();
  myTimerFunctionStart();
  int x = 1;
  int y;
  extern void myTimerFunctionEnd();
  myTimerFunctionEnd();
// I think that this case fails currently
// if (x) y = 1; else y = 2;
}
```

Figure 15.20: Output of input to instrumenting translator.

```
/*! \brief  test instrumentation right before the end of a function
*/
#include "rose.h"
#include <iostream>
using namespace SageInterface;
using namespace SageBuilder;

int main (int argc, char *argv[])
{
  SgProject *project = frontend (argc, argv);

  // Find all function definitions we want to instrument
  std::vector<SgNode* > funcDefList =
    NodeQuery::querySubTree (project, V_SgFunctionDefinition);

 std::vector<SgNode*>::iterator iter;
 for (iter = funcDefList.begin(); iter!= funcDefList.end(); iter++)
 {
   SgFunctionDefinition* cur_def = isSgFunctionDefinition(*iter);
   ROSE_ASSERT(cur_def);
  SgBasicBlock* body = cur_def->get_body();
  // Build the call statement for each place
  SgExprStatement* callStmt1 = buildFunctionCallStmt("call1",
               buildIntType(),buildExprListExp() ,body);

  // instrument the function
  int i= instrumentEndOfFunction(cur_def->get_declaration(), callStmt1);
  std::cout<<"Instrumented "<<i<<" places."<<std::endl;

 }  // end of instrumentation

  AstTests::runAllTests(project);
  // translation only
   project->unparse();
}
```

Figure 15.21: Example source code instrumenting end of functions

```
/* Example code:
 *  a function with multiple returns
 *     some returns have expressions with side effects
 *  a function without any return
 */
extern int foo();
extern int call1();
int main(int argc, char* argv[])
{
  if (argc>1)
    return foo();
  else
     return foo();
  return 0;
}

void bar()
{
  int i;
}
```

Figure 15.22: Example input code of the instrumenting translator for end of functions.

```
/* Example code:
 *  a function with multiple returns
 *     some returns have expressions with side effects
 *  a function without any return
 */
extern int foo();
extern int call1();

int main(int argc,char *argv[])
{
  if (argc > 1) {
    int rose_temp__1 = foo();
    call1();
    return rose_temp__1;
  }
  else {
    int rose_temp__2 = foo();
    call1();
    return rose_temp__2;
  }
  call1();
  return 0;
}


void bar()
{
  int i;
  call1();
}
```

Figure 15.23: Output of instrumenting translator for end of functions.

Chapter 16

Loop Optimization

This section is specific to loop optimization and show several tutorial examples using the optimization mechanisms within ROSE.

FIXME: *We might want to reference Qing's work explicitly since this is really just showing off here work.*

16.1 Example Loop Optimizer

Simple example translator showing use of pre-defined loop optimizations.

FIXME: *We are not running performance tests within this tutorial, but perhaps we could later.*

Figure 16.1 shows the code required to call some loop optimizations within ROSE. The translator that we build for this tutorial is simple and takes the following command line options to control which optimizations are done.

```
-ic1 :loop interchange for more reuses
-bk1/2/3 <blocksize> :block outer/inner/all loops
-fs1/2 :single/multi-level loop fusion for more reuses
-cp <copydim> :copy array
-fs0 : loop fission
-splitloop: loop splitting
-unroll [locond] [nvar] <unrollsize> : loop unrolling
-bs <stmtsize> : break up statements in loops
-annot <filename>:
     Read annotation from a file which defines side effects of functions
-arracc <funcname> :
     Use special function to denote array access (the special function can be replaced
     with macros after transformation). This option is for circumventing complex
     subscript expressions for linearized multi-dimensional arrays.
-opt <level=0> : The level of loop optimizations to apply (By default, only the outermost
              level is optimized).
-ta <int> : Max number of nodes to split for transitive dependence analysis (to limit the
         overhead of transitive dep. analysis)
-clsize <int> : set cache line size in evaluating spatial locality (affect decisions in
             applying loop optimizations)
-reuse_dist <int> : set maximum distance of reuse that can exploit cache (used to evaluate
```

```
                temporal locality of loops)
```

16.2 Matrix Multiply Example

Using the matrix multiply example code shown in figure 16.2, we run the loop optimizer in figure 16.1 and generate the code shown in figure 16.3.

16.3 Loop Fusion Example

Using the loop fusion example code shown in figure 16.4, we run the loop optimizer in figure 16.1 and generate the code shown in figure 16.5.

16.4 Example Loop Processor (LoopProcessor.C)

This section contains a more detail translator which uses the command-line for input of specific loop processing options and is more sophisticated than the previous translator used to handle the previous two examples.

Figure 16.6 shows the code required to call the loop optimizations within ROSE. The translator that we build for this tutorial is simple and takes command line parameters to control which optimizations are done.

16.5 Matrix Multiplication Example (mm.C)

Using the matrix multiplication example code shown in figure 16.8, we run the loop optimizer in figure 16.6 and generate the code shown in figure 16.9.

16.6 Matrix Multiplication Example Using Linearized Matrices (dgemm.C)

Using the matrix multiplication example code shown in figure 16.10, we run the loop optimizer in figure 16.6 and generate the code shown in figure 16.11.

16.7 LU Factorization Example (lufac.C)

Using the LU factorization example code shown in figure 16.12, we run the loop optimizer in figure 16.6 and generate the code shown in figure 16.13.

16.8 Loop Fusion Example (tridvpk.C)

Using the loop fusion example code shown in figure 16.14, we run the loop optimizer in figure 16.6 and generate the code shown in figure 16.15.

```
// LoopProcessor:
//   Assume no aliasing
//   apply loop opt to the bodies of all function definitions

//=====================================

#include "rose.h"

#include <AstInterface_ROSE.h>
#include "LoopTransformInterface.h"
#include "CommandOptions.h"

using namespace std;

int
main ( int argc,  char * argv[] )
   {
     vector<string> argvList(argv, argv + argc);
     CmdOptions::GetInstance()->SetOptions(argvList);
     SetLoopTransformOptions(argvList);
     AssumeNoAlias aliasInfo;

     SgProject* project = new SgProject(argvList);

  // Loop over the number of files in the project
     int filenum = project->numberOfFiles();
     for (int i = 0; i < filenum; ++i)
        {
          SgSourceFile* file = isSgSourceFile(project->get_fileList()[i]);
          SgGlobal *root = file->get_globalScope();
          SgDeclarationStatementPtrList& declList = root->get_declarations ();

       // Loop over the declaration in the global scope of each file
          for (SgDeclarationStatementPtrList::iterator p = declList.begin(); p != declList.end(); ++p)
             {
               SgFunctionDeclaration *func = isSgFunctionDeclaration(*p);
               if (func == NULL)
                    continue;
               SgFunctionDefinition *defn = func->get_definition();
               if (defn == NULL)
                    continue;

               SgBasicBlock *stmts = defn->get_body();
               AstInterfaceImpl faImpl(stmts);
               AstInterface fa(&faImpl);

            // This will do as much fusion as possible (finer grained
            // control over loop optimizations uses a different interface).
               LoopTransformTraverse( fa, AstNodePtrImpl(stmts), aliasInfo);

            // JJW 10-29-2007 Adjust for iterator invalidation and possible
            // inserted statements
               p = std::find(declList.begin(), declList.end(), func);
               assert (p != declList.end());
             }
        }

  // Generate source code from AST and call the vendor's compiler
     return backend(project);
   }
```

Figure 16.1: Example source code showing use of loop optimization mechanisms.

```
// Example program showing matrix multiply
// (for use with loop optimization tutorial example)

#define N 50

int main()
   {
     int i,j, k;
     double a[N][N], b[N][N], c[N][N];

     for (i = 0; i <= N-1; i+=1)
        {
          for (j = 0; j <= N-1; j+=1)
              {
                for (k = 0; k <= N-1; k+=1)
                    {
                      c[i][j] = c[i][j] + a[i][k] * b[k][j];
                    }
              }
        }

     return 0;
   }
```

Figure 16.2: Example source code used as input to loop optimization processor.

```
int min2(int a0,int a1)
{
  return a0 < a1?a0:a1;
}

// Example program showing matrix multiply
// (for use with loop optimization tutorial example)
#define N 50

int main()
{
  int i;
  int j;
  int k;
  double a[50UL][50UL];
  double b[50UL][50UL];
  double c[50UL][50UL];
  int _var_0;
  int _var_1;
  for (_var_1 = 0; _var_1 <= 49; _var_1 += 16) {
    for (_var_0 = 0; _var_0 <= 49; _var_0 += 16) {
      for (k = 0; k <= 49; k += 1) {
        for (i = _var_1; i <= min2(49,_var_1 + 15); i += 1) {
          for (j = _var_0; j <= min2(49,_var_0 + 15); j += 1) {
            (c[i])[j] = (((c[i])[j]) + (((a[i])[k]) * ((b[k])[j])));
          }
        }
      }
    }
  }
  return 0;
}
```

Figure 16.3: Output of loop optimization processor showing matrix multiply optimization (using options: `-bk1 -fs0`).

```
main() {

 int x[30], i;

  for (i = 1; i <= 10; i += 1) {
    x[2 * i] = x[2 * i + 1] + 2;
  }
  for (i = 1; i <= 10; i += 1) {
    x[2 * i + 3] = x[2 * i] + i;
  }

}
```

Figure 16.4: Example source code used as input to loop optimization processor.

```
int main()
{
  int x[30UL];
  int i;
  for (i = 1; i <= 11; i += 1) {
    if (i <= 10) {
      x[2 * i] = ((x[(2 * i) + 1]) + 2);
    }
    else {
    }
    if (i >= 2) {
      x[(2 * (-1 + i)) + 3] = ((x[2 * (-1 + i)]) + (-1 + i));
    }
    else {
    }
  }
  return 0;
}
```

Figure 16.5: Output of loop optimization processor showing loop fusion (using options: `-fs2`).

```
#include <general.h>

#include "pre.h"
#include "finiteDifferencing.h"
#include "rose.h"

// DQ (1/2/2008): I think this is no longer used!
// #include "copy_unparser.h"

#include "rewrite.h"
#include <CommandOptions.h>
#include <AstInterface_ROSE.h>
#include <LoopTransformInterface.h>
#include <AnnotCollect.h>
#include <OperatorAnnotation.h>

using namespace std;

#ifdef USE_OMEGA
#include <DepTestStatistics.h>

extern DepTestStatistics DepStats;
#endif

extern bool DebugAnnot();
extern void FixFileInfo(SgNode* n);
class UnparseFormatHelp;
class UnparseDelegate;
void unparseProject( SgProject* project, UnparseFormatHelp* unparseHelp /*= NULL*/, UnparseDelegate *repl  /*=

void PrintUsage( char* name)
{
  cerr << name << " <options>" << "<program name>" << "\n";
  cerr << "-gobj: generate object file\n";
  cerr << "-orig: copy non-modified statements from original file\n";
  cerr << "-splitloop: applying loop splitting to remove conditionals inside loops\n";
  cerr << ReadAnnotation::get_inst()->OptionString() << endl;
  cerr << "-pre:  apply partial redundancy elimination\n";
  cerr << "-fd:  apply finite differencing to array index expressions\n";
  PrintLoopTransformUsage( cerr );
}
```

Figure 16.6: Detailed example source code showing use of loop optimization mechanisms (loopProcessor.C part 1).

```
bool GenerateObj()
{
  return CmdOptions::GetInstance()->HasOption("-gobj");
}


int
main ( int argc,  char * argv[] )
{

  if (argc <= 1) {
      PrintUsage(argv[0]);
      return -1;
  }
#if 0
  CmdOptions::GetInstance()->SetOptions(argc, argv);
  SetLoopTransformOptions(argc, argv);

  OperatorSideEffectAnnotation *funcInfo =
         OperatorSideEffectAnnotation::get_inst();
  funcInfo->register_annot();
  ReadAnnotation::get_inst()->read();
  AssumeNoAlias aliasInfo;

  vector<string> argvList(argv, argv + argc);
  SgProject sageProject ( argvList);
#else
// DQ (2/10/2008): Using command-line support similar to that in tests/roseTests/loopProcessor
  vector<string> argvList(argv, argv + argc);
  SetLoopTransformOptions(argvList);
  CmdOptions::GetInstance()->SetOptions(argvList);

#ifdef USE_OMEGA
  DepStats.SetFileName(buffer.str());
#endif

  OperatorSideEffectAnnotation *funcInfo =
         OperatorSideEffectAnnotation::get_inst();
  funcInfo->register_annot();
  ReadAnnotation::get_inst()->read();
  if (DebugAnnot())
```

Figure 16.7: loopProcessor.C source code (Part 2).

```
#define N 50

void printmatrix( double x[][N]);
void initmatrix( double x[][N], double s);

main()
{
  int i,j, k;
  double a[N][N], b[N][N], c[N][N];

  double s;
  s = 235.0;
  initmatrix(a, s);
  s = 321.0;
  initmatrix(b, s);

  printmatrix(a);
  printmatrix(b);
  for (i = 0; i <= N-1; i+=1) {
    for (j = 0; j <= N-1; j+=1) {
       for (k = 0; k <= N-1; k+=1) {
          c[i][j] = c[i][j] + a[i][k] * b[k][j];
       }
    }
  }

  printmatrix(c);
}
```

Figure 16.8: Example source code used as input to loopProcessor, show in figure 16.6.

```
int min2(int a0,int a1)
{
  return a0 < a1?a0:a1;
}

#define N 50
extern void printmatrix(double x[][50UL]);
extern void initmatrix(double x[][50UL],double s);

int main()
{
  int i;
  int j;
  int k;
  double a[50UL][50UL];
  double b[50UL][50UL];
  double c[50UL][50UL];
  double s;
  int _var_0;
  int _var_1;
  s = 235.0;
  initmatrix(a,s);
  s = 321.0;
  initmatrix(b,s);
  printmatrix(a);
  printmatrix(b);
  for (_var_1 = 0; _var_1 <= 49; _var_1 += 16) {
    for (_var_0 = 0; _var_0 <= 49; _var_0 += 16) {
      for (k = 0; k <= 49; k += 1) {
        for (i = _var_1; i <= min2(49,_var_1 + 15); i += 1) {
          for (j = _var_0; j <= min2(49,_var_0 + 15); j += 1) {
            (c[i])[j] = (((c[i])[j]) + (((a[i])[k]) * ((b[k])[j])));
          }
        }
      }
    }
  }
  printmatrix(c);
  return 0;
}
```

Figure 16.9: Output of loopProcessor using input from figure 16.8 (using options: `-bk1 -fs0`).

```
// Function prototype
void dgemm(double *a, double *b, double *c, int n);

// Function definition
void dgemm(double *a, double *b, double *c, int n)
   {
     int i, j, k;

  // int n;

     for(k=0;k<n;k+=1){
       for(j=0;j<n;j+=1){
         for(i=0;i<n;i+=1){
           c[j*n+i]=c[j*n+i]+a[k*n+i]*b[j*n+k];
         }
       }
     }
   }
```

Figure 16.10: Example source code used as input to loopProcessor, show in figure 16.6.

```
int min2(int a0,int a1)
{
  return a0 < a1?a0:a1;
}

// Function prototype
extern void dgemm(double *a,double *b,double *c,int n);
// Function definition

void dgemm(double *a,double *b,double *c,int n)
{
  int i;
  int j;
  int k;
  int _var_0;
  int _var_1;
  for (_var_1 = 0; _var_1 <= -1 + n; _var_1 += 16) {
    for (_var_0 = 0; _var_0 <= -1 + n; _var_0 += 16) {
      for (i = 0; i <= -1 + n; i += 1) {
        for (k = _var_1; k <= min2(-1 + n,_var_1 + 15); k += 1) {
          for (j = _var_0; j <= -15 + min2(-1 + n,_var_0 + 15); j += 16) {
            int _var_2 = j;
            c[(j * n) + i] = ((c[(j * n) + i]) + ((a[(k * n) + i]) * (b[(j * n) + k])));
            _var_2 = 1 + _var_2;
            c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])));
            _var_2 = 1 + _var_2;
            c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])));
            _var_2 = 1 + _var_2;
            c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])));
            _var_2 = 1 + _var_2;
            c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])));
            _var_2 = 1 + _var_2;
            c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])));
            _var_2 = 1 + _var_2;
            c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])));
            _var_2 = 1 + _var_2;
            c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])));
            _var_2 = 1 + _var_2;
            c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])));
            _var_2 = 1 + _var_2;
            c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])));
            _var_2 = 1 + _var_2;
            c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])));
            _var_2 = 1 + _var_2;
            c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])));
            _var_2 = 1 + _var_2;
            c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])));
            _var_2 = 1 + _var_2;
            c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])));
            _var_2 = 1 + _var_2;
            c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])));
            _var_2 = 1 + _var_2;
            c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])));
          }
          for (; j <= min2(-1 + n,_var_0 + 15); j += 1) {
            c[(j * n) + i] = ((c[(j * n) + i]) + ((a[(k * n) + i]) * (b[(j * n) + k])));
          }
        }
      }
    }
  }
}
```

Figure 16.11: Output of loopProcessor using input from figure 16.10 (using options: `-bk1 -unroll nvar 16`).

```
double abs(double x) { if (x < 0) return -x; else return x; }

#define n 50
void printmatrix( double x[][n]);
void initmatrix( double x[][n], double s);

main(int argc, char* argv[]) {
int p[n], i, j, k;
double a[n][n], mu, t;

initmatrix(a, 5.0);
printmatrix(a);

for (k = 0; k<=n-2; k+=1) {
    p[k] = k;
    mu = abs(a[k][k]);
    for (i = k+1; i <= n-1; i+=1) {
      if (mu < abs(a[i][ k])) {
        mu = abs(a[i][k]);
        p[k] = i;
      }
    }

    for (j = k; j <= n-1; j+=1) {
       t = a[k][j];
       a[k][j] = a[p[k]][j];
       a[p[k]][j] = t;
    }

    for (i = k+1; i <= n-1; i+=1) {
          a[i][k] = a[i][k]/a[k][k];
    }
    for (j = k+1; j <=n-1; j+=1) {
       for (i = k+1; i <=n-1; i+=1) {
            a[i][j] = a[i][j] - a[i][k]*a[k][j];
       }
    }
 }

printmatrix(a);
}
```

Figure 16.12: Example source code used as input to loopProcessor, show in figure 16.6.

```
double abs(double x)
{
  if (x < (0))
    return -x;
  else
    return x;
}

#define n 50
extern void printmatrix(double x[][50UL]);
extern void initmatrix(double x[][50UL],double s);

int main(int argc,char *argv[])
{
  int p[50UL];
  int i;
  int j;
  int k;
  double a[50UL][50UL];
  double mu;
  double t;
  initmatrix(a,5.0);
  printmatrix(a);
  for (k = 0; k <= 48; k += 1) {
    p[k] = k;
    mu = abs(((a[k])[k]));
    for (i = 1 + k; i <= 49; i += 1) {
      if (mu < abs(((a[i])[k]))) {
        mu = abs(((a[i])[k]));
        p[k] = i;
      }
    }
    for (j = k; j <= 49; j += 1) {
      t = ((a[k])[j]);
      (a[k])[j] = ((a[p[k]])[j]);
      (a[p[k]])[j] = t;
    }
    for (i = 1 + k; i <= 49; i += 1) {
      (a[i])[k] = (((a[i])[k]) / ((a[k])[k]));
    }
    for (j = 1 + k; j <= 49; j += 1) {
      for (i = 1 + k; i <= 49; i += 1) {
        (a[i])[j] = (((a[i])[j]) - (((a[i])[k]) * ((a[k])[j])));
      }
    }
  }
  printmatrix(a);
  return 0;
}
```

Figure 16.13: Output of loopProcessor using input from figure 16.12 (using options: `-bk1 -fs0 -splitloop -annotation`).

```
#define n 100

  double a[n], b[n], c[n], d[n], e[n];
  double tot[n][ n];
  double dux[n][n][n], duy[n][n][n], duz[n][n][n];

main()
{
  int i,j,k;
      for ( j=0; j <=n-1; j+=1)
      for ( i=0; i<=n-1; i+=1)
         duz[i][j][0] = duz[i][j][0]*b[0];

      for (k=1; k<=n-2; k+=1)
      for (j = 0; j <= n-1; j+=1)
      for (i=0; i<=n-1; i+=1)
         duz[i][j][k]=(duz[i][j][k]-a[k]*duz[i][j][k-1])*b[k];

      for ( j=0; j <=n-1; j+=1)
      for ( i=0; i<=n-1; i+=1)
         tot[i][j] = 0;

      for ( k=0; k<=n-2; k+=1)
      for ( j=0; j <=n-1; j+=1)
      for ( i=0; i<=n-1; i+=1)
           tot[i][j] = tot[i][j] + d[k]*duz[i][j][k];

      for ( j=0; j <=n-1; j+=1)
      for ( i=0; i<=n-1; i+=1)
         duz[i][j][n-1] = (duz[i][j][n-1] - tot[i][j])*b[n-1];

      for ( j=0; j <=n-1; j+=1)
      for ( i=0; i<=n-1; i+=1)
         duz[i][j][n-2]=duz[i][j][n-2] - e[n-2]*duz[i][j][n-1];

      for (k=n-3; k>=0; k+=-1)
      for (j = 0; j <= n-1; j+=1)
      for (i=0; i<=n-1; i+=1)
         duz[i][j][k] = duz[i][j][k] - c[k]*duz[i][j][k+1] - e[k]*duz[i][j][n-1];

}
```

Figure 16.14: Example source code used as input to loopProcessor, show in figure 16.6.

```
#define n 100
double a[100UL];
double b[100UL];
double c[100UL];
double d[100UL];
double e[100UL];
double tot[100UL][100UL];
double dux[100UL][100UL][100UL];
double duy[100UL][100UL][100UL];
double duz[100UL][100UL][100UL];

int main()
{
  int i;
  int j;
  int k;
  for (i = 0; i <= 99; i += 1) {
    for (j = 0; j <= 99; j += 1) {
      (tot[i])[j] = (0);
      ((duz[i])[j])[0] = ((((duz[i])[j])[0]) * (b[0]));
      for (k = 0; k <= 98; k += 1) {
        if (k >= 1) {
          ((duz[i])[j])[k] = (((((duz[i])[j])[k]) - ((a[k]) * (((duz[i])[j])[k - 1]))) * (b[k]));
        }
        else {
        }
        (tot[i])[j] = (((tot[i])[j]) + ((d[k]) * (((duz[i])[j])[k])));
      }
      ((duz[i])[j])[(100 - 1)] = (((((duz[i])[j])[(100 - 1)]) - ((tot[i])[j])) * (b[(100 - 1)]));
      ((duz[i])[j])[(100 - 2)] = ((((duz[i])[j])[(100 - 2)]) - ((e[(100 - 2)]) * (((duz[i])[j])[(100 - 1)])));
      for (k = 97; k >= 0; k += -1) {
        ((duz[i])[j])[k] = (((((duz[i])[j])[k]) - ((c[k]) * (((duz[i])[j])[k + 1]))) - ((e[k]) * (((duz[i])[j])[(100 - 1
      }
    }
  }
  return 0;
}
```

Figure 16.15: Output of loopProcessor input from figure 16.14 (using options: `-fs2 -ic1 -opt 1`).

Part III

Experimental Parts of ROSE

These parts of ROSE are either new or not well tested on real applications.

FIXME: *Lay out rules and criteria for the classification of different parts of ROSE.*

Chapter 17

Generating Control Flow Graphs

The control flow of a program is broken into *basic blocks* as nodes with control flow forming edges between the basic blocks. Thus the control flow forms a graph which often labeled edges (true and false), and basic blocks representing sequentially executed code. This chapter presents the Control Flow Graph (CFG) and the ROSE application code for generating such graphs for any function in an input code. The CFG forms a fundamental building block for more complex forms of program analysis.

Figure 17.1 shows the code required to generate the control flow graph for each function of an application. Using the input code shown in figure 17.2 the first function's control flow graph is shown in figure 17.3.

Figure 17.3 shows the control flow graph for the function in the input code in figure 17.2.

```
// Example ROSE Translator: used within ROSE/tutorial

#include "rose.h"
#include <GraphUpdate.h>
#include "CFGImpl.h"
#include "GraphDotOutput.h"
#include "controlFlowGraph.h"
#include "CommandOptions.h"

using namespace std;

// Use the ControlFlowGraph is defined in both PRE
// and the DominatorTreesAndDominanceFrontiers namespaces.
// We want to use the one in the PRE namespace.
using namespace PRE;

class visitorTraversal : public AstSimpleProcessing
   {
     public:
          virtual void visit(SgNode* n);
   };

void visitorTraversal::visit(SgNode* n)
   {
     SgFunctionDeclaration* functionDeclaration = isSgFunctionDeclaration(n);
     if (functionDeclaration != NULL)
        {
          SgFunctionDefinition* functionDefinition = functionDeclaration->get_definition();
          if (functionDefinition != NULL)
             {
               SgBasicBlock* functionBody = functionDefinition->get_body();
               ROSE_ASSERT(functionBody != NULL);

               ControlFlowGraph controlflow;

             // The CFG can only be called on a function definition (at present)
               makeCfg(functionDefinition,controlflow);
               string fileName = functionDeclaration->get_name().str();
               fileName += ".dot";
               ofstream dotfile(fileName.c_str());
               printCfgAsDot(dotfile, controlflow);
             }
        }
   }

int main( int argc, char * argv[] )
   {
  // Build the AST used by ROSE
     SgProject* project = frontend(argc,argv);

     CmdOptions::GetInstance()->SetOptions(argc,argv);

  // Build the traversal object
     visitorTraversal exampleTraversal;

  // Call the traversal starting at the project node of the AST
     exampleTraversal.traverseInputFiles(project,preorder);

     return 0;
   }
```

Figure 17.1: Example source code showing visualization of control flow graph.

```
#include <stdio.h>
#include <string.h>
#include <assert.h>

int main(int argc, char *argv[])
{
  int i;
  char buffer[10];
  for (i=0; i < strlen(argv[1]); i++)
  {
 // Buffer overflow for strings of over 9 characters
    assert(i < 10);
    buffer[i] = argv[1][i];
  }
  return 0;
}

#if 0
void
bar(int& w)
   {
     ++w;
   }

int
main(int, char**)
   {
     int z = 3;
     int a = 5 + z + 9;
     int b = (6 - z) * (a + 2) - 3;
     bar(b);
     while (b - 7 > 0)
        {
          b-=5;
          --b;
        }

     do {
          --b; LLL: if (b <= -999) return 0;
        }
     while (b > 2);

     for (b = 0; b < 10; ++b)
          ++z;

     for (int z2 = 7 + z * 5; z2 + 9 < b % 10; ++*(&z2 + 5 - 5))
        {
          (a += 7) += 7;
          ++(++a);
        }
     b = -999;
     goto LLL;
   }
#endif
```

Figure 17.2: Example source code used as input to build control flow graph.

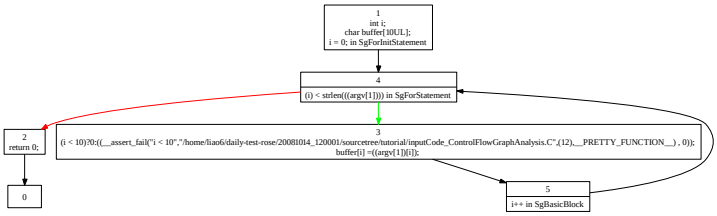

Figure 17.3: Control flow graph for function in input code file: inputCode_1.C.

Chapter 18

Runtime Error Check

This chapter demonstrates the capabilities of ROSE to check for runtime errors. Instead of a segmentation fault, a valid error message with line number and associated file is printed.

For this, the original code is instrumented (transformed) with calls to a function that checks whether a variable is NULL at any given location in the code. In particular, this example searches (amongst other things) for all VarRefExp expressions in a program. When found, a function call to check_var(SgNode* n, std::string desc) is inserted before the statement containing that expression. The first parameter to check_var contains the variable to check and the second parameter the line number and file description.

18.1 Interface

The runtime check is performed by the following lines of code:

```
SgProject* project  = frontend(argc,argv);
RuntimeInstrumentation* runtime = new RuntimeInstrumentation();
runtime->run(project);
return backend(project);
```

18.2 Example 1

As an example, consider the input program in Figure 18.1.

ROSE traverses the program and inserts check_var() statements at locations before the variables are read. Besides VarRefExp, the runtime error checker looks also for isSgPointerDerefExp and isSgPntrArrRefExp expressions. The output program can be found in Figure 18.2.

Finally, if the new code is executed – the obvious segmentation fault is caught and the following error message is printed to the screen, cf. Figure 18.3.

18.3 Example 2

The following demonstrates a second example. The input code is illustrated in Figure 18.4.

```
using namespace std;

int main(int argc, char** argv) {
  int* x= new int[2];
  int* k ;
  x[1]=2;
  x[2]=3;
  int y=x[5];
  int a=x[0]/y;

  x[0]=*k;
}
```

Figure 18.1: Example source code.

```
#include <iostream>
void check_var(void *n,char* desc);
using namespace std;

int main(int argc,char **argv)
{
  int *x = new int [2UL];
  int *k;
  x[1] = 2;
  x[2] = 3;
  check_var(((void *)(x[5])),"(x[5])==NULL on line: 8\
  in File : /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/input_runtimeErrorCheck.C");
  int y = (x[5]);
  check_var(((void *)y),"y==NULL on line: 9\
  in File : /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/input_runtimeErrorCheck.C");
  check_var(((void *)(x[0])),"(x[0])==NULL on line: 9\
  in File : /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/input_runtimeErrorCheck.C");
  int a = ((x[0]) / y);
  check_var(((void *)k),"k==NULL on line: 11\
  in File : /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/input_runtimeErrorCheck.C");
  check_var(((void *)( *k)),"*k==NULL on line: 11\
  in File : /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/input_runtimeErrorCheck.C");
  x[0] =  *k;
  return 0;
}


void check_var(void *n,char* desc)
{
  if (n == 0) {
    printf("\n\nERROR detected: %s \n",desc);
    printf("Aborting this program. Goodbye. \n");
    exit(0);
  }
  else {
  }
}
```

Figure 18.2: Transformed source code.

Figure 18.3: Execution of transformed source code.

```
using namespace std;

class Thomas {
public:
  int x;
};

int main(int argc, char** argv) {
  Thomas* thomas = new Thomas();
  thomas->x=5;
  Thomas tom;
  tom.x=6;

  Thomas** hans = &thomas;
  int x = (*hans)->x;
  thomas=0;
  x = (*hans)->x;
  x = thomas->x;
}
```

Figure 18.4: Example source code.

The transformed code is shown in Figure 18.5.
The new program when executed prints the error message shows in Figure 18.6.

```
#include <iostream>
void check_var(void *n,char* desc);
using namespace std;

class Thomas
{
  public: int x;
}

;

int main(int argc,char **argv)
{
  class Thomas *thomas = ::new Thomas ;
  check_var(((void *)thomas),"thomas -> Thomas::x==NULL on line: 10\
  in File : /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/input_runtimeErrorCheck2.C");
  thomas -> Thomas::x = 5;
  class Thomas tom;
  check_var(((void *)(&tom)),"tom.Thomas::x==NULL on line: 12\
  in File : /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/input_runtimeErrorCheck2.C");
  tom.Thomas::x = 6;
  check_var(((void *)thomas),"thomas==NULL on line: 14\
  in File : /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/input_runtimeErrorCheck2.C");
  class Thomas **hans = &thomas;
  check_var(((void *)(&hans)),"(( *( *hans)).Thomas::x)==NULL on line: 15\
  in File : /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/input_runtimeErrorCheck2.C");
  check_var(((void *)(&( *hans))),"( *hans)==NULL on line: 15\
  in File : /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/input_runtimeErrorCheck2.C");
  check_var(((void *)(&( *( *hans)))),"( *( *hans))==NULL on line: 15\
  in File : /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/input_runtimeErrorCheck2.C");
  int x = (( *( *hans)).Thomas::x);
  thomas = ((0));
  check_var(((void *)(&hans)),"(( *( *hans)).Thomas::x)==NULL on line: 17\
  in File : /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/input_runtimeErrorCheck2.C");
  check_var(((void *)(&( *hans))),"( *hans)==NULL on line: 17\
  in File : /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/input_runtimeErrorCheck2.C");
  check_var(((void *)(&( *( *hans)))),"( *( *hans))==NULL on line: 17\
  in File : /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/input_runtimeErrorCheck2.C");
  x = (( *( *hans)).Thomas::x);
  check_var(((void *)thomas),"(thomas -> Thomas::x)==NULL on line: 18\
  in File : /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/input_runtimeErrorCheck2.C");
  x = (thomas -> Thomas::x);
  return 0;
}


void check_var(void *n,char* desc)
{
  if (n == 0) {
    printf("\n\nERROR detected: %s \n",desc);
    printf("Aborting this program. Goodbye. \n");
    exit(0);
  }
  else {
  }
}
```

Figure 18.5: Transformed source code.

```
ERROR detected: (*(*hans))==NULL on line: 17  in File : /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial
Aborting this program. Goodbye.
```

Figure 18.6: Execution of transformed source code.

Chapter 19

Dataflow Analysis

The dataflow analysis in Rose is based on the control flow graph (CFG). One type of dataflow analysis is called def-use analysis, which is explained next.

19.1 Def-Use Analysis

The definition-usage (def-use) analysis allows to query the definition and usage for each *control flow node* (CFN). Any statement or expression within ROSE is represented as a sequence of CFN's. For instance, the CFG for the following program

```
int main()
{
  int x = 9;
  x = x + 1;
}
```

Figure 19.1: Example source code.

is illustrated in Figure 19.2. For each CFN in the CFG, the definition and usage for variable references can be determined with the public function calls:

```
vector <SgNode*> getDefFor(SgNode*, SgInitializedName*)
vector <SgNode*> getUseFor(SgNode*, SgInitializedName*)
```

where SgNode* represents any control flow node and SgInitializedName any variable (being used or defined at that CFN). The result is a vector of possible CFN's that either define (getDefFor) or use (getUseFor) a specific variable.

Figure 19.2 shows how the variable x is being declared and defined in CFN's between node 1 and 6. Note that the definition is annotated along the edge. For instance at node 6, the edge reads *(6) DEF: x (3) = 5*. This means that variable x was declared at CFN 3 but defined at CFN 5.

The second statement x=x+1 is represented by CFN's from 7 to 12. One can see in the figure that x is being re-defined at CFN 11. However, the definition of x within the same statement

happens at CFN 8. Hence, the definition of the right hand side x in the statement is at CFN 5 : *(8) DEF: x (3) = 5.*

Another usage of the def-use analysis is to determine which variables actually are defined at each CFN. The following function allows to query a CFN for all its variables (SgInitializedNames) and the positions those variables are defined (SgNode):

```
std::multimap <SgInitializedName*, SgNode*> getDefMultiMapFor(SgNode*)
std::multimap <SgInitializedName*, SgNode*> getUseMultiMapFor(SgNode*)
```

All public functions are described in *DefuseAnalysis.h*. To use the def-use analysis, one needs to create an object of the class DefUseAnalysis and execute the run function. After that, the described functions above help to evaluate definition and usage for each CFN.

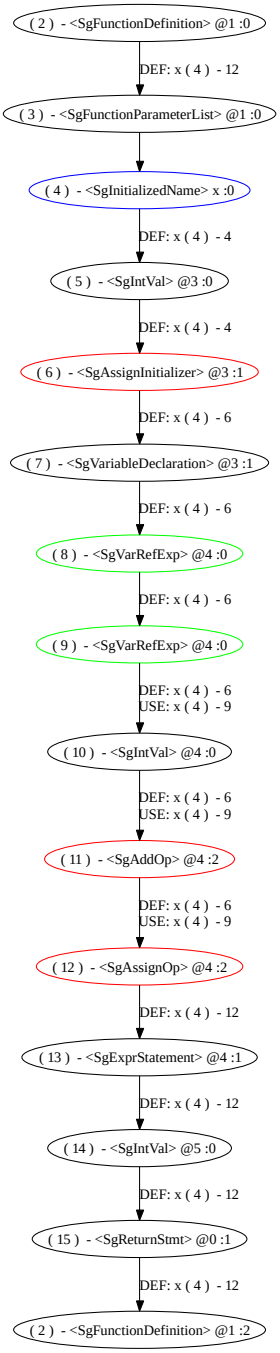

Figure 19.2: Def-Use graph for example program.

Chapter 20

Binary Analysis

This chapter discusses the capabilities of ROSE to read and analyze software binaries.

In the following we use a small example that demonstrates various features of Binary-Rose. The source code of our binary example is:

```
#include <stdlib.h>

int main(int argc, char** argv) {
  int* arr = (int*) malloc( sizeof(int)*10);

  int i=0;
  for (i=0; i<10;++i) {
    arr[i]=5;
  }
  int x = arr[12];
}
```

Figure 20.1: Example source code.

20.1 Loading binaries

Binary-ROSE is currently based on two front-ends: *objdump* and *IdaPro-mysql*.

20.1.1 objdump

The following code reads in a binary and creates a binary ROSE AST:

```
SgProject* project = frontend(argc,argv);
ROSE_ASSERT (project != NULL);
```

20.1.2 IdaPro-mysql

A binary processed by IdaPro needs to be processed into a MySQL database (DB). With that DB, the following lines create a binary ROSE AST:

```
// create RoseBin object with DB specific information
RoseBin* roseBin = new RoseBin(def_host_name,
                               def_user_name,
                               def_password,
                               def_db_name);

RoseBin_Arch::arch=RoseBin_Arch::bit32;
RoseBin_OS::os_sys=RoseBin_OS::linux_op;
RoseBin_OS_VER::os_ver=RoseBin_OS_VER::linux_26;
RoseBin_Def::RoseAssemblyLanguage=x86;
// connect to the DB
roseBin->connect_DB(socket);
// query the DB to retrieve all data
SgAsmNode* globalBlock = roseBin->retrieve_DB_IDAPRO();
// close the DB
roseBin->close_DB();

// traverse the AST and test it
roseBin->test();
```

20.2 The AST

The binary AST can now be unparsed to represent assembly instructions, cf. Figure 20.2:

```
SgAsmFile* file = project->get_file(0).get_binaryFile();
RoseBin_unparse* unparser = new RoseBin_unparse();
unparser->init(file->get_global_block(), fileName);
unparser->unparse();
```

To visualize the binary AST, we can use the following lines to write the AST out to a .dot format:

```
string filename=''_binary_tree.dot'';
AST_BIN_Traversal* trav = new AST_BIN_Traversal();
trav->run(file->get_global_block(), filename);
```

20.3 The ControlFlowGraph

Based on a control flow traversal of the binary AST, a separate control flow graph is created that can be used for further analyses:

```
// forward or backward analysis ?
bool forward = true;
// when creating a visual representation, visualized edges?
bool edges = true;
// visualize multiple edges or merge edges between same nodes to one edge?
bool mergedEdges = false;
// create DotGraph Object or/and GmlGraph Object
RoseBin_DotGraph* dotGraph = new RoseBin_DotGraph();
RoseBin_GMLGraph* gmlGraph = new RoseBin_GMLGraph();
char* cfgFileNameDot = ``cfg.dot'';
char* cfgFileNameGml = ``gml.dot'';
RoseBin_ControlFlowAnalysis* cfganalysis = new RoseBin_ControlFlowAnalysis(
                                           file->get_global_block(), forward, NULL, edges);
cfganalysis->run(dotGraph, cfgFileNameDot, mergedEdges);
cfganalysis->run(gmlGraph, cfgFileNameGml, mergedEdges);
```

The control flow graph of our example is represented in Figure 20.3. The graph shows different functions represented purple, containing various instructions. Instructions within a

```
/* Interpretation ELF */
/* Block 0x0 */
/* Function _init at 0x8048278 */
/* Block 0x8048278 */
0x08048278 55                |U      | :: 0x8048278:push   ebp
0x08048279 89 e5             |..     | :: 0x8048279:mov    ebp, esp
0x0804827b 83 ec 08          |...    | :: 0x804827b:sub    esp, 0x8
0x0804827e e8 61 00 00 00    |.a...  | :: 0x804827e:call   0x80482e4
/* Block 0x8048283 */
0x08048283 e8 b8 00 00 00    |.....  | :: 0x8048283:call   0x8048340
/* Block 0x8048288 */
0x08048288 e8 d3 01 00 00    |.....  | :: 0x8048288:call   0x8048460
/* Block 0x804828d */
0x0804828d c9                |.      | :: 0x804828d:leave
0x0804828e c3                |.      | :: 0x804828e:ret
/* Block 0x804828f */
0x0804828f 00 ff             |..     | :: 0x804828f:add    bh, bh
0x08048291 35 a0 95 04 08    |5....  | :: 0x8048291:xor    eax, 0x80495a0
/* Block 0x8048290 */
0x08048290 ff 35 a0 95 04 08 |.5....| :: 0x8048290:push   DWORD PTR ds:[0x80495a0]
0x08048296 ff 25 a4 95 04 08 |.%....| :: 0x8048296:jmp    DWORD PTR ds:[0x80495a4]
/* Block 0x804829c */
0x0804829c 00 00             |..     | :: 0x804829c:add    BYTE PTR ds:[eax], al
0x0804829e 00 00             |..     | :: 0x804829e:add    BYTE PTR ds:[eax], al
/* Block 0x80482a0 */
0x080482a0 ff 25 a8 95 04 08 |.%....| :: 0x80482a0:jmp    DWORD PTR ds:[0x80495a8]
/* Block 0x80482a6 */
0x080482a6 68 00 00 00 00    |h....  | :: 0x80482a6:push   0x0
0x080482ab e9 e0 ff ff ff    |.....  | :: 0x80482ab:jmp    0x8048290
/* Block 0x80482b0 */
0x080482b0 ff 25 ac 95 04 08 |.%....| :: 0x80482b0:jmp    DWORD PTR ds:[0x80495ac]
/* Block 0x80482b6 */
0x080482b6 68 08 00 00 00    |h....  | :: 0x80482b6:push   0x8
0x080482bb e9 d0 ff ff ff    |.....  | :: 0x80482bb:jmp    0x8048290
/* Function _start at 0x80482c0 */
/* Block 0x80482c0 */
0x080482c0 31 ed             |1.     | :: 0x80482c0:xor    ebp, ebp
0x080482c2 5e                |^      | :: 0x80482c2:pop    esi
0x080482c3 89 e1             |..     | :: 0x80482c3:mov    ecx, esp
0x080482c5 83 e4 f0          |...    | :: 0x80482c5:and    esp, 0xfffffff0
0x080482c8 50                |P      | :: 0x80482c8:push   eax
0x080482c9 54                |T      | :: 0x80482c9:push   esp
```

Figure 20.2: Assembly code.

function are represented within the same box. Common instructions are represented yellow. Green instructions are jumps and pink instructions calls and returns. Respectively, blue edges are call relationships and red edges return relationships. Black edges represent plain controlflow from one instruction to the next.

20.4 DataFlow Analysis

Based on the control flow many forms of dataflow analysis may be performed. The code to perform the dataflow analysis looks as follows:

```
string dfgFileName = "dfg.dot";
forward = true;
bool printEdges = true;
// choose between a interprocedural and intraprocedural dataflow analysis
```

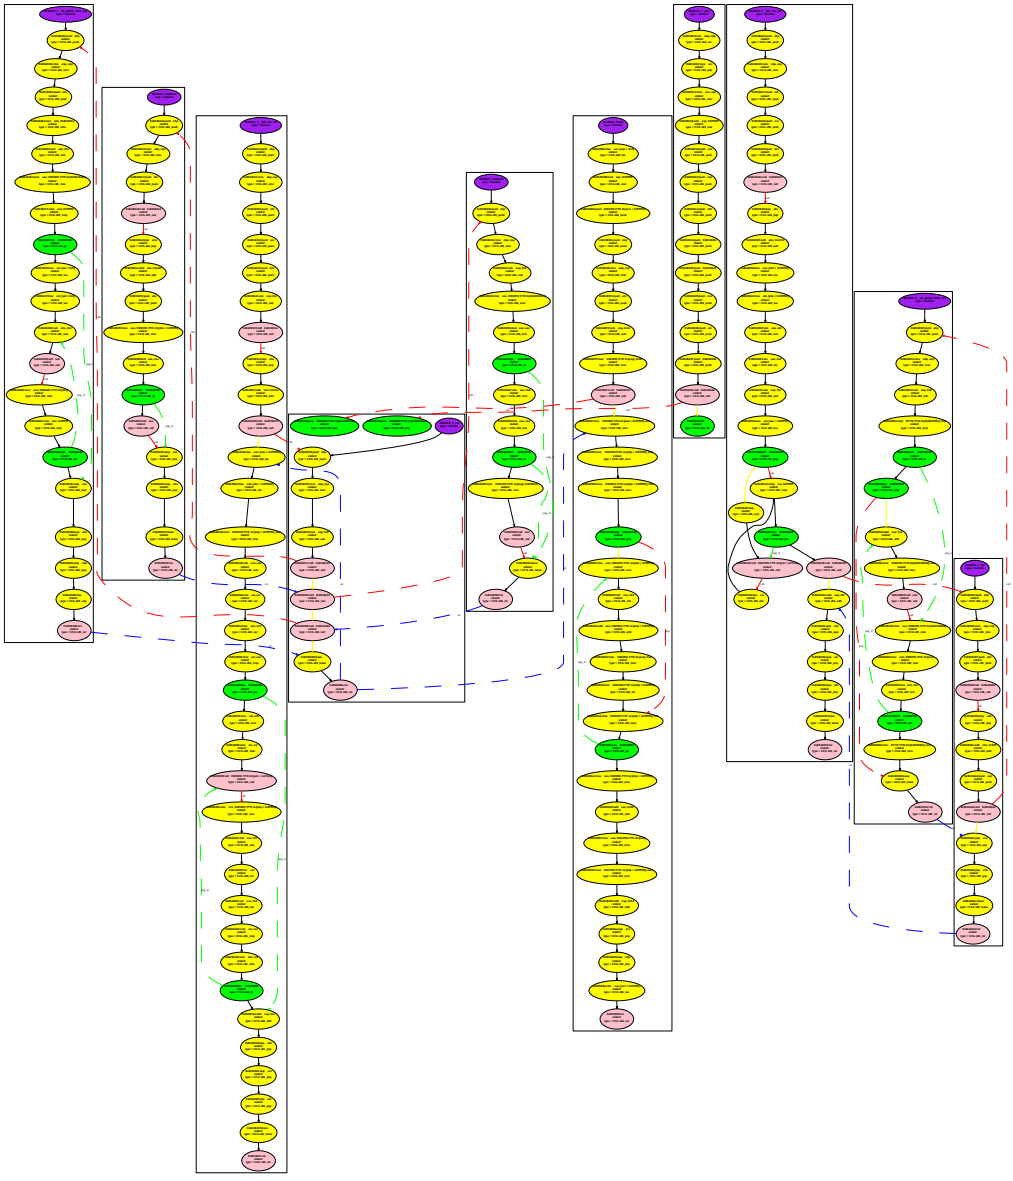

Figure 20.3: Controlflow graph for example program.

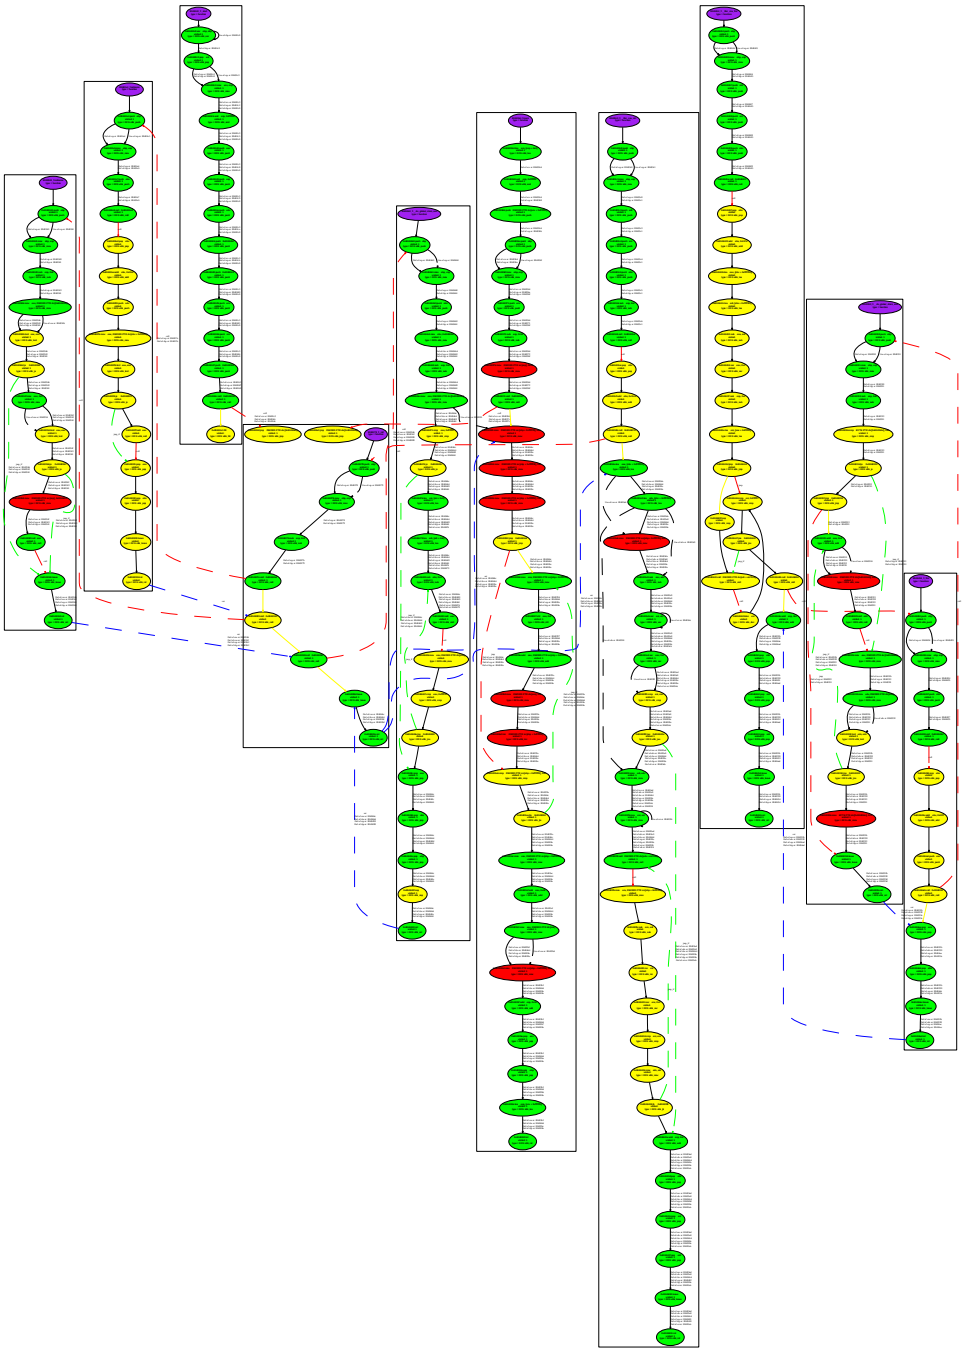

Figure 20.4: Dataflowflow graph for example program.

```
bool interprocedural = true;
RoseBin_DataFlowAnalysis* dfanalysis = new RoseBin_DataFlowAnalysis(
                                    file->get_global_block(), forward, NULL);
dfanalysis->init(interprocedural, printEdges);
dfanalysis->run(dotGraph, dfgFileName, mergedEdges);
```

Dataflow analyses available are:

20.4.1 Def-Use Analysis

Definition-Usage is one way to compute dataflow information about a binary program. Figure 20.4 shows a typical dataflow graph with additional definition and usage information visualized at the edges.

20.4.2 Variable Analysis

This analysis helps to detect different types within a binary. Currently, we use this analysis to detect interrupt calls and their parameters together with the def-use analysis.

This allows us to track back the value of parameters to the calls, such as eax and therefore determine whether a interrupt call is for instance a write or read.

Another feature is the buffer overflow analysis. By traversing the CFG, we can detect buffer overflows. The black node in 20.4 shows such a buffer overlow, cf. source code above.

Chapter 21

Generating the Call Graph (CG)

The formal definition of a call graph is:

'A diagram that identifies the modules in a system or computer program and shows which modules call one another.' IEEE

A call graph shows all function call paths of an arbitrary code. These paths are found by following all function calls in a function, where a function in the graph is represented by a node and each possible function call by an edge (arrow). To make a call graph this process is redone for every called function until all edges are followed and there are no ungraphed functions. ROSE has an in-build mechanism for generating call graphs.

Figure 21 shows the code required to generate the call graph for each function of an application. Using the input code shown in figure 21 the first function's call graph is shown in figure 21.3.

Figure 21.3 shows the call graph for the function in the input code in figure 21.

```
#include "rose.h"
#include <CallGraph.h>
#include <GraphUpdate.h>
using namespace std;

struct filterNodes : public unary_function<CallGraphNode*,bool>{
    public:
      bool operator()(CallGraphNode* test){
              bool returnValue = false;
              SgFunctionDeclaration* CallGraphNode2 = test->functionDeclaration;
              string filename = CallGraphNode2->get_file_info()->get_filename();
              if( filename.find("g++_HEADERS")!=string::npos ||
                  filename.find("/usr/include")!=string::npos){
                returnValue= true;
              }
              if(test->toString().find(string("_"))!=string::npos)
                returnValue = true;
              if(filename.find("rose_edg_macros_and_functions_required_for_gnu.h")!=string::npos)
                 returnValue = true;
         if(CallGraphNode2->get_file_info()->isCompilerGenerated()==true)
            returnValue=true;
         return returnValue;
      }
};

int
main( int argc, char * argv[] ) {
   RoseTestTranslator test;
   SgProject* project = new SgProject(argc, argv);
   CallGraphBuilder CGBuilder( project );
   CGBuilder.buildCallGraph();

   cout << "Classifying...\n"; // Classify subgraphs within call graph
   CGBuilder.classifyCallGraph();
   cout << "Done classifying\n";
  //  GenerateDotGraph(CGBuilder.getGraph(),"callgraph.dot");

   ClassHierarchyWrapper hier( project );
#ifdef HAVE_MYSQL
   hier.setDBName( "ClassHierarchy" );
   hier.createHierarchySchema();
   hier.writeHierarchyToDB();
#endif
   // Use the information in the graph to output a dot file for the call graph
   CallGraphDotOutput output( *(CGBuilder.getGraph()) );

#ifdef HAVE_MYSQL
   output.writeToDB( 1, "DATABASE" );
   output.filterNodesByDirectory( "DATABASE", "/export" );
   output.filterNodesByDB( "DATABASE", "_filter.db" );

#ifdef SOLVE_FUNCTION_CALLS_IN_DB
   output.solveVirtualFunctions( "DATABASE", "ClassHierarchy" );
   output.solveFunctionPointers( "DATABASE" );
#endif // SOLVE...

   cout << "Loading from DB...\n";
   CallGraphCreate *newGraph = output.loadGraphFromDB( "DATABASE" );
   cout << "Loaded\n";
#else
   // Not SQL Database case
   printf ("Not using the SQLite Database ... \n");
   CallGraphCreate *newGraph = CGBuilder.getGraph();
#endif // USE_ROSE...

   ostringstream st;
   st << "DATABASE.dot";
   cout << "Generating DOT...\n";
   filterGraph(*newGraph, filterNodes());
   generateDOT( *project );
   cout << "Done with DOT\n";
   GenerateDotGraph(newGraph, st.str());
   printf ("\nLeaving main program ... \n");
   return 0; // backend(project);
}
```

Figure 21.1: Example source code showing visualization of call graph.

```
// #include <iostream>
// #include <sstream>
// #include <stdio.h>
// #include <stdlib.h>

class A
{
public:
        int f1() {}
        int f2() { pf = &A::f1; return (this->*pf)(); }
        int (A::*pf) ();
};

// simple (trivial) example code used to demonstrate the call graph generatation

// #include <sstream>

void foo1();
void foo2();
void foo3();
void foo4();


void foo1()
   {
     foo1();
     foo2();
     foo3();
     foo4();

     void (*pointerToFunction)();

   }

void foo2()
   {
     foo1();
     foo2();
     foo3();
     foo4();
   }

void foo3()
   {
     foo1();
     foo2();
     foo3();
     foo4();
   }

void foo4()
   {
     foo1();
     foo2();
     foo3();
     foo4();
   }

int main()
   {
     foo1();
     foo2();
     foo3();
     foo4();

     return 0;
   }
```

Figure 21.2: Example source code used as input to build call graph.

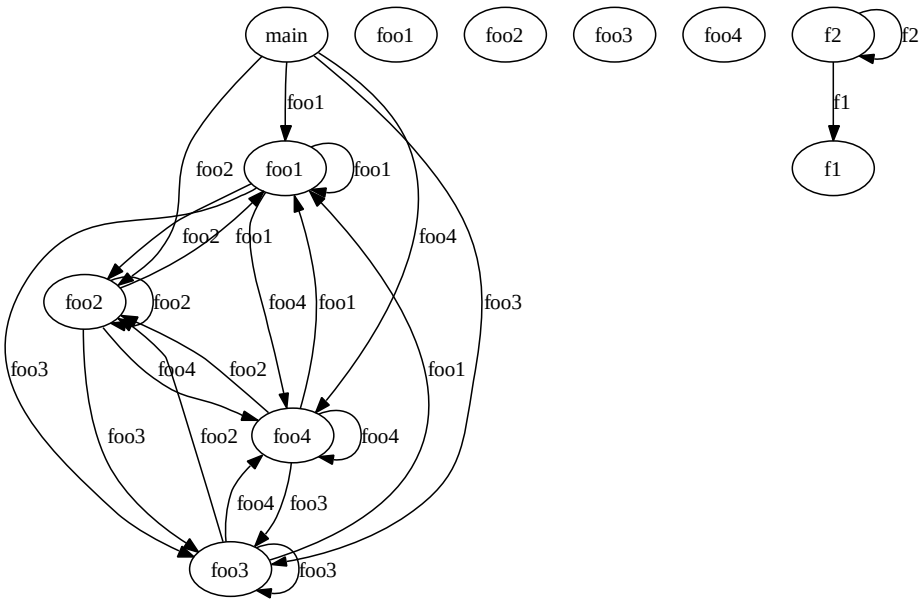

Figure 21.3: Call graph for function in input code file: inputCode_BuildCG.C.

Chapter 22

Generating the Class Hierarchy Graph

```
// Example ROSE Preprocessor
// used for testing ROSE infrastructure

// #include <assert.h>
// #include <string>
#include "rose.h"
// #include <iomanip>
// #include "AstConsistencyTests.h"

#include <CallGraph.h>

using namespace std;

  int
main( int argc, char * argv[] )
   {
     SgProject* project = new SgProject(argc, argv);

     ClassHierarchyWrapper hier( project );

     ClassHierarchy* classHier = hier.getClassHierarchyGraph();
     GraphDotOutput<ClassHierarchy> outputHier (*classHier);
     outputHier.writeToDOTFile ("classHier.dot", "Class Hierarchy");

     return 0;
   }
```

Figure 22.1: Example source code showing visualization of class hierarchy graph.

For C++, because of multiple inheritance, a class hierarchy graph is a directed graph with pointers from a class to a superclass. A superclass is a class which does not inherit from any other class. A class may inherit from a superclass by inheriting from another class which does rather than by a direct inheritance.

Figure 22 shows the code required to generate the class hierarchy graph for each class of an application. Using the input code shown in figure 22 the first function's call graph is shown in

figure 22.3.

```
class A{};
class B : public A{};
class C : public B{};
```

Figure 22.2: Example source code used as input to build class hierarchy graph.

Figure 22.3 shows the class hierarchy graph for the classes in the input code in figure 22.

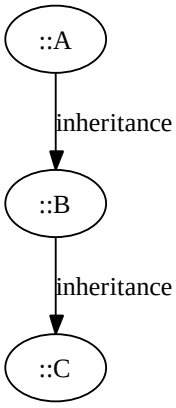

Figure 22.3: Class hierarchy graph in input code file: inputCode_ClassHierarchyGraph.C.

Chapter 23

Calling the Inliner

Figure 23.1 shows an example of how to use the inline mechanism. This chapter presents an example translator to to inlining of function calls where they are called. Such transformations are quite complex in a number of cases (one case is shown in the input code; a function call in a for loop conditional test). The details of functionality are hidden from the user and a high level interface is provided.

23.1 Source Code for Inliner

Figure 23.1 shows an example translator which calls the inliner mechanism. The code is designed to only inline up to ten functions. the list of function calls is recomputed after any function call is successfully inlined.

The input code is shown in figure 23.2, the output of this code is shown in figure 23.3.

23.2 Input to Demonstrate Function Inlining

Figure 23.2 shows the example input used for demonstration of an inlining transformation.

23.3 Final Code After Function Inlining

Figure 23.3 shows the results from the inlining of three function calls. The first two function calls are the same, and trivial. The second function call appears in the test of a for loop and is more complex.

```
// Example demonstrating function inlining (maximal inlining, up to preset number of inlinings).

#include "rose.h"

using namespace std;

// This is a function in Qing's AST interface
void FixSgProject(SgProject& proj);

int main (int argc, char* argv[])
   {
  // Build the project object (AST) which we will fill up with multiple files and use as a
  // handle for all processing of the AST(s) associated with one or more source files.
     SgProject* project = new SgProject(argc,argv);

  // DQ (7/20/2004): Added internal consistancy tests on AST
     AstTests::runAllTests(project);

     bool modifiedAST = true;
     int  count    = 0;

  // Inline one call at a time until all have been inlined.  Loops on recursive code.
     do {
          modifiedAST = false;

       // Build a list of functions within the AST
          Rose_STL_Container<SgNode*> functionCallList = NodeQuery::querySubTree (project,V_SgFunctionCallExp)

       // Loop over all function calls
       // for (list<SgNode*>::iterator i = functionCallList.begin(); i != functionCallList.end(); i++)
          Rose_STL_Container<SgNode*>::iterator i = functionCallList.begin();
          while (modifiedAST == false && i != functionCallList.end())
             {
               SgFunctionCallExp* functionCall = isSgFunctionCallExp(*i);
               ROSE_ASSERT(functionCall != NULL);

             // Not all function calls can be inlined in C++, so report if successful.
                bool sucessfullyInlined = doInline(functionCall);

                if (sucessfullyInlined == true)
                   {
                  // As soon as the AST is modified recompute the list of function
                  // calls (and restart the iterations over the modified list)
                     modifiedAST = true;
                   }
                  else
                   {
                     modifiedAST = false;
                   }

             // Increment the list iterator
                i++;
              }

       // Quite when we have ceased to do any inline transformations
       // and only do a predefined number of inline transformations
          count++;
        }
     while(modifiedAST == true && count < 10);

  // Call function to postprocess the AST and fixup symbol tables
     FixSgProject(*project);

  // Rename each variable declaration
     renameVariables(project);

  // Fold up blocks
     flattenBlocks(project);

  // Clean up inliner-generated code
     cleanupInlinedCode(project);

  // Change members to public
     changeAllMembersToPublic(project);

  // DQ (3/11/2006): This fails so the inlining, or the AST Interface
  // support, needs more work even though it generated good code.
  // AstTests::runAllTests(project);

     return backend(project);
   }
```

```
// This test code is a combination of pass1 and pass7, selected somewhat randomly
// from Jeremiah's test code of his inlining transformation from summer 2004.

int x = 0;

// Function it increment "x"
void incrementX()
   {
     x++;
   }

int foo()
   {
     int a = 0;
     while (a < 5)
        {
          ++a;
        }

     return a + 3;
   }

int main(int, char**)
   {
  // Two trival function calls to inline
     incrementX();
     incrementX();

  // Somthing more interesting to inline
     for (; foo() < 7;)
        {
          x++;
        }

     return x;
   }
```

Figure 23.2: Example source code used as input to program to the inlining transformation.

```
// This test code is a combination of pass1 and pass7, selected somewhat randomly
// from Jeremiah's test code of his inlining transformation from summer 2004.
int x = 0;
// Function it increment "x"

void incrementX()
{
  x++;
}


int foo()
{
  int a__0 = 0;
  while(a__0 < 5){
    ++a__0;
  }
  return a__0 + 3;
}


int main(int ,char **)
{
  x++;
  x++;
// Somthing more interesting to inline
  for (; true; ) {
    int a__1 = 0;
    while(a__1 < 5){
      ++a__1;
    }
    int rose_temp__7__0 = a__1 + 3;
    bool rose__temp__2 = (bool )(rose_temp__7__0 < 7);
    if (!rose__temp__2) {
      break;
    }
    else {
    }
    x++;
  }
  return x;
}
```

Figure 23.3: Output of input code after inlining transformations.

Chapter 24

Function Outlining using the AST Outliner

Function outlining is the process of replacing a block of consecutive statements with a function call to a new function containing those statements. Conceptually, outlining the inverse of inlining (Chapter 23). This chapter shows how to use the basic experimental outliner implementation included in the ROSE projects directory.

There are two basic ways to use the outliner. The first is a "user-level" method, in which you may use a special pragma to mark outline targets in the input program, and then call a high-level driver routine to process these pragmas. The second method is to call "low-level" outlining routines that operate directly on AST nodes. After a brief example of what the outliner can do and a discussion of its limits (Sections 24.1–24.2), we discuss each of these methods in Sections 24.3 and 24.4, respectively.

24.1 An Outlining Example

Figure 24.1 shows a small program with a pragma marking the outline target, a nested for loop, and Figure 24.2 shows the result. The outliner extracts the loop and inserts it into the body of a new function, and inserts a call to that function. The outlined code's input and output variables are wrapped up as parameters to this function. We make the following observations about this output.

Placement and forward declarations. The function itself is placed, by default, at the end of the input file to guarantee that it has access to all of the same declarations that were available at the outline target site. The outliner inserts any necessary forward declarations as well, including any necessary `friend` declarations if the outline target appeared in a class member function.

Calling convention. The outliner generates a C-callable function (`extern ''C''`, with pointer arguments). This design choice is motivated by our need to use the outliner to extract code into external, dynamically loadable library modules.

24.2 Limitations of the Outliner

The main limitation of the outliner implementation is that it can only outline single SgStatement nodes. However, since an SgStatement node may be a block (*i.e.*, an SgBasicBlock node), a "single statement" may actually comprise a sequence of complex statements.

The rationale for restricting to single SgStatement nodes is to avoid subtly changing the program's semantics when outlining code. Consider the following example, in which we wish to outline the middle 3 lines of executable code.

```
  int x = 5;
// START outlining here.
  foo (x);
  Object y (x);
  y.foo ();
// STOP outlining here.
  y.bar ();
```

This example raises a number of issues. How should an outliner handle the declaration of `y`, which constructs an object in local scope? It cannot just cut-and-paste the declaration of `y` to the body of the new outlined function because that will change its scope and lifetime, rendering the call to `y.bar()` impossible. Additionally, it may be unsafe to move the declaration of `y` so that it precedes the outlined region because the constructor call may have side-effects that could affect the execution of `foo(x)`. It is possible to heap-allocate `y` inside the body of the

```
namespace N
{
  class A
  {
    int foo (void) const { return 7; }
    int bar (void) const { return foo () / 2; }
  public:
    int biz (void) const
    {
      int result = 0;
#pragma rose_outline
      for (int i = 1; i <= foo (); i++)
        for (int j = 1; j <= bar (); j++)
          result += i * j;
      return result;
    }
  };
}

extern "C" int printf (const char* fmt, ...);

int main ()
{
  N::A x;
  printf ("%d\n", x.biz ()); // Prints '168'
  return 0;
}
```

Figure 24.1: inputCode_OutlineLoop.cc: Sample input program. The `#pragma` directive marks the nested for loop for outlining.

```
extern "C" void OUT__1__9204__ (int *resultp__, const void *this__ptr__p__);
extern "C" void OUT__1__9204__ (int *resultp__, const void *this__ptr__p__);
namespace N
{

  class A
  {
  public:friend void::OUT__1__9204__ (int *resultp__,
                                      const void *this__ptr__p__);

  private:  inline int foo () const
    {
      return 7;
    }


    inline int bar () const
    {
      return (this)->foo () / 2;
    }


  public:  inline int biz () const
    {
      int result = 0;
      {
        const class A *this__ptr__ = this;
          OUT__1__9204__ (&result, &this__ptr__);
      }
      return result;
    }

  }

   ;
}

extern "C"
{
  int printf (const char *fmt, ...);
}

int main ()
{
  class N::A x;
// Prints '168'
  printf ("%d\n", x.biz ());
  return 0;
}


extern "C" void OUT__1__9204__ (int *resultp__, const void *this__ptr__p__)
{
  int &result = *((int *) resultp__);
  const class N::A * &this__ptr__ = *((const class N::A **) this__ptr__p__);
  for (int i = 1; i <= this__ptr__->foo (); i++)
    for (int j = 1; j <= this__ptr__->bar (); j++)
      result += (i * j);
}
```

Figure 24.2: rose_outlined-inputCode_OutlineLoop.cc: The nested for loop of Figure 24.1 has been outlined.

outlined function so that it can be returned to the caller and later freed, but it is not clear if changing `y` from a stack-allocated variable to a heap-allocated one will always be acceptable, particularly if the developer of the original program has, for instance, implemented a customized memory allocator. Restricting outlining to well-defined SgStatement objects avoids these issues. It is possible to build a "higher-level" outliner that extends the outliner's basic infrastructure to handle these and other issues.

The outliner cannot outline all possible SgStatement nodes. However, the outliner interface provides a routine, `outliner::isOutlineable(s)`, for testing whether an SgStatement object `s` is known to satisfy the outliner's preconditions (see Section 24.4 for details).

24.3 User-Directed Outlining *via* Pragmas

Figure 24.3 shows the basic translator, `outline`, that produces Figure 24.2 from Figure 24.1. This translator extends the identity translator with an include directive on line 5 of Figure 24.3, and a call to the outliner on line 16. All outliner routines live in the `Outliner` namespace. Here, the call to `Outliner::outlineAll (proj)` on line 16 traverses the AST, looks for `#pragma rose_outline` directives, outlines the SgStatement objects to which each pragma is attached, and returns the number of outlined objects.

A slightly lower-level outlining primitive. The `Outliner::outlineAll()` routine is a wrapper around calls to a simpler routine, `Outliner::outline()`, that operates on pragmas:

```
Outliner::Result  Outliner::outline (SgPragmaDeclaration* s);
```

Given a pragma statement AST node `s`, this routine checks if `s` is a `rose_outline` directive, and if so, outlines the statement with which `s` is associated. It returns a `Outliner::Result` object, which is simply a structure that points to (a) the newly generated outlined function and (b) the statement corresponding to the new function call (*i.e.*, the outlined function call-site). See `Outliner.hh` or the ROSE Programmer's Reference for more details.

The `Outliner::outlineAll()` wrapper. The advantage of using the wrapper instead of the lower-level primitive is that the wrapper processes the pragmas in an order that ensures the outlining can be performed correctly in-place. This order is neither a preorder nor a postorder traversal, but in fact a "reverse" preorder traversal; refer to the wrapper's documentation for an explanation.

24.4 Calling Outliner Directly on AST Nodes

The preceding examples rely on the outliner's `#pragma` interface to identify outline targets. In this section, we show how to call the outliner directly on SgStatement nodes from within your translator.

Figure 24.4 shows an example translator that finds all `if` statements and outlines them. A sample input appears in Figure 24.5, with the corresponding output shown in Figure 24.6. Notice that valid preprocessor control structure is accounted for and preserved in the output.

The translator has two distinct phases. The first phase selects all *outlineable* if-statements, using the `CollectOutlineableIfs` helper class. This class produces a list that stores the targets in an order appropriate for outlining them in-place. The second phase iterates over the list of

```
// outline.cc: Demonstrates the pragma-interface of the Outliner.

#include <iostream>
#include <rose.h>
#include <Outliner.hh>

using namespace std;

int
main (int argc, char* argv[])
{
  SgProject* proj = frontend (argc, argv);
  ROSE_ASSERT (proj);

#if 1
  cerr << "[Outlining...]" << endl;
  size_t count = Outliner::outlineAll (proj);
  cerr << "  [Processed " << count << " outline directives.]" << endl;
#else
  printf ("Skipping outlining due to recent move from std::list to std::vector in ROSE\n");
#endif

  cerr << "[Unparsing...]" << endl;
  return backend (proj);
}
```

Figure 24.3: outline.cc: A basic outlining translator, which generates Figure 24.2 from Figure 24.1. This outliner relies on the high-level driver, `Outliner::outlineAll()`, which scans the AST for outlining pragma directives (`#pragma rose_outline`) that mark outline targets.

statements and outlines each one. The rest of this section explains these phases, as well as various aspects of the sample input and output.

24.4.1 Selecting the *outlineable* `if` statements

Line 45 of Figure 24.4 builds a list, `ifs` (declared on line 44), of outlineable if-statements. The helper class, `CollectOutlineableIfs` in lines 12–35, implements a traversal to build this list. Notice that a node is inserted into the target list only if it satisfies the outliner's preconditions; this check is the call to `Outliner::isOutlineable()` on line 28.

The function `Outliner::isOutlineable()` also accepts an optional second boolean parameter (not shown). When this parameter is true and the statement cannot be outlined, the check will print an explanatory message to standard error. Such messages are useful for discovering why the outliner will not outline a particular statement. The default value of this parameter is `false`.

24.4.2 Properly ordering statements for in-place outlining

Each call to `Outliner::outline(*i)` on line 50 of Figure 24.4 outlines a target if-statement `*i` in `if_targets`. However, in order for these statements to be outlined in-place, it is essential to outline the statements in the proper order.

The postorder traversal implemented by the helper class, `CollectOutlineableIfs`, produces the correct ordering. To see why, consider the following example code:

```
// outlineIfs.cc: Calls Outliner directly to outline if statements.

#include <iostream>
#include <set>
#include <list>
#include <rose.h>
#include <Outliner.hh>

using namespace std;

// Traversal to gather all outlineable SgIfStmt nodes.
class CollectOutlineableIfs : public AstSimpleProcessing
{
public:
  // Container of list statements in ``outlineable'' order.
  typedef list<SgIfStmt *> IfList_t;

  // Call this routine to gather the outline targets.
  static void collect (SgProject* p, IfList_t& final)
  {
    CollectOutlineableIfs collector (final);
    collector.traverseInputFiles (p, postorder);
  }

  virtual void visit (SgNode* n)
  {
    SgIfStmt* s = isSgIfStmt (n);
    if (Outliner::isOutlineable (s))
      final_targets_.push_back (s);
  }

private:
  CollectOutlineableIfs (IfList_t& final) : final_targets_ (final) {}
  IfList_t& final_targets_; // Final list of outline targets.
};

//====================================================================
int main (int argc, char* argv[])
{
  SgProject* proj = frontend (argc, argv);
  ROSE_ASSERT (proj);

#if 1
  // Build a set of outlineable if statements.
  CollectOutlineableIfs::IfList_t ifs;
  CollectOutlineableIfs::collect (proj, ifs);

  // Outline them all.
  for (CollectOutlineableIfs::IfList_t::iterator i = ifs.begin ();
       i != ifs.end (); ++i)
    Outliner::outline (*i);
#else
  printf ("Skipping outlining due to recent move from std::list to std::vector in ROSE\n");
#endif

  // Unparse
  return backend (proj);
}
```

Figure 24.4: outlineIfs.cc: A lower-level outlining translator, which calls `Outliner::outline()` directly on SgStatement nodes. This particular translator outlines all SgIfStmt nodes.

```
if (a) // [1]
```

```
#include <iostream>

using namespace std;

#define LEAP_YEAR 0

int main (int argc, char* argv[])
{
  for (int i = 1; i < argc; ++i)
    {
      string month (argv[i]);
      size_t days = 0;
      if (month == "January"
          || month == "March"
          || month == "May"
          || month == "July"
          || month == "August"
          || month == "October"
          || month == "December")
        days = 31;
#if LEAP_YEAR
      else if (month == "February")
        days = 29;
#else
      else if (month == "February")
        days = 28;
#endif
      else if (month == "April"
               || month == "June"
               || month == "September"
               || month == "November")
        days = 30;
      cout << argv[i] << " " << days << endl;
    }
  return 0;
}
```

Figure 24.5: inputCode_Ifs.cc: Sample input program, without explicit outline targets specified using `#pragma rose_outline`, as in Figures 24.1 and 24.9.

```
{
  if (b) foo (); // [2]
}
else if (c) // [3]
{
  if (d) bar (); // [4]
}
```

The corresponding AST is (roughly)

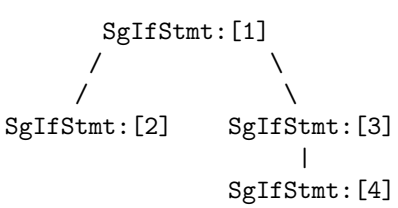

The postorder traversal—2, 4, 3, 1—ensures that child if-statements are outlined before their

parents.

```
#include <iostream>
using namespace std;
#define LEAP_YEAR 0
extern "C" void OUT__1__7996__(void *monthp__,size_t *daysp__);
extern "C" void OUT__2__7996__(void *monthp__,size_t *daysp__);
extern "C" void OUT__3__7996__(void *monthp__,size_t *daysp__);

int main(int argc,char *argv[])
{
  for (int i = 1; i < argc; ++i) {
    std::string month(((argv[i])));
    size_t days = (0);
    OUT__3__7996__(&month,&days);
    (*(&((*(&( *((&std::cout))<<((argv[i]))<<" ")))<<days)))<<std::endl;
  }
  return 0;
}


extern "C" void OUT__1__7996__(void *monthp__,size_t *daysp__)
{
  std::string &month =  *((std::string *)monthp__);
  size_t &days =  *((size_t *)daysp__);
  if (((month=="April" || month=="June") || month=="September") || month=="November")
    days = (30);
}


extern "C" void OUT__2__7996__(void *monthp__,size_t *daysp__)
{
  std::string &month =  *((std::string *)monthp__);
  size_t &days =  *((size_t *)daysp__);
#if 1 /* #if LEAP_YEAR ... #else */
  if (month=="February")
    days = (28);
  else
#endif
    OUT__1__7996__(&month,&days);
}

extern "C" void OUT__3__7996__(void *monthp__,size_t *daysp__)
{
  std::string &month =  *((std::string *)monthp__);
  size_t &days =  *((size_t *)daysp__);
  if ((((((month=="January" || month=="March") || month=="May") || month=="July") || month=="August") || month
    days = (31);
  else
#if LEAP_YEAR
#else
{
#endif /* #if LEAP_YEAR ... #else  */
    OUT__2__7996__(&month,&days);
  }
}
```

Figure 24.6: rose_inputCode_Ifs.cc: Figure 24.5, after outlining using the translator in Figure 24.4.

24.5 Outliner's Preprocessing Phase

Internally, the outliner implementation itself has two distinct phases. The first is a *preprocessing phase*, in which an arbitrary outlineable target is placed into a canonical form that is relatively simple to extract. The second phase then creates the outlined function, replacing the original target with a call to the outlined function. It is possible to run just the preprocessing phase, which is useful for understanding or even debugging the outliner implementation.

```
// outlinePreproc.cc: Shows the outliner's preprocessor-only phase.

#include <iostream>
#include <rose.h>
#include <Outliner.hh>

using namespace std;

int
main (int argc, char* argv[])
{
  SgProject* proj = frontend (argc, argv);
  ROSE_ASSERT (proj);

#if 1
  cerr << "[Running outliner's preprocessing phase only...]" << endl;
  size_t count = Outliner::preprocessAll (proj);
  cerr << "  [Processed " << count << " outline directives.]" << endl;
#else
  printf ("Skipping outlining due to recent move from std::list to std::vector in ROSE\n");
#endif

  cerr << "[Unparsing...]" << endl;
  return backend (proj);
}
```

Figure 24.7: outlinePreproc.cc: The basic translator of Figure 24.3, modified to execute the Outliner's preprocessing phase only. In particular, the original call to `Outliner::outlineAll()` has been replaced by a call to `Outliner::preprocessAll()`.

To call just the preprocessor, simply replace a call to `Outliner::outlineAll(s)` or `Outliner::outline(s)` with a call to `Outliner::preprocessAll(s)` or `Outliner::preprocess(s)`, respectively. The translator in Figure 24.7 modifies the translator in Figure 24.3 in this way to create a preprocessing-only translator.

The preprocessing phase consists of a sequence of initial analyses and transformations that the outliner performs in order to put the outline target into a particular canonical form. Roughly speaking, this form is an enclosing SgBasicBlock node, possibly preceded or followed by start-up and tear-down code. Running just the preprocessing phase on Figure 24.1 produces the output in Figure 24.8. In this example, the original loop is now enclosed in two additional SgBasicBlocks (Figure 24.8, lines 24–35), the outermost of which contains a declaration that shadows the object's `this` pointer, replacing all local references to `this` with the new shadow pointer. In this case, this initial transformation is used by the main underlying outliner implementation to explicitly identify all references to the possibly implicit references to `this`.

The preprocessing phase is more interesting in the presence of non-local control flow outside the outline target. Consider Figure 24.9, in which the outline target contains two `break`

```
namespace N
{

  class A
  {

  private:inline int foo () const
    {
      return 7;
    }


    inline int bar () const
    {
      return (this)->foo () / 2;
    }


  public:  inline int biz () const
    {
      int result = 0;

#pragma rose_outline
      {
        const class A *this__ptr__ = this;
        {
          for (int i = 1; i <= this__ptr__->foo (); i++)
            for (int j = 1; j <= this__ptr__->bar (); j++)
              result += (i * j);
        }
      }
      return result;
    }

  }

   ;
}

extern "C"
{
  int printf (const char *fmt, ...);
}

int main ()
{
  class N::A x;
// Prints '168'
  printf ("%d\n", x.biz ());
  return 0;
}
```

Figure 24.8: rose_outlined_pp-inputCode_OutlineLoop.cc: Figure 24.1 after outline preprocessing only, *i.e.*, specifying `-rose:outline:preproc-only` as an option to the translator of Figure 24.3.

statements, which require jumping to a regions of code outside the target. We show the preprocessed code in Figure 24.10. The original non-local jumps are first transformed into assignments to a flag, `EXIT_TAKEN__` (lines 18–20 and 26–29), and then relocated to a subsequent block of code (lines 38–53) with their execution controlled by the value of the flag. The final outlined result appears in Figure 24.11; the initial preprocessing simplifies this final step of extracting the outline target.

```
#include <iostream>

size_t factorial (size_t n)
{
  size_t i = 1;
  size_t r = 1;
  while (1)
    {
#pragma rose_outline
      if (i <= 1)
        break; // Non-local jump #1
      else if (i >= n)
        break; // Non-local jump #2
      else
        r *= ++i;
    }
  return r;
}

int main (int argc, char* argv[])
{
  std::cout << "7! == " << factorial (7) << std::endl; // Prints 5040
  return 0;
}
```

Figure 24.9: inputCode_OutlineNonLocalJumps.cc: Sample input program, with an outlining target that contains two non-local jumps (here, `break` statements).

```
#include <iostream>

size_t factorial (size_t n)
{
  size_t i = (1);
  size_t r = (1);
  while ((1))
    {

#pragma rose_outline
      {
        int EXIT_TAKEN__ = 0;
        {
          if (i <= (1))
            {
              EXIT_TAKEN__ = 1;
              goto NON_LOCAL_EXIT__;
            }
          else if (i >= n)
            {
              EXIT_TAKEN__ = 2;
              goto NON_LOCAL_EXIT__;
            }
          else
            r *= ++i;
        NON_LOCAL_EXIT__:;
        }
        if (EXIT_TAKEN__ == 1)
          {
// Non-local jump #1
            break;
          }
        else
          {
            if (EXIT_TAKEN__ == 2)
              {
// Non-local jump #2
                break;
              }
            else
              {
              }
          }
      }
    }
  return r;
}

int main (int argc, char *argv[])
{
// Prints 5040
  (*(&((*(&(*((&std::cout)) << "7!==")))) << factorial ((7))))) << std::endl;
  return 0;
}
```

Figure 24.10: rose_outlined_pp-inputCode_OutlineNonLocalJumps.cc: The non-local jump example of Figure 24.9 after outliner preprocessing, but before the actual outlining. The non-local jump is handled by an additional flag, `EXIT_TAKEN__`, which indicates what non-local jump is to be taken.

```
#include <iostream>
extern "C" void OUT__1__10111__ (size_t * np__, size_t * ip__, size_t * rp__,
                                 int *EXIT_TAKEN__p__);

size_t factorial (size_t n)
{
  size_t i = (1);
  size_t r = (1);
  while ((1))
    {
      {
        int EXIT_TAKEN__ = 0;
        OUT__1__10111__ (&n, &i, &r, &EXIT_TAKEN__);
        if (EXIT_TAKEN__ == 1)
          {
// Non-local jump #1
            break;
          }
        else
          {
            if (EXIT_TAKEN__ == 2)
              {
// Non-local jump #2
                break;
              }
            else
              {
              }
          }
      }
    }
  return r;
}

int main (int argc, char *argv[])
{
// Prints 5040
  (*(&((*(&(*((&std::cout)) << "7!==")))) << factorial ((7))))) << std::endl;
  return 0;
}


extern "C" void OUT__1__10111__ (size_t * np__, size_t * ip__, size_t * rp__,
                                 int *EXIT_TAKEN__p__)
{
  size_t & n = *((size_t *) np__);
  size_t & i = *((size_t *) ip__);
  size_t & r = *((size_t *) rp__);
  int &EXIT_TAKEN__ = *((int *) EXIT_TAKEN__p__);
  if (i <= (1))
    {
      EXIT_TAKEN__ = 1;
      goto NON_LOCAL_EXIT__;
    }
  else if (i >= n)
    {
      EXIT_TAKEN__ = 2;
      goto NON_LOCAL_EXIT__;
    }
  else
    r *= ++i;
NON_LOCAL_EXIT__:;
}
```

Figure 24.11: rose_outlined-inputCode_OutlineNonLocalJumps.cc: Figure 24.9 after outlining.

Chapter 25

Partial Redundancy Elimination (PRE)

Figure 25.1 shows an example of how to call the Partial Redundancy Elimination (PRE) implemented by Jeremiah Willcock. This transformation is useful for cleaning up code generated from other transformations (used in Qing's loop optimizations).

25.1 Source Code for example using PRE

Figure 25.1 shows an example translator which calls the PRE mechanism.

The input code is shown in figure 25.2, the output of this code is shown in figure 25.3.

```
// Example translator demontrating Partial Redundancy Elimination (PRE).

#include "rose.h"
#include "CommandOptions.h"

int main (int argc, char* argv[])
   {
  // Build the project object (AST) which we will fill up with multiple files and use as a
  // handle for all processing of the AST(s) associated with one or more source files.
     std::vector<std::string> l = CommandlineProcessing::generateArgListFromArgcArgv(argc, argv);

     CmdOptions::GetInstance()->SetOptions(argc, argv);
     SgProject* project = frontend(l);

     PRE::partialRedundancyElimination(project);

     return backend(project);
   }
```

Figure 25.1: Example source code showing how use Partial Redundancy Elimination (PRE).

25.2 Input to Example Demonstrating PRE

Figure 25.2 shows the example input used for demonstration of Partial Redundancy Elimination (PRE) transformation.

```
// Program, based on example in Knoop et al ("Optimal code motion: theory and
// practice", ACM TOPLAS 16(4), 1994, pp. 1117-1155, as cited in Paleri et al
// (see pre.C)), converted to C++

int unknown(); // ROSE bug: including body "return 0;" here doesn't work

void foo() {
  int a, b, c, x, y, z, w;

  if (unknown()) {
    y = a + b;
    a = c;
 // Added by Jeremiah Willcock to test local PRE
    w = a + b;
    a = b;
    x = a + b;
    w = a + b;
    a = c;
 // End of added part
    x = a + b;
  }

  if (unknown()) {
    while (unknown()) {y = a + b;}
  } else if (unknown()) {
    while (unknown()) {}
    if (unknown()) {y = a + b;} else {goto L9;} // FIXME: the PRE code crashes if this isn't in a block
  } else {
    goto L10;
  }

  z = a + b;
  a = c;

  L9: x = a + b;

  L10: 0; // ROSE bug: using return; here doesn't work
}

int unknown() {
  0; // Works around ROSE bug
  return 0;
}

int main(int, char**) {
  foo();
  return 0;
}
```

Figure 25.2: Example source code used as input to program to the Partial Redundancy Elimination (PRE) transformation.

25.3 Final Code After PRE Transformation

Figure 25.3 shows the results from the use of PRE on an the example input code.

```
// Program, based on example in Knoop et al ("Optimal code motion: theory and
// practice", ACM TOPLAS 16(4), 1994, pp. 1117-1155, as cited in Paleri et al
// (see pre.C)), converted to C++
// ROSE bug: including body "return 0;" here doesn't work
extern int unknown();

void foo()
{
// Partial redundancy elimination: cachevar__1 is a cache of (a + b)
  int cachevar__1;
  int a;
  int b;
  int c;
  int x;
  int y;
  int z;
  int w;
  if ((unknown())) {
    y = (a + b);
    a = c;
// Added by Jeremiah Willcock to test local PRE
    w = (a + b);
    a = b;
    cachevar__1 = (a + b);
    x = cachevar__1;
    w = cachevar__1;
    a = c;
// End of added part
    x = (a + b);
  }
  else {
  }
  if ((unknown())) {
    cachevar__1 = (a + b);
    while((unknown())){
      y = cachevar__1;
    }
  }
  else if ((unknown())) {
    while((unknown())){
    }
// FIXME: the PRE code crashes if this isn't in a block
    if ((unknown())) {
      cachevar__1 = (a + b);
      y = cachevar__1;
    }
    else {
      goto L9;
    }
  }
  else {
    goto L10;
  }
  z = cachevar__1;
  a = c;
  L9:
  x = (a + b);
// ROSE bug: using return; here doesn't work
  L10:
  0;
}


int unknown()
{
// Works around ROSE bug
  0;
  return 0;
}


int main(int ,char **)
{
  foo();
  return 0;
}
```

Figure 25.3: Output of input code after Partial Redundancy Elimination (PRE) transformation.

Chapter 26

AST File I/O

Figure 26.1 shows an example of how to use the AST File I/O mechanism. This chapter presents an example translator to write out an AST to a file and then read it back in.

26.1 Source Code for File I/O

Figure 26.1 shows an example translator which reads an input application, forms the AST, writes out the AST to a file, then deletes the AST and reads the AST from the previously written file.

The input code is shown in figure 26.2, the output of this code is shown in figure 26.3.

26.2 Input to Demonstrate File I/O

Figure 26.2 shows the example input used for demonstration of the AST file I/O. In this case we are reusing the example used in the inlining example.

26.3 Output from File I/O

Figure 26.3 shows the output from the example file I/O tutorial example.

26.4 Final Code After Passing Through File I/O

Figure 26.4 shows the same file as the input demonstrating that the file I/O didn't change the resulting generated code. *Much more sophisticated tests are applied internally to verify the correctness of the AST after AST file I/O.*

```
// Example demonstrating function inlining (maximal inlining, up to preset number of inlinings).

#include "rose.h"

using namespace std;

// This is a function in Qing's AST interface
void FixSgProject(SgProject& proj);

int main (int argc, char* argv[])
   {
  // Build the project object (AST) which we will fill up with multiple files and use as a
  // handle for all processing of the AST(s) associated with one or more source files.
     SgProject* project = new SgProject(argc,argv);

  // DQ (7/20/2004): Added internal consistancy tests on AST
     AstTests::runAllTests(project);

     bool modifiedAST = true;
     int  count    = 0;

  // Inline one call at a time until all have been inlined.  Loops on recursive code.
     do {
          modifiedAST = false;

       // Build a list of functions within the AST
          Rose_STL_Container<SgNode*> functionCallList = NodeQuery::querySubTree (project,V_SgFunctionCallExp);

       // Loop over all function calls
       // for (list<SgNode*>::iterator i = functionCallList.begin(); i != functionCallList.end(); i++)
          Rose_STL_Container<SgNode*>::iterator i = functionCallList.begin();
          while (modifiedAST == false && i != functionCallList.end())
             {
               SgFunctionCallExp* functionCall = isSgFunctionCallExp(*i);
               ROSE_ASSERT(functionCall != NULL);

             // Not all function calls can be inlined in C++, so report if successful.
               bool sucessfullyInlined = doInline(functionCall);

               if (sucessfullyInlined == true)
                  {
                 // As soon as the AST is modified recompute the list of function
                 // calls (and restart the iterations over the modified list)
                    modifiedAST = true;
                  }
                 else
                  {
                    modifiedAST = false;
                  }

             // Increment the list iterator
                i++;
              }

       // Quite when we have ceased to do any inline transformations
       // and only do a predefined number of inline transformations
          count++;
        }
     while(modifiedAST == true && count < 10);

  // Call function to postprocess the AST and fixup symbol tables
     FixSgProject(*project);

  // Rename each variable declaration
     renameVariables(project);

  // Fold up blocks
     flattenBlocks(project);

  // Clean up inliner-generated code
     cleanupInlinedCode(project);

  // Change members to public
     changeAllMembersToPublic(project);

  // DQ (3/11/2006): This fails so the inlining, or the AST Interface
  // support, needs more work even though it generated good code.
  // AstTests::runAllTests(project);

     return backend(project);
   }
```

```
// This test code is a combination of pass1 and pass7, selected somewhat randomly
// from Jeremiah's test code of his inlining transformation from summer 2004.

int x = 0;

// Function it increment "x"
void incrementX()
   {
     x++;
   }

int foo()
   {
     int a = 0;
     while (a < 5)
        {
          ++a;
        }

     return a + 3;
   }

int main(int, char**)
   {
  // Two trival function calls to inline
     incrementX();
     incrementX();

  // Somthing more interesting to inline
     for (; foo() < 7;)
        {
          x++;
        }

     return x;
   }
```

Figure 26.2: Example source code used as input to demonstrate the AST file I/O support.

```
Extending memory pools ... done
 Setting data of AST #0
```

Figure 26.3: Output of input code after inlining transformations.

```
// This test code is a combination of pass1 and pass7, selected somewhat randomly
// from Jeremiah's test code of his inlining transformation from summer 2004.
int x = 0;
// Function it increment "x"

void incrementX()
{
  x++;
}


int foo()
{
  int a__0 = 0;
  while(a__0 < 5){
    ++a__0;
  }
  return a__0 + 3;
}


int main(int ,char **)
{
  x++;
  x++;
// Somthing more interesting to inline
  for (; true; ) {
    int a__1 = 0;
    while(a__1 < 5){
      ++a__1;
    }
    int rose_temp__7__0 = a__1 + 3;
    bool rose__temp__2 = (bool )(rose_temp__7__0 < 7);
    if (!rose__temp__2) {
      break;
    }
    else {
    }
    x++;
  }
  return x;
}
```

Figure 26.4: Output of input code after file I/O.

Chapter 27

Generating Unique Names for Declarations

There are many instances where a unique name must be generated for either a function or variable declaration. ROSE defines a mechanism to make the generation of unique names from all SgDeclarationStatment IR nodes and the SgInitializedName IR node. This simplifies ROSE based applications that require this sort of mechanism. Our experience has found that a significant number of tools require such a mechanism and that it correct implementation can have subtle complex points, thus we have provided one as part of ROSE.

The specific translator described in this chapter traverses an AST and outputs the unique names that can be generated for each declaration showing the use of the unique name generation mechanism. This tool is intended as an example of how to generate unique names using ROSE. Not all IR nodes can be used to generate a unique name. The generated names are unique under the following rules:

1. Any two generated names are the same if the declarations are the same.
   Declaration can be the same across files or within the same file. Declarations that are the same can have different location in the same file (be represented multiple times) or be in different files. Language constructs that are the same must follow the One-time Definition Rule (ODR) across files.

2. Declarations in different unnamed scopes (e.g. for loop bodies) will generate different names.

3. Names are the same when generated buy different ROSE tools.
   Pointer values could be used to generate unique names of all IR nodes, but this would work only within a single invocation of the ROSE based tool. Generated names are not based on internal pointer values and are thus insensitive to pointer values. Generated names of the same declaration are thus the same even if generated from different tools. This allows multiple ROSE tools to inter-operate.

This unique name generation mechanism is only applicable to specific IR nodes, specifically:

- SgInitializedName

- SgDeclarationStatement IR nodes:
  - Obvious IR nodes supported:
    - SgClassDeclaration
    - SgFunctionDeclaration
    - SgEnumDeclaration
    - SgNamespaceDeclarationStatement
    - SgTypedefDeclaration
  - Less obvious IR nodes not supported (support for these would not make sense):
    - SgAsmStmt
    - SgCtorInitializerList
    - SgFunctionParameterList
    - SgNamespaceAliasDeclarationStatement
    - SgPragmaDeclaration
    - SgTemplateDeclaration (can this have a mangled name?)
    - SgTemplateInstantiationDirectiveStatement
    - SgUsingDeclarationStatement
    - SgUsingDirectiveStatement
    - SgVariableDeclaration
      Note that the SgVariableDeclaration contains a list of SgInitializedName nodes and the mangled names are best queried from each SgInitializedName instead of the SgVariableDeclaration.
    - SgVariableDefinition
- Un-named scopes
  A number of scopes are un-names and so there is an opportunity to generate non-unique names from declarations in such scopes. To fix this we generate names for each un-named scope to guarantee uniqueness. Nodes handled are:
  - SgForStatement
  - SgBasicBlock
  - SgIfStmt
  - get the complete list ...

Other language constructs can generate unique names as well, but their name could be invalid after certain transformation that move it structurally within the generated source code.

27.1 Example Code Showing Generation of Unique Names

27.2 Input For Examples Showing Unique Name Generation for Variables

Figure 27.1, shows an example translator demonstrating the generation of unique names from declarations in the AST. For each SgInitializedName we generate the mangled name. Figure 27.2

shows the input code and figure 27.3 shows the generated output from the translator (the mangled names from the AST associated with the input application).

27.3 Example Output Showing Unique Variable Names

27.4 Input For Examples Showing Unique Name Generation for Functions

Figure 27.1, shows an example translator demonstrating the generation of unique names from declarations in the AST. For each SgInitializedName we generate the mangled name. Figure 27.4 shows the input code and figure 27.5 shows the generated output from the translator (the mangled names from the AST associated with the input application).

27.5 Example Output Showing Unique Function Names

```
// This example shows the generation of unique names from declarations.

// Mangled name demo
//
// This translator queries the AST for all SgInitializedNames and
// SgFunctionDeclarations, and for each one prints (a) the source
// location, (b) the source name of the object, and (c) the mangled
// name.

#include <rose.h>

using namespace std;

// Returns a Sg_File_Info object as a display-friendly string, "[source:line]".
static string toString (const Sg_File_Info* info)
   {
     ostringstream info_str;
     if (info)
     info_str << '['
              << info->get_raw_filename ()
              << ":" << info->get_raw_line ()
              << ']';
     return info_str.str ();
   }

// Displays location and mangled name of an SgInitializedName object.
static void printInitializedName (const SgNode* node)
   {
     const SgInitializedName* name = isSgInitializedName (node);
     ROSE_ASSERT (name != NULL);

     if (name->get_file_info()->isCompilerGenerated() == false)
          cout // << toString (name->get_file_info ())
               // << " "
               << name->get_name ().str ()
               << " ---> " << name->get_mangled_name ().str ()
               << endl;
   }

// Displays location and mangled name of an SgFunctionDeclaration object.
static void printFunctionDeclaration (const SgNode* node)
   {
     const SgFunctionDeclaration* decl = isSgFunctionDeclaration (node);
     ROSE_ASSERT (decl != NULL);

     if (decl->get_file_info()->isCompilerGenerated() == false)
          cout // << toString (decl->get_startOfConstruct ())
               // << " "
               << decl->get_qualified_name ().str ()
               << " ---> " << decl->get_mangled_name ().str ()
               << endl;
   }

int main (int argc, char** argv)
   {
     SgProject* proj = frontend (argc, argv);

     cout << endl << "***** BEGIN initialized names *****" << endl;
     Rose_STL_Container<SgNode *> init_names = NodeQuery::querySubTree (proj, V_SgInitializedName);
     for_each (init_names.begin (), init_names.end (), printInitializedName);
     cout << "***** END initialized names *****" << endl;

     cout << endl << "***** BEGIN function declarations *****" << endl;
     Rose_STL_Container<SgNode *> func_decls = NodeQuery::querySubTree (proj, V_SgFunctionDeclaration);
     for_each (func_decls.begin (), func_decls.end (), printFunctionDeclaration);
     cout << "***** END function declarations *****" << endl;

     return backend (proj);
   }
```

Figure 27.1: Example source code showing the output of mangled name. The string represents the code associated with the subtree of the target IR node.

```
// Input file to test mangling of SgInitializedName objects.

int x;

// Global class
class A
   {
     private:
          int x;
       // Nested class
          class B
             {
               private:
                    int x;
               public:
                    void foo (int x_arg) { int x; }
             };
   };

template <typename T>
void
foo (T x_arg)
   {
     T x;
     for (x = 0; x < 10; x++)
        {
          T x = 0;
          do {
            // Local class
               class A
                  {
                    private:
                      // Nested class
                         class B
                            {
                              T x;
                            };
                    public:
                         void foo (T x) {}
                  };
               T x = 0;
             }
          while (x > 0);

          do {
               T x = 0;
             }
          while (x > 0);

       // Nested scope
             {
               T x = 0;
             }
        }
   }

template void foo<int> (int x);
template void foo<double> (double x);

void bar (void)
   {
     for (int x = 0; x != 0; x++)
          for (int x = 0; x != 0; x++)
               for (long x = 0; x != 0; x++)
                    ;
     try {
       for (int x = 0; x != 0; x++) ;
     }
     catch (int) {}
     catch (char x) {}
   }
```

Figure 27.2: Example source code used as input to program in codes showing debugging techniques shown in this section.

```
2   ***** BEGIN initialized names *****
    x --> x
4   x --> A__scope__x
    x --> A__scope__B__scope__x
6   x_arg --> L2R__ARG1
    x --> L2R__scope____SgSS2____scope__x
8   x_arg --> foo__tas__i__tae_____Fb_v_Gb_i_Fe___ARG1
    x_arg --> foo__tas__i__tae_____Fb_v_Gb_i_Fe___ARG1
10  x_arg --> foo__tas__d__tae_____Fb_v_Gb_d_Fe___ARG1
    x_arg --> foo__tas__d__tae_____Fb_v_Gb_d_Fe___ARG1
12  x --> bar___Fb_v_Gb__Fe___scope____SgSS2____scope____SgSS3____scope__x
    x --> bar___Fb_v_Gb__Fe___scope____SgSS2____scope____SgSS3____scope____SgSS4____scope__x
14  x --> bar___Fb_v_Gb__Fe___scope____SgSS2____scope____SgSS3____scope____SgSS4____scope____SgSS5____scope__x
    x --> bar___Fb_v_Gb__Fe___scope____SgSS2____scope____SgSS7____scope____SgSS8____scope__x
16   --> bar___Fb_v_Gb__Fe___scope____SgSS2____scope____SgSS10____scope__CATCHARG
    x --> bar___Fb_v_Gb__Fe___scope____SgSS2____scope____SgSS12____scope__x
18  x_arg --> foo__tas__i__tae_____Fb_v_Gb_i_Fe___ARG1
    x --> foo__tas__i__tae_____Fb_v_Gb_i_Fe___scope____SgSS2____scope__x
20  x --> foo__tas__i__tae_____Fb_v_Gb_i_Fe___scope____SgSS2____scope____SgSS3____scope____SgSS4____scope__x
    x --> L0R__scope__B__scope__x
22  x --> L3R__ARG1
    x --> foo__tas__i__tae_____Fb_v_Gb_i_Fe___scope____SgSS2____scope____SgSS3____scope____SgSS4____scope____SgSS5____scope____SgSS6_
24  x --> foo__tas__i__tae_____Fb_v_Gb_i_Fe___scope____SgSS2____scope____SgSS3____scope____SgSS4____scope____SgSS11____scope____SgSS1
    x --> foo__tas__i__tae_____Fb_v_Gb_i_Fe___scope____SgSS2____scope____SgSS3____scope____SgSS4____scope____SgSS13____scope__x
26  x_arg --> foo__tas__d__tae_____Fb_v_Gb_d_Fe___ARG1
    x --> foo__tas__d__tae_____Fb_v_Gb_d_Fe___scope____SgSS2____scope__x
28  x --> foo__tas__d__tae_____Fb_v_Gb_d_Fe___scope____SgSS2____scope____SgSS3____scope____SgSS4____scope__x
    x --> L1R__scope__B__scope__x
30  x --> L4R__ARG1
    x --> foo__tas__d__tae_____Fb_v_Gb_d_Fe___scope____SgSS2____scope____SgSS3____scope____SgSS4____scope____SgSS5____scope____SgSS6_
32  x --> foo__tas__d__tae_____Fb_v_Gb_d_Fe___scope____SgSS2____scope____SgSS3____scope____SgSS4____scope____SgSS11____scope____SgSS1
    x --> foo__tas__d__tae_____Fb_v_Gb_d_Fe___scope____SgSS2____scope____SgSS3____scope____SgSS4____scope____SgSS13____scope__x
34  ***** END initialized names *****

36  ***** BEGIN function declarations *****
    ::A::B::foo --> L2R
38  ::foo < int >  --> foo__tas__i__tae_____Fb_v_Gb_i_Fe_
    ::foo < int >  --> foo__tas__i__tae_____Fb_v_Gb_i_Fe_
40  ::foo < double >  --> foo__tas__d__tae_____Fb_v_Gb_d_Fe_
    ::foo < double >  --> foo__tas__d__tae_____Fb_v_Gb_d_Fe_
42  ::bar --> bar___Fb_v_Gb__Fe_
    ::foo < int >  --> foo__tas__i__tae_____Fb_v_Gb_i_Fe_
44  A::foo --> L3R
    ::foo < double >  --> foo__tas__d__tae_____Fb_v_Gb_d_Fe_
46  A::foo --> L4R
    ***** END function declarations *****
```

Figure 27.3: Output of input code using generatingUniqueNamesFromDeclaration.C

```
// Input file to test mangling of SgFunctionDeclaration objects.

long foobar();
long foobar(int);
long foobar(int y);
long foobar(int x);
long foobar(int x = 0);
long foobar(int xyz)
   {
     return xyz;
   }

char foobarChar(char);
char foobarChar(char c);

// Input file to test mangling of SgFunctionDeclaration objects.

typedef int value0_t;
typedef value0_t value_t;
namespace N
   {
     typedef struct { int a; } s_t;
     class A { public: A () {} virtual void foo (int) {} };
     class B { public: B () {} void foo (value_t) const {} };
     class C : public A { public: C () {} void foo (int) {} void foo (const s_t&) {} };
     void foo (const s_t*) {}
   }

typedef N::s_t s2_t;
void foo (value_t);
void foo (s2_t) {}
void foo (float x[]) {}
void foo (value_t, s2_t);

template <typename T>
void foo (T) {}

namespace P
   {
     typedef long double type_t;
     namespace Q
        {
          template <typename T>
          void foo (T) {}

          class R
             {
               public:
                    R () {}
                    template <typename T>
                    void foo (T) {}
                    void foo (P::type_t) {}
                    template <typename T, int x>
                    int foo (T) { return x; }
             };
        }
   }

template <typename T, int x>
int foo (T) { return x; }

template void foo<char> (char);
template void foo<const value_t *> (const value_t *);
template void P::Q::foo<long> (long);
template void P::Q::R::foo<value_t> (value_t);
```

Figure 27.4: Example source code used as input to program in codes showing debugging techniques shown in this section.

```
2   ***** BEGIN initialized names *****
     --> foobar___Fb_l_Gb_i_Fe___ARG1
4   y --> foobar___Fb_l_Gb_i_Fe___ARG1
    x --> foobar___Fb_l_Gb_i_Fe___ARG1
6   x --> foobar___Fb_l_Gb_i_Fe___ARG1
    xyz --> foobar___Fb_l_Gb_i_Fe___ARG1
8    --> foobarChar___Fb_c_Gb_c_Fe___ARG1
    c --> foobarChar___Fb_c_Gb_c_Fe___ARG1
10  a --> L0R__scope__a
     --> L3R__ARG1
12   --> L4R__ARG1
     --> L5R__ARG1
14   --> L6R__ARG1
     --> L7R__ARG1
16   --> foo___Fb_v_Gb_L1R_Fe___ARG1
     --> foo___Fb_v_Gb_L2R_Fe___ARG1
18  x --> L8R__ARG1
     --> foo___Fb_v_Gb_L1R__sep__L2R_Fe___ARG1
20   --> foo___Fb_v_Gb_L1R__sep__L2R_Fe___ARG2
     --> L9R__ARG1
22   --> foo__tas__c__tae_____Fb_v_Gb_c_Fe___ARG1
     --> foo__tas__c__tae_____Fb_v_Gb_c_Fe___ARG1
24   --> L10R__ARG1
     --> L10R__ARG1
26   --> L11R__ARG1
     --> L11R__ARG1
28   --> L12R__ARG1
     --> L12R__ARG1
30   --> foo__tas__c__tae_____Fb_v_Gb_c_Fe___ARG1
     --> L10R__ARG1
32   --> L11R__ARG1
     --> L12R__ARG1
34  ***** END initialized names *****

36  ***** BEGIN function declarations *****
    ::foobar --> foobar___Fb_l_Gb__Fe_
38  ::foobar --> foobar___Fb_l_Gb_i_Fe_
    ::foobar --> foobar___Fb_l_Gb_i_Fe_
40  ::foobar --> foobar___Fb_l_Gb_i_Fe_
    ::foobar --> foobar___Fb_l_Gb_i_Fe_
42  ::foobar --> foobar___Fb_l_Gb_i_Fe_
    ::foobarChar --> foobarChar___Fb_c_Gb_c_Fe_
44  ::foobarChar --> foobarChar___Fb_c_Gb_c_Fe_
    ::N::A::A --> L13R
46  ::N::A::foo --> L3R
    ::N::B::B --> L14R
48  ::N::B::foo --> L4R
    ::N::C::C --> L15R
50  ::N::C::foo --> L5R
    ::N::C::foo --> L6R
52  ::N::foo --> L7R
    ::foo --> foo___Fb_v_Gb_L1R_Fe_
54  ::foo --> foo___Fb_v_Gb_L2R_Fe_
    ::foo --> L8R
56  ::foo --> foo___Fb_v_Gb_L1R__sep__L2R_Fe_
    ::P::Q::R::R --> L16R
58  ::P::Q::R::foo --> L9R
    ::foo < char >  --> foo__tas__c__tae_____Fb_v_Gb_c_Fe_
60  ::foo < char >  --> foo__tas__c__tae_____Fb_v_Gb_c_Fe_
    ::foo < const value_t * >  --> L10R
62  ::foo < const value_t * >  --> L10R
    ::P::Q::foo < long >  --> L11R
64  ::P::Q::foo < long >  --> L11R
    ::P::Q::R::foo < value_t >  --> L12R
66  ::P::Q::R::foo < value_t >  --> L12R
    ::foo < char >  --> foo__tas__c__tae_____Fb_v_Gb_c_Fe_
68  ::foo < const value_t * >  --> L10R
    ::P::Q::foo < long >  --> L11R
70  ::P::Q::R::foo < value_t >  --> L12R
    ***** END function declarations *****
```

Figure 27.5: Output of input code using generatingUniqueNamesFromDeclaration.C

Chapter 28

Scopes of Declarations

The scope of a IR nodes may be either stored explicitly in the IR node or obtained through computation through it's parent information in the AST. Figure X shows an example where the variable definition for a variable is the scope of namespace X. The declaration for variable a is in the namespace X. In a more common way the function foo is a member function of B which a declaration appearing in class B, but with a function definition in global scope.

```
namespace X{
   extern int a;
}
int X::a = 0;
class B
  {
    void foo();
  };
void B::foo() {}
```

In C++, using name qualification the scope of a declaration can be independent of it structural location in the AST. The `get_parent()` member function (available on most IR nodes) communicates the structural information of the original source code (also represented in the AST). The scope information must at times be stored explicitly when it can not be interpreted structurally.

The example in this chapter show how to find the scope of each C++ construct were it make sense. Note that SgExpression IR nodes can take there scope from that of the statement where they are found. SgStatement and SgInitializedName IR nodes are the interesting IR nodes from the point of scope.

The SgInitializedName and all SgStatement IR nodes have a member function `get_scope()` which returns the scope of the associated IR nodes. The example code in this chapter traverses the AST and reports the scope of any SgInitializedName and all SgStatement IR nodes. It is intended to provide so simple intuition about what the scope can be expected to be in an application. The example code is also useful as a simple means of exploring the scopes of any other input application.

28.1 Input For Examples Showing Scope Information

Figure 28.1 shows the input example code form this tutorial example.

```
2  int xyz;

4  void foo (int x)
       {
6        int y;
         for (int i=0; i < 10; i++)
8            {
               int z;
10             z = 42;
             }
12     }
```

Figure 28.1: Example source code used as input to program in codes used in this chapter.

28.2 Generating the code representing any IR node

The following code traverses each IR node and for a SgInitializedName of SgStatement output the scope information. The input code is shown in figure 28.1, the output of this code is shown in figure 28.3.

```
// This example shows the scope of each statement and name (variable names, base class names, etc.).

#include "rose.h"

class visitorTraversal : public AstSimpleProcessing
   {
     public:
          virtual void visit(SgNode* n);
   };

void visitorTraversal::visit(SgNode* n)
   {
  // There are three types ir IR nodes that can be queried for scope:
  //   - SgStatement, and
  //   - SgInitializedName
     SgStatement* statement = isSgStatement(n);
     if (statement != NULL)
        {
          SgScopeStatement* scope = statement->get_scope();
          ROSE_ASSERT(scope != NULL);
          printf ("SgStatement ======= %12p = %30s has scope = %12p = %s (total number = %d) \n",
               statement,statement->class_name().c_str(),
               scope,scope->class_name().c_str(),(int)scope->numberOfNodes());
        }

     SgInitializedName* initializedName = isSgInitializedName(n);
     if (initializedName != NULL)
        {
          SgScopeStatement* scope = initializedName->get_scope();
          ROSE_ASSERT(scope != NULL);
          printf ("SgInitializedName = %12p = %30s has scope = %12p = %s (total number = %d)\n",
               initializedName,initializedName->get_name().str(),
               scope,scope->class_name().c_str(),(int)scope->numberOfNodes());
        }
   }

int
main ( int argc, char* argv[] )
   {
     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT (project != NULL);

  // Build the traversal object
     visitorTraversal exampleTraversal;

  // Call the traversal starting at the project node of the AST
     exampleTraversal.traverseInputFiles(project,preorder);

     printf ("Number of scopes (SgScopeStatement) = %d \n",(int)SgScopeStatement::numberOfNodes());
     printf ("Number of scopes (SgBasicBlock)  ===== = %d \n",(int)SgBasicBlock::numberOfNodes());

#if 0
     printf ("\n\n");
     printf ("Now output all the symbols in each symbol table \n");
     SageInterface::outputLocalSymbolTables(project);
     printf ("\n\n");
#endif

     return 0;
   }
```

Figure 28.2: Example source code showing how to get scope information for each IR node.

```
    SgStatement        =    0x987da58 =                    SgGlobal has scope =    0x987da58 = SgGlobal (total n
 2  SgStatement        =   0xb7e5e008 =          SgVariableDeclaration has scope =    0x987da58 = SgGlobal (total n
    SgInitializedName  =    0x9934e08 =                       xyz has scope =    0x987da58 = SgGlobal (total n
 4  SgStatement        =   0xb7f1eab0 =        SgFunctionDeclaration has scope =    0x987da58 = SgGlobal (total n
    SgStatement        =   0xb7ed8c78 =      SgFunctionParameterList has scope =    0x987da58 = SgGlobal (total n
 6  SgInitializedName  =    0x9934e60 =                         x has scope =   0xb7df9008 = SgFunctionDefinit
    SgStatement        =   0xb7df9008 =         SgFunctionDefinition has scope =    0x987da58 = SgGlobal (total n
 8  SgStatement        =    0x99c8f68 =                SgBasicBlock has scope =   0xb7df9008 = SgFunctionDefinit
    SgStatement        =   0xb7e5e140 =          SgVariableDeclaration has scope =    0x99c8f68 = SgBasicBlock (tot
10  SgInitializedName  =    0x9934eb8 =                         y has scope =    0x99c8f68 = SgBasicBlock (tot
    SgStatement        =    0x99e7658 =              SgForStatement has scope =    0x99c8f68 = SgBasicBlock (tot
12  SgStatement        =    0x9a07ad8 =           SgForInitStatement has scope =    0x99e7658 = SgForStatement (t
    SgStatement        =   0xb7e5e278 =          SgVariableDeclaration has scope =    0x99e7658 = SgForStatement (t
14  SgInitializedName  =    0x9934f10 =                         i has scope =    0x99e7658 = SgForStatement (t
    SgStatement        =    0x9a446c0 =            SgExprStatement has scope =    0x99e7658 = SgForStatement (t
16  SgStatement        =    0x99c8fe4 =                SgBasicBlock has scope =    0x99e7658 = SgForStatement (t
    SgStatement        =   0xb7e5e3b0 =          SgVariableDeclaration has scope =    0x99c8fe4 = SgBasicBlock (tot
18  SgInitializedName  =    0x9934f68 =                         z has scope =    0x99c8fe4 = SgBasicBlock (tot
    SgStatement        =    0x9a446e8 =            SgExprStatement has scope =    0x99c8fe4 = SgBasicBlock (tot
20  Number of scopes (SgScopeStatement) = 0
    Number of scopes (SgBasicBlock)     = 2
```

Figure 28.3: Output of input code using scopeInformation.C

Chapter 29

Type and Declaration Modifiers

Most languages support the general concept of modifiers to types, declarations, etc. The keyword *volatile* for example is a modifier to the type where it is used in a declaration. Searching for the modifiers for types and declarations can however be confusing. They are often not where one would expect, and most often because of corner cases in the language that force them to be handled in specific ways.

This example tutorial code is a demonstration of a how to access the *volatile* modifier used in the declaration of types for variables. We demonstrate that the modifier is not present in the SgVariableDeclaration or the SgVariableDefinitoon, but is located in the SgModifierType used to wrap the type returned from the SgInitializedName (the variable in the variable declaration).

29.1 Input For Example Showing use of *Volatile* type modifier

Figure 29.1 shows the example input used for demonstration of test for the *volatile* type modifier.

```
// Input example of use of "volatile" type modifier

volatile int a,*b;

void foo()
   {
     for (volatile int y = 0; y < 10; y++)
        {
        }
   }
```

Figure 29.1: Example source code used as input to program in codes used in this chapter.

29.2 Generating the code representing the seeded bug

Figure 29.2 shows a code that traverses each IR node and for and SgInitializedName IR node checks it type. The input code is shown in figure 29.1, the output of this code is shown in figure 29.3.

```
#include "rose.h"

using namespace std;

class visitorTraversal : public AstSimpleProcessing
   {
     public:
          void visit(SgNode* n);
   };

void visitorTraversal::visit(SgNode* n)
   {
  // The "volatile" madifier is in the type of the SgInitializedName
     SgInitializedName* initializedName = isSgInitializedName(n);
     if (initializedName != NULL)
        {
          printf ("Found a SgInitializedName = %s \n",initializedName->get_name().str());
          SgType* type = initializedName->get_type();

          printf ("   initializedName: type = %p = %s \n",type,type->class_name().c_str());
          SgModifierType* modifierType = isSgModifierType(type);
          if (modifierType != NULL)
             {
               bool isVolatile = modifierType->get_typeModifier().get_constVolatileModifier().isVolatile();
               printf ("   initializedName: SgModifierType: isVolatile = %s \n",(isVolatile == true) ? "true" : "false");
             }

          SgModifierNodes* modifierNodes = type->get_modifiers();
          printf ("   initializedName: modifierNodes = %p \n",modifierNodes);
          if (modifierNodes != NULL)
             {
               SgModifierTypePtrVector modifierList = modifierNodes->get_nodes();
               for (SgModifierTypePtrVector::iterator i = modifierList.begin(); i != modifierList.end(); i++)
                  {
                    printf ("initializedName: modifiers: i = %s \n",(*i)->class_name().c_str());
                  }
             }
        }

  // Note that the "volatile" madifier is not in the SgVariableDeclaration nor the SgVariableDefinition
     SgVariableDeclaration* variableDeclaration = isSgVariableDeclaration(n);
     if (variableDeclaration != NULL)
        {
          bool isVolatile = variableDeclaration->get_declarationModifier().get_typeModifier().get_constVolatileModifier().isVola
          printf ("SgVariableDeclaration: isVolatile = %s \n",(isVolatile == true) ? "true" : "false");
          SgVariableDefinition* variableDefinition = variableDeclaration->get_definition();
        // printf ("variableDefinition = %p \n",variableDefinition);
          if (variableDefinition != NULL)
             {
               bool isVolatile = variableDefinition->get_declarationModifier().get_typeModifier().get_constVolatileModifier().isV
               printf ("SgVariableDefinition: isVolatile = %s \n",(isVolatile == true) ? "true" : "false");
             }
        }
   }

// must have argc and argv here!!
int main(int argc, char * argv[])
   {
     SgProject *project = frontend (argc, argv);

     visitorTraversal myvisitor;
     myvisitor.traverseInputFiles(project,preorder);

     return backend(project);
   }
```

Figure 29.2: Example source code showing how to detect *volatile* modifier.

```
// Input example of use of "volatile" type modifier
volatile int a;
volatile int *b;

void foo()
{
  for (volatile int y = 0; y < 10; y++) {
  }
}
```

Figure 29.3: Output of input code using volatileTypeModifier.C

Chapter 30

Debugging Techniques

There are numerous methods ROSE provides to help debug the development of specialized source-to-source translators. This section shows some of the techniques for getting information from IR nodes and displaying it. Show how to use the PDF generator for ASTs. This section may contain several subsections. More information about generation of specialized AST graphs to support debugging can be found in 35 and custom graph generation in 34.

30.1 Input For Examples Showing Debugging Techniques

Figure 30.1 shows the input code used for the example translators that report useful debugging information in this chapter.

```
// Example program showing matrix multiply
// (for use with loop optimization tutorial example)

#define N 50

int main()
   {
     int i,j, k;
     double a[N][N], b[N][N], c[N][N];

     for (i = 0; i <= N-1; i+=1)
        {
          for (j = 0; j <= N-1; j+=1)
             {
               for (k = 0; k <= N-1; k+=1)
                  {
                    c[i][j] = c[i][j] + a[i][k] * b[k][j];
                  }
             }
        }

     return 0;
   }
```

Figure 30.1: Example source code used as input to program in codes showing debugging techniques shown in this section.

30.2 Generating the code from any IR node

Any IR node may be converted to the string that represents its subtree within the AST. It is is a type then it the string will be the value of the type, if it is a statement the value will be the source code associated with that statement, including any sub-statements. To support the generation for strings from IR nodes we use the `unparseToString()` member function. This function strips comments and preprocessor control structure. The resulting string is useful for both debugging and when forming larger strings associated with the specification of transformations using the string-based rewrite mechanism. Using ROSE IR nodes may be converted to string and string converted to AST fragments of IR nodes.

Note that unparsing associated with generating source code for the backend vendor compiler in however as more than just calling the unparseToString member function since it introduces comments, preprocessor control structure and formating.

Figure 30.2 shows a translator which generates a string for a number of predefined IR nodes. Figure 30.1 shows the sample input code and figure 30.5 shows the output from the translator when using the example input application.

30.3 Displaying the source code position of any IR node

This example shows how to obtain information about the position of any IR node relative to where it appeared in the original source code. New IR nodes (or subtrees) that are added to the AST as part of a transformation will be marked as part of a transformation and have no position in the source code. Shared IR nodes (as generated by the AST merge mechanism are marked as shared explicitly (other IR nodes that are shared by definition don't have a SgFileInfo object and are thus not marked explicitly as shared.

The example translator to output the source code position is shown in figure 30.4. Using the input code in figure 30.1 the output code is shown in figure 30.5.

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.
// rose.C: Example (default) ROSE Preprocessor: used for testing ROSE infrastructure

#include "rose.h"

using namespace std;

int
main ( int argc, char* argv[] )
   {
     ios::sync_with_stdio();      // Syncs C++ and C I/O subsystems!

     if (SgProject::get_verbose() > 0)
          printf ("In preprocessor.C: main()\n");

     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT (project != NULL);

  // AST diagnostic tests
     AstTests::runAllTests(const_cast<SgProject*>(project));

  // test statistics
     if (project->get_verbose() > 1)
        {
          cout << AstNodeStatistics::traversalStatistics(project);
          cout << AstNodeStatistics::IRnodeUsageStatistics();
        }

     if (project->get_verbose() > 0)
          printf ("Generate the pdf output of the SAGE III AST\n");
     generatePDF ( *project );

     if (project->get_verbose() > 0)
          printf ("Generate the DOT output of the SAGE III AST\n");
     generateDOT ( *project );

     Rose_STL_Container<SgNode*> nodeList;
  // nodeList = NodeQuery::querySubTree (project,V_SgType,NodeQuery::ExtractTypes);
     nodeList = NodeQuery::querySubTree (project,V_SgForStatement);
     printf ("\nnodeList.size() = %zu \n",nodeList.size());

     Rose_STL_Container<SgNode*>::iterator i = nodeList.begin();
     while (i != nodeList.end())
        {
          printf ("Query node = %p = %s = %s \n",*i,(*i)->sage_class_name(),(*i)->unparseToString().c_str());
          i++;
        }

     return 0;
   }
```

Figure 30.2: Example source code showing the output of the string from an IR node. The string represents the code associated with the subtree of the target IR node.

```
nodeList.size() = 3
Query node = 0x9741110 = SgForStatement = for(i = 0;i <=(50 - 1);i += 1) {for(j = 0;j <=(50 - 1);j += 1) {for(k = 0;k <=
Query node = 0x9741194 = SgForStatement = for(j = 0;j <=(50 - 1);j += 1) {for(k = 0;k <=(50 - 1);k += 1) {(c[i])[j] =(((c
Query node = 0x9741218 = SgForStatement = for(k = 0;k <=(50 - 1);k += 1) {(c[i])[j] =(((c[i])[j]) +(((a[i])[k]) *((b[k])[
```

Figure 30.3: Output of input code using debuggingIRnodeToString.C

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.
// rose.C: Example (default) ROSE Preprocessor: used for testing ROSE infrastructure

#include "rose.h"

using namespace std;

int
main ( int argc, char* argv[] )
   {
     if (SgProject::get_verbose() > 0)
          printf ("In preprocessor.C: main() \n");

     SgProject* project = frontend(argc,argv);
     ROSE_ASSERT (project != NULL);

     Rose_STL_Container<SgNode*> nodeList;
     nodeList = NodeQuery::querySubTree (project,V_SgForStatement);
     printf ("\nnodeList.size() = %zu \n",nodeList.size());

     Rose_STL_Container<SgNode*>::iterator i = nodeList.begin();
     while (i != nodeList.end())
        {
          Sg_File_Info & fileInfo = *((*i)->get_file_info());
          printf ("Query node = %p = %s in %s \n ----- at line %d on column %d \n",
              *i,(*i)->sage_class_name(),fileInfo.get_filename(),
               fileInfo.get_line(), fileInfo.get_col());
          i++;
        }

     if (project->get_verbose() > 0)
          printf ("Calling the backend() \n");

     return 0;
   }
```

Figure 30.4: Example source code showing the output of the string from an IR node. The string represents the code associated with the subtree of the target IR node.

```
nodeList.size() = 3
Query node = 0x99bd130 = SgForStatement in /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inp
 ----- at line 11 on column 6
Query node = 0x99bd1b4 = SgForStatement in /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inp
 ----- at line 13 on column 11
Query node = 0x99bd238 = SgForStatement in /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inp
 ----- at line 15 on column 16
```

Figure 30.5: Output of input code using debuggingSourceCodePositionInformation.C

Chapter 31

Handling Comments, Preprocessor Directives, And Adding Arbitrary Text to Generated Code

What To Learn From These Examples Learn how to access existing comments and CPP directives and modify programs to include new ones. Where such comments can be automated they can be used to explain transformations or for more complex transformations using other tools designed to read the generated comments. Also included is how to add arbitrary text to the generated code (often useful for embedded system programming to support back-end compiler specific functionality).

This chapter deals with comments and preprocessor directives. These are often dropped from compiler infrastructures and ignored by make tools. ROSE takes great care to preserve all comments and preprocessor directives. Where they exist in the input code we save them (note that EDG drops them from their AST) and weave them back into the AST.

Note that *#pragma* is not a CPP directive and is part of the C and C++ grammar, thus represented explicitly in the AST (see SgPragmaDeclaration).

31.1 How to Access Comments and Preprocessor Directives

Comments and CPP directives are treated identically within ROSE and are saved as special preprocessor attributes to IR nodes within the AST. Not all IR nodes can have these specific type of attributes, only SgLocatedNodes can be associated with such preprocessor attributes. The more general *persistent attribute mechanism* within ROSE is separate from this preprocessor attribute mechanism and is available on a wider selection of IR nodes.

31.1.1 Source Code Showing How to Access Comments and Preprocessor Directives

Figure 31.1 shows an example translator which access the comments and preprocessor directives on each statement. Note that in this example the AST is traversed twice, first header files are ignored and then the full AST (including header files) are traversed (generated additional comments).

The input code is shown in figure 31.2, the output of this code is shown in figure 31.3 for the source file only. Figure 31.4 shows the same input code processed to output comments and preprocessor directives assembled from the source file *and* all header files.

31.1.2 Input to example showing how to access comments and CPP directives

Figure 31.2 shows the example input used for demonstration of how to collect comments and CPP directives.

31.1.3 Comments and CPP Directives collected from source file (skipping headers)

Figure 31.3 shows the results from the collection of comments and CPP directives within the input source file only (without -rose:collectAllCommentsAndDirectives).

31.1.4 Comments and CPP Directives collected from source file and all header files

Figure 31.4 shows the results from the collection of comments and CPP directives within the input source file and all headers (with -rose:collectAllCommentsAndDirectives).

31.2 Collecting #define C Preprocessor Directives

This example shows how to collect the #define directives as a list for later processing.

31.2.1 Source Code Showing How to Collect #define Directives

Figure 31.5 shows an example translator which access the comments and preprocessor directives on each statement. Note that in this example the AST is traversed twice, first header files are ignored and then the full AST (including header files) are traversed (generated additional comments).

The input code is shown in figure 31.6, the output of this code is shown in Figure 31.7 shows the same input code processed to output comments and preprocessor directives assembled from the source file *and* all header files.

```
// Example ROSE Translator: used within ROSE/tutorial

#include "rose.h"

using namespace std;

// Class declaration
class visitorTraversal : public AstSimpleProcessing
   {
     public:
          virtual void visit(SgNode* n);
   };

void visitorTraversal::visit(SgNode* n)
   {
  // On each node look for any comments of CPP directives
     SgLocatedNode* locatedNode = isSgLocatedNode(n);
     if (locatedNode != NULL)
        {
          AttachedPreprocessingInfoType* comments = locatedNode->getAttachedPreprocessingInfo();

          if (comments != NULL)
             {
               printf ("Found attached comments (to IR node at %p of type: %s): \n",locatedNode,locatedNode->class_name()
               int counter = 0;
               AttachedPreprocessingInfoType::iterator i;
               for (i = comments->begin(); i != comments->end(); i++)
                  {
                    printf ("           Attached Comment #%d in file %s (relativePosition=%s): classification %s: \n%s\n",
                         counter++,(*i)->get_file_info()->get_filenameString().c_str(),
                         ((*i)->getRelativePosition() == PreprocessingInfo::before) ? "before" : "after",
                         PreprocessingInfo::directiveTypeName((*i)->getTypeOfDirective()).c_str(),
                         (*i)->getString().c_str());
                  }
             }
            else
             {
               printf ("No attached comments (at %p of type: %s): \n",locatedNode,locatedNode->sage_class_name());
             }
        }
   }

int main( int argc, char * argv[] )
   {
  // Build the AST used by ROSE
     SgProject* project = frontend(argc,argv);

  // Build the traversal object
     visitorTraversal exampleTraversal;

  // Call the traversal starting at the project node of the AST
  // Traverse all header files and source file (the -rose:collectAllCommentsAndDirectives
  // commandline option controls if comments and CPP directives are separately extracted
  // from header files).
     exampleTraversal.traverse(project,preorder);

     return 0;
   }
```

Figure 31.1: Example source code showing how to access comments.

```
2  // #include<stdio.h>

4  #define SOURCE_CODE_BEFORE_INCLUDE_A
   #define SOURCE_CODE_BEFORE_INCLUDE_B
6  #include<inputCode_collectComments.h>
   #define SOURCE_CODE_AFTER_INCLUDE_A
8  #define SOURCE_CODE_AFTER_INCLUDE_B

10 // main program: collectComments input test code
   int main()
12    {
        return 0;
14    }
```

Figure 31.2: Example source code used as input to collection of comments and CPP directives.

31.2.2 Input to example showing how to access comments and CPP directives

Figure 31.6 shows the example input used for demonstration of how to collect comments and CPP directives.

31.2.3 Comments and CPP Directives collected from source file and all header files

Figure 31.7 shows the results from the collection of comments and CPP directives within the input source file and all headers (with -rose:collectAllCommentsAndDirectives).

31.3 Automated Generation of Comments

Figure 31.8 shows an example of how to introduce comments into the AST which will then show up in the generated source code. The purpose for this is generally to add comments to where transformations are introduced. If the code is read by the use the generated comments can be useful in identifying, marking, and/or explaining the transformation.

This chapter presents an example translator which just introduces a comment at the top of each function. The comment includes the name of the function and indicates that the comment is automatically generated.

Where appropriate such techniques could be used to automate the generation of documentation templates in source code that would be further filled in by the used. In this case the automatically generated templates would be put into the generated source code and a patch formed between the generated source and the original source. The patch could be easily inspected and applied to the original source code to place the documentation templates into the original source. The skeleton of the documentation in the source code could been be filled in by the use. The template would have all relevant information obtained by analysis (function parameters, system functions used, security information, side-effects, anything that could come from an analysis of the source code using ROSE).

31.3.1 Source Code Showing Automated Comment Generation

Figure 31.8 shows an example translator which calls the mechanism to add a comment to the IR node representing a function declaration (SgFunctionDeclaration).

The input code is shown in figure 31.9, the output of this code is shown in figure 31.10.

31.3.2 Input to Automated Addition of Comments

Figure 31.9 shows the example input used for demonstration of an automated commenting.

31.3.3 Final Code After Automatically Adding Comments

Figure 31.10 shows the results from the addition of comments to the generated source code.

31.4 Addition of Arbitrary Text to Unparsed Code Generation

This section is different from the comment generation (section 31.3) because it is more flexible and does not introduce any formatting. It also does not use the same internal mechanism, this mechanism supports the addition of new strings or the replacement of the IR node (where the string is attached) with the new string. It is fundamentally lower level and a more powerful mechanism to support generation of tailored output where more than comments, CPP directives, or AST transformation are required. It is also much more dangerous to use.

This mechanism is expected to be used rarely and sparingly since no analysis of the AST is likely to leverage this mechanism and search for code that introduced as a transformation here. Code introduced using this mechanism is for the most part unanalyzable since it would have to be reparsed in the context of the location in the AST were it is attached. (Technically this is possible and is the subject of the existing ROSE AST Rewrite mechanism, but that is a different subject).

Figure 31.11 shows an example of how to introduce arbitrary text into the AST for output by the unparser which will then show up in the generated source code. The purpose for this is generally to add backend compiler or tool specific code generation which don't map to any formal language constructs and so cannot be represented in the AST. However, since most tools that require specialized annotations read them as comments, the mechanism in the previous section 31.3 may be more appropriate. It is because this is not always that case that we have provide this more general mechanism (often useful for embedded system compilers).

31.4.1 Source Code Showing Automated Arbitrary Text Generation

Figure 31.11 shows an example translator which calls the mechanism to add a arbitrary text to the IR node representing a function declaration (SgFunctionDeclaration).

The input code is shown in figure 31.12, the output of this code is shown in figure 31.13.

31.4.2 Input to Automated Addition of Arbitrary Text

Figure 31.12 shows the example input used for demonstration of the automated introduction of text via the unparser.

31.4.3 Final Code After Automatically Adding Arbitrary Text

Figure 31.13 shows the results from the addition of arbitrary text to the generated source code.

```
    No attached comments (at 0x8cc8b20 of type: SgGlobal):
 2  No attached comments (at 0xb7ecb008 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e88008 of type: SgFunctionParameterList):
 4  No attached comments (at 0xb7ecb1a0 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e88118 of type: SgFunctionParameterList):
 6  No attached comments (at 0xb7ecb338 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e88228 of type: SgFunctionParameterList):
 8  No attached comments (at 0xb7ecb4d0 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e88338 of type: SgFunctionParameterList):
10  No attached comments (at 0xb7ecb668 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e88448 of type: SgFunctionParameterList):
12  No attached comments (at 0xb7ecb800 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e88558 of type: SgFunctionParameterList):
14  No attached comments (at 0xb7ecb998 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e88668 of type: SgFunctionParameterList):
16  No attached comments (at 0xb7ecbb30 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e88778 of type: SgFunctionParameterList):
18  No attached comments (at 0xb7ecbcc8 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e88888 of type: SgFunctionParameterList):
20  No attached comments (at 0xb7ecbe60 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e88998 of type: SgFunctionParameterList):
22  No attached comments (at 0xb7ecbff8 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e88aa8 of type: SgFunctionParameterList):
24  No attached comments (at 0xb7ecc190 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e88bb8 of type: SgFunctionParameterList):
26  No attached comments (at 0xb7ecc328 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e88cc8 of type: SgFunctionParameterList):
28  No attached comments (at 0xb7ecc4c0 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e88dd8 of type: SgFunctionParameterList):
30  No attached comments (at 0xb7ecc658 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e88ee8 of type: SgFunctionParameterList):
32  No attached comments (at 0xb7ecc7f0 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e88ff8 of type: SgFunctionParameterList):
34  No attached comments (at 0xb7ecc988 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e89108 of type: SgFunctionParameterList):
36  No attached comments (at 0xb7eccb20 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e89218 of type: SgFunctionParameterList):
38  No attached comments (at 0xb7ecccb8 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e89328 of type: SgFunctionParameterList):
40  No attached comments (at 0xb7ecce50 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e89438 of type: SgFunctionParameterList):
42  No attached comments (at 0xb7eccfe8 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e89548 of type: SgFunctionParameterList):
44  No attached comments (at 0xb7ecd180 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e89658 of type: SgFunctionParameterList):
46  No attached comments (at 0xb7ecd318 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e89768 of type: SgFunctionParameterList):
48  No attached comments (at 0xb7ecd4b0 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e89878 of type: SgFunctionParameterList):
50  No attached comments (at 0xb7ecd648 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e89988 of type: SgFunctionParameterList):
52  No attached comments (at 0xb7ecd7e0 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e89a98 of type: SgFunctionParameterList):
54  No attached comments (at 0xb7ecd978 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e89ba8 of type: SgFunctionParameterList):
56  No attached comments (at 0xb7ecdb10 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e89cb8 of type: SgFunctionParameterList):
58  No attached comments (at 0xb7ecdca8 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e89dc8 of type: SgFunctionParameterList):
60  No attached comments (at 0xb7ecde40 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e89ed8 of type: SgFunctionParameterList):
62  No attached comments (at 0xb7ecdfd8 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e89fe8 of type: SgFunctionParameterList):
64  No attached comments (at 0xb7ece170 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e8a0f8 of type: SgFunctionParameterList):
66  No attached comments (at 0xb7ece308 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e8a208 of type: SgFunctionParameterList):
68  No attached comments (at 0xb7ece4a0 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e8a318 of type: SgFunctionParameterList):
70  No attached comments (at 0xb7ece638 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e8a428 of type: SgFunctionParameterList):
72  No attached comments (at 0xb7ece7d0 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e8a538 of type: SgFunctionParameterList):
74  No attached comments (at 0xb7ece968 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e8a648 of type: SgFunctionParameterList):
76  No attached comments (at 0xb7eceb00 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e8a758 of type: SgFunctionParameterList):
78  No attached comments (at 0xb7ecec98 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e8a868 of type: SgFunctionParameterList):
80  No attached comments (at 0xb7ecee30 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e8a978 of type: SgFunctionParameterList):
82  No attached comments (at 0xb7ecefc8 of type: SgFunctionDeclaration):
    No attached comments (at 0xb7e8aa88 of type: SgFunctionParameterList):
```

```
Collect all comments and CPP directives into the AST (from header files)
AttachAllPreprocessingInfoTreeTrav constructor: This code is now redundant with the include/exclude support fo
AttachAllPreprocessingInfoTreeTrav::evaluateInheritedAttribute(): Calling getPreprocessorDirectives for file =
AttachAllPreprocessingInfoTreeTrav::evaluateInheritedAttribute(): Calling getPreprocessorDirectives for file =
No attached comments (at 0x9491bb0 of type: SgGlobal):
No attached comments (at 0xb7f02008 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ebf008 of type: SgFunctionParameterList):
No attached comments (at 0xb7f021a0 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ebf118 of type: SgFunctionParameterList):
No attached comments (at 0xb7f02338 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ebf228 of type: SgFunctionParameterList):
No attached comments (at 0xb7f024d0 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ebf338 of type: SgFunctionParameterList):
No attached comments (at 0xb7f02668 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ebf448 of type: SgFunctionParameterList):
No attached comments (at 0xb7f02800 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ebf558 of type: SgFunctionParameterList):
No attached comments (at 0xb7f02998 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ebf668 of type: SgFunctionParameterList):
No attached comments (at 0xb7f02b30 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ebf778 of type: SgFunctionParameterList):
No attached comments (at 0xb7f02cc8 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ebf888 of type: SgFunctionParameterList):
No attached comments (at 0xb7f02e60 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ebf998 of type: SgFunctionParameterList):
No attached comments (at 0xb7f02ff8 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ebfaa8 of type: SgFunctionParameterList):
No attached comments (at 0xb7f03190 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ebfbb8 of type: SgFunctionParameterList):
No attached comments (at 0xb7f03328 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ebfcc8 of type: SgFunctionParameterList):
No attached comments (at 0xb7f034c0 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ebfdd8 of type: SgFunctionParameterList):
No attached comments (at 0xb7f03658 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ebfee8 of type: SgFunctionParameterList):
No attached comments (at 0xb7f037f0 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ebfff8 of type: SgFunctionParameterList):
No attached comments (at 0xb7f03988 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec0108 of type: SgFunctionParameterList):
No attached comments (at 0xb7f03b20 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec0218 of type: SgFunctionParameterList):
No attached comments (at 0xb7f03cb8 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec0328 of type: SgFunctionParameterList):
No attached comments (at 0xb7f03e50 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec0438 of type: SgFunctionParameterList):
No attached comments (at 0xb7f03fe8 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec0548 of type: SgFunctionParameterList):
No attached comments (at 0xb7f04180 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec0658 of type: SgFunctionParameterList):
No attached comments (at 0xb7f04318 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec0768 of type: SgFunctionParameterList):
No attached comments (at 0xb7f044b0 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec0878 of type: SgFunctionParameterList):
No attached comments (at 0xb7f04648 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec0988 of type: SgFunctionParameterList):
No attached comments (at 0xb7f047e0 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec0a98 of type: SgFunctionParameterList):
No attached comments (at 0xb7f04978 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec0ba8 of type: SgFunctionParameterList):
No attached comments (at 0xb7f04b10 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec0cb8 of type: SgFunctionParameterList):
No attached comments (at 0xb7f04ca8 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec0dc8 of type: SgFunctionParameterList):
No attached comments (at 0xb7f04e40 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec0ed8 of type: SgFunctionParameterList):
No attached comments (at 0xb7f04fd8 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec0fe8 of type: SgFunctionParameterList):
No attached comments (at 0xb7f05170 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec10f8 of type: SgFunctionParameterList):
No attached comments (at 0xb7f05308 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec1208 of type: SgFunctionParameterList):
No attached comments (at 0xb7f054a0 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec1318 of type: SgFunctionParameterList):
No attached comments (at 0xb7f05638 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec1428 of type: SgFunctionParameterList):
No attached comments (at 0xb7f057d0 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec1538 of type: SgFunctionParameterList):
No attached comments (at 0xb7f05968 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec1648 of type: SgFunctionParameterList):
No attached comments (at 0xb7f05b00 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec1758 of type: SgFunctionParameterList):
No attached comments (at 0xb7f05c98 of type: SgFunctionDeclaration):
No attached comments (at 0xb7ec1868 of type: SgFunctionParameterList):
```

```
// Example ROSE Translator: used within ROSE/tutorial

#include "rose.h"

using namespace std;

// Build a synthesized attribute for the tree traversal
class SynthesizedAttribute
   {
     public:
       // List of #define directives (save the PreprocessingInfo objects
       // so that we have all the source code position information).
          list<PreprocessingInfo*> accumulatedList;

          void display() const;
   };

void
SynthesizedAttribute::display() const
   {
     list<PreprocessingInfo*>::const_iterator i = accumulatedList.begin();
     while (i != accumulatedList.end())
        {
          printf ("CPP_define_directive_=_%s_\n",(*i)->getString().c_str());
          i++;
        }
   }

class visitorTraversal : public AstBottomUpProcessing<SynthesizedAttribute>
   {
     public:
       // virtual function must be defined
          virtual SynthesizedAttribute evaluateSynthesizedAttribute (
                       SgNode* n, SynthesizedAttributesList childAttributes );
   };

SynthesizedAttribute
visitorTraversal::evaluateSynthesizedAttribute ( SgNode* n, SynthesizedAttributesList childAttributes )
   {
     SynthesizedAttribute localResult;

  // printf ("In evaluateSynthesizedAttribute(n = %p = %s) \n",n,n->class_name().c_str());

  // Build the list from children (in reverse order to preserve the final ordering)
     for (SynthesizedAttributesList::reverse_iterator child = childAttributes.rbegin(); child != childAttributes.rend();
        {
          localResult.accumulatedList.splice(localResult.accumulatedList.begin(),child->accumulatedList);
        }

  // Add in the information from the current node
     SgLocatedNode* locatedNode = isSgLocatedNode(n);
     if (locatedNode != NULL)
        {
          AttachedPreprocessingInfoType* commentsAndDirectives = locatedNode->getAttachedPreprocessingInfo();

          if (commentsAndDirectives != NULL)
             {
            // printf ("Found attached comments (to IR node at %p of type: %s): \n",locatedNode,locatedNode->class_name()
            // int counter = 0;

            // Use a reverse iterator so that we preserve the order when using push_front to add each directive to the ac
               AttachedPreprocessingInfoType::reverse_iterator i;
               for (i = commentsAndDirectives->rbegin(); i != commentsAndDirectives->rend(); i++)
                  {
                 // The different classifications of comments and directives are in ROSE/src/frontend/SageIII/rose_attrib
                    if ((*i)->getTypeOfDirective() == PreprocessingInfo::CpreprocessorDefineDeclaration)
                       {
#if 0
                         printf ("___________Attached_Comment_#%d_in_file_%s_(relativePosition=%s):_classification_%s_:\n%
                            counter++,(*i)->get_file_info()->get_filenameString().c_str(),
                            ((*i)->getRelativePosition() == PreprocessingInfo::before) ? "before" : "after",
                            PreprocessingInfo::directiveTypeName((*i)->getTypeOfDirective()).c_str(),
                            (*i)->getString().c_str());
#endif
                       // use push_front() to end up with source ordering of final list of directives
                          localResult.accumulatedList.push_front(*i);
                       }
                  }
             }
        }

  // printf ("localResult after adding current node info \n");
  // localResult.display();
```

```
2 #define JUST_A_MACRO just_a_macro

4 #define ANOTHER_MACRO another_macro
```

Figure 31.6: Example source code used as input to collection of comments and CPP directives.

```
  Collect all comments and CPP directives into the AST (from header files)
2 AttachAllPreprocessingInfoTreeTrav constructor: This code is now redundant with the include/exclude support fo
  AttachAllPreprocessingInfoTreeTrav::evaluateInheritedAttribute(): Calling getPreprocessorDirectives for file =
4 AttachAllPreprocessingInfoTreeTrav::evaluateInheritedAttribute(): Calling getPreprocessorDirectives for file =
  CPP define directive = #define max(a,b) ((a) > (b) ? (a) : (b))
6
  CPP define directive = #define maxint(a,b) ({int _a = (a), _b = (b); _a > _b ? _a : _b; })
8
  CPP define directive = #define SOURCE_CODE_BEFORE_INCLUDE_A
10
  CPP define directive = #define SOURCE_CODE_BEFORE_INCLUDE_B
12
  CPP define directive = #define SOURCE_CODE_AFTER_INCLUDE_A
14
  CPP define directive = #define SOURCE_CODE_AFTER_INCLUDE_B
```

Figure 31.7: Output from collection of comments and CPP directives on the input source file and all header files.

```
// Example ROSE Translator: used within ROSE/tutorial

#include "rose.h"

using namespace std;

class visitorTraversal : public AstSimpleProcessing
   {
     public:
          virtual void visit(SgNode* n);
   };

void visitorTraversal::visit(SgNode* n)
   {
     SgFunctionDeclaration* functionDeclaration = isSgFunctionDeclaration(n);
     if (functionDeclaration != NULL)
        {
          string comment = string("Auto-comment function name: ") +
                           functionDeclaration->get_name().str() +
                           " is now a commented function";

       // Note that this function will add the "//" or "/* */" comment syntax as required for C or C++, or Fortran.
          SageInterface::attachComment(functionDeclaration,comment);
        }

     SgValueExp* valueExp = isSgValueExp(n);
     if (valueExp != NULL)
        {
       // Check if there is an expression tree from the original unfolded expression.
       // This is a trivial example ouf the output of an analysis result.
          string comment = string("Auto-comment value: ") +
               ((valueExp->get_originalExpressionTree() != NULL) ?
                    " this IS a constant folded value" : " this is NOT a constant folded value");

          SageInterface::attachComment(valueExp,comment);
        }

   }

// Typical main function for ROSE translator
int main( int argc, char * argv[] )
   {
  // Build the AST used by ROSE
     SgProject* project = frontend(argc,argv);

  // Build the traversal object
     visitorTraversal exampleTraversal;

  // Call the traversal starting at the project node of the AST
     exampleTraversal.traverseInputFiles(project,preorder);

     return backend(project);
   }
```

Figure 31.8: Example source code showing how automate comments.

```
2  int
   foo()
4      {
         int x = 2;
6        return x;
       }
```

Figure 31.9: Example source code used as input to automate generation of comments.

```
   // Auto-comment function name: foo is now a commented function
2
   int foo()
4  {
     int x =
6  // Auto-comment value:  this is NOT a constant folded value
   2;
8    return x;
   }
```

Figure 31.10: Output of input code after automating generation of comments.

```
// Example ROSE Translator: used within ROSE/tutorial

#include "rose.h"

using namespace std;

class visitorTraversal : public AstSimpleProcessing
   {
     public:
          virtual void visit(SgNode* n);
   };

void visitorTraversal::visit(SgNode* n)
   {
     SgFunctionDeclaration* functionDeclaration = isSgFunctionDeclaration(n);
     if (functionDeclaration != NULL)
        {
       // This is an example of a XYZ tool specific annotation
          string compilerSpecificDirective = "\n#if XYZ_TOOL \n   \" builtin \"\n#endif\n";
          SageInterface::addTextForUnparser(functionDeclaration,compilerSpecificDirective,AstUnparseAttribute::e_before);
        }

     SgValueExp* valueExp = isSgValueExp(n);
     if (valueExp != NULL)
        {
       // Add a backend specific compiler directive
          string compilerSpecificDirective = "\n#if CRAY \n   cray_specific_attribute \n#endif\n";
          SageInterface::addTextForUnparser(valueExp,compilerSpecificDirective,AstUnparseAttribute::e_before);
        }

   }

// Typical main function for ROSE translator
int main( int argc, char * argv[] )
   {
  // Build the AST used by ROSE
     SgProject* project = frontend(argc,argv);

  // Build the traversal object
     visitorTraversal exampleTraversal;

  // Call the traversal starting at the project node of the AST
     exampleTraversal.traverseInputFiles(project,preorder);

     return backend(project);
   }
```

Figure 31.11: Example source code showing how automate the introduction of arbitrary text.

```
int
foo()
   {
     int x = 42;
     return x;
   }
```

Figure 31.12: Example source code used as input to automate generation of arbitrary text.

```
#if XYZ_TOOL
    "builtin"
#endif

int foo()
{
  int x =
#if CRAY
    cray_specific_attribute
#endif
 42;
  return x;
}
```

Figure 31.13: Output of input code after automating generation of arbitrary text.

Chapter 32

Tailoring The Code Generation Format

Figure 32.1 shows an example of how to use the mechanisms in ROSE to tailor the format and style of the generated code. This chapter presents an example translator that modifies the formatting of the code that is generated within ROSE.

The details of functionality are hidden from the user and a high level interface is provided that permits key parameters to be specified. This example will be made more sophisticated later, for now it just modifies the indentation of nested code blocks (from 2 spaces/block to 5 spaces/block).

32.1 Source Code for Example that Tailors the Code Generation

Figure 32.1 shows an example translator which calls the inliner mechanism. The code is designed to only inline up to ten functions. the list of function calls is recomputed after any function call is successfully inlined.

The input code is shown in figure 32.2, the output of this code is shown in figure 32.3.

32.2 Input to Demonstrate Tailoring the Code Generation

Figure 32.2 shows the example input used for demonstration of how to control the formatting of generated code.

32.3 Final Code After Tailoring the Code Generation

Figure 32.3 shows the results from changes to the formatting of generated code.

```
// This example will be made more sophisticated later, for now it just
// modifies the indentation of nested code blocks (from 2 spaces/block
// to 5 spaces/block).

#include "rose.h"
#include "unparseFormatHelp.h"

class CustomCodeFormat : public UnparseFormatHelp
   {
     public:
          CustomCodeFormat();
         ~CustomCodeFormat();

          virtual int getLine( SgLocatedNode*, SgUnparse_Info& info, FormatOpt opt);
          virtual int getCol ( SgLocatedNode*, SgUnparse_Info& info, FormatOpt opt);

       // return the value for indentation of code (part of control over style)
          virtual int tabIndent ();

       // return the value for where line wrapping starts (part of control over style)
          virtual int maxLineLength ();

     private:
          int defaultLineLength;
          int defaultIndentation;
   };


CustomCodeFormat::CustomCodeFormat()
   {
  // default values here!
     defaultLineLength = 20;
     defaultIndentation = 5;
   }

CustomCodeFormat::~CustomCodeFormat()
   {}

// return: > 0: start new lines; == 0: use same line; < 0:default
int
CustomCodeFormat::getLine( SgLocatedNode*, SgUnparse_Info& info, FormatOpt opt)
   {
  // Use default mechanism to select the line where to output generated code
     return -1;
   }

// return starting column. if < 0, use default
int
CustomCodeFormat::getCol( SgLocatedNode*, SgUnparse_Info& info, FormatOpt opt)
   {
  // Use default mechanism to select the column where to output generated code
     return -1;
   }

int
CustomCodeFormat::tabIndent()
   {
  // Modify the indentation of the generated code (trival example of tailoring code generation)
     return defaultIndentation;
   }

int
CustomCodeFormat::maxLineLength()
   {
     return defaultLineLength;
   }


int main (int argc, char* argv[])
   {
  // Build the project object (AST) which we will fill up with multiple files and use as a
  // handle for all processing of the AST(s) associated with one or more source files.
     SgProject* project = new SgProject(argc,argv);

     CustomCodeFormat* formatControl = new CustomCodeFormat();

     return backend(project,formatControl);
   }
```

Figure 32.1: Example source code showing how to tailor the code generation format.

```c
extern int min(int ,int );

void dgemm(double *a,double *b,double *c,int n)
{
  int _var_1;
  int _var_0;
  int i;
  int j;
  int k;
  for (_var_1 = 0; _var_1 <= -1 + n; _var_1 += 16) {
    for (_var_0 = 0; _var_0 <= -1 + n; _var_0 += 16) {
      for (i = 0; i <= -1 + n; i += 1) {
        for (k = _var_1; k <= min(-1 + n,_var_1 + 15); k += 1) {
          int dummy_1 = k * n + i;
          for (j = _var_0; j <= min(n + -16,_var_0); j += 16) {
            int _var_2 = (j);
            c[j * n + i] = c[j * n + i] + a[k * n + i] * b[j * n + k];
            _var_2 = 1 + _var_2;
            c[_var_2 * n + i] = c[_var_2 * n + i] + a[k * n + i] * b[_var_2 * n + k];
            _var_2 = 1 + _var_2;
            c[_var_2 * n + i] = c[_var_2 * n + i] + a[k * n + i] * b[_var_2 * n + k];
            _var_2 = 1 + _var_2;
            c[_var_2 * n + i] = c[_var_2 * n + i] + a[k * n + i] * b[_var_2 * n + k];
            _var_2 = 1 + _var_2;
            c[_var_2 * n + i] = c[_var_2 * n + i] + a[k * n + i] * b[_var_2 * n + k];
            _var_2 = 1 + _var_2;
            c[_var_2 * n + i] = c[_var_2 * n + i] + a[k * n + i] * b[_var_2 * n + k];
            _var_2 = 1 + _var_2;
            c[_var_2 * n + i] = c[_var_2 * n + i] + a[k * n + i] * b[_var_2 * n + k];
            _var_2 = 1 + _var_2;
            c[_var_2 * n + i] = c[_var_2 * n + i] + a[k * n + i] * b[_var_2 * n + k];
            _var_2 = 1 + _var_2;
            c[_var_2 * n + i] = c[_var_2 * n + i] + a[k * n + i] * b[_var_2 * n + k];
            _var_2 = 1 + _var_2;
            c[_var_2 * n + i] = c[_var_2 * n + i] + a[k * n + i] * b[_var_2 * n + k];
            _var_2 = 1 + _var_2;
            c[_var_2 * n + i] = c[_var_2 * n + i] + a[k * n + i] * b[_var_2 * n + k];
            _var_2 = 1 + _var_2;
            c[_var_2 * n + i] = c[_var_2 * n + i] + a[k * n + i] * b[_var_2 * n + k];
            _var_2 = 1 + _var_2;
            c[_var_2 * n + i] = c[_var_2 * n + i] + a[k * n + i] * b[_var_2 * n + k];
            _var_2 = 1 + _var_2;
            c[_var_2 * n + i] = c[_var_2 * n + i] + a[k * n + i] * b[_var_2 * n + k];
            _var_2 = 1 + _var_2;
            c[_var_2 * n + i] = c[_var_2 * n + i] + a[k * n + i] * b[_var_2 * n + k];
            _var_2 = 1 + _var_2;
            c[_var_2 * n + i] = c[_var_2 * n + i] + a[k * n + i] * b[_var_2 * n + k];
          }
          for (; j <= min(-1 + n,_var_0 + 15); j += 1) {
            c[j * n + i] = c[j * n + i] + a[k * n + i] * b[j * n + k];
          }
        }
      }
    }
  }
}
```

Figure 32.2: Example source code used as input to program to the tailor the code generation.

```
extern int min(int ,int );

void dgemm(double *a,double *b,double *c,int n)
{
    int _var_1;
    int _var_0;
    int i;
    int j;
    int k;
    for (_var_1 = 0; _var_1 <= ((-1) + n); _var_1 += 16) {
        for (_var_0 = 0; _var_0 <= ((-1) + n); _var_0 += 16) {
            for (i = 0; i <= ((-1) + n); i += 1) {
                for (k = _var_1; k <= min(((-1) + n),(_var_1 + 15)); k += 1) {
                int dummy_1 = ((k * n) + i);
                for (j = _var_0; j <= min((n + (-16)),_var_0); j += 16) {
                int _var_2 = j;
                c[(j * n) + i] = ((c[(j * n) + i]) + ((a[(k * n) + i]) * (b[(j * n) + k])));
                _var_2 = (1 + _var_2);
                c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])))
                _var_2 = (1 + _var_2);
                c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])))
                _var_2 = (1 + _var_2);
                c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])))
                _var_2 = (1 + _var_2);
                c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])))
                _var_2 = (1 + _var_2);
                c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])))
                _var_2 = (1 + _var_2);
                c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])))
                _var_2 = (1 + _var_2);
                c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])))
                _var_2 = (1 + _var_2);
                c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])))
                _var_2 = (1 + _var_2);
                c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])))
                _var_2 = (1 + _var_2);
                c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])))
                _var_2 = (1 + _var_2);
                c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])))
                _var_2 = (1 + _var_2);
                c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])))
                _var_2 = (1 + _var_2);
                c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])))
                _var_2 = (1 + _var_2);
                c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])))
                _var_2 = (1 + _var_2);
                c[(_var_2 * n) + i] = ((c[(_var_2 * n) + i]) + ((a[(k * n) + i]) * (b[(_var_2 * n) + k])))
                }
                for (; j <= min(((-1) + n),(_var_0 + 15)); j += 1) {
                c[(j * n) + i] = ((c[(j * n) + i]) + ((a[(k * n) + i]) * (b[(j * n) + k])));
                }
                }
            }
        }
    }
}
```

Figure 32.3: Output of input code after changing the format of the generated code.

Chapter 33

Command-line Processing Within Translators

ROSE includes mechanism to simplify the processing of command-line arguments so that translators using ROSE can trivially replace compilers within makefiles. This example shows some of the many command-line handling options within ROSE and the ways in which customized options may be added for specific translators.

33.1 Commandline Selection of Files

Overview This example shows the optional processing of specific files selected after the call to the frontend to build the project. First the SgProject if build and *then* the files are selected for processing via ROSE or the backend compiler directly.

This example demonstrates the separation of the construction of a SgProject with valid SgFile objects for each file on the command line, but with an empty SgGlobal scope, and the call to the frontend, called for each SgFile in a separate loop over all the SgFile objects.

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.
// rose.C: Example (default) ROSE Preprocessor: used for testing ROSE infrastructure

#include "rose.h"

using namespace std;

int
main ( int argc, char* argv[] )
   {
     Rose_STL_Container<string> l = CommandlineProcessing::generateArgListFromArgcArgv (argc,argv);
     printf ("Preprocessor (before): argv = \n%s \n",StringUtility::listToString(l).c_str());

  // Remove certain sorts of options from the command line
     CommandlineProcessing::removeArgs (l,"-edg:");
     CommandlineProcessing::removeArgs (l,"--edg:");
     CommandlineProcessing::removeArgsWithParameters (l,"-edg_parameter:");
     CommandlineProcessing::removeArgsWithParameters (l,"--edg_parameter:");

  // Add a test for a custom command line option
     int integerOptionForVerbose = 0;
     if ( CommandlineProcessing::isOptionWithParameter(l,"-myTranslator:","(v|verbose)",integerOptionForVerbos
        {
          printf ("Turning on my translator's verbose mode (set to %d) \n",integerOptionForVerbose);
        }

  // Adding a new command line parameter (for mechanisms in ROSE that take command lines)

     // printf ("argc = %zu \n",l.size());
     // l = CommandlineProcessing::generateArgListFromArgcArgv (argc,argv);
     printf ("l.size() = %zu \n",l.size());
     printf ("Preprocessor (after): argv = \n%s \n",StringUtility::listToString(l).c_str());

  // SgProject* project = frontend(argc,argv);
  // ROSE_ASSERT (project != NULL);
  // Generate the source code and compile using the vendor's compiler
  // return backend(project);

  // Build the AST, generate the source code and call the backend compiler ...
     frontend(l);
     return 0;
   }
```

Figure 33.1: Example source code showing simple command-line processing within ROSE translator.

```
Preprocessor (before): argv =
/home/liao6/daily-test-rose/20081014_120001/build/tutorial/.libs/lt-commandlineProcessing --edg:no_warnings -w
Turning on my translator's verbose mode (set to 42)
l.size() = 4
Preprocessor (after): argv =
/home/liao6/daily-test-rose/20081014_120001/build/tutorial/.libs/lt-commandlineProcessing -w -c ../../sourcetr
```

Figure 33.2: Output of input code using commandlineProcessing.C

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.
// rose.C: Example (default) ROSE Preprocessor: used for testing ROSE infrastructure

#include "rose.h"

using namespace std;

int
main ( int argc, char* argv[] )
   {
     Rose_STL_Container<string> l = CommandlineProcessing::generateArgListFromArgcArgv (argc,argv);
     printf ("Preprocessor (before): argv = \n%s \n",StringUtility::listToString(l).c_str());

  // Remove certain sorts of options from the command line
     CommandlineProcessing::removeArgs (l,"-edg:");
     CommandlineProcessing::removeArgs (l,"--edg:");
     CommandlineProcessing::removeArgsWithParameters (l,"-edg_parameter:");
     CommandlineProcessing::removeArgsWithParameters (l,"--edg_parameter:");

  // Add a test for a custom command line option
     int integerOptionForVerbose = 0;
     if ( CommandlineProcessing::isOptionWithParameter(l,"-myTranslator:","(v|verbose)",integerOptionForVerbose,true) )
        {
          printf ("Turning on my translator's verbose mode (set to %d) \n",integerOptionForVerbose);
        }

  // Adding a new command line parameter (for mechanisms in ROSE that take command lines)

     // printf ("argc = %zu \n",l.size());
     // l = CommandlineProcessing::generateArgListFromArgcArgv (argc,argv);
     printf ("l.size() = %zu \n",l.size());
     printf ("Preprocessor (after): argv = \n%s \n",StringUtility::listToString(l).c_str());

  // SgProject* project = frontend(argc,argv);
  // ROSE_ASSERT (project != NULL);
  // Generate the source code and compile using the vendor's compiler
  // return backend(project);

  // Build the AST, generate the source code and call the backend compiler ...
     frontend(l);
     return 0;
   }
```

Figure 33.3: Example source code showing simple command-line processing within ROSE translator.

```
Preprocessor (before): argv =
/home/liao6/daily-test-rose/20081014_120001/build/tutorial/.libs/lt-commandlineProcessing --edg:no_warnings -w -c ../../s
Turning on my translator's verbose mode (set to 42)
l.size() = 4
Preprocessor (after): argv =
/home/liao6/daily-test-rose/20081014_120001/build/tutorial/.libs/lt-commandlineProcessing -w -c ../../sourcetree/tutorial
```

Figure 33.4: Output of input code using commandlineProcessing.C

Chapter 34

Building Custom Graphs

What To Learn From This Example This example shows how to generate custom graphs (typically used for analysis results).

The mechanisms used internally to build different graphs of program data are also made externally available. This section shows how new graphs of program information can be built or existing graphs customized. More information about generation of specialized AST graphs to support debugging can be found in 35.

Figure 34.1 shows a graph representing the highest levels of the ROSE directory tree built using the code shown in figure 34.2. This example is build explicitly in the example code, the script *lsdot* in the *ROSE/scripts/* directory builds such a graph of the directory tree automatically for any directory structure and is used to present the ROSE source directory tree in the *ROSE User Manual*.

Figure 34.3 shows the same graph but with filtering to tailor the graph by removing nodes. The removal of any node from the graph automatically removes all edged pointing to that node removed and all edges pointing away from the node being removed. This mechanism is provided as a way to customized graphs automatically generated using ROSE (e.g. call graphs, control flow graphs, etc.).

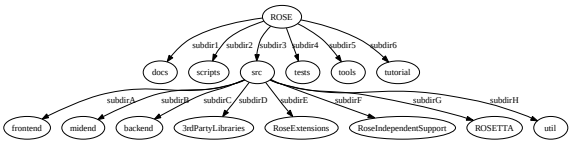

Figure 34.1: Graph of top level of ROSE directory tree.

```
// Example of using Qing's graph interface for an arbitrary graph.

#include "rose.h"
#include <GraphDotOutput.h>
#include <VirtualGraphCreate.h>
using namespace std;

class Node : public MultiGraphElem {
     public:
          std::string name;
          Node( std::string n ) : MultiGraphElem( NULL ), name( n ) {}
          virtual std::string toString() const { return name; }
   };

class Edge : public MultiGraphElem {
     public:
          std::string label;
          Edge ( std::string label = "default edge" ) : MultiGraphElem( NULL ), label( label ) {};
          void Dump() const { printf ("EDGE: label = %s \n",label.c_str()); }
          virtual std::string toString() const { return label;}
   };

template <class NodeType, class EdgeType>
class GraphBuilder : public VirtualGraphCreateTemplate<NodeType, EdgeType> {
     public:
          void addNode ( NodeType* node );
          void addEdge ( NodeType* src, NodeType* snk, EdgeType* edge );
          GraphBuilder () : VirtualGraphCreateTemplate<NodeType, EdgeType> (NULL) {}
         ~GraphBuilder() { printf ("Inside of ~GraphBuilder() \n"); }
   };

template <class NodeType, class EdgeType>
void GraphBuilder<NodeType, EdgeType>::addNode ( NodeType* node )
   { VirtualGraphCreateTemplate<NodeType,EdgeType>::AddNode ( node ); }

template <class NodeType, class EdgeType>
void GraphBuilder<NodeType, EdgeType>::addEdge ( NodeType* src, NodeType* snk, EdgeType* edge )
   { VirtualGraphCreateTemplate<NodeType,EdgeType>::AddEdge ( src, snk, edge ); }

int main( int argc, char * argv[] )
   {
     GraphBuilder<Node, Edge> graph;

     Node* rose                   = new Node("ROSE");                   graph.addNode(rose);
     Node* docs                   = new Node("docs");                   graph.addNode(docs);
     Node* scripts                = new Node("scripts");                graph.addNode(scripts);
     Node* src                    = new Node("src");                    graph.addNode(src);
     Node* tests                  = new Node("tests");                  graph.addNode(tests);
     Node* tools                  = new Node("tools");                  graph.addNode(tools);
     Node* tutorial               = new Node("tutorial");               graph.addNode(tutorial);
     Node* frontend               = new Node("frontend");               graph.addNode(frontend);
     Node* midend                 = new Node("midend");                 graph.addNode(midend);
     Node* backend                = new Node("backend");                graph.addNode(backend);
     Node* thirdPartyLibraries    = new Node("3rdPartyLibraries");      graph.addNode(thirdPartyLibraries);
     Node* roseExtensions         = new Node("RoseExtensions");         graph.addNode(roseExtensions);
     Node* roseIndependentSupport = new Node("RoseIndependentSupport"); graph.addNode(roseIndependentSupport);
     Node* rosetta                = new Node("ROSETTA");                graph.addNode(rosetta);
     Node* util                   = new Node("util");                   graph.addNode(util);

     graph.addEdge(rose,docs,                   new Edge("subdir1"));
     graph.addEdge(rose,scripts,                new Edge("subdir2"));
     graph.addEdge(rose,src,                    new Edge("subdir3"));
     graph.addEdge(rose,tests,                  new Edge("subdir4"));
     graph.addEdge(rose,tools,                  new Edge("subdir5"));
     graph.addEdge(rose,tutorial,               new Edge("subdir6"));
     graph.addEdge(src,frontend,                new Edge("subdirA"));
     graph.addEdge(src,midend,                  new Edge("subdirB"));
     graph.addEdge(src,backend,                 new Edge("subdirC"));
     graph.addEdge(src,thirdPartyLibraries,     new Edge("subdirD"));
     graph.addEdge(src,roseExtensions,          new Edge("subdirE"));
     graph.addEdge(src,roseIndependentSupport, new Edge("subdirF"));
     graph.addEdge(src,rosetta,                 new Edge("subdirG"));
     graph.addEdge(src,util,                    new Edge("subdirH"));

   // Build a wrapper for outputing DOT graph
     GraphDotOutput<GraphBuilder<Node, Edge> > output(graph);

   // Write out the DOT graph
     output.writeToDOTFile("customGraph.dot", "custom graph");

     return 0;
   }
```

Figure 34.3: Graph of top level of ROSE directory tree with filtering of subtree.

Chapter 35

General AST Graph Generation

What To Learn From This Example This example shows a maximally complete representation of the AST (often in more detail that is useful).

Where chapter 3 presented a ROSE-based translator which presented the AST as a tree, this chapter presents the more general representation of the graph in which the AST is embedded. The AST may be thought of as a subset of a more general graph or equivalently as an AST (a tree in a formal sense) with annotations (extra edges and information), sometimes referred to as a '*decorated* AST.

We present tools for seeing all the IR nodes in the graph containing the AST, including all types (SgType nodes), symbols (SgSymbol nodes), compiler generated IR nodes, and supporting IR nodes. In general it is a specific filtering of this larger graph which is more useful to communicating how the AST is designed and internally connected. We use these graphs for internal debugging (typically on small problems where the graphs are reasonable in size). The graphs presented using these graph mechanism present all back-edges, and demonstrate what IR nodes are shared internally (typically SgType IR nodes).

First a few names, we will call the AST those nodes in the IR that are specified by a traversal using the ROSE traversal (SgSimpleTraversal, etc.). We will call the graph of all IR nodes the *Graph of all IR nodes*. the AST is embedded in the *Graph of all IR nodes*. The AST *is* a tree, while the *graph of all IR nodes* typically not a tree (in a Graph Theory sense) since it typically contains cycles.

We cover the visualization of both the AST and the *Graph of all IR nodes*.

- AST graph
  These techniques define ways of visualizing the AST and filtering IR nodes from being represented.
  - Simple AST graphs
  - Colored AST graphs
  - Filtering the graph
    The AST graph may be generated for any subtree of the AST (not possible for the graphs of all IR nodes). Additionally runtime options permit null pointers to be ignored. .

FIXME: *Is this true?*

- *Graph of all IR nodes*
  These techniques define the ways of visualizing the whole graph of IR nodes and is based on the memory pool traversal as a means to access all IR nodes. Even disconnected portions of the AST will be presented.
  - Simple graphs
  - Colored graphs
  - Filtering the graph

35.1 Whole Graph Generation

This example shows how to generate and visualize the AST from any input program. Each node of the graph in figure 35.3 shows a node of the Intermediate Representation (IR). Each edge shows the connection of the IR nodes in memory. The generated graph shows the connection of different IR nodes to form the AST.

The program in figure 35.1 calls an internal ROSE function that traverses the AST and generates an ASCII file in `dot` format. Figure 35.2 shows an input code which is processed to generate a graph of the AST, generating a `dot` file. The `dot` file is then processed using `dot` to generate a postscript file 35.3 (within the `Makefile`). Note that a similar utility program already exists within ROSE/exampleTranslators (and includes a utility to output an alternative PDF representation (suitable for larger ASTs) as well). Figure 35.3 (../../..//tutorial/test.ps) can be found in the compile tree (in the tutorial directory) and viewed directly using ghostview or any postscript viewer to see more detail.

Figure 35.3 displays the individual C++ nodes in ROSE's intermediate representation (IR). Each circle represents a single IR node, the name of the C++ construct appears in the center of the circle, with the edge numbers of the traversal on top and the number of child nodes appearing below. Internal processing to build the graph generates unique values for each IR node, a pointer address, which is displays at the bottom of each circle. The IR nodes are connected for form a tree, and abstract syntax tree (AST). Each IR node is a C++ class, see SAGE III reference for details, the edges represent the values of data members in the class (pointers which connect the IR nodes to other IR nodes). The edges are labeled with the names of the data members in the classes representing the IR nodes.

```
/*****************************************************************************
 * This example explains how to generate a DOT graph using the
 *       - whole AST traversal
 *       - memory pool traversal
 * In order to reduce the size of the graphs it is possible to filter on both
 * nodes and edges. Coloring of nodes and edges is also possible.
 *****************************************************************************/

#include "rose.h"

using namespace std;

// This functor is derived from the STL functor mechanism

/*****************************************************************************************************
 * The unary functional
 *      struct filterOnNodes
 * returns an AST_Graph::FunctionalReturnType and takes a std::pair<SgNode*,std::string> as a paramater.
 * The pair represents a variable of type std::string and a variablename.
 * The type AST_Graph::FunctionalReturnType contains the variables
 *       * addToGraph : if false do not graph node or edge, else graph
 *       * DOTOptions : a std::string which sets the color of the node etc.
 *
 * PS!!! The std::string is currently not set to anything useful. DO NOT USE THE STRING.
 *        Maybe it does not make sence in this case and should be removed.
 *****************************************************************************************************/
struct filterOnNodes: public std::unary_function< pair< SgNode*, std::string >,AST_Graph::FunctionalReturnType >
   {
  // This functor filters SgFileInfo objects and IR nodes from the GNU compatability file
     result_type operator() ( argument_type x ) const;
   };


//The argument to the function is
filterOnNodes::result_type
filterOnNodes::operator()(filterOnNodes::argument_type x) const
   {
     AST_Graph::FunctionalReturnType returnValue;
     //Determine if the node is to be added to the graph. true=yes
     returnValue.addToGraph = true;
     //set colors etc. for the graph Node
     returnValue.DOTOptions = "shape=polygon,regular=0,URL=\"\\N\",tooltip=\"more info at \\N\",sides=4,peripheries=1,colo

     if( isSgProject(x.first) != NULL )
        returnValue.DOTOptions = "shape=ellipse,regular=0,URL=\"\\N\",tooltip=\"more info at \\N\",sides=4,peripheries=1,c



     //Filter out SgSymbols from the graph
     if ( isSgSymbol(x.first) != NULL )
          returnValue.addToGraph = false;

     if ( isSgType(x.first) != NULL )
          returnValue.addToGraph = false;

     //Filter out compiler generated nodes
     SgLocatedNode* locatedNode = isSgLocatedNode(x.first);
     if ( locatedNode != NULL )
        {
          Sg_File_Info* fileInfo = locatedNode->get_file_info();
          std::string filename(ROSE::stripPathFromFileName(fileInfo->get_filename()));


          if (filename.find("rose_edg_macros_and_functions_required_for_gnu.h") != std::string::npos)
             {
               returnValue.addToGraph = false;
             }

          if (fileInfo->isCompilerGenerated()==true)
             {
            // std::cout << "Is compiler generated\n";
               returnValue.addToGraph = false;
             }
        }

     return returnValue;
   }


/*****************************************************************************************************
 * The binary functional
 *       struct filterOnEdges
```

```
int x;

#if 0
int
main()
   {
     int x = 0;

     return 0;
   }
#endif
```

Figure 35.2: Example source code used as input to generate the AST graph.

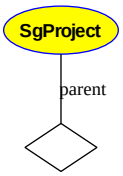

Figure 35.3: AST representing the source code file: inputCode_wholeGraphAST.C.

Chapter 36

Required Makefile for Tutorial Examples

This section shows an example makefile 36.1 required for the compilation of many of the tutorial example programs using the installed libraries (assumed to be generated from `make install`). The build process can be tested by running `make installcheck` from within the ROSE compile tree. This `makefile` can be found in the compile tree *(not the source tree)* for ROSE in the `tutorial` directory.

FIXME: *The exampleMakefile needs to be modified to include the support for compiling all of the tutorial examples.*

```
# Example Makefile for ROSE users
# This makefile is provided as an example of how to use ROSE when ROSE is
# installed (using "make install").  This makefile is tested as part of the
# "make distcheck" rule (run as part of tests before any SVN checkin).
# The test of this makefile can also be run by using the "make installcheck"
# rule (run as part of "make distcheck").

# Location of include directory after "make install"
ROSE_INCLUDE_DIR = /home/liao6/daily-test-rose/20081014_120001/install/include

# Location of Boost include directory
BOOST_CPPFLAGS = -pthread -I/home/liao6/opt/boost_1_35_0/include

# Location of library directory after "make install"
ROSE_LIB_DIR = /home/liao6/daily-test-rose/20081014_120001/install/lib

CPPFLAGS              =  -I/usr/apps/java/jdk1.5.0_11/include -I/usr/apps/java/jdk1.5.0_11/include/linux
CXXFLAGS              =  -g -Wall

ROSE_LIBS = $(ROSE_LIB_DIR)/librose.la

# Location of source code
ROSE_SOURCE_DIR = \
   ../../sourcetree/tutorial

executableFiles = identityTranslator ASTGraphGenerator \
               visitorTraversal inheritedAttributeTraversal \
               synthesizedAttributeTraversal \
               inheritedAndSynthesizedAttributeTraversal \
               accumulatorAttributeTraversal persistantAttributes \
               queryLibraryExample nestedTraversal \
               loopRecognition \
               typeInfoFromFunctionParameters \
               resolveOverloadedFunction templateParameter \
               instrumentationExample addVariableDeclaration \
               addFunctionDeclaration loopOptimization \
               buildCFG debuggingIRnodeToString \
               debuggingSourceCodePositionInformation \
               commandlineProcessing \
               loopNestingInfoProcessing

# Default make rule to use
all: $(executableFiles)
        @if [ x$${ROSE_IN_BUILD_TREE:+present} = xpresent ]; then echo "ROSE_IN_BUILD_TREE should not be set" >&2; exit 1; fi

# Example of how to use ROSE (linking to dynamic library, which is must faster
# and smaller than linking to static libraries).  Dynamic linking requires the
# use of the "-L$(ROSE_LIB_DIR) -Wl,-rpath" syntax if the LD_LIBRARY_PATH is not
# modified to use ROSE_LIB_DIR.  We provide two example of this; one using only
# the "-lrose -ledg" libraries, and one using the many separate ROSE libraries.
$(executableFiles):
#       g++ -I$(ROSE_INCLUDE_DIR) -o $@ $(ROSE_SOURCE_DIR)/$@.C -L$(ROSE_LIB_DIR) -Wl,-rpath $(ROSE_LIB_DIR) $(ROSE_LIBS)
#       g++ -I$(ROSE_INCLUDE_DIR) -o $@ $(ROSE_SOURCE_DIR)/$@.C $(LIBS_WITH_RPATH) $(ROSE_LIBS)
        ..//libtool --mode=link g++ $(CPPFLAGS) $(CXXFLAGS) -I$(ROSE_INCLUDE_DIR) $(BOOST_CPPFLAGS) -o $@ $(ROSE_SOURCE_DIR)/$@.C $(ROSE_L
```

Figure 36.1: Example `Makefile` showing how to use an installed version of ROSE (generated by `make install`).

Chapter 37

Tutorial Wrap-up

This tutorial has shown the construction and simple manipulation of the AST as part of the construction of the source-to-source translators using ROSE. Much more complex translators are possible using ROSE, but they are not such that they present well as part of a tutorial with short example programs. The remaining chapters of the tutorial include examples of translators built using ROSE as part of active collaborations with external research groups.

FIXME: *Reference the User Manual, HTML Doxygen generated documentation, unresolved issues, etc. Reference other work currently using ROSE (ANL, Cornell in the future), academic collaborations.*

Chapter 38

Code Coverage

This translator is part of ongoing collaboration with IBM on the support of code coverage analysis tools for C, C++ and F90 applications. the subject of code coverage is much more complex than this example code would cover. The following web site: *http://www.bullseye.com/coverage.html* contains more information and is the source for the descriptions below. Code coverage can include:

- Statement Coverage
  This measure reports whether each executable statement is encountered.

- Decision Coverage
  This measure reports whether boolean expressions tested in control structures (such as the if-statement and while-statement) evaluated to both true and false. The entire boolean expression is considered one true-or-false predicate regardless of whether it contains logical-and or logical-or operators. Additionally, this measure includes coverage of switch-statement cases, exception handlers, and interrupt handlers.

- Condition Coverage
  Condition coverage reports the true or false outcome of each boolean sub-expression, separated by logical-and and logical-or if they occur. Condition coverage measures the sub-expressions independently of each other.

- Multiple Condition Coverage
  Multiple condition coverage reports whether every possible combination of boolean sub-expressions occurs. As with condition coverage, the sub-expressions are separated by logical-and and logical-or, when present. The test cases required for full multiple condition coverage of a condition are given by the logical operator truth table for the condition.

- Condition/Decision Coverage
  Condition/Decision Coverage is a hybrid measure composed by the union of condition coverage and decision coverage. This measure was created at Boeing and is required for aviation software by RCTA/DO-178B.

- Modified Condition/Decision Coverage
  This measure requires enough test cases to verify every condition can affect the result of its encompassing decision.

- Path Coverage
  This measure reports whether each of the possible paths in each function have been followed. A path is a unique sequence of branches from the function entry to the exit.

- Function Coverage
  This measure reports whether you invoked each function or procedure. It is useful during preliminary testing to assure at least some coverage in all areas of the software. Broad, shallow testing finds gross deficiencies in a test suite quickly.

- Call Coverage
  This measure reports whether you executed each function call. The hypothesis is that faults commonly occur in interfaces between modules.

- Linear Code Sequence and Jump (LCSAJ) Coverage
  This variation of path coverage considers only sub-paths that can easily be represented in the program source code, without requiring a flow graph. An LCSAJ is a sequence of source code lines executed in sequence. This "linear" sequence can contain decisions as long as the control flow actually continues from one line to the next at run-time. Sub-paths are constructed by concatenating LCSAJs. Researchers refer to the coverage ratio of paths of length n LCSAJs as the test effectiveness ratio (TER) n+2.

- Data Flow Coverage
  This variation of path coverage considers only the sub-paths from variable assignments to subsequent references of the variables.

- Object Code Branch Coverage
  This measure reports whether each machine language conditional branch instruction both took the branch and fell through.

- Loop Coverage
  This measure reports whether you executed each loop body zero times, exactly once, and more than once (consecutively). For do-while loops, loop coverage reports whether you executed the body exactly once, and more than once. The valuable aspect of this measure is determining whether while-loops and for-loops execute more than once, information not reported by others measure.

- Race Coverage
  This measure reports whether multiple threads execute the same code at the same time. It helps detect failure to synchronize access to resources. It is useful for testing multi-threaded programs such as in an operating system.

- Relational Operator Coverage
  This measure reports whether boundary situations occur with relational operators (¡, ¡=, ¿, ¿=). The hypothesis is that boundary test cases find off-by-one errors and mistaken uses of wrong relational operators such as ¡ instead of ¡=.

- Weak Mutation Coverage
  This measure is similar to relational operator coverage but much more general [Howden1982]. It reports whether test cases occur which would expose the use of wrong operators and also wrong operands. It works by reporting coverage of conditions derived by substituting (mutating) the program's expressions with alternate operators, such as "-" substituted for "+", and with alternate variables substituted.

- Table Coverage
  This measure indicates whether each entry in a particular array has been referenced. This is useful for programs that are controlled by a finite state machine.

The rest of this text must be changed to refer to the code coverage example within ROSE/-tutorial.

Figure 38 shows the low level construction of a more complex AST fragment (a function declaration) and its insertion into the AST at the top of each block. Note that the code does not handle symbol table issues, yet.

Building a function in global scope.

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.
// Specifically it shows the design of a transformation to instrument source code, placing source code
// at the top and bottome of each basic block.

#include "rose.h"

using namespace std;

/*
   Design of this code.
      Inputs: source code (file.C)
      Outputs: instrumented source code (rose_file.C and file.o)

   Properties of instrumented source code:
       1) added declaration for coverage support function
          (either forward function declaration or a #include to include a header file).
       2) Each function in the source program is instrumented to include a call to the
          coverage support function/
*/


// Global variables so that the global function declaration can be reused to build
// each function call expression in the AST traversal to instrument all functions.
SgFunctionDeclaration* globalFunctionDeclaration = NULL;
SgFunctionType*        globalFunctionType        = NULL;
SgFunctionSymbol* functionSymbol = NULL;

// Simple ROSE traversal class: This allows us to visit all the functions and add
// new code to instrument/record their use.
class SimpleInstrumentation : public SgSimpleProcessing
   {
     public:
       // required visit function to define what is to be done
          void visit ( SgNode* astNode );
   };

// Code to build function declaration: This declares Shmuel's function call which
// will be inserted (as a function call) into each function body of the input
// application.
void buildFunctionDeclaration(SgProject* project)
   {
   // ****************************************************
   // Create the functionDeclaration
   // ****************************************************

   // SgGlobal* globalScope = project->get_file(0).get_root();
      SgSourceFile* sourceFile = isSgSourceFile(project->get_fileList()[0]);
      ROSE_ASSERT(sourceFile != NULL);
      SgGlobal* globalScope = sourceFile->get_globalScope();
      ROSE_ASSERT(globalScope != NULL);

      Sg_File_Info * file_info             = Sg_File_Info::generateDefaultFileInfoForTransformationNode();
      SgType * function_return_type        = new SgTypeVoid();

      SgName function_name                          = "coverageTraceFunc1";
      SgFunctionType * function_type                = new SgFunctionType(function_return_type,false);
      SgFunctionDeclaration * functionDeclaration = new SgFunctionDeclaration(file_info, function_name, function_type);

   // DQ (9/8/2007): Fixup the defining and non-defining declarations
      ROSE_ASSERT(functionDeclaration->get_definingDeclaration() == NULL);
      functionDeclaration->set_definingDeclaration(functionDeclaration);
      ROSE_ASSERT(functionDeclaration->get_definingDeclaration()          != NULL);
      ROSE_ASSERT(functionDeclaration->get_firstNondefiningDeclaration() != functionDeclaration);

   // DQ (9/8/2007): We have not build a non-defining declaration, so this should be NULL.
      ROSE_ASSERT(functionDeclaration->get_firstNondefiningDeclaration() == NULL);

   // DQ (9/8/2007): Need to add function symbol to global scope!
      printf ("Fixing up the symbol table in scope = %p = %s for function = %p = %s \n",globalScope,globalScope->class_name().c_st
      functionSymbol = new SgFunctionSymbol(functionDeclaration);
      globalScope->insert_symbol(functionDeclaration->get_name(),functionSymbol);
      ROSE_ASSERT(globalScope->lookup_function_symbol(functionDeclaration->get_name()) != NULL);

   // ****************************************************************
   // Create the InitializedName for a parameter within the parameter list
   // ****************************************************************
      SgName var1_name = "textString";
```

Figure 38.1: Example source code shows instrumentation to call a test function from the top of each function body in the application (part 1).

```
        SgTypeChar * var1_type             = new SgTypeChar();
        SgPointerType *pointer_type        = new SgPointerType(var1_type);
        SgInitializer * var1_initializer   = NULL;
        SgInitializedName *var1_init_name = new SgInitializedName(var1_name, pointer_type, var1_initializer, NULL);
        var1_init_name->set_file_info(Sg_File_Info::generateDefaultFileInfoForTransformationNode());

    // Insert argument in function parameter list
        ROSE_ASSERT(functionDeclaration != NULL);
        ROSE_ASSERT(functionDeclaration->get_parameterList() != NULL);

        ROSE_ASSERT(functionDeclaration->get_parameterList() != NULL);
        functionDeclaration->get_parameterList()->append_arg(var1_init_name);

    // Set the parent node in the AST (this could be done by the AstPostProcessing
        functionDeclaration->set_parent(globalScope);

    // Set the scope explicitly (since it could be different from the parent?)
    // This can't be done by the AstPostProcessing (unless we relax some constraints)
        functionDeclaration->set_scope(globalScope);

    // If it is not a forward declaration then the unparser will skip the ";" at the end (need to fix this better)
        functionDeclaration->setForward();
        ROSE_ASSERT(functionDeclaration->isForward() == true);

    // Mark function as extern "C"
        functionDeclaration->get_declarationModifier().get_storageModifier().setExtern();
        functionDeclaration->set_linkage("C");  // This mechanism could be improved!

        globalFunctionType = function_type;
        globalFunctionDeclaration = functionDeclaration;

    // Add function declaration to global scope!
        globalScope->prepend_declaration(globalFunctionDeclaration);

    // functionSymbol = new SgFunctionSymbol(globalFunctionDeclaration);
    // All any modifications to be fixed up (parents etc)
    // AstPostProcessing(project); // This is not allowed and should be fixed!
        AstPostProcessing(globalScope);

     }

#if 0
// DQ (12/1/2005): This version of the visit function handles the special case of
// instumentation at each function (a function call at the top of each function).
// At IBM we modified this to be a version which instrumented every block.

void
SimpleInstrumentation::visit ( SgNode* astNode )
   {
     SgFunctionDeclaration* functionDeclaration = isSgFunctionDeclaration(astNode);
     SgFunctionDefinition*  functionDefinition  = functionDeclaration != NULL ? functionDeclaration->get_definition() : NULL;
     if (functionDeclaration != NULL && functionDefinition != NULL)
        {
       // It is up to the user to link the implementations of these functions link time
          string functionName = functionDeclaration->get_name().str();
          string fileName     = functionDeclaration->get_file_info()->get_filename();

       // Build a source file location object (for construction of each IR node)
       // Note that we should not be sharing the same Sg_File_Info object in multiple IR nodes.
          Sg_File_Info * file_info = Sg_File_Info::generateDefaultFileInfoForTransformationNode();

          SgFunctionSymbol* functionSymbol        = new SgFunctionSymbol(globalFunctionDeclaration);
          SgFunctionRefExp* functionRefExpression = new SgFunctionRefExp(file_info,functionSymbol,globalFunctionType);
          SgExprListExp* expressionList           = new SgExprListExp(file_info);

          string converageFunctionInput = functionName + string ("_in_") + fileName;
          SgStringVal* functionNameStringValue = new SgStringVal(file_info,(char*)converageFunctionInput.c_str());
          expressionList->append_expression(functionNameStringValue);
          SgFunctionCallExp* functionCallExp    = new SgFunctionCallExp(file_info,functionRefExpression,expressionList,globalFunctionType);

       // create an expression type
          SgTypeVoid* expr_type = new SgTypeVoid();

       // create an expression root
       // SgExpressionRoot * expr_root = new SgExpressionRoot(file_info,functionCallExp,expr_type);

       // create an expression statement
          SgExprStatement* new_stmt = new SgExprStatement(file_info,functionCallExp);

       // expr_root->set_parent(new_stmt);
```

Figure 38.2: Example source code shows instrumentation to call a test function from the top of each function body in the application (part 2).

```
        // new_stmt->set_expression_root(expr_root);
2       // functionCallExp->set_parent(new_stmt->get_expression_root());

4           functionCallExp->set_parent(new_stmt);

6       // insert a statement into the function body
            functionDefinition->get_body()->prepend_statement(new_stmt);
8
#if 0
10      // This shows the alternative use of the ROSE Rewrite Mechanism to do the same thing!
        // However, the performance is not as good as the version written above (more directly
12      // building the IR nodes).

14      // string codeAtTopOfBlock = "void printf(char*); printf (\"FUNCTION_NAME in FILE_NAME \\n\");";
            string codeAtTopOfBlock = "coverageTraceFunc1(\"FUNCTION_NAME_in_FILE_NAME\");";
16
            string functionTarget   = "FUNCTION_NAME";
18          string fileTarget       = "FILE_NAME";

20          codeAtTopOfBlock.replace(codeAtTopOfBlock.find(functionTarget),functionTarget.size(),functionName);
            codeAtTopOfBlock.replace(codeAtTopOfBlock.find(fileTarget),fileTarget.size(),fileName);
22
        // printf ("codeAtTopOfBlock = %s \n",codeAtTopOfBlock.c_str());
24          printf ("%s_in_%s_\n",functionName.c_str(),fileName.c_str());

26      // Insert new code into the scope represented by the statement (applies to SgScopeStatements)
            MiddleLevelRewrite::ScopeIdentifierEnum scope = MidLevelCollectionTypedefs::StatementScope;
28
            SgBasicBlock* functionBody = functionDefinition->get_body();
30          ROSE_ASSERT(functionBody != NULL);

32      // Insert the new code at the top of the scope represented by block
            MiddleLevelRewrite::insert(functionBody,codeAtTopOfBlock,scope,MidLevelCollectionTypedefs::TopOfCurrentScope);
34  #endif
         }
36  }
#endif
38
void
40  SimpleInstrumentation::visit ( SgNode* astNode ) {
    SgBasicBlock *block = NULL;
42  block = isSgBasicBlock(astNode);
    if (block != NULL) {
44      // It is up to the user to link the implementations of these functions link time
        Sg_File_Info *fileInfo = block->get_file_info();
46      string fileName = fileInfo->get_filename();
        int lineNum = fileInfo->get_line();
48
        // Build a source file location object (for construction of each IR node)
50      // Note that we should not be sharing the same Sg_File_Info object in multiple IR nodes.
        Sg_File_Info * newCallfileInfo = Sg_File_Info::generateDefaultFileInfoForTransformationNode();
52
        ROSE_ASSERT(functionSymbol != NULL);
54      SgFunctionRefExp* functionRefExpression = new SgFunctionRefExp(newCallfileInfo,functionSymbol,globalFunctionType);
        SgExprListExp* expressionList           = new SgExprListExp(newCallfileInfo);
56
        string codeLocation = fileName + "_" + StringUtility::numberToString(lineNum);
58      SgStringVal* functionNameStringValue = new SgStringVal(newCallfileInfo,(char*)codeLocation.c_str());
        expressionList->append_expression(functionNameStringValue);
60      SgFunctionCallExp* functionCallExp    = new SgFunctionCallExp(newCallfileInfo,functionRefExpression,expressionList,globalFu

62    // create an expression type
      // SgTypeVoid* expr_type = new SgTypeVoid();
64    // create an expression root
      // SgExpressionRoot * expr_root = new SgExpressionRoot(newCallfileInfo,functionCallExp,expr_type);
66
      // create an expression statement
68      SgExprStatement* new_stmt = new SgExprStatement(newCallfileInfo,functionCallExp);

70    // expr_root->set_parent(new_stmt);
      // new_stmt->set_expression(expr_root);
72    // functionCallExp->set_parent(new_stmt->get_expression());
        functionCallExp->set_parent(new_stmt);
74
      // insert a statement into the function body
76      block->prepend_statement(new_stmt);

78  }
}
```

Figure 38.3: Example source code shows instrumentation to call a test function from the top of each function body in the application (part 3).

```
void foo()
   {
  // Should detect that foo IS called
  if (true) {
      int x = 3;
  }
  else {
      int x = 4;
   }
}

void foobar()
   {
      int y =4;
      switch (y) {
         case 1:
            //hello world
            break;
         case 2:

         case 3:

         default: {
            //
         }
      }

  // Should detect that foobar is NOT called
   }

int main()
   {
      if (true) {
      }
     foo();
     return 0;
   }
```

Figure 38.4: Example source code used as input to translator adding new function.

```
extern "C" void coverageTraceFunc1(char *textString);

void foo()
{
  coverageTraceFunc1("/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleCodeCo
// Should detect that foo IS called
  if (true) {
    coverageTraceFunc1("/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleCodeC
    int x = 3;
  }
  else {
    coverageTraceFunc1("/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleCodeC
    int x = 4;
  }
}

void foobar()
{
  coverageTraceFunc1("/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleCodeCo
  int y = 4;
  switch(y){
    coverageTraceFunc1("/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleCodeC
    case 1:
{
      coverageTraceFunc1("/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleCod
//hello world
      break;
    }
    default:
{
      coverageTraceFunc1("/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleCod
{
        coverageTraceFunc1("/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleC
//
      }
    }
  }
// Should detect that foobar is NOT called
}

int main()
{
  coverageTraceFunc1("/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleCodeCo
  if (true) {
    coverageTraceFunc1("/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleCodeC
  }
  foo();
  return 0;
}
```

Figure 38.5: Output of input to translator adding new function.

Chapter 39

Creating a 'struct' for Global Variables

This is an example written to support the Charm++ tool. This translator extracts global variables from the program and builds a structure to hold them. The support is part of a number of requirements associated with using Charm++ and AMPI.

Figure 39.1 shows repackaging of global variables within an application into a struct. All reference to the global variables are also transformed to reference the original variable indirectly through the structure. This processing is part of preprocessing to use Charm++.

This example shows the low level handling directly at the level of the IR.

```
// ROSE is a tool for building preprocessors, this file is an example preprocessor built with ROSE.
// Specifically it shows the design of a transformation to do a transformation specific for Charm++.

#include "rose.h"

using namespace std;

Rose_STL_Container<SgInitializedName*>
buildListOfGlobalVariables ( SgSourceFile* file )
   {
  // This function builds a list of global variables (from a SgFile).
     assert(file != NULL);

     Rose_STL_Container<SgInitializedName*> globalVariableList;

     SgGlobal* globalScope = file->get_globalScope();
     assert(globalScope != NULL);
     Rose_STL_Container<SgDeclarationStatement*>::iterator i = globalScope->get_declarations().begin();
     while(i != globalScope->get_declarations().end())
        {
          SgVariableDeclaration *variableDeclaration = isSgVariableDeclaration(*i);
          if (variableDeclaration != NULL)
             {
               Rose_STL_Container<SgInitializedName*> & variableList = variableDeclaration->get_variables();
               Rose_STL_Container<SgInitializedName*>::iterator var = variableList.begin();
               while(var != variableList.end())
                  {
                    globalVariableList.push_back(*var);
                    var++;
                  }
             }
          i++;
        }

     return globalVariableList;
   }

// This function is not used, but is useful for
// generating the list of all global variables
Rose_STL_Container<SgInitializedName*>
buildListOfGlobalVariables ( SgProject* project )
   {
  // This function builds a list of global variables (from a SgProject).

     Rose_STL_Container<SgInitializedName*> globalVariableList;

     const SgFilePtrList& fileList = project->get_fileList();
     SgFilePtrList::const_iterator file = fileList.begin();

  // Loop over the files in the project (multiple files exist
  // when multiple source files are placed on the command line).
     while(file != fileList.end())
        {
          Rose_STL_Container<SgInitializedName*> fileGlobalVariableList = buildListOfGlobalVariables(isSgSourceFile(*file));

        // DQ (9/26/2007): Moved from std::list to std::vector
        // globalVariableList.merge(fileGlobalVariableList);
          globalVariableList.insert(globalVariableList.begin(),fileGlobalVariableList.begin(),fileGlobalVariableList.end());

          file++;
        }

     return globalVariableList;
   }

Rose_STL_Container<SgVarRefExp*>
buildListOfVariableReferenceExpressionsUsingGlobalVariables ( SgNode* node )
   {
  // This function builds a list of "uses" of variables (SgVarRefExp IR nodes) within the AST.

  // return variable
     Rose_STL_Container<SgVarRefExp*> globalVariableUseList;

  // list of all variables (then select out the global variables by testing the scope)
     Rose_STL_Container<SgNode*> nodeList = NodeQuery::querySubTree ( node, V_SgVarRefExp );

     Rose_STL_Container<SgNode*>::iterator i = nodeList.begin();
     while(i != nodeList.end())
        {
          SgVarRefExp *variableReferenceExpression = isSgVarRefExp(*i);
```

Figure 39.1: Example source code shows repackaging of global variables to a struct (part 1).

```
            assert(variableReferenceExpression != NULL);

            assert(variableReferenceExpression->get_symbol() != NULL);
            assert(variableReferenceExpression->get_symbol()->get_declaration() != NULL);
            assert(variableReferenceExpression->get_symbol()->get_declaration()->get_scope() != NULL);

        // Note that variableReferenceExpression->get_symbol()->get_declaration() returns the
        // SgInitializedName (not the SgVariableDeclaration where it was declared)!
            SgInitializedName* variable = variableReferenceExpression->get_symbol()->get_declaration();

            SgScopeStatement* variableScope = variable->get_scope();

        // Check if this is a variable declared in global scope, if so, then save it
            if (isSgGlobal(variableScope) != NULL)
               {
                 globalVariableUseList.push_back(variableReferenceExpression);
               }
            i++;
          }

       return globalVariableUseList;
     }

SgClassDeclaration*
buildClassDeclarationAndDefinition (string name, SgScopeStatement* scope)
   {
  // This function builds a class declaration and definition
  // (both the defining and nondefining declarations as required).

  // Build a file info object marked as a transformation
     Sg_File_Info* fileInfo = Sg_File_Info::generateDefaultFileInfoForTransformationNode();
     assert(fileInfo != NULL);

  // This is the class definition (the fileInfo is the position of the opening brace)
     SgClassDefinition* classDefinition    = new SgClassDefinition(fileInfo);
     assert(classDefinition != NULL);

  // Set the end of construct explictly (where not a transformation this is the location of the closing brace)
     classDefinition->set_endOfConstruct(fileInfo);

  // This is the defining declaration for the class (with a reference to the class definition)
     SgClassDeclaration* classDeclaration = new SgClassDeclaration(fileInfo,name.c_str(),SgClassDeclaration::e_struct,NULL,classDefinition);
     assert(classDeclaration != NULL);

  // Set the defining declaration in the defining declaration!
     classDeclaration->set_definingDeclaration(classDeclaration);

  // Set the non defining declaration in the defining declaration (both are required)
     SgClassDeclaration* nondefiningClassDeclaration = new SgClassDeclaration(fileInfo,name.c_str(),SgClassDeclaration::e_struct,NULL,NULL);
     assert(classDeclaration != NULL);

  // Set the internal reference to the non-defining declaration
     classDeclaration->set_firstNondefiningDeclaration(nondefiningClassDeclaration);

  // Set the defining and no-defining declarations in the non-defining class declaration!
     nondefiningClassDeclaration->set_firstNondefiningDeclaration(nondefiningClassDeclaration);
     nondefiningClassDeclaration->set_definingDeclaration(classDeclaration);

  // Set the nondefining declaration as a forward declaration!
     nondefiningClassDeclaration->setForward();

  // Don't forget the set the declaration in the definition (IR node constructors are side-effect free!)!
     classDefinition->set_declaration(classDeclaration);

  // set the scope explicitly (name qualification tricks can imply it is not always the parent IR node!)
     classDeclaration->set_scope(scope);
     nondefiningClassDeclaration->set_scope(scope);

  // some error checking
     assert(classDeclaration->get_definingDeclaration() != NULL);
     assert(classDeclaration->get_firstNondefiningDeclaration() != NULL);
     assert(classDeclaration->get_definition() != NULL);

  // DQ (9/8/2007): Need to add function symbol to global scope!
     printf ("Fixing up the symbol table in scope = %p = %s for class = %p = %s \n",scope,scope->class_name().c_str(),classDeclaration,classDe
     SgClassSymbol* classSymbol = new SgClassSymbol(classDeclaration);
     scope->insert_symbol(classDeclaration->get_name(),classSymbol);
     ROSE_ASSERT(scope->lookup_class_symbol(classDeclaration->get_name()) != NULL);
```

Figure 39.2: Example source code shows repackaging of global variables to a struct (part 2).

```
     return classDeclaration;
   }


SgVariableSymbol*
putGlobalVariablesIntoClass (Rose_STL_Container<SgInitializedName*> & globalVariables, SgClassDeclaration* classDeclaration )
   {
  // This function iterates over the list of global variables and inserts them into the iput class definition

     SgVariableSymbol* globalClassVariableSymbol = NULL;

     for (Rose_STL_Container<SgInitializedName*>::iterator var = globalVariables.begin(); var != globalVariables.end(); var++)
        {
       // printf ("Appending global variable = %s to new globalVariableContainer \n",(*var)->get_name().str());
          SgVariableDeclaration* globalVariableDeclaration = isSgVariableDeclaration((*var)->get_parent());
          assert(globalVariableDeclaration != NULL);

       // Get the global scope from the global variable directly
          SgGlobal* globalScope = isSgGlobal(globalVariableDeclaration->get_scope());
          assert(globalScope != NULL);

          if (var == globalVariables.begin())
             {
            // This is the first time in this loop, replace the first global variable with
            // the class declaration/definition containing all the global variables!
            // Note that initializers in the global variable declarations require modification
            // of the preinitialization list in the class's constructor!  I am ignoring this for now!
               globalScope->replace_statement(globalVariableDeclaration,classDeclaration);

            // Build source position informaiton (marked as transformation)
               Sg_File_Info* fileInfo = Sg_File_Info::generateDefaultFileInfoForTransformationNode();
               assert(fileInfo != NULL);

            // Add the variable of the class type to the global scope!
               SgClassType* variableType = new SgClassType(classDeclaration->get_firstNondefiningDeclaration());
               assert(variableType != NULL);
               SgVariableDeclaration* variableDeclaration = new SgVariableDeclaration(fileInfo,"AMPI_globals",variableType);
               assert(variableDeclaration != NULL);

               globalScope->insert_statement(classDeclaration,variableDeclaration,false);

               assert(variableDeclaration->get_variables().empty() == false);
               SgInitializedName* variableName = *(variableDeclaration->get_variables().begin());
               assert(variableName != NULL);

            // DQ (9/8/2007): Need to set the scope of the new variable.
               variableName->set_scope(globalScope);

            // build the return value
               globalClassVariableSymbol = new SgVariableSymbol(variableName);

            // DQ (9/8/2007): Need to add the symbol to the global scope (new testing requires this).
               globalScope->insert_symbol(variableName->get_name(),globalClassVariableSymbol);
               ROSE_ASSERT(globalScope->lookup_variable_symbol(variableName->get_name()) != NULL);
             }
            else
             {
            // for all other iterations of this loop ...
            // remove variable declaration from the global scope
               globalScope->remove_statement(globalVariableDeclaration);
             }

       // add the variable declaration to the class definition
          classDeclaration->get_definition()->append_member(globalVariableDeclaration);
         }

     return globalClassVariableSymbol;
   }

void
fixupReferencesToGlobalVariables ( Rose_STL_Container<SgVarRefExp*> & variableReferenceList, SgVariableSymbol* globalClassVariable
   {
  // Now fixup the SgVarRefExp to reference the global variables through a struct
     for (Rose_STL_Container<SgVarRefExp*>::iterator var = variableReferenceList.begin(); var != variableReferenceList.end(); var
        {
          assert(*var != NULL);
       // printf ("Variable reference for %s \n",(*var)->get_symbol()->get_declaration()->get_name().str());
```

Figure 39.3: Example source code shows repackaging of global variables to a struct (part 3).

```
            SgNode* parent = (*var)->get_parent();
2           assert(parent != NULL);

4         // If this is not an expression then is likely a meaningless statement such as ("x;")
            SgExpression* parentExpression = isSgExpression(parent);
6           assert(parentExpression != NULL);

8         // Build the reference through the global class variable ("x" --> "AMPI_globals.x")

10        // Build source position informaiton (marked as transformation)
            Sg_File_Info* fileInfo = Sg_File_Info::generateDefaultFileInfoForTransformationNode();
12          assert(fileInfo != NULL);

14        // Build "AMPI_globals"
            SgExpression* lhs = new SgVarRefExp(fileInfo,globalClassVariableSymbol);
16          assert(lhs != NULL);
          // Build "AMPI_globals.x" from "x"
18          SgDotExp* globalVariableReference = new SgDotExp(fileInfo,lhs,*var);
            assert(globalVariableReference != NULL);
20
            if (parentExpression != NULL)
22              {
                // Introduce reference to *var through the data structure
24
                // case of binary operator
26                SgUnaryOp* unaryOperator = isSgUnaryOp(parentExpression);
                  if (unaryOperator != NULL)
28                    {
                      unaryOperator->set_operand(globalVariableReference);
30                    }
                    else
32                    {
                    // case of binary operator
34                    SgBinaryOp* binaryOperator = isSgBinaryOp(parentExpression);
                      if (binaryOperator != NULL)
36                        {
                        // figure out if the *var is on the lhs or the rhs
38                        if (binaryOperator->get_lhs_operand() == *var)
                            {
40                            binaryOperator->set_lhs_operand(globalVariableReference);
                            }
42                          else
                            {
44                            assert(binaryOperator->get_rhs_operand() == *var);
                              binaryOperator->set_rhs_operand(globalVariableReference);
46                            }
                        }
48                      else
                        {
50                      // ignore these cases for now!
                        switch(parentExpression->variantT())
52                          {
                          // Where the variable appers in the function argument list the parent is a SgExprListExp
54                          case V_SgExprListExp:
                              {
56                              printf ("Sorry not implemented, case of global variable in function argument list ...\n");
                                assert(false);
58                              break;
                              }
60                          case V_SgInitializer:
                            case V_SgRefExp:
62                          case V_SgVarArgOp:
                            default:
64                              {
                                printf ("Error: default reached in switch  parentExpression = %p = %s \n",parentExpression,parentExpression
66                              assert(false);
                              }
68                          }
                        }
70                  }
                }
72          }
        }
74
    #define OUTPUT_NAMES_OF_GLOBAL_VARIABLES            0
76  #define OUTPUT_NAMES_OF_GLOBAL_VARIABLE_REFERENCES  0

78  void
    transformGlobalVariablesToUseStruct ( SgSourceFile *file )
80      {
```

Figure 39.4: Example source code shows repackaging of global variables to a struct (part 4).

```
     assert(file != NULL);

  // These are the global variables in the input program (provided as helpful information)
     Rose_STL_Container<SgInitializedName*> globalVariables = buildListOfGlobalVariables(file);

#if OUTPUT_NAMES_OF_GLOBAL_VARIABLES
     printf ("global variables (declared in global scope): \n");
     for (Rose_STL_Container<SgInitializedName*>::iterator var = globalVariables.begin(); var != globalVariables.end(); var++)
        {
          printf ("   %s \n",(*var)->get_name().str());
        }
     printf ("\n");
#endif

  // get the global scope within the first file (currently ignoring all other files)
     SgGlobal* globalScope = file->get_globalScope();
     assert(globalScope != NULL);

  // Build the class declaration
     SgClassDeclaration* classDeclaration = buildClassDeclarationAndDefinition("AMPI_globals_t",globalScope);

  // Put the global variables intothe class
     SgVariableSymbol* globalClassVariableSymbol = putGlobalVariablesIntoClass(globalVariables,classDeclaration);

  // Their associated symbols will be located within the project's AST
  // (where they occur in variable reference expressions).
     Rose_STL_Container<SgVarRefExp*> variableReferenceList = buildListOfVariableReferenceExpressionsUsingGlobalVariables(file);

#if OUTPUT_NAMES_OF_GLOBAL_VARIABLE_REFERENCES
     printf ("global variables appearing in the application: \n");
     for (Rose_STL_Container<SgVarRefExp*>::iterator var = variableReferenceList.begin(); var != variableReferenceList.end(); var
        {
          printf ("   %s \n",(*var)->get_symbol()->get_declaration()->get_name().str());
        }
     printf ("\n");
#endif

  // Fixup all references to global variable to access the variable through the class ("x" --> "AMPI_globals.x")
     fixupReferencesToGlobalVariables(variableReferenceList,globalClassVariableSymbol);
   }


void
transformGlobalVariablesToUseStruct ( SgProject *project )
   {
  // Call the transformation of each file (there are multiple SgFile
  // objects when multiple files are specfied on the command line!).
     assert(project != NULL);

     const SgFilePtrList& fileList = project->get_fileList();
     SgFilePtrList::const_iterator file = fileList.begin();
     while(file != fileList.end())
        {
          transformGlobalVariablesToUseStruct(isSgSourceFile(*file));
          file++;
        }
   }

// *****************************************
//              MAIN PROGRAM
// *****************************************
int
main( int argc, char * argv[] )
   {
  // Build the AST used by ROSE
     SgProject* project = frontend(argc,argv);
     assert(project != NULL);

  // transform application as required
     transformGlobalVariablesToUseStruct(project);

  // Code generation phase (write out new application "rose_<input file name>")
     return backend(project);
   }
```

Figure 39.5: Example source code shows repackaging of global variables to a struct (part 5).

```
int x;
int y;
long z;
float pressure;

int main()
   {
     int a = 0;
     int b = 0;
     float density = 1.0;

     x++;
     b++;

     x = a + y;

     return 0;
   }
```

Figure 39.6: Example source code used as input to translator adding new function.

```
struct AMPI_globals_t
{
  int x;
  int y;
  long z;
  float pressure;
}

;
struct AMPI_globals_t AMPI_globals;

int main()
{
  int a = 0;
  int b = 0;
  float density = (1.0);
  AMPI_globals.x++;
  b++;
  AMPI_globals.x = (a + AMPI_globals.y);
  return 0;
}
```

Figure 39.7: Output of input to translator adding new function.

Chapter 40

TAU Instrumentation

Tau is a performance analysis tool from University of Oregon. They have mechanisms for automating instrumentation of the source code's text file directly, but these can be problematic in the present of macros. We present an example of instrumentation combined with code generation to provide a more robust means of instrumentation for source code. This work is preliminary and depends upon two separate mechanisms for the rewrite of the AST (one high level using strings and one low level representing a more direct handling of the AST at the IR level).

40.1 Input For Examples Showing Information using Tau

Figure 40.1 shows the example input used for demonstration of Tau instrumentation.

40.2 Generating the code representing any IR node

Figure 40.2 shows a code that traverses each IR node and for a SgInitializedName of SgStatement output the scope information. The input code is shown in figure 40.2, the output of this code is shown in figure 40.3.

```
// #include <math.h>
// #include <stdlib.h>

double foo(double x)
   {
     double theValue = x;
     theValue*= x;
     return theValue;
   }

int main(int argc, char* argv[])
   {
     int j,i;
     double tSquared,t;
     t = 1.0;

     tSquared = t*t;

     i = 1000;
     for( j=1; j < i; j += 2)
        {
          tSquared += 1.0;
          tSquared += foo(2.2);
        }

     return 0;
   }
```

Figure 40.1: Example source code used as input to program in codes used in this chapter.

```
// Demonstration of instrumentation using the TAU performance monitoring tools (University of Oregon)

#include "rose.h"

using namespace std;

#define NEW_FILE_INFO Sg_File_Info::generateDefaultFileInfoForTransformationNode()

SgClassDeclaration*
getProfilerClassDeclaration (SgProject* project)
   {
  // Note that it would be more elegant to look this up in the Symbol table (do this next)

     SgClassDeclaration* returnClassDeclaration = NULL;
     Rose_STL_Container<SgNode*> classDeclarationList = NodeQuery::querySubTree (project,V_SgClassDeclaration);
     for (Rose_STL_Container<SgNode*>::iterator i = classDeclarationList.begin(); i != classDeclarationList.end(); i++)
        {
       // Need to cast *i from SgNode to at least a SgStatement
          SgClassDeclaration* classDeclaration = isSgClassDeclaration(*i);
          ROSE_ASSERT (classDeclaration != NULL);

       // printf ("In getProfilerClassDeclaration(): classDeclaration->get_name() = %s \n",classDeclaration->get_name().s

          if (classDeclaration->get_name() == "Profiler")
               returnClassDeclaration = classDeclaration;
        }

     ROSE_ASSERT(returnClassDeclaration != NULL);
     return returnClassDeclaration;
   }


int
main( int argc, char * argv[] )
   {
  // This test code tests the AST rewrite mechanism to add TAU Instrumention to the AST.

     SgProject* project = frontend(argc,argv);

  // Output the source code file (as represented by the SAGE AST) as a PDF file (with bookmarks)
  // generatePDF(*project);

  // Output the source code file (as represented by the SAGE AST) as a DOT file (graph)
  // generateDOT(*project);

  // Allow ROSE translator options to influence if we transform the AST
     if ( project->get_skip_transformation() == false )
        {
       // NOTE: There can be multiple files on the command line and each file has a global scope

          MiddleLevelRewrite::ScopeIdentifierEnum    scope           = MidLevelCollectionTypedefs::StatementScope;
          MiddleLevelRewrite::PlacementPositionEnum locationInScope = MidLevelCollectionTypedefs::TopOfCurrentScope;

       // Add a TAU include directive to the top of the global scope
          Rose_STL_Container<SgNode*> globalScopeList = NodeQuery::querySubTree (project,V_SgGlobal);
          for (Rose_STL_Container<SgNode*>::iterator i = globalScopeList.begin(); i != globalScopeList.end(); i++)
             {
            // Need to cast *i from SgNode to at least a SgStatement
               SgGlobal* globalScope = isSgGlobal(*i);
               ROSE_ASSERT (globalScope != NULL);

            // TAU does not seem to compile using EDG or g++ (need to sort this out with Brian)
            // MiddleLevelRewrite::insert(globalScope,"#define PROFILING_ON \n#include<TAU.h> \n",scope,locationInScope);
            // MiddleLevelRewrite::insert(globalScope,"#include<TAU.h> \n",scope,locationInScope);
               MiddleLevelRewrite::insert(globalScope,"#define PROFILING_ON 1 \n#define TAU_STDCXXLIB 1 \n#include<TAU.h>
             }
#if 1
       // Now get the class declaration representing the TAU type with which to build variable declarations
          SgClassDeclaration* tauClassDeclaration = getProfilerClassDeclaration(project);
          ROSE_ASSERT(tauClassDeclaration != NULL);
          SgClassType* tauType = tauClassDeclaration->get_type();
          ROSE_ASSERT(tauType != NULL);

       // Get a constructor to use with the variable declaration (anyone will due for code generation)
          SgMemberFunctionDeclaration* memberFunctionDeclaration = SageInterface::getDefaultConstructor(tauClassDeclaratio
          ROSE_ASSERT(memberFunctionDeclaration != NULL);

       // Add the instrumentation to each function
          Rose_STL_Container<SgNode*> functionDeclarationList = NodeQuery::querySubTree (project,V_SgFunctionDeclaration)
          for (Rose_STL_Container<SgNode*>::iterator i = functionDeclarationList.begin(); i != functionDeclarationList.en
             {
               SgFunctionDeclaration* functionDeclaration = isSgFunctionDeclaration(*i);
               ROSE_ASSERT(functionDeclaration != NULL);
```

```
Test Failed!
```

Figure 40.3: Output of input code using tauInstrumenter.C

Chapter 41

ROSE-HPCToolkit Interface

The HPCToolkit (`http://www.hipersoft.rice.edu/hpctoolkit`) is a set of tools for analyzing the dynamic performance behavior of applications. It includes a tool that instruments a program's binary, in order to observe CPU hardware counters during execution; additional post-processing tools attribute the observed data to statements in the original source code. HPCToolkit stores this data and the source attributions in XML files. In this chapter, we give an overview of simple interfaces in ROSE that can read this data and attach it to the AST.

ROSE-HPCT is included in the ROSE distribution as an optional module. To enable it, specify the `--enable-rosehpct` option when running configure, and be sure that you have an existing installation of the Gnome XML library, libxml2 (`http://xmlsoft.org`).

41.1 An HPCToolkit Example Run

Consider the sample source program shown in Figures 41.1–41.2. This program takes an integer n on the command line, and has a number of loops whose flop and memory-operation complexity are either $\Theta(n)$ or $\Theta(n^2)$. For this example, we would expect the loop nest at line 56, which has $O(n^2)$ cost, to be the most expensive loop in the program for large n.

Suppose we use the HPCToolkit to profile this program, collecting cycle counts and floating-point instructions.[1] HPCToolkit will generate one XML file for each metric.

A schema specifying the format of these XML files appears in Figure 41.3. In essence, this schema specifies that the XML file will contain a structured, abstract representation of the program in terms of abstract program entities such as "modules," "procedures," "loops," and "statements." Each of these entities may have line number information and a "metric" value. (Refer to the HPCToolkit documentation for more information.) This schema is always the first part of an HPCToolkit-generated XML profile data file.

We ran HPCToolkit on the program in Figures 41.1–41.2, and collected cycle and flop counter data. The actual XML files storing this data appear in Figures 41.4 and 41.5. By convention, these metrics are named according to their PAPI symbolic name, as shown on line 67 in both Figures 41.4 and 41.5. According to the cycle data on line 90 of Figure 41.4, the most expensive statement in profiled.c is line 62 of Figure 41.1, as expected.

[1] In this example, HPCToolkit uses the PAPI to read CPU counters (`http://icl.cs.utk.edu/papi`).

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef double val_t;

static val_t *
gen (size_t n)
{
  val_t* x = (val_t *) malloc (sizeof(val_t) * n);
  if (x != NULL)
    {
      size_t i;
      for (i = 0; i < n; i++)
          x[i] = (val_t)1.0 / n;
    }
  return x;
}

static val_t
dot (size_t n, const val_t* x, const val_t* y)
{
  size_t i;
  val_t sum;
  for (i = 0, sum = 0.0; i < n; i++)
    sum += x[i] * y[i];
  return sum;
}

static size_t
max (size_t n, const val_t* x)
{
  size_t i;
  size_t i_max;
  val_t v_max;

  if (n <= 0)
    return 0;

  v_max = x[0];
  i_max = 0;
  for (i = 1; i < n; i++)
    if (x[i] > v_max)
      {
        v_max = x[i];
        i_max = i;
      }
  return i_max;
}

static void
mv (size_t n, const val_t* A, const val_t* x, val_t* y)
{
  size_t j;
  memset (y, 0, sizeof(val_t) * n);
  for (j = 0; j < n; j++)
    {
      const val_t* Ap;
      register val_t xj = x[j];
      size_t i;
      for (i = 0, Ap = A + j*n; i < n; i++, Ap++)
        y[i] += Ap[0] * xj;
    }
}
```

Figure 41.1: profiled.c (part 1 of 2): Sample input program, profiled using the HPCToolkit.

```
int
main (int argc, char* argv[])
{
  size_t n;
  val_t* x;
  val_t* y;
  val_t* A;
  size_t i;

  /* outputs */
  val_t sum;
  size_t i_max;
  val_t y_max;

  if (argc != 2)
    {
      fprintf (stderr, "usage: %s <n>\n", argv[0]);
      return 1;
    }
  n = atoi (argv[1]);
  if (n <= 0)
    return 1;

  A = gen (n * n);
  x = gen (n);
  y = gen (n);

  if (A == NULL || x == NULL || y == NULL)
    {
      fprintf (stderr, "*** Out of memory ***\n");
      return 1;
    }

  sum = 0;
  for (i = 0; i < n; i++)
    sum += dot (n, x, y);
  mv (n, A, x, y);
  i_max = max (n, y);
  y_max = y[i_max];

  printf ("%g %lu %g\n", sum, (unsigned long)i_max, y_max);

  return 0;
}

/* eof */
```

Figure 41.2: profiled.c (part 2 of 2): Sample input program, profiled using the HPCToolkit.

```
<?xml version="1.0"?>
<!DOCTYPE PROFILE [
<!-- Profile: correlates profiling info with program structure. -->
<!ELEMENT PROFILE (PROFILEHDR, PROFILEPARAMS, PROFILESCOPETREE)>
<!ATTLIST PROFILE
        version CDATA #REQUIRED>
  <!ELEMENT PROFILEHDR (#PCDATA)>
  <!ELEMENT PROFILEPARAMS (TARGET, METRICS)>
    <!ELEMENT TARGET EMPTY>
    <!ATTLIST TARGET
        name CDATA #REQUIRED>
    <!ELEMENT METRICS (METRIC)+>
    <!ELEMENT METRIC EMPTY>
    <!ATTLIST METRIC
        shortName   CDATA #REQUIRED
        nativeName  CDATA #REQUIRED
        period      CDATA #REQUIRED
        units       CDATA #IMPLIED
        displayName CDATA #IMPLIED
        display     (true|false) #IMPLIED>
  <!ELEMENT PROFILESCOPETREE (PGM)*>
    <!-- This is essentially the PGM dtd with M element added. -->
    <!ELEMENT PGM (G|LM|F|M)+>
    <!ATTLIST PGM
        n CDATA #REQUIRED>
    <!-- Groups create arbitrary sets of other elements except PGM. -->
    <!ELEMENT G (G|LM|F|P|L|S|M)*>
    <!ATTLIST G
        n CDATA #IMPLIED>
    <!-- Runtime load modules for PGM (e.g., DSOs, exe) -->
    <!ELEMENT LM (G|F|M)*>
    <!ATTLIST LM
        n CDATA #REQUIRED>
    <!-- Files contain procedures and source line info -->
    <!ELEMENT F (G|P|L|S|M)*>
    <!ATTLIST F
        n CDATA #REQUIRED>
    <!-- Procedures contain source line info
         n: processed name; ln: link name -->
    <!ELEMENT P (G|L|S|M)*>
    <!ATTLIST P
        n  CDATA #REQUIRED
        ln CDATA #IMPLIED
        b CDATA #IMPLIED
        e CDATA #IMPLIED>
    <!-- Loops -->
    <!ELEMENT L (G|L|S|M)*>
    <!ATTLIST L
        b CDATA #IMPLIED
        e CDATA #IMPLIED>
    <!-- Statement/Statement range -->
    <!ELEMENT S (M)*>
    <!ATTLIST S
        b CDATA #REQUIRED
        e CDATA #IMPLIED
        id CDATA #IMPLIED>
    <!ELEMENT M EMPTY>
    <!ATTLIST M
        n CDATA #REQUIRED
        v CDATA #REQUIRED>
]>
```

Figure 41.3: XML schema for HPCToolkit data files: This schema, prepended to each of the HPCToolkit-generated XML files, describes the format of the profiling data. This particular schema was generated by HPCToolkit 1.0.4.

```
<PROFILE version="3.0">
<PROFILEHDR></PROFILEHDR>
<PROFILEPARAMS>
  <TARGET name="example"/>
  <METRICS>
    <METRIC shortName="0" nativeName="PAPI_TOT_CYC" displayName="PAPI_TOT_CYC" period="32767" units="PAPI_events"/>
  </METRICS>
</PROFILEPARAMS>
<PROFILESCOPETREE>
<PGM n="example">
  <LM n="/home/vuduc2/projects/ROSE/tmp/xml/xerces-c/hpctif/examples/data/01/example">
    <F n="./profiled.c">
      <P n="main">
        <S b="15">
          <M n="0" v="163835"/>
        </S>
        <S b="25">
          <M n="0" v="65534"/>
        </S>
        <S b="26">
          <M n="0" v="65534"/>
        </S>
      </P>
      <P n="mv">
        <S b="61">
          <M n="0" v="65534"/>
        </S>
        <S b="62">
          <M n="0" v="327670"/>
        </S>
      </P>
    </F>
  </LM>
</PGM>
</PROFILESCOPETREE>
</PROFILE>
```

Figure 41.4: PAPI_TOT_CYC.xml: Sample cycle counts observed during profiling, generated from running the HPCToolkit on profiled.c (Figures 41.1–41.2.) These lines would appear after the schema shown in Figure 41.3.

```
<PROFILE version="3.0">
<PROFILEHDR></PROFILEHDR>
<PROFILEPARAMS>
  <TARGET name="example"/>
  <METRICS>
    <METRIC shortName="0" nativeName="PAPI_FP_OPS" displayName="PAPI_FP_OPS" period="32767" units="PAPI_events
  </METRICS>
</PROFILEPARAMS>
<PROFILESCOPETREE>
<PGM n="example">
  <LM n="/home/vuduc2/projects/ROSE/tmp/xml/xerces-c/hpctif/examples/data/01/example">
    <F n="./profiled.c">
      <P n="main">
        <S b="25">
          <M n="0" v="32767"/>
        </S>
        <S b="26">
          <M n="0" v="98301"/>
        </S>
      </P>
      <P n="mv">
        <S b="61">
          <M n="0" v="131068"/>
        </S>
        <S b="62">
          <M n="0" v="262136"/>
        </S>
      </P>
    </F>
  </LM>
</PGM>
</PROFILESCOPETREE>
</PROFILE>
```

Figure 41.5: PAPI_FP_OPS.xml: Sample flop counts observed during profiling, generated from running the HPCToolkit on profiled.c (Figures 41.1–41.2.) These lines would appear after the schema shown in Figure 41.3.

41.2 Attaching HPCToolkit Data to the ROSE AST

To attach the data of Figures 41.4 and 41.5 to the AST, we augment a basic ROSE translator with two additional calls, as shown in Figure 41.6, lines 47–48 and 54. We describe these calls below.

41.2.1 Calling ROSE-HPCT

We must first include rosehpct/rosehpct.hh, as shown on line 6 of Figure 41.6. All ROSE-HPCT routines and intermediate data structures reside in the RoseHPCT namespace.

Next, lines 47–48 of Figure 41.6 store the contents of the raw XML file into an intermediate data structure of type RoseHPCT::ProgramTreeList_t. The RoseHPCT::loadHPCTProfiles() routine processes command-line arguments, extracting ROSE-HPCT-specific options that specify the files. We discuss these options in Section 41.2.4.

Line 54 of Figure 41.6, attaches the intermediate profile data structure to the ROSE AST. The RoseHPCT::attachMetrics() routine creates new persistent attributes that store the counter data.[2] The attributes are named using the metric name taken from the XML file (see lines 67 of Figures 41.4–41.5); in this example, the attributes are named PAPI_TOT_CYC and PAPI_FP_OPS. Following the conventions of persistent attribute mechanism as described in Chapter 8, the attributes themselves are of type RoseHPCT::MetricAttr, which derives from the AstAttribute type.

41.2.2 Retrieving the attribute values

We retrieve the attribute values as described in Chapter 8. In particular, given a located node with cycle and flop attribute values, the printFlopRate() routine defined in lines 11–42 of Figure 41.6 prints the source position, AST node type, and estimated flops per cycle. We call printFlopRate() for each expression statement (SgExpressionStmt), for-initializer (SgForInitStatement), and for-statement (SgForStatement) in lines 59–66 of Figure 41.6. The output is shown in Figure 41.7.

Inspecting the output carefully, you may notice seeming discrepancies between the values shown and the values that appear in the XML files, or other values which seem unintuitive. We explain how these values are derived in Section 41.2.3.

This example dumps the AST as a PDF file, as shown on line 68 of Figure 41.6. You can inspect this file to confirm where attributes have been attached. We show an example of such a page in Figure 41.8. This page is the SgExprStatement node representing the sum-accumulate on line 26 of Figure 41.1.

41.2.3 Metric propagation

The example program in Figure 41.6 dumps metric values at each expression statement, for-initializer, and for-statement, but the input XML files in Figure 41.4–41.5 only attribute the profile data to "statements" that are not loop constructs. (The `<S ...>` XML tags refer to statements, intended to be "simple" non-scoping executable statements; a separate `<L ...>` tag

[2] The last parameter to RoseHPCT::attachMetrics() is a boolean that, when true, enables verbose (debugging) messages to standard error.

would refer to a loop.) Since the XML file specifies statements only by source line number, RoseHPCT::attachMetrics() attributes measurements to AST nodes in a heuristic way.

For example, lines 78–80 of Figure 41.4 indicate that all executions of the "simple statements" of line 25 of the original source (Figure 41.1) accounted for 65534 observed cycles, and that line 26 accounted for an additional 65534 cycles. In the AST, there are multiple "statement" and expression nodes that occur on line 25; indeed, Figure 41.7 lists 4 such statements.

The ROSE-HPCT modules uses a heuristic which only assigns `<S ...>` metric values to non-scoping nodes derived from SgStatement. When multiple SgStatement nodes occur at a particular source line, ROSE-HPCT simply divides the observed metric equally among all the SgStatement nodes on that line.

How is the measurement of 65534 cycles attributed to all of the AST nodes corresponding to line 25 of Figure 41.1? Indeed, line 25 actually "contains" four different SgStatement nodes: an SgForStatement representing the whole loop on lines 25–26, an SgForInitStatement (initializer), and two SgExprStatements (one which is a child of the SgForInitStatement, and another for the for-loop's test expression). The loop's increment is stored in the SgForStatement node as an SgExpression, not an SgStatement. The SgForStatement node is a scoping statement, and so it "receives" none of the 65534 cycles. Since the increment is not a statement and one of the SgExprStatements is a child of the initializer, this leaves only two direct descendants of the SgForStatement—the initializer and the test expression statement—among which to divide the 65534 cycles. Thus, each receives 32767 cycles. The initializer's SgExprStatement child gets the same 32767 as its parent, since the two nodes are equivalent (see first two cases of Figure 41.7).

For the entire loop on lines 25–26 of Figure 41.1, the original XML files attribute 65534 cycles to line 25, and another 65534 cycles to line 26 (see Figure 41.4). Moreover, the XML files do not attribute any costs to this loop *via* an explicit `<L ...>` tag. Thus, the best we can infer is that the entire for-statement's costs is the sum of its immediate child costs; in this case, 131068 cycles. The RoseHPCT::attachMetrics() routine will heuristically accumulate and propagate metrics in this way to assign higher-level scopes approximate costs.

The RoseHPCT::attachMetrics() routine automatically propagates metric values through parent scopes. A given metric attribute, RoseHPCT::MetricAttr* x, is "derived" through propagation if x->isDerived() returns true. In fact, if you call x->toString() to obtain a string representation of the metric's value, two asterisks will be appended to the string as a visual indicator that the metric is derived. We called RoseHPCT::MetricAttr::toString() on lines 27 and 29 of Figure 41.6, and all of the SgForStatement nodes appearing in the output in Figure 41.7 are marked as derived.

Alternatively, you cann call RoseHPCT::attachMetricsRaw(), rather than calling RoseHPCT::attachMetrics(). The "raw" routine takes the same arguments but only attaches the raw data, *i.e.*, without attempting to propagate metric values through parent scopes.

41.2.4 ROSE-HPCT command-line options

The call to RoseHPCT::loadHPCTProfiles() on line 48 of Figure 41.6 processes and extracts ROSE-HPCT-specific command-line options. To generate the output in this chapter, we invoked Figure 41.6 with the following command-line:

```
./attachMetrics \
  -rose:hpctprof ../../../sourcetree/tutorial/roseHPCT/PAPI_TOT_CYC.xml \
```

```
-rose:hpctprof ../../../sourcetree/tutorial/roseHPCT/PAPI_FP_OPS.xml \
-rose:hpcteqpath .=/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/roseHPCT\
-c ../../../sourcetree/tutorial/roseHPCT/profiled.c
```

The main option is `-rose:hpct:prof <file>`, which specifies the HPCToolkit-generated XML file containing metric data. Here, we use this option twice to specify the names of the cycle and flop data files (Figures 41.4–41.5).

We need the other option, `-rose:hpct:eqpath <A>=<B>`, to specify how hard-coded paths in the HPCToolkit XML files map to file paths in the ROSE AST. This option is necessary because analyzing the HPCToolkit-generated data need not occur on the same machine or environment in which the data was collected. In this example, the XML files specify the source file as, "./profiled.c" (line 73 of Figures 41.4 and 41.5); the "eqpath" command-line option above remaps the relative path "." to an absolute path as it would appear in the ROSE AST.

```
// attachMetrics.cc -- Sample translator showing how to attach
// HPCToolkit data to the ROSE AST.

#include <iostream>
#include <list>
#include <rosehpct/rosehpct.hh>

using namespace std;

// Prints the estimated flops/cycle at a located node.
static void printFlopRate (const SgNode* n)
{
  const SgLocatedNode* n_loc = isSgLocatedNode (n);
  if (n_loc
      && n_loc->attributeExists ("PAPI_TOT_CYC")
      && n_loc->attributeExists ("PAPI_FP_OPS"))
    {
      // Extract attributes.
      const RoseHPCT::MetricAttr* cycles_attr =
        dynamic_cast<const RoseHPCT::MetricAttr *> (n->getAttribute ("PAPI_TOT_CYC"));
      const RoseHPCT::MetricAttr* flops_attr =
        dynamic_cast<const RoseHPCT::MetricAttr *> (n->getAttribute ("PAPI_FP_OPS"));
      ROSE_ASSERT (cycles_attr && flops_attr);

      // Get the metric values, as doubles and strings.
      double cycles = cycles_attr->getValue ();
      string cycles_s = const_cast<RoseHPCT::MetricAttr *> (cycles_attr)->toString ();
      double flops = flops_attr->getValue ();
      string flops_s = const_cast<RoseHPCT::MetricAttr *> (flops_attr)->toString ();

      // Print node pointer/type, parent, estimated flop rate, and source position.
      const SgNode* n_par = n_loc->get_parent ();
      cout << (const void *)n_loc << ":<" << n_loc->class_name () << ">"
           << " (Par=" << (const void *)n_par << ":<" << n_par->class_name () << ">)"
           << " = (" << flops_s << " flops)"
           << " / (" << cycles_s << " cy)"
           << " ~= " << flops / cycles << " flops/cy" << endl
           << "  [" << n_loc->get_startOfConstruct ()->get_raw_filename ()
           << ':' << n_loc->get_startOfConstruct ()->get_raw_line () << "]"
           << endl;
    }
}

int main (int argc, char* argv[])
{
  vector<string> argvList(argv, argv+argc);
  cerr << "[Loading HPCToolkit profiling data...]" << endl;
  RoseHPCT::ProgramTreeList_t profiles
    = RoseHPCT::loadHPCTProfiles (argvList);

  cerr << "[Building the ROSE AST...]" << endl;
  SgProject* proj = frontend (argvList);

  cerr << "[Attaching HPCToolkit metrics to the AST...]" << endl;
  RoseHPCT::attachMetrics (profiles, proj, proj->get_verbose () > 0);

  cerr << "[Estimating flop execution rates...]" << endl;
  //typedef list<SgNode *> NodeList_t;
  //Liao (10/1/2007): Commented out as part of move from std::list to std::vector
  typedef Rose_STL_Container<SgNode *> NodeList_t;

  NodeList_t estmts = NodeQuery::querySubTree (proj, V_SgExprStatement);
  for_each (estmts.begin (), estmts.end (), printFlopRate);

  NodeList_t for_inits = NodeQuery::querySubTree (proj, V_SgForInitStatement);
  for_each (for_inits.begin (), for_inits.end (), printFlopRate);

  NodeList_t fors = NodeQuery::querySubTree (proj, V_SgForStatement);
  for_each (fors.begin (), fors.end (), printFlopRate);

  cerr << "[Dumping a PDF...]" << endl;
  generatePDF (*proj);

// DQ (1/2/2008): This output appears to have provided enough synchronization in
// the output to allow -j32 to work. Since I can't debug the problem further for
// now I will leave it.
  cerr << "[Calling backend...]" << endl;

  return backend (proj);
}
```

Figure 41.6: attachMetrics.cc: Sample translator to attach HPCToolkit metrics to the AST

```
0xa0ee358:<SgExprStatement> (Par=0xa0f80f8:<SgForInitStatement>) = (21844.7 flops) / (21844.7 cy) ~= 1 flops/cy
  [/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/roseHPCT/profiled.c:25]
0xa0ee380:<SgExprStatement> (Par=0xb7bd208c:<SgForStatement>) = (21844.7 flops) / (21844.7 cy) ~= 1 flops/cy
  [/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/roseHPCT/profiled.c:25]
0xa0ee3a8:<SgExprStatement> (Par=0xb7bd208c:<SgForStatement>) = (32767 flops) / (21844.7 cy) ~= 1.5 flops/cy
  [/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/roseHPCT/profiled.c:26]
0xa0ee560:<SgExprStatement> (Par=0xb7bd2194:<SgForStatement>) = (98301 flops) / (98301 cy) ~= 1 flops/cy
  [/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/roseHPCT/profiled.c:56]
0xa0ee588:<SgExprStatement> (Par=0xa0f8188:<SgForInitStatement>) = (65534 flops) / (65534 cy) ~= 1 flops/cy
  [/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/roseHPCT/profiled.c:61]
0xa0ee5b0:<SgExprStatement> (Par=0xb7bd2218:<SgForStatement>) = (65534 flops) / (65534 cy) ~= 1 flops/cy
  [/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/roseHPCT/profiled.c:61]
0xa0ee5d8:<SgExprStatement> (Par=0xb7bd2218:<SgForStatement>) = (87378.7 flops) / (109223 cy) ~= 0.8 flops/cy
  [/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/roseHPCT/profiled.c:62]
0xa0f80f8:<SgForInitStatement> (Par=0xb7bd208c:<SgForStatement>) = (21844.7 flops) / (21844.7 cy) ~= 1 flops/cy
  [/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/roseHPCT/profiled.c:25]
0xa0f8158:<SgForInitStatement> (Par=0xb7bd2194:<SgForStatement>) = (98301 flops) / (98301 cy) ~= 1 flops/cy
  [/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/roseHPCT/profiled.c:56]
0xa0f8188:<SgForInitStatement> (Par=0xb7bd2218:<SgForStatement>) = (65534 flops) / (65534 cy) ~= 1 flops/cy
  [/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/roseHPCT/profiled.c:61]
0xb7bd208c:<SgForStatement> (Par=0x9fe2e04:<SgBasicBlock>) = (76456.3** flops) / (65534** cy) ~= 1.16667 flops/cy
  [/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/roseHPCT/profiled.c:25]
0xb7bd2194:<SgForStatement> (Par=0x9fe2f78:<SgBasicBlock>) = (415049** flops) / (436893** cy) ~= 0.95 flops/cy
  [/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/roseHPCT/profiled.c:56]
0xb7bd2218:<SgForStatement> (Par=0x9fe2ff4:<SgBasicBlock>) = (218447** flops) / (240291** cy) ~= 0.909091 flops/cy
  [/home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/roseHPCT/profiled.c:61]
```

Figure 41.7: Sample output, when running attachMetrics.cc (Figure 41.6) with the XML inputs in Figures 41.4–41.5. Here, we only show the output sent to standard output (*i.e.*, cout and not cerr).

Figure 41.8: Sample PDF showing attributes.

Chapter 42

Bug Seeding

Bug seeding is a technique used to construct example codes from existing codes which can be used to evaluate tools for the finding bugs in randomly selected existing applications. The idea is to seed an existing application with a known number of known bugs and evaluate the bug finding or security tool based on the percentage of the number of bugs found by the tool. If the bug finding tool can identify all known bugs then there can be some confidence that the tool detects all bugs of that type used to seed the application.

This example tutorial code is a demonstration of a more complete technique being developed in collaboration with NIST to evaluate tools for finding security flaws in applications. It will in the future be a basis for testing of tools built using ROSE, specifically Compass, but the techniques are not in any way specific to ROSE or Compass.

42.1 Input For Examples Showing Bug Seeding

Figure **??** shows the example input used for demonstration of bug seeding as a transformation.

```
2  void foobar()
     {
4    // Static array declaration
       float array[10];
6
       array[0] = 0;
8
       for (int i=0; i < 5; i++)
10       {
           array[i] = 0;
12       }
   }
```

Figure 42.1: Example source code used as input to program in codes used in this chapter.

42.2 Generating the code representing the seeded bug

Figure 42.2 shows a code that traverses each IR node and for and modifies array reference index expressions to be out of bounds. The input code is shown in figure 42.1, the output of this code is shown in figure 42.3.

```
// This example demonstrates the seeding of a specific type
// of bug (buffer overflow) into any existing application to
// test bug finding tools.

#include "rose.h"
using namespace SageBuilder;
using namespace SageInterface;

namespace SeedBugsArrayIndexing {

class InheritedAttribute
   {
     public:
          bool isLoop;
          bool isVulnerability;
          InheritedAttribute() : isLoop(false),isVulnerability(false) {}
          InheritedAttribute(const InheritedAttribute & X) : isLoop(X.isLoop),isVulnerability(X.isVulnerability) {}
   };

class BugSeeding : public SgTopDownProcessing<InheritedAttribute>
   {
     public:
          InheritedAttribute evaluateInheritedAttribute (
             SgNode* astNode,
             InheritedAttribute inheritedAttribute );
   };

InheritedAttribute
BugSeeding::evaluateInheritedAttribute (
     SgNode* astNode,
     InheritedAttribute inheritedAttribute )
   {
  // Use this if we only want to seed bugs in loops
     bool isLoop = inheritedAttribute.isLoop           ||
                   (isSgForStatement(astNode) != NULL) ||
                   (isSgWhileStmt(astNode) != NULL)    ||
                   (isSgDoWhileStmt(astNode) != NULL);
  // Add Fortran support
     isLoop = isLoop || (isSgFortranDo(astNode) != NULL);

  // Mark future noes in this subtree as being part of a loop
     inheritedAttribute.isLoop = isLoop;

  // To test this on simple codes, optionally allow it to be applied everywhere
     bool applyEveryWhere = true;

     if (isLoop == true || applyEveryWhere == true)
        {
       // The inherited attribute is true iff we are inside a loop and this is a SgPntrArrRefExp.
          SgPntrArrRefExp *arrayReference = isSgPntrArrRefExp(astNode);
          if (arrayReference != NULL)
             {
            // Mark as a vulnerability
               inheritedAttribute.isVulnerability = true;

            // Now change the array index (to seed the buffer overflow bug)
               SgVarRefExp* arrayVarRef = isSgVarRefExp(arrayReference->get_lhs_operand());
               ROSE_ASSERT(arrayVarRef != NULL);
               ROSE_ASSERT(arrayVarRef->get_symbol() != NULL);
               SgInitializedName* arrayName = isSgInitializedName(arrayVarRef->get_symbol()->get_declaration());
               ROSE_ASSERT(arrayName != NULL);
               SgArrayType* arrayType = isSgArrayType(arrayName->get_type());
               ROSE_ASSERT(arrayType != NULL);
               SgExpression* arraySize = arrayType->get_index();

               SgTreeCopy copyHelp;
            // Make a copy of the expression used to hold the array size in the array declaration.
               SgExpression* arraySizeCopy = isSgExpression(arraySize->copy(copyHelp));
               ROSE_ASSERT(arraySizeCopy != NULL);

            // This is the existing index expression
               SgExpression* indexExpression = arrayReference->get_rhs_operand();
               ROSE_ASSERT(indexExpression != NULL);

            // Build a new expression: "array[n]" --> "array[n+arraySizeCopy]", where the arraySizeCopy is a size of "array"
               SgExpression* newIndexExpression = buildAddOp(indexExpression,arraySizeCopy);

            // Substitute the new expression for the old expression
               arrayReference->set_rhs_operand(newIndexExpression);
             }
        }

     return inheritedAttribute;
   }
}

int
main (int argc, char *argv[])
   {
     SgProject *project = frontend (argc, argv);
     ROSE_ASSERT (project != NULL);

     SeedBugsArrayIndexing::BugSeeding treeTraversal;
     SeedBugsArrayIndexing::InheritedAttribute inheritedAttribute;

     treeTraversal.traverseInputFiles (project,inheritedAttribute);

  // Running internal tests (optional)
     AstTests::runAllTests (project);

  // Output the new code seeded with a specific form of bug.
     return backend (project);
   }
```

Figure 42.2: Example source code showing how to seed bugs.

```
2  void foobar()
   {
4  // Static array declaration
     float array[10UL];
6    array[0 + 10UL] = (0);
     for (int i = 0; i < 5; i++) {
8      array[i + 10UL] = (0);
     }
10 }
```

Figure 42.3: Output of input code using seedBugsExample_arrayIndexing.C

Chapter 43

Shared-Memory Parallel Traversals

Besides the traversal classes introduced in Chapter 8, ROSE also includes classes to run multi-threaded traversals to make use of multi-CPU systems with shared memory (such as typical multicore desktop systems). These shared memory traversals are like the combined traversal classes in that they run several small traversal classes simultaneously; the difference is that here different visit functions may be executed concurrently on different CPUs, while the combined traversals always execute visit functions sequentially.

Because of this similarity, the public interface for the parallel traversals is a superset of the combined traversal interface. For each Ast*Processing class there is an AstSharedMemoryParallel*Processing class that provides an interface for adding traversals to its internal list, getting a reference to the internal list, and for starting the combined traversal. The `traverse()` method performs the same combined traversal as in the corresponding AstCombined*Processing class, and the new `traverseInParallel()` method (with the same signature as `traverse()`) must be used to start a parallel traversal. (We currently do not provide `traverseWithinFileInParallel()` and `traverseInputFilesInParallel()` that would be needed to make the parallel processing classes a fully-featured drop-in replacement for other classes.)

A example of how to use the parallel traversal classes is given in Figure 43.1 (note the similarity to Figure 8.24 on page 60). A group of traversal objects is executed first in combined mode and then in parallel threads.

It is the user's responsibility to make sure that the actions executed in the parallel traversal are thread-safe. File or terminal I/O may produce unexpected results if several threads attempt to write to the same stream at once. Allocation of dynamic memory (including the use of ROSE or standard C++ library calls that allocate memory) may defeat the purpose of multi-threading as such calls will typically be serialized by the library.

Two member functions in each AstSharedMemoryParallel*Processing class are available to tune the performance of the parallel traversals. The first is `void set_numberOfThreads(size_t threads)` which can be used to set the number of parallel threads. The default value for this parameter is 2. Our experiments suggest that even on systems with more than two CPUs, running more than two traversal threads in parallel does not typically increase performance

```
#include <rose.h>

class NodeTypeTraversal: public AstSimpleProcessing {
public:
    NodeTypeTraversal(enum VariantT variant, std::string typeName)
        : myVariant(variant), typeName(typeName) {
    }

protected:
    virtual void visit(SgNode *node) {
        if (node->variantT() == myVariant) {
            std::cout << "Found " << typeName;
            if (SgLocatedNode *loc = isSgLocatedNode(node)) {
                Sg_File_Info *fi = loc->get_startOfConstruct();
                if (fi->isCompilerGenerated()) {
                    std::cout << ": compiler generated";
                } else {
                    std::cout << ": " << fi->get_filenameString()
                              << ":" << fi->get_line();
                }
            }
            std::cout << std::endl;
        }
    }

private:
    enum VariantT myVariant;
    std::string typeName;
};

int main(int argc, char **argv) {
    SgProject *project = frontend(argc, argv);

    std::cout << "combined execution of traversals" << std::endl;
    AstSharedMemoryParallelSimpleProcessing parallelTraversal(5);
    parallelTraversal.addTraversal(new NodeTypeTraversal(V_SgForStatement, "for loop"));
    parallelTraversal.addTraversal(new NodeTypeTraversal(V_SgIntVal, "int constant"));
    parallelTraversal.addTraversal(new NodeTypeTraversal(V_SgVariableDeclaration, "variable declaration"));
    parallelTraversal.traverse(project, preorder);
    std::cout << std::endl;

    std::cout << "shared-memory parallel execution of traversals" << std::endl;
    parallelTraversal.traverseInParallel(project, preorder);
}
```

Figure 43.1: Example source showing the shared-memory parallel execution of traversals.

because the memory bandwidth is saturated.

The second function is `void set_synchronizationWindowSize(size_t windowSize)`. This sets a parameter that corresponds to the size of a 'window' of AST nodes that the parallel threads use to synchronize. The value is, in effect, the number of AST nodes that are visited by each thread before synchronizing. Smaller values may in theory result in more locality and therefore better cache utilization at the expense of more time spent waiting for other threads. In practice, synchronization overhead appears to dominate caching effects, so making this parameter too small inhibits performance. The default value is 100000; any large values will result in comparable execution times.

```
combined execution of traversals
Found variable declaration: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_1
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:17
Found variable declaration: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_1
Found for loop: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:20
Found variable declaration: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_1
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:20
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:20
Found variable declaration: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_1
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:27
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:29
Found variable declaration: compiler generated
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:34
Found variable declaration: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_1
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:38
Found for loop: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:41
Found variable declaration: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_1
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:41
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:41
Found variable declaration: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_1
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:46

shared-memory parallel execution of traversals
Found for loop: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:20
Found for loop: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:41
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:17
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:20
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:20
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:27
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:29
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:34
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:38
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:41
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:41
Found int constant: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_18.C:46
Found variable declaration: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_1
Found variable declaration: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_1
Found variable declaration: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_1
Found variable declaration: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_1
Found variable declaration: compiler generated
Found variable declaration: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_1
Found variable declaration: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_1
Found variable declaration: /home/liao6/daily-test-rose/20081014_120001/sourcetree/tutorial/inputCode_ExampleTraversals_1
```

Figure 43.2: Output of input file to the shared-memory parallel traversals. Output may be garbled depending on the multi-threaded behavior of the underlying I/O libraries.

Chapter 44

Distributed-Memory Parallel Traversals

ROSE provides an experimental distributed-memory AST traversal mechanism meant for very large scale program analysis. It allows you to distribute expensive program analyses among a distributed-memory system consisting of many processors; this can be a cluster or a network of workstations. Different processes in the distributed system will get different parts of the AST to analyze: Each process is assigned a number of defining function declarations in the AST, and a method implemented by the user is invoked on each of these. The parts of the AST outside of function definitions are shared among all processes, but there is no guarantee that all function definitions are visible to all processes.

The distributed memory analysis framework uses the MPI message passing library for communicating attributes among processes. You will need an implementation of MPI to be able to build and run programs using distributed memory traversals; consult your documentation on how to run MPI programs. (This is often done using a program named `mpirun`, `mpiexecute`, or similar.)

Distributed memory analyses are performed in three phases:

1. A top-down traversal (the 'pre-traversal') specified by the user runs on the shared AST outside of function definitions. The inherited attributes this traversal computes for defining function declaration nodes in the AST are saved by the framework for use in the next phase.

2. For every defining function declaration, the user-provided `analyzeSubtree()` method is invoked; these calls run concurrently, but on different function declarations, on all processors. It takes as arguments the AST node for the function declaration and the inherited attribute computed for this node by the pre-traversal. Within `analyzeSubtree()` any analysis features provided by ROSE can be used. This method returns the value that will be used as the synthesized attribute for this function declaration in the bottom-up traversal (the 'post-traversal').

   However, unlike normal bottom-up traversals, the synthesized attribute is not simply copied in memory as the AST is distributed. The user must therefore provide the methods `serializeAttribute()` and `deserializeAttribute()`. These compute a serialized

representation of a synthesized attribute, and convert such a representation back to the user's synthesized attribute type, respectively. A serialized attribute is a pair of an integer specifying the size of the attribute in bytes and a pointer to a region of memory of that size that will be copied byte by byte across the distributed system's communication network. Attributes from different parts of the AST may have different sizes. As serialization of attributes will often involve dynamic memory allocation, the user can also implement the `deleteSerializedAttribute()` method to such dynamic memory after the serialized data has been copied to the communication subsystem's internal buffer.

Within the `analyzeSubtree()` method the methods `numberOfProcesses()` and `myID()` can be called. These return the total number of concurrent processes, and an integer uniquely identifying the currently running process, respectively. The ID ranges from 0 to one less than the number of processes, but has no semantics other than that it is different for each process.

3. Finally, a bottom-up traversal is run on the shared AST outside of function definitions. The values returned by the distributed analyzers in the previous phase are used as synthesized attributes for function definition nodes in this traversal.

After the bottom-up traversal has finished, the `getFinalResults()` method can be invoked to obtain the final synthesized attribute. The `isRootProcess()` method returns true on exactly one designated process and can be used to perform output, program transformations, or other tasks that are not meant to be run on each processor.

Figure 44.1 gives a complete example of how to use the distributed memory analysis framework. It implements a trivial analysis that determines for each function declaration at what depth in the AST it can be found and what its name is. Figure 44.2 shows the output produced by this program when running using four processors on some input files.

```
// This is a small example of how to use the distributed memory traversal mechanism. It computes a list of function
// definitions in a program and outputs their names, their depth in the AST, and the ID of the process that found it.

#include <rose.h>
#include "DistributedMemoryAnalysis.h"

// The pre-traversal runs before the distributed part of the analysis and is used to propagate context information down
// to the individual function definitions in the AST. Here, it just computes the depth of nodes in the AST.
class FunctionNamesPreTraversal: public AstTopDownProcessing<int>
{
protected:
    int evaluateInheritedAttribute(SgNode *, int depth)
    {
        return depth + 1;
    }
};

// The post-traversal runs after the distributed part of the analysis and is used to collect the information it
// computed. Here, the synthesized attributes computed by the distributed analysis are strings representing information
// about functions. These strings are concatenated by the post-traversal (and interleaved with newlines where necessary).
class FunctionNamesPostTraversal: public AstBottomUpProcessing<std::string>
{
protected:
    std::string evaluateSynthesizedAttribute(SgNode *node, SynthesizedAttributesList synAttributes)
    {
        std::string result = "";
        SynthesizedAttributesList::iterator s;
        for (s = synAttributes.begin(); s != synAttributes.end(); ++s)
        {
            std::string &str = *s;
            result += str;
            if (str.size() > 0 && str[str.size()-1] != '\n')
                result += "\n";
        }
        return result;
    }

    std::string defaultSynthesizedAttribute()
    {
        return "";
    }
};

// This is the distributed part of the analysis. The DistributedMemoryTraversal base class is a template taking an
// inherited and a synthesized attribute type as template parameters; these are the same types used by the pre- and
// post-traversals.
class FunctionNames: public DistributedMemoryTraversal<int, std::string>
{
protected:
 // The analyzeSubtree() method is called for every defining function declaration in the AST. Its second argument is the
 // inherited attribute computed for this node by the pre-traversal, the value it returns becomes the synthesized
 // attribute used by the post-traversal.
    std::string analyzeSubtree(SgFunctionDeclaration *funcDecl, int depth)
    {
        std::string funcName = funcDecl->get_name().str();
        std::stringstream s;
        s << "process " << myID() << ": at depth " << depth << ": function " << funcName;
        return s.str();
    }

 // The user must implement this method to pack a synthesized attribute (a string in this case) into an array of bytes
 // for communication. The first component of the pair is the number of bytes in the buffer.
    std::pair<int, void *> serializeAttribute(std::string attribute) const
    {
        int len = attribute.size() + 1;
        char *str = strdup(attribute.c_str());
        return std::make_pair(len, str);
    }

 // This method must be implemented to convert the serialized data to the application's synthesized attribute type.
    std::string deserializeAttribute(std::pair<int, void *> serializedAttribute) const
    {
        return std::string((const char *) serializedAttribute.second);
    }

 // This method is optional (the default implementation is empty). Its job is to free memory that may have been
 // allocated by the serializeAttribute() method.
    void deleteSerializedAttribute(std::pair<int, void *> serializedAttribute) const
    {
        std::free(serializedAttribute.second);
    }
};
```

```
----- found the following functions: ------
process 0: at depth 3: function il
process 0: at depth 5: function head
process 0: at depth 5: function eq
process 1: at depth 3: function headhead
process 1: at depth 3: function List
process 1: at depth 3: function find
process 1: at depth 3: function head
process 2: at depth 3: function operator!=
process 2: at depth 3: function find
process 2: at depth 3: function head
process 2: at depth 3: function fib
process 3: at depth 3: function xform
process 3: at depth 3: function func
process 3: at depth 3: function f
process 3: at depth 3: function g
process 3: at depth 3: function deref
-------------------------------------------
```

Figure 44.2: Example output of a distributed-memory analysis running on four processors.

Chapter 45

Parallel Checker

This Chapter is about the project *DistributedMemoryAnalysisCompass*, which runs Compass Checkers in Parallel, i.e. shared, combined and distributed.

45.1 Different Implementations

The project contains the following files:

- parallel_functionBased_ASTBalance contains the original implementation, which is based on an AST traversal that is balanced based on the number of nodes in each function. Then the functions are distributed over all processors. It contains as well the original interfaces to the shared and combined traversal work.
- parallel_file_compass distributed on the granularity level of files.
- parallel_functionBased_dynamicBalance is the implementation of dynamically scheduling functions across processors. In addition, this algorithm weights the functions first and then sorts them in descending order according to their weight.
- parallel_compass performs dynamic scheduling based on nodes. The nodes are weighted and then sorted. This algorithm allows the greatest scalability.

45.2 Running through PSUB

The following represents a typical script to run parallel_compass on 64 processors using CXX_Grammer. CXX_Grammar is a binary ROSE AST representation of a previously parsed program. We specify 65 processors because processor 0 does only communication and no computation. Furthermore, we ask for 17 nodes of which each has 8 processors giving us a total of 136 possible processes. We only need 65 but still want to use this configuration. This will average out our 65 processes over 17 nodes, resulting in about 4 processors per node. This trick is used because the AST loaded into memory takes about 400 MB per process. We end up with 1600MB per node.

```
#!/bin/bash
#  Sample LCRM script to be submitted with psub
#PSUB -r ncxx65 # sets job name (limit of 7 characters)
#PSUB -b nameofbank  # sets bank account
#PSUB -ln 17 # == defines the amount of nodes needed
#PSUB -o ~/log.log
#PSUB -e ~/log.err
#PSUB -tM 0:05  # Max time 5 min runtime
#PSUB -x  # export current env var settings
#PSUB -nr # do NOT rerun job after system reboot
#PSUB -ro # send output log directly to file
#PSUB -re # send err log directly to file
#PSUB -mb # send email at execution start
#PSUB -me # send email at execution finish
#PSUB -c zeus
#PSUB # no more psub commands
# job commands start here
set echo
echo LCRM job id = $PSUB_JOBID
cd ~/new-src/build-rose/projects/DistributedMemoryAnalysisCompass/
srun -n 65 ./parallel_compass -load ~/CXX_Grammar.ast
echo "ALL DONE"
```

There are a few tricks that could be considered. Prioritization is based on the amount of time and nodes requested. If less time is specified, it is more likely that a job runs very soon, as processing time becomes available.

To submit the job above, use *psub file-name*. To check the job in the queue, use *squeue* and to cancel the job use *mjobctl -c job-number*.

Chapter 46

Abstract Handles to Language Constructs

This chapter describes a reference design and its corresponding implementation for supporting abstract handles to language constructs in source code and optimization phases. It can be used to facilitate the interoperability between compilers and tools.

The idea is to define unique identifiers for statements, loops, functions, and other language constructs in source code. Given the diverse user requirements, an ideal specification should include multiple forms to specify a language construct. Current, we are interested in the following forms for specifying language constructs:

- Source file position information including path, filename, line and column number etc. GNU standard source position from `http://www.gnu.org/prep/standards/html_node/Errors.html` presents some examples.
- Global or local numbering of specified language construct in source file (e.g. 2nd "do" loop in a global scope). The file is itself specified using an abstract handle (typically generated from the file name).
- Global or local names of constructs. Some language constructs, such as files, function definitions and namespace, have names which can be used as their handle within a context.
- Language-specific label mechanisms. They include named constructs in Fortran, numbered labels in Fortran, and statement labels in C and C++ and so on.

In addition to human-readable forms, compilers and tools can generate internal IDs for language constructs. It is up to compiler/tool developers to provide a way to convert their internal representations into human-readable formats.

We define an abstract handle as a unique representation for a language construct. It can have any of the human-readable or machine-generated forms. A handle can be used alone or combined with other handles to specify a language construct. A handle can also be converted from one form to another. Abstract handles can have different lifetimes depending on their use and implementation. An abstract handle might be required to be persistent if it is used to

reference a language construct that would be optimized over multiple executions of one or more different tools. Where as an abstract-handle might be internally generated only for purposes of optimizations used in a single execution (e.g. optimization within a compiler).

46.1 Syntax

A possible specification of language handles can have the following syntax:

```
/* a handle is a single handle item or a link of them separated by ::, or
other delimiters */
handle ::= handle_item | handle '::' handle_item

/* Each handle item consists of construct_type and a specifier.
Or it can be a compiler generated id of any forms. */

handle_item ::= construct_type specifier | compiler_generated_handle

/*
Construct types are implementation dependent.
An implementation can support a subset of legal constructs or all of them.
For ROSE, they are the type names of IR nodes.
*/
construct_type ::= SgProject | SgFile | SgForStmt|SgVariableDeclaration | ...

/* A specifier is used to locate a particular construct
  e.g: <name, "foo">
*/

specifier::= '<' specifier_type ',' specifier_value '>'

/* tokens for specifier types could be name, position,numbering, label, etc.
specifier type is necessary to avoid ambiguity for specifier values,
because a same value could be interpreted in different specifier types otherwise
*/

specifier_type::= name | position | numbering | label

/* Possible values for a specifier */

specifier_value::= string_lit|int_lit|position_value| label_value

/*A label could be either integer or string */
label_value::= int_lit | string_lit

/* Start and end source line and column information
```

```
e.g.: 13.5-55.4,  13,  13.5 , 13.5-55 */
position_value:: = line_number[ '.' column_number][ '-' line_number[ '.' column_number]]

/* Integer value: one or more digits */
int_lit ::= [0-9]+

/* String value: start with a letter, followed by zero or more letters or digits */
string_lit ::= [a-z][a-z0-9]*
```

46.2 Examples

We give some examples of language handles using the syntax mentioned above. ROSE AST's node type names are used as the construct type names. Other implementations can use their own construct type names.

- A file handle consisting of only one handle item:

  ```
  SgFile<name,"/home/PERI/test111.f">
  ```

- A function handle using a named handle item, combined with a parent handle using a name also:

  ```
  SgFile<name,"/home/PERI/test111.f">::SgFunctionDefiningDeclaration<name,"foo">
  ```

- A function handle using source position(A function starting at line 12, column 1 till line 30, column 1 within a file):

  ```
  SgFile<name,"/home/PERI/test111.f">::SgFunctionDefiningDeclaration<position,"12.1-30.1">
  ```

- A function handle using numbering(The first function definition in a file):

  ```
  SgFile<name,/home/PERI/test111.f">::SgFunctionDefiningDeclaration<numbering,1>
  ```

- A return statement using source position (A return statement at line 100):

  ```
  SgFile<name,/home/PERI/test222.c>::SgReturnStatement<position,"100">
  ```

- A loop using numbering information (The second loop in function main()):

  ```
  SgFile<name,"/home/PERI/test222.c">::SgFunctionDefiningDeclaration<name,"main">::
  SgForStatement<numbering,2>
  ```

46.3 Reference Implementation

We provide a reference implementation of the abstract handle concept using ROSE. The source files are located in *src/midend/abstractHandle*. A generic interface (abstract_handle.h and abstract_handle.cpp) provides data structures and operations for manipulating abstract handles using source file positions, numbering, or names. A ROSE adapter (roseAdapter.h and roseAdapter.cpp) using the interface is provided as a concrete implementation. Other compilers and tools can have their own implementations using the same interface.

Figure 46.1 shows the code to generate abstract handles for loops in an input source file (as in Figure 46.2). Abstract handle constructors generate handles from abstract nodes, which are implemented using ROSE AST nodes. Source position is used by default to generate a handle item. Names or numbering are used instead when source position information is not available. The Constructor can also be used to generate a handle item using a specified handle type (numbering handles in the example). Figure 46.3 is the output showing the generated handles for the loops.

Another example (shown in Figure 46.4) demonstrates how to create handles using user-specified strings representing handle items for language constructs within a source file (shown in Figure 46.5). This is particularly useful to grab internal language constructs from handles provided by external software tools. The output of the example is given in Figure 46.6.

```
/*
Example code to generate abstract handles for language constructs

by Liao, 10/6/2008
*/
#include "rose.h"
#include <iostream>
#include "abstract_handle.h"
#include "roseAdapter.h"
#include <string.h>

using namespace std;
using namespace AbstractHandle;

// a global handle for the current file
static abstract_handle* file_handle = NULL;

class visitorTraversal : public AstSimpleProcessing
{
  protected:
    virtual void visit(SgNode* n);
};

void visitorTraversal::visit(SgNode* n)
{
  SgForStatement* forloop = isSgForStatement(n);

  if (forloop)
  {
    //Create an abstract node
    abstract_node* anode= new roseNode(forloop);

    //Create an abstract handle from the abstract node
    //Using source position specifiers by default
    abstract_handle * ahandle = new abstract_handle(anode);
    cout<<ahandle->toString()<<endl;

    // Create handles based on numbering specifiers within the file
    abstract_handle * bhandle = new abstract_handle(anode,e_numbering,file_handle);
    cout<<bhandle->toString()<<endl;
  }
}

int main(int argc, char * argv[])
{
  SgProject *project = frontend (argc, argv);

  //Generate a file handle
  abstract_node * file_node = new roseNode((project->get_fileList())[0]);
  file_handle = new abstract_handle(file_node);

  //Generate handles for language constructs
  visitorTraversal myvisitor;
  myvisitor.traverseInputFiles(project,preorder);

// Generate source code from AST and call the vendor's compiler
  return backend(project);
}
```

Figure 46.1: Generated handles for loops: using constructors with or without a specified handle type

```
/* test input for generated abstract handles */
int a[100][100][100];

void foo()
{
  int i,j,k;
  for (i=0;i++;i<100)
    for (j=0;j++;j<100)
      for (k=0;k++;k<100)
        a[i][j][k]=i+j+k;

  for (i=0;i++;i<100)
    for (j=0;j++;j<100)
      for (k=0;k++;k<100)
        a[i][j][k]+=5;
}
```

Figure 46.2: Source code with some loops

```
SgProject<numbering,1>::SgSourceFile<name,/home/liao6/daily-test-rose/20081014_1
20001/sourcetree/tutorial/inputCode_AbstractHandle1.cpp>::SgForStatement<positio
n,7.3-10.25>
SgProject<numbering,1>::SgSourceFile<name,/home/liao6/daily-test-rose/20081014_1
20001/sourcetree/tutorial/inputCode_AbstractHandle1.cpp>::SgForStatement<numberi
ng,1>
SgProject<numbering,1>::SgSourceFile<name,/home/liao6/daily-test-rose/20081014_1
20001/sourcetree/tutorial/inputCode_AbstractHandle1.cpp>::SgForStatement<positio
n,8.5-10.25>
SgProject<numbering,1>::SgSourceFile<name,/home/liao6/daily-test-rose/20081014_1
20001/sourcetree/tutorial/inputCode_AbstractHandle1.cpp>::SgForStatement<numberi
ng,2>
SgProject<numbering,1>::SgSourceFile<name,/home/liao6/daily-test-rose/20081014_1
20001/sourcetree/tutorial/inputCode_AbstractHandle1.cpp>::SgForStatement<positio
n,9.7-10.25>
SgProject<numbering,1>::SgSourceFile<name,/home/liao6/daily-test-rose/20081014_1
20001/sourcetree/tutorial/inputCode_AbstractHandle1.cpp>::SgForStatement<numberi
ng,3>
SgProject<numbering,1>::SgSourceFile<name,/home/liao6/daily-test-rose/20081014_1
20001/sourcetree/tutorial/inputCode_AbstractHandle1.cpp>::SgForStatement<positio
n,12.3-15.22>
SgProject<numbering,1>::SgSourceFile<name,/home/liao6/daily-test-rose/20081014_1
20001/sourcetree/tutorial/inputCode_AbstractHandle1.cpp>::SgForStatement<numberi
ng,4>
SgProject<numbering,1>::SgSourceFile<name,/home/liao6/daily-test-rose/20081014_1
20001/sourcetree/tutorial/inputCode_AbstractHandle1.cpp>::SgForStatement<positio
n,13.5-15.22>
SgProject<numbering,1>::SgSourceFile<name,/home/liao6/daily-test-rose/20081014_1
20001/sourcetree/tutorial/inputCode_AbstractHandle1.cpp>::SgForStatement<numberi
ng,5>
SgProject<numbering,1>::SgSourceFile<name,/home/liao6/daily-test-rose/20081014_1
20001/sourcetree/tutorial/inputCode_AbstractHandle1.cpp>::SgForStatement<positio
n,14.7-15.22>
SgProject<numbering,1>::SgSourceFile<name,/home/liao6/daily-test-rose/20081014_1
20001/sourcetree/tutorial/inputCode_AbstractHandle1.cpp>::SgForStatement<numberi
ng,6>
```

Figure 46.3: Handles generated for loops

```
/*
Example code to generate language handles from input strings about
* source position information
* numbering information

by Liao, 10/9/2008
*/
#include "rose.h"
#include <iostream>
#include <string.h>
#include "abstract_handle.h"
#include "roseAdapter.h"

using namespace std;
using namespace AbstractHandle;

int main(int argc, char * argv[])
{
  SgProject *project = frontend (argc, argv);

  // Generate a file handle from the first file of the project
  abstract_node* file_node= new roseNode((project->get_fileList())[0]);
  abstract_handle* handle1 = new abstract_handle(file_node);
  cout<<"Created_file_handle:\n"<<handle1->toString()<<endl<<endl;;

  //Create a handle to a namespace given its name and parent handle
  string input1("SgNamespaceDeclarationStatement<name,space1>");
  abstract_handle* handle0 = new abstract_handle(handle1,input1);
  cout<<"Created_a_handle:\n"<<handle0->toString()<<endl<<endl;
  cout<<"It_points_to:\n"<<handle0->getNode()->toString()<<endl;

  // Create a handle within the file, given a string specifying
  // its construct type (class declaration) and source position
  string input("SgClassDeclaration<position,4.3-9.2>");
  abstract_handle* handle2 = new abstract_handle(handle1,input);

  cout<<"Created_a_handle:\n"<<handle2->toString()<<endl;
  cout<<"It_points_to:\n"<<handle2->getNode()->toString()<<endl<<endl;;

  // find the second function declaration within handle2
  abstract_handle handle3(handle2,"SgFunctionDeclaration<numbering,2>");

  cout<<"Created_a_handle:\n"<<handle3.toString()<<endl;
  cout<<"It_points_to:\n"<<handle3.getNode()->toString()<<endl;

// Generate source code from AST and call the vendor's compiler
  return backend(project);
}
```

Figure 46.4: Generated handles from strings representing handle items

```
void bar(int x);
namespace space1
{
  class A
  {
  public:
    void bar(int x);
    void foo();
  };
}
```

Figure 46.5: Source code with some language constructs

```
#! /bin/sh

# abstractHandle2 - temporary wrapper script for .libs/abstractHandle2
# Generated by ltmain.sh - GNU libtool 1.5.6 (1.1220.2.95 2004/04/11 05:50:42)
#
# The abstractHandle2 program cannot be directly executed until all the libtool
# libraries that it depends on are installed.
#
# This wrapper script should never be moved out of the build directory.
# If it is, it will not operate correctly.

# Sed substitution that helps us do robust quoting.  It backslashifies
# metacharacters that are still active within double-quoted strings.
Xsed='/bin/sed -e 1s/^X//'
sed_quote_subst='s/\([\\`\"$\\]\)/\\\1/g'

# The HP-UX ksh and POSIX shell print the target directory to stdout
# if CDPATH is set.
if test "${CDPATH+set}" = set; then CDPATH=:; export CDPATH; fi

relink_command="(cd /home/liao6/daily-test-rose/20081014_120001/build/tutorial;{ test -z \"\${LIBRARY_PATH+set}\" || unset LIBRA

# This environment variable determines our operation mode.
if test "$libtool_install_magic" = "%%%MAGIC variable%%%"; then
  # install mode needs the following variable:
  notinst_deplibs='../src/librose.la /home/liao6/daily-test-rose/20081014_120001/build/src/3rdPartyLibraries/libharu-2.1.0/src/l
else
  # When we are sourced in execute mode, $file and $echo are already set.
  if test "$libtool_execute_magic" != "%%%MAGIC variable%%%"; then
    echo="echo"
    file="$0"
    # Make sure echo works.
    if test "X$1" = X--no-reexec; then
      # Discard the --no-reexec flag, and continue.
      shift
    elif test "X`($echo '\t') 2>/dev/null`" = 'X\t'; then
      # Yippee, $echo works!
      :
    else
      # Restart under the correct shell, and then maybe $echo will work.
      exec /bin/sh "$0" --no-reexec ${1+"$@"}
    fi
  fi

  # Find the directory that this script lives in.
  thisdir=`$echo "X$file" | $Xsed -e 's%/[^/]*$%%'`
  test "x$thisdir" = "x$file" && thisdir=.

  # Follow symbolic links until we get to the real thisdir.
  file=`ls -ld "$file" | /bin/sed -n 's/.*-> //p'`
  while test -n "$file"; do
    destdir=`$echo "X$file" | $Xsed -e 's%/[^/]*$%%'`

    # If there was a directory component, then change thisdir.
    if test "x$destdir" != "x$file"; then
      case "$destdir" in
      [\\/]* | [A-Za-z]:[\\/]*) thisdir="$destdir" ;;
      *) thisdir="$thisdir/$destdir" ;;
      esac
    fi

    file=`$echo "X$file" | $Xsed -e 's%^.*/%%'`
    file=`ls -ld "$thisdir/$file" | /bin/sed -n 's/.*-> //p'`
  done

  # Try to get the absolute directory name.
  absdir=`cd "$thisdir" && pwd`
  test -n "$absdir" && thisdir="$absdir"

  program=lt-'abstractHandle2'
  progdir="$thisdir/.libs"

  if test ! -f "$progdir/$program" || \
     { file=`ls -1dt "$progdir/$program" "$progdir/../$program" 2>/dev/null | /bin/sed 1q`; \
       test "X$file" != "X$progdir/$program"; }; then

    file="$$-$program"

    if test ! -d "$progdir"; then
      mkdir "$progdir"
    else
      rm -f "$progdir/$file"
    fi

    # relink executable if necessary
    if test -n "$relink_command"; then
      if relink_command_output=`eval $relink_command 2>&1`; then :
      else
        echo "$relink_command_output" >&2
        rm -f "$progdir/$file"
        exit 1
      fi
    fi

    mv -f "$progdir/$file" "$progdir/$program" 2>/dev/null ||
    { rm -f "$progdir/$program";
      mv -f "$progdir/$file" "$progdir/$program"; }
    rm -f "$progdir/$file"
  fi

  if test -f "$progdir/$program"; then
    if test "$libtool_execute_magic" != "%%%MAGIC variable%%%"; then
      # Run the actual program with our arguments.

      exec $progdir/$program ${1+"$@"}

      $echo "$0: cannot exec $program ${1+"$@"}"
      exit 1
    fi
  else
    # The program doesn't exist.
```

Chapter 47

Making your Contributions to ROSE

To make the process of contributing to ROSE go more smoothly, the following procedure has been established. Steps 1-3 should be followed for contributions of new code, and step 2 for changes to existing code.

1. First of all, work outside of the ROSE tree. You can use your own Makefile to control the build process. An example of such a Makefile is shown in Figure 36.1 .

2. If you find you need to make changes to ROSE, make them directly to the ROSE source tree, and contribute patches as you work. If you have access to the CVS repository, use the '`cvs diff`' command to create patches. If you are starting with a heavily patched version of ROSE, or don't have access to CVS, make a pristine copy of ROSE before you start and use the plain '`diff`' command.

   In order to meet the standard for ROSE patches, which include:

   - Context information
   - Omission of machine-generated files

   please supply the following flags to your '`diff`' or '`cvs diff`' command:

   ```
   -u --exclude 'aclocal.m4' --exclude 'autom4te.cache' --exclude 'Makefile.in'
   --exclude 'configure'
   ```

   If using '`diff`' rather than '`cvs diff`', the '`-r`' flag is also required.

   The syntax for the `diff` command is as follows:

   ```
   diff [flags] <dir1> <dir2> > <patchfile>
   ```

For our purposes, the `diff` command should be issued in the directory directly above the ROSE source directory. Be sure to give the patch file a suitable name such as `enhanced-template-support.patch`. `<dir1>` should be the name of your pristine source directory; `<dir2>` that of your modified sources. For example, if your pristine source tree is located at `/home/me/research/ROSE-orig` and your modified sources are in `/home/me/research/ROSE`, you can issue the commands:

```
$ cd /home/me/research
$ diff -ur --exclude 'aclocal.m4' --exclude 'autom4te.cache' --exclude 'Makefile.in'
  --exclude 'configure' ROSE-orig ROSE > my-changes.patch
```

The syntax for the `cvs diff` command is as follows:

```
cvs diff [flags] <dir> > <patchfile>
```

This command should be given in the directory directly above that to which you wish to contribute changes. Then the `<dir>` option shall be the name of this directory. For example, to create a patch for the `src/frontend` directory, issue the following commands:

```
$ cd src
$ cvs diff -u --exclude 'aclocal.m4' --exclude 'autom4te.cache' --exclude 'Makefile.in'
  --exclude 'configure' frontend > my-changes.patch
```

You will need to supply information to the maintainer on how to apply your patch. You will need to look at the paths following the '`---`' and '`+++`' markers in the patch file. If you used '`cvs diff`' to create the patch, the paths will likely be relative to the directory you were in when you created the patch, and you'll need to apply the patch with '`patch -p0`' in the same directory. If you used '`diff`' the path will probably include the name of your source tree root directory, and you'll need to use '`patch -p1`' in the root of the source tree.

3. When you are ready to contribute the work, it should first be integrated into the ROSE tree. Create a new directory either under `projects` or a subdirectory of `src`. If you are unsure where to create your directory, contact the maintainer. Once you have created the directory copy your files there, excluding the Makefile.

   You will need to create a Makefile.am which integrates with the ROSE build process. The best way to learn about these is by looking at other Makefile.am files in sibling directories to yours. Note the differences between a '`project`' and a '`src`' Makefile.am, e.g. only the latter will create shared libraries. Remember to add your new directory to the `SUBDIRS` list in the parent directory's Makefile.am, and your Makefile.am to the list of Makefiles in `configure.in` in the root directory.

   If your project is in `src`, your test cases should go into a subdirectory of `tests/roseTests`, which mirrors the `src` directory structure. Otherwise they can go into your project directory. Again, look at other Makefiles for information on how test cases work.

   While making these changes, keep track of all the new files you create within the ROSE tree. These will typically be all the files in your new directory as well as your test codes

and test input codes in the `tests` directory, if applicable. Once you are done, '`tar`' up all these files and send them to the maintainer, along with patches to existing files, created using '`diff`' or '`cvs diff`' as described in step 2.

The current maintainer is Dan Quinlan `<dquinlan@llnl.gov>`.

Appendix

This appendix includes information useful in supporting the ROSE Tutorial.

47.1 Location of To Do List

This was an older location for the Tutorial Tod List. We now keep the Tod list in the `ROSE/docs/testDoxygen/ProjectToDoList.docs` in the section called: `ROSE Tutorial Todo List`.

47.2 Abstract Grammar

In this section we show an abstract grammar for the ROSE AST. The grammar generates the set of all ASTs. On the left hand side of a production we have a non-terminal that corresponds to an inner node of the class hierarchy. On the right hand side of a production we have either one non-terminal or one terminal. The terminal corresponds to a leaf-node where the children of the respective node are listed as double-colon separated pairs, consisting of an access name (= name for get function) and a name that directly corresponds to the class of the child. Details like pointers are hidden. The asterisk shows where lists of children (containers) are used in the ROSE AST. For each terminal, a name followed by '(' and ')', a variant exists in ROSE with the prefix `V_` that can be obtained by using the function variantT() on a node. Note, that concrete classes of AST nodes directly correspond to terminals and base classes to non-terminals.

```
START:SgNode
SgNode : SgSupport
 | SgLocatedNode
 | SgSymbol
 ;

SgSupport : SgName()
 | SgSymbolTable()
 | SgInitializedName ( initptr:SgInitializer )
 | SgFile ( root:SgGlobal ( declarations:SgDeclarationStatement* ) )
 | SgProject ( fileList:SgFile ( root:SgGlobal ( declarations:SgDeclarationStatement* ) ) )
 | SgOptions()
 | SgBaseClass ( base_class:SgClassDeclaration )
 | SgTemplateParameter ( expression:SgExpression, defaultExpressionParameter:SgExpression,
```

```
                        templateDeclaration:SgTemplateDeclaration(),
                        defaultTemplateDeclarationParameter:SgTemplateDeclaration() )
 | SgTemplateArgument ( expression:SgExpression,
                       templateInstantiation:SgTemplateInstantiationDecl
                          ( definition:SgClassDefinition ) )
 | SgFunctionParameterTypeList()
 | SgAttribute
 | SgModifier
 ;

SgAttribute : SgPragma()
 | SgBitAttribute
 ;

SgBitAttribute : SgFuncDecl_attr()
 | SgClassDecl_attr()
 ;

SgModifier : SgModifierNodes()
 | SgConstVolatileModifier()
 | SgStorageModifier()
 | SgAccessModifier()
 | SgFunctionModifier()
 | SgUPC_AccessModifier()
 | SgSpecialFunctionModifier()
 | SgElaboratedTypeModifier()
 | SgLinkageModifier()
 | SgBaseClassModifier()
 | SgDeclarationModifier()
 ;

SgLocatedNode : SgStatement
 | SgExpression
 ;

SgStatement : SgExprStatement ( expression_root:SgExpressionRoot ( operand_i:SgExpression ) )
 | SgLabelStatement()
 | SgCaseOptionStmt ( key_root:SgExpressionRoot ( operand_i:SgExpression ),
                      body:SgBasicBlock ( statements:SgStatement* ) )
 | SgTryStmt ( body:SgBasicBlock ( statements:SgStatement* ),
               catch_statement_seq_root:SgCatchStatementSeq ( catch_statement_seq:SgStatement* ) )
 | SgDefaultOptionStmt ( body:SgBasicBlock ( statements:SgStatement* ) )
 | SgBreakStmt()
 | SgContinueStmt()
 | SgReturnStmt ( expression_root:SgExpressionRoot ( operand_i:SgExpression ) )
 | SgGotoStatement()
```

```
 | SgSpawnStmt ( the_func_root:SgExpressionRoot ( operand_i:SgExpression ) )
 | SgForInitStatement ( init_stmt:SgStatement* )
 | SgCatchStatementSeq ( catch_statement_seq:SgStatement* )
 | SgClinkageStartStatement()
 | SgDeclarationStatement
 | SgScopeStatement
 ;

SgDeclarationStatement :
   SgVariableDeclaration ( variables:SgInitializedName ( initptr:SgInitializer ) )
 | SgVariableDefinition ( vardefn:SgInitializedName ( initptr:SgInitializer ),
                          bitfield:SgUnsignedLongVal() )
 | SgEnumDeclaration()
 | SgAsmStmt ( expr_root:SgExpressionRoot ( operand_i:SgExpression ) )
 | SgTemplateDeclaration()
 | SgNamespaceDeclarationStatement ( definition:SgNamespaceDefinitionStatement
                                        ( declarations:SgDeclarationStatement* )
                                   )
 | SgNamespaceAliasDeclarationStatement()
 | SgUsingDirectiveStatement()
 | SgUsingDeclarationStatement()
 | SgFunctionParameterList ( args:SgInitializedName ( initptr:SgInitializer ) )
 | SgCtorInitializerList ( ctors:SgInitializedName ( initptr:SgInitializer ) )
 | SgPragmaDeclaration ( pragma:SgPragma() )
 | SgClassDeclaration
 | SgFunctionDeclaration
 ;

SgClassDeclaration : SgTemplateInstantiationDecl ( definition:SgClassDefinition )
 ;

SgFunctionDeclaration :
   SgTemplateInstantiationFunctionDecl ( parameterList:SgFunctionParameterList
                                           ( args:SgInitializedName ( initptr:SgInitializer ) ),
                                         definition:SgFunctionDefinition
                                            ( body:SgBasicBlock ( statements:SgStatement* ) )
                                       )
 | SgMemberFunctionDeclaration
 ;

SgMemberFunctionDeclaration :
    SgTemplateInstantiationMemberFunctionDecl ( parameterList:SgFunctionParameterList
                                                  ( args:SgInitializedName ( initptr:SgInitializer ) ),
                                                definition:SgFunctionDefinition
                                                  ( body:SgBasicBlock ( statements:SgStatement* ) ),
                                                CtorInitializerList:SgCtorInitializerList
```

```
                                                    ( ctors:SgInitializedName ( initptr:SgInitializer
                                                  )
 ;

SgScopeStatement : SgGlobal ( declarations:SgDeclarationStatement* )
 | SgBasicBlock ( statements:SgStatement* )
 | SgIfStmt ( conditional:SgStatement,
              true_body:SgBasicBlock ( statements:SgStatement* ),
              false_body:SgBasicBlock ( statements:SgStatement* ) )
 | SgForStatement ( for_init_stmt:SgForInitStatement ( init_stmt:SgStatement* ),
                    test_expr_root:SgExpressionRoot ( operand_i:SgExpression ),
                    increment_expr_root:SgExpressionRoot ( operand_i:SgExpression ),
                    loop_body:SgBasicBlock ( statements:SgStatement* ) )
 | SgFunctionDefinition ( body:SgBasicBlock ( statements:SgStatement* ) )
 | SgWhileStmt ( condition:SgStatement, body:SgBasicBlock ( statements:SgStatement* ) )
 | SgDoWhileStmt ( condition:SgStatement, body:SgBasicBlock ( statements:SgStatement* ) )
 | SgSwitchStatement ( item_selector_root:SgExpressionRoot ( operand_i:SgExpression ),
                       body:SgBasicBlock ( statements:SgStatement* ) )
 | SgCatchOptionStmt ( condition:SgVariableDeclaration
                                     ( variables:SgInitializedName ( initptr:SgInitializer ) ),
                                       body:SgBasicBlock ( statements:SgStatement* )
                     )
 | SgNamespaceDefinitionStatement ( declarations:SgDeclarationStatement* )
 | SgClassDefinition
 ;

SgClassDefinition : SgTemplateInstantiationDefn ( members:SgDeclarationStatement* )
 ;

SgExpression : SgExprListExp ( expressions:SgExpression* )
 | SgVarRefExp()
 | SgClassNameRefExp()
 | SgFunctionRefExp()
 | SgMemberFunctionRefExp()
 | SgFunctionCallExp ( function:SgExpression, args:SgExprListExp ( expressions:SgExpression* ) )
 | SgSizeOfOp ( operand_expr:SgExpression )
 | SgConditionalExp ( conditional_exp:SgExpression,
                      true_exp:SgExpression,
                      false_exp:SgExpression )
 | SgNewExp ( placement_args:SgExprListExp ( expressions:SgExpression* ),
              constructor_args:SgConstructorInitializer(
                                       args:SgExprListExp(expressions:SgExpression*)
                                                  ),
              builtin_args:SgExpression )
 | SgDeleteExp ( variable:SgExpression )
 | SgThisExp()
```

```
 | SgRefExp()
 | SgVarArgStartOp ( lhs_operand:SgExpression, rhs_operand:SgExpression )
 | SgVarArgOp ( operand_expr:SgExpression )
 | SgVarArgEndOp ( operand_expr:SgExpression )
 | SgVarArgCopyOp ( lhs_operand:SgExpression, rhs_operand:SgExpression )
 | SgVarArgStartOneOperandOp ( operand_expr:SgExpression )
 | SgInitializer
 | SgValueExp
 | SgBinaryOp
 | SgUnaryOp
 ;

SgInitializer :
   SgAggregateInitializer ( initializers:SgExprListExp ( expressions:SgExpression* ) )
 | SgConstructorInitializer ( args:SgExprListExp ( expressions:SgExpression* ) )
 | SgAssignInitializer ( operand_i:SgExpression )
 ;

SgValueExp : SgBoolValExp()
 | SgStringVal()
 | SgShortVal()
 | SgCharVal()
 | SgUnsignedCharVal()
 | SgWcharVal()
 | SgUnsignedShortVal()
 | SgIntVal()
 | SgEnumVal()
 | SgUnsignedIntVal()
 | SgLongIntVal()
 | SgLongLongIntVal()
 | SgUnsignedLongLongIntVal()
 | SgUnsignedLongVal()
 | SgFloatVal()
 | SgDoubleVal()
 | SgLongDoubleVal()
 ;

SgBinaryOp : SgArrowExp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgDotExp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgDotStarOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgArrowStarOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgEqualityOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgLessThanOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgGreaterThanOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgNotEqualOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgLessOrEqualOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
```

```
 | SgGreaterOrEqualOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgAddOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgSubtractOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgMultiplyOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgDivideOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgIntegerDivideOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgModOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgAndOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgOrOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgBitXorOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgBitAndOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgBitOrOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgCommaOpExp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgLshiftOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgRshiftOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgPntrArrRefExp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgScopeOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgAssignOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgPlusAssignOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgMinusAssignOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgAndAssignOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgIorAssignOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgMultAssignOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgDivAssignOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgModAssignOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgXorAssignOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgLshiftAssignOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 | SgRshiftAssignOp ( lhs_operand_i:SgExpression, rhs_operand_i:SgExpression )
 ;

SgUnaryOp : SgExpressionRoot ( operand_i:SgExpression )
 | SgMinusOp ( operand_i:SgExpression )
 | SgUnaryAddOp ( operand_i:SgExpression )
 | SgNotOp ( operand_i:SgExpression )
 | SgPointerDerefExp ( operand_i:SgExpression )
 | SgAddressOfOp ( operand_i:SgExpression )
 | SgMinusMinusOp ( operand_i:SgExpression )
 | SgPlusPlusOp ( operand_i:SgExpression )
 | SgBitComplementOp ( operand_i:SgExpression )
 | SgCastExp ( operand_i:SgExpression )
 | SgThrowOp ( operand_i:SgExpression )
 ;

SgSymbol : SgVariableSymbol()
 | SgClassSymbol ( declaration:SgClassDeclaration )
 | SgTemplateSymbol ( declaration:SgTemplateDeclaration() )
```

```
 | SgEnumSymbol ( declaration:SgEnumDeclaration() )
 | SgEnumFieldSymbol()
 | SgLabelSymbol ( declaration:SgLabelStatement() )
 | SgDefaultSymbol()
 | SgNamespaceSymbol ( declaration:SgNamespaceDeclarationStatement
                       ( definition:SgNamespaceDefinitionStatement
                          ( declarations:SgDeclarationStatement* )
                       )
                     )
 | SgFunctionSymbol
 ;

SgFunctionSymbol : SgMemberFunctionSymbol ( declaration:SgFunctionDeclaration )
 ;

SgPartialFunctionType :
   SgPartialFunctionModifierType ( ref_to:SgReferenceType, ptr_to:SgPointerType,
                                   modifiers:SgModifierNodes(), typedefs:SgTypedefSeq,
                                   return_type:SgType, orig_return_type:SgType )
 ;
```

This grammar was generated with GRATO, a grammar transformation tool, written by Markus Schordan. The input is a representation generated by ROSETTA. Several other versions of the grammar can be generated as well, such as eliminating nested tree nodes by introducing auxiliary non-terminals, introducing base types as non-terminals etc. Additionally from that grammar we can also generate grammars that can be used with yacc/bison, Coco, and other attribute grammar tools, as well as tree grammar based tools such as burg (requires a transformation to a binary tree).

Glossary

We define terms used in the ROSE manual which might otherwise be unclear.

FIXME: *Define the following terms: IR node, Inherited Attribute, Synthesized Attribute, Accumulator Attribute, AST Traversal*

- **AST** Abstract Syntax Tree. A very basic understanding of an AST is the entry level into ROSE.
- **Attribute** User defined information (objects) associated with IR nodes. Forms of attributes include: accumulator, inherited, persistent, and synthesized. Both inherited and synthesized attributes are managed automatically on the stack within a traversal. Accumulator attributes are typically something semantically equivalent to a global variable (often a static data member of a class). Persistent attributes are explicitly added to the AST and are managed directly by the user. As a result, they can persist across multiple traversals of the AST. Persistent attributes are also saved in the binary file I/O, but only if the user provides the attribute specific `pack()` and `unpack()` virtual member functions. See the ROSE User Manual for more information, and the ROSE Tutorial for examples.
- **CFG** As used in ROSE, this is the Control Flow Graph, not Context Free Grammar or anything else.
- **EDG** Edison Design Group (the commercial company that produces the C and C++ front-end that is used in ROSE).
- **IR** Intermediate Representation (IR). The IR is the set of classes defined within SAGE III that allow an AST to be built to define any application in C, C++, and Fortran application.
- **Query** (as in AST Query) Operations on the AST that return answers to questions posed about the content or context in the AST.
- **ROSE** A project that covers both research in optimization and a specific infrastructure for handling large scale C, C++, and Fortran applications.
- **Rosetta** A tool (written by the ROSE team) used within ROSE to automate the generation of the SAGE III IR.
- **SAGE++ and SAGE II** An older object-oriented IR upon which the API of SAGE III IR is based.
- **Semantic Information** What abstractions mean (short answer). (This might be better as a description of what kind of semantic information ROSE could take advantage, not a definition.)

- **Telescoping Languages** A research area that defines a process to generate domain-specific languages from a general purpose languages.
- **Transformation** The process of automating the editing (either reconfiguration, addition, or deletion; or some combination) of input application parts to build a new application. In the context of ROSE, all transformations are source-to-source.
- **Translator** An executable program (in our context built using ROSE) that performs source-to-source translation on an existing input application source to generate a second (generated) source code file. The second (generated) source code is then typically provided as input to a vendor provided compiler (which generates object code or an executable program).
- **Traversal** The process of operating on the AST in some order (usually pre-order, post-order, out of order [randomly], depending on the traversal that is used). The ROSE user builds a traversal from base classes that do the traversal and execute a function, or a number of functions, provided by the user.